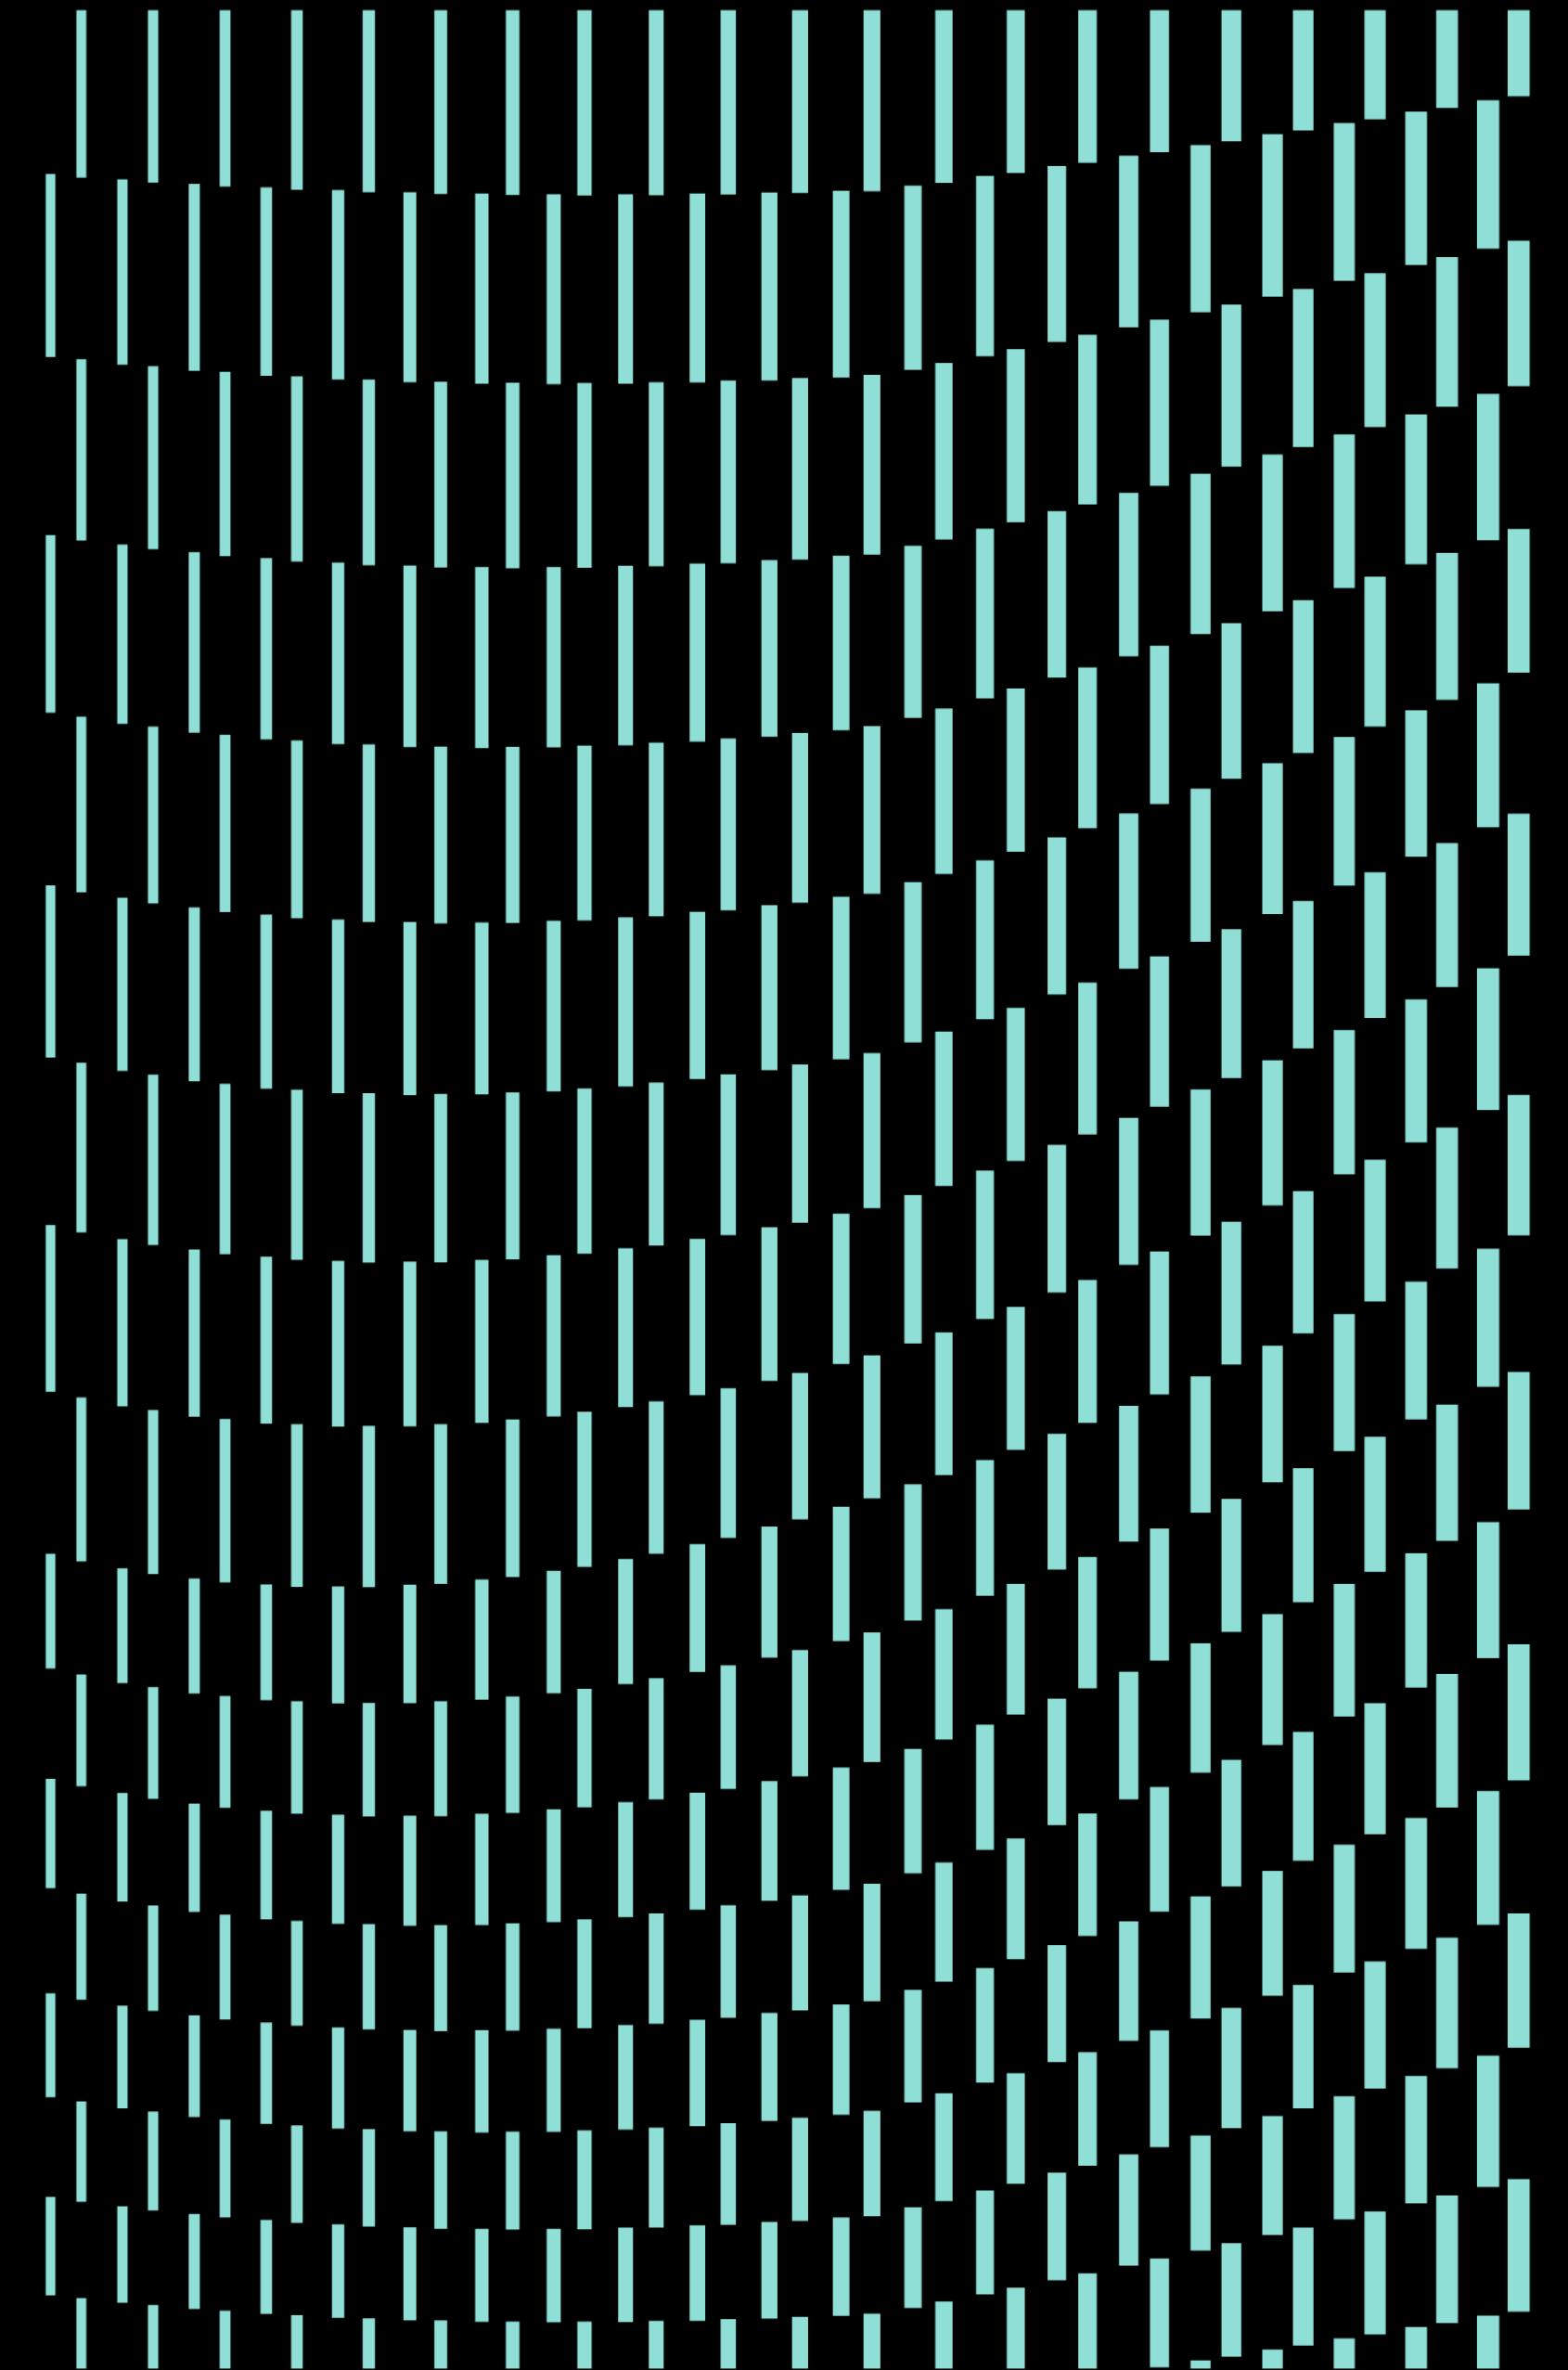


# Deep Dive into Spark

Eric Xiao, Data Developer | Storage and Query Technologies



# Motivation

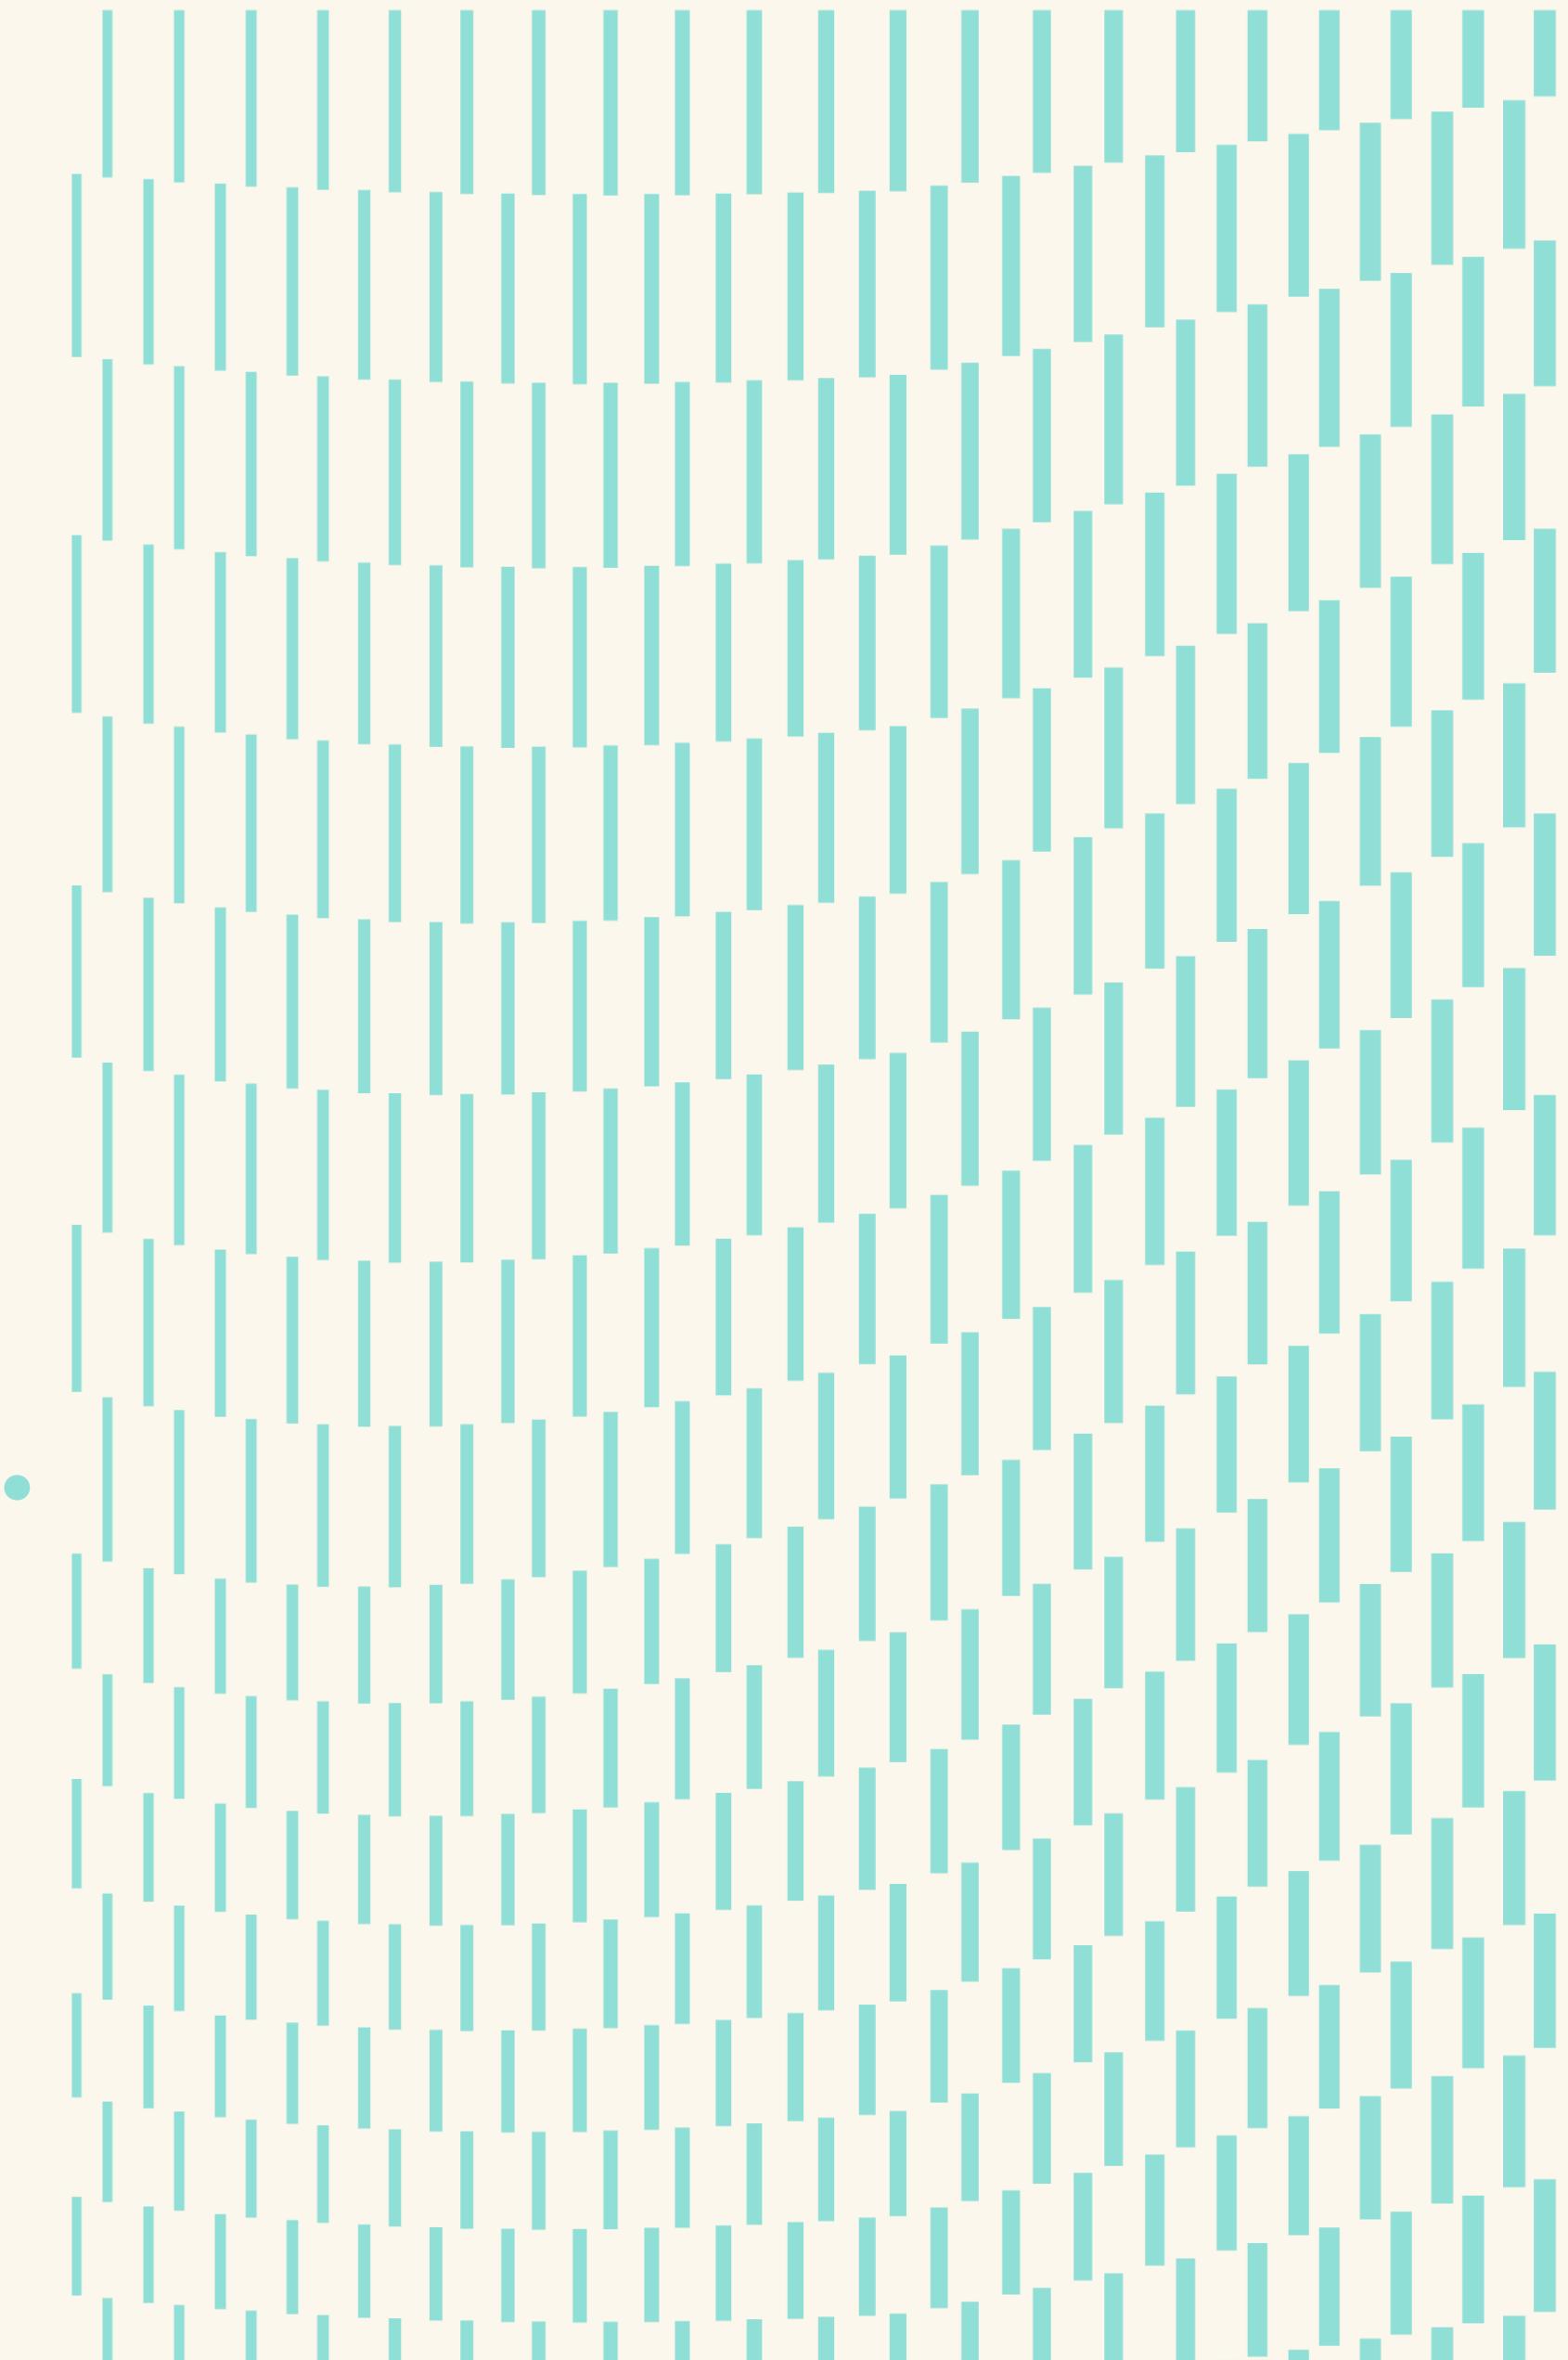
We'll be learning how to analyze a spark application through the sparkUI.

Then implement 2 optimizations that will address 2 main problems that are common in spark applications.



## Overview:

1. **Concrete Example Query.**
2. **From Code to Execution.**
3. **Optimization 1: Improving joins.**
4. **Optimization 2: Eliminate data spill.**
5. **Recap.**
6. **Questions.**



# **Section 1:**

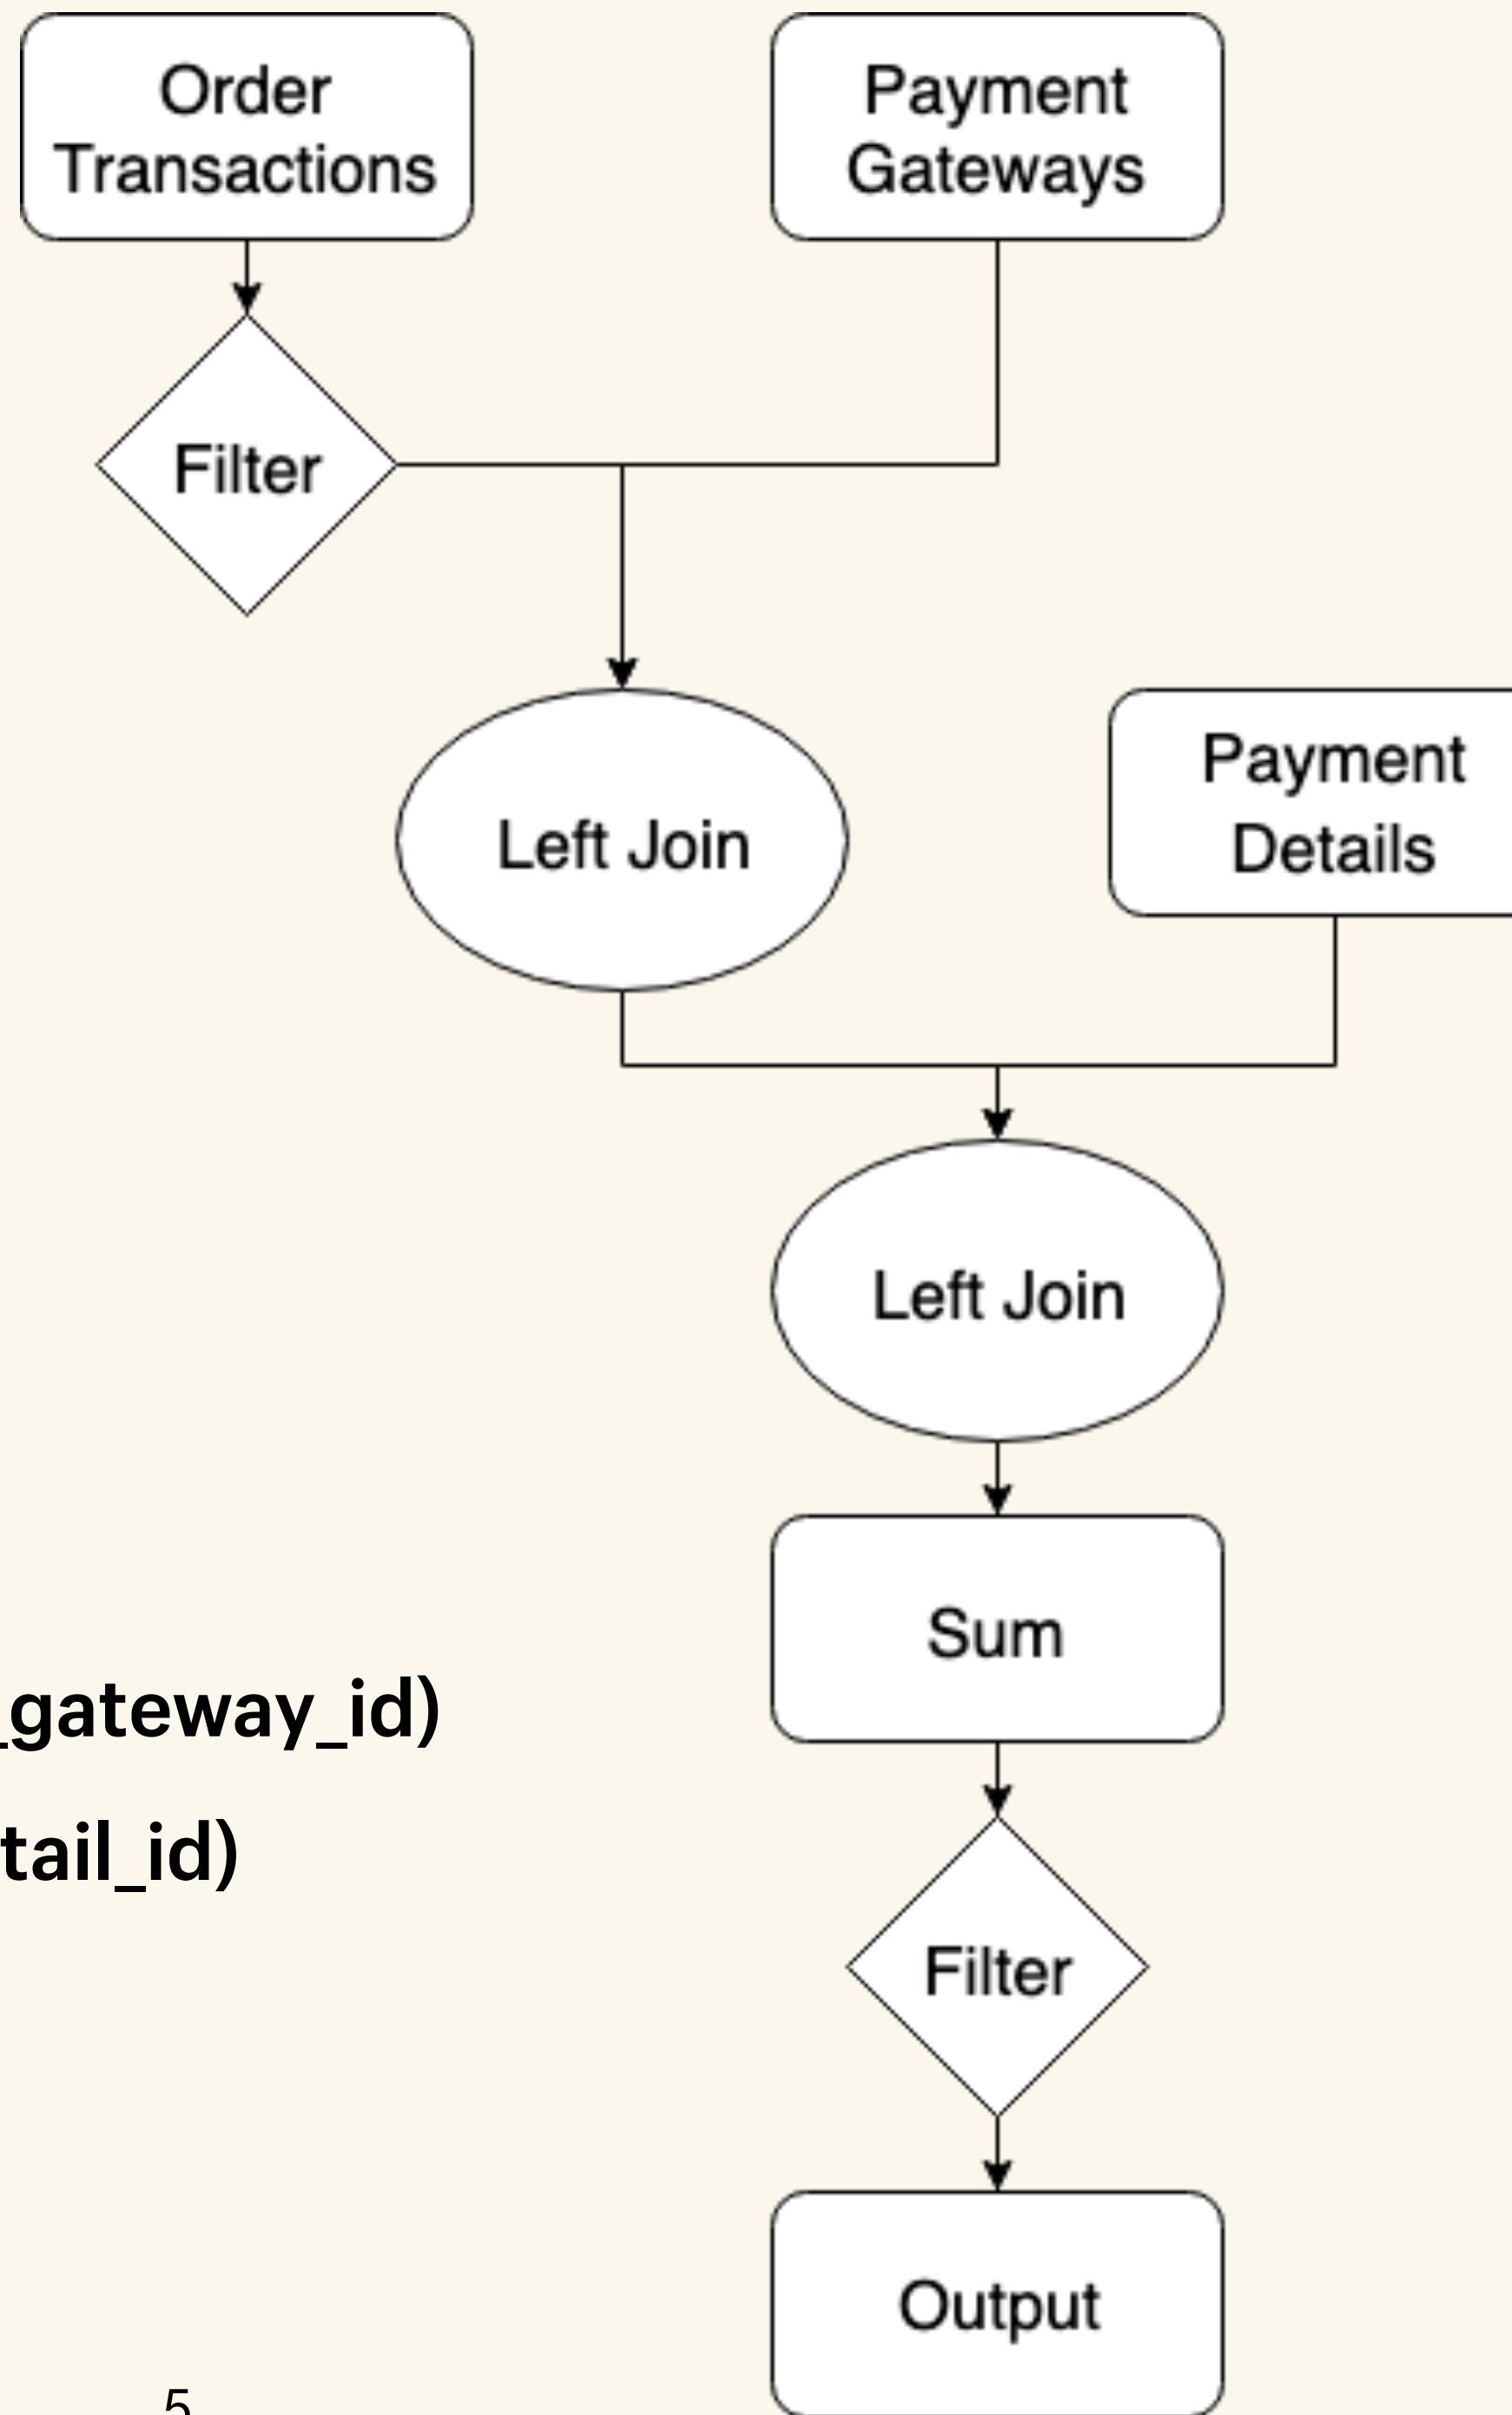
## **Concrete Query**

### **Example**

# Example Query

I want to know the count of order transactions per credit card types.

```
SELECT  
    card_type,  
    count(*)  
FROM order_transactions  
LEFT JOIN payment_gateways USING (payment_gateway_id)  
LEFT JOIN payment_details USING (payment_detail_id)  
GROUP BY 1  
WHERE ....
```



# **Section 2:** **From Code to** **Execution**

**Lazy Evaluation**

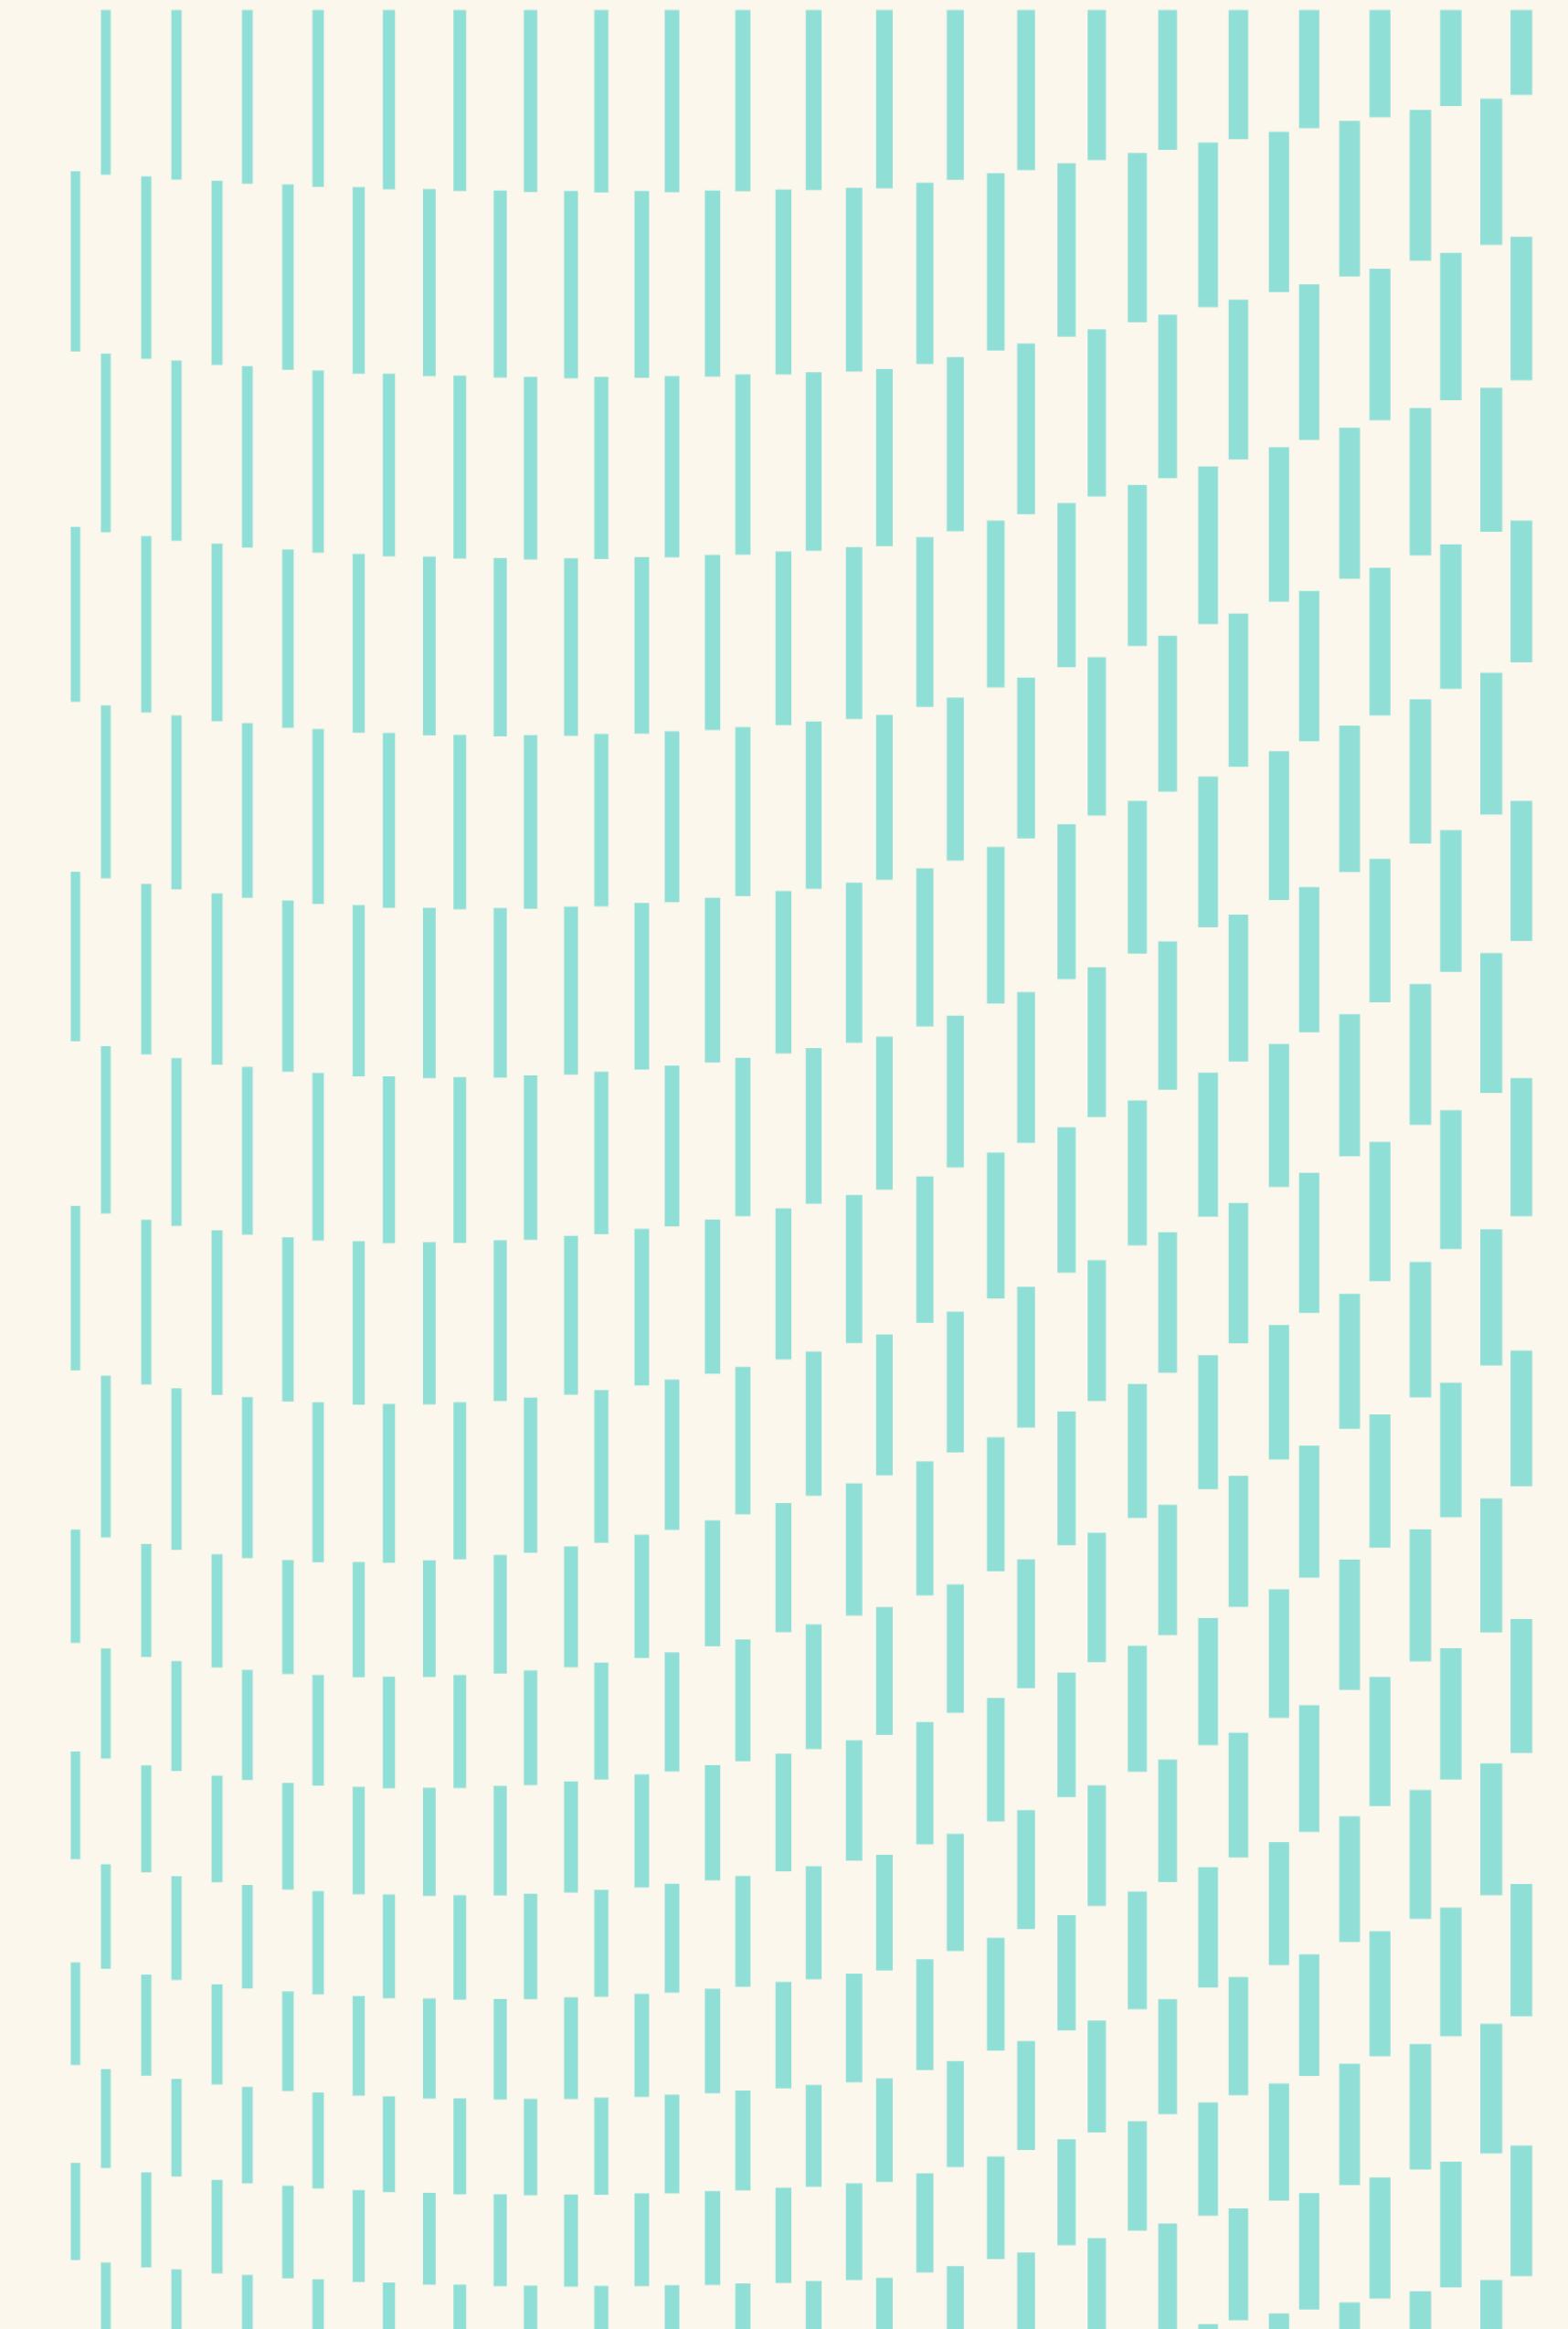
**Actions**

**Spark Plans**

**Logical Plan**

**Optimized Logical Plan**

**Physical Plans**



# Lazy Evaluation

Spark doesn't perform any transformations until an “action” is called.

```
1 df
2 > DataFrame[card_type: string, count:
3 bigint]
4
5
6
7
```

# Actions

An “action” triggers Spark execution.

```
1 df.write()  
2  
3 df.head(n)  
4 df.take(n)  
5  
6 df.collect()  
7 df.show()  
df.toPandas()
```

# Spark Plans

`df.explain(True)`

- Spark generates different execution plans from the spark DataFrames/RDD code:
  - Logical Plan
  - Optimized Logical Plan
  - Physical Plan
- Read bottom up, opposite of the plan in SparkUI



# Spark Plans

**== Physical Plan ==**

```
*(10) HashAggregate(keys=[card_type#229], functions=[sum(1)], output=[card_type#229, count#289L])
```

```
+ Exchange hashpartitioning(card_type#229, 200)
```

```
  +- *(9) HashAggregate(keys=[card_type#229], functions=[partial_sum(1)], output=[card_type#229, sum#296L])
```

```
  +- *(9) Project [card_type#229]
```

```
    +- *(9) SortMergeJoin [order_transaction_id#202L], [order_transaction_id#160L], Inner
```

```
      :- *(6) Sort [order_transaction_id#202L ASC NULLS FIRST], false, 0
```

...

# Logical Plan

- Is a set of abstract expressions that represent the spark code.



Order  
Transactions

Payment  
Gateways

Filter

Left Join

Payment  
Details

Left Join

Sum

Filter

Output

**SQL QUERY**

Relation

Relation

Filter

Join LeftOuter

Relation

Join LeftOuter

Aggregate

Filter

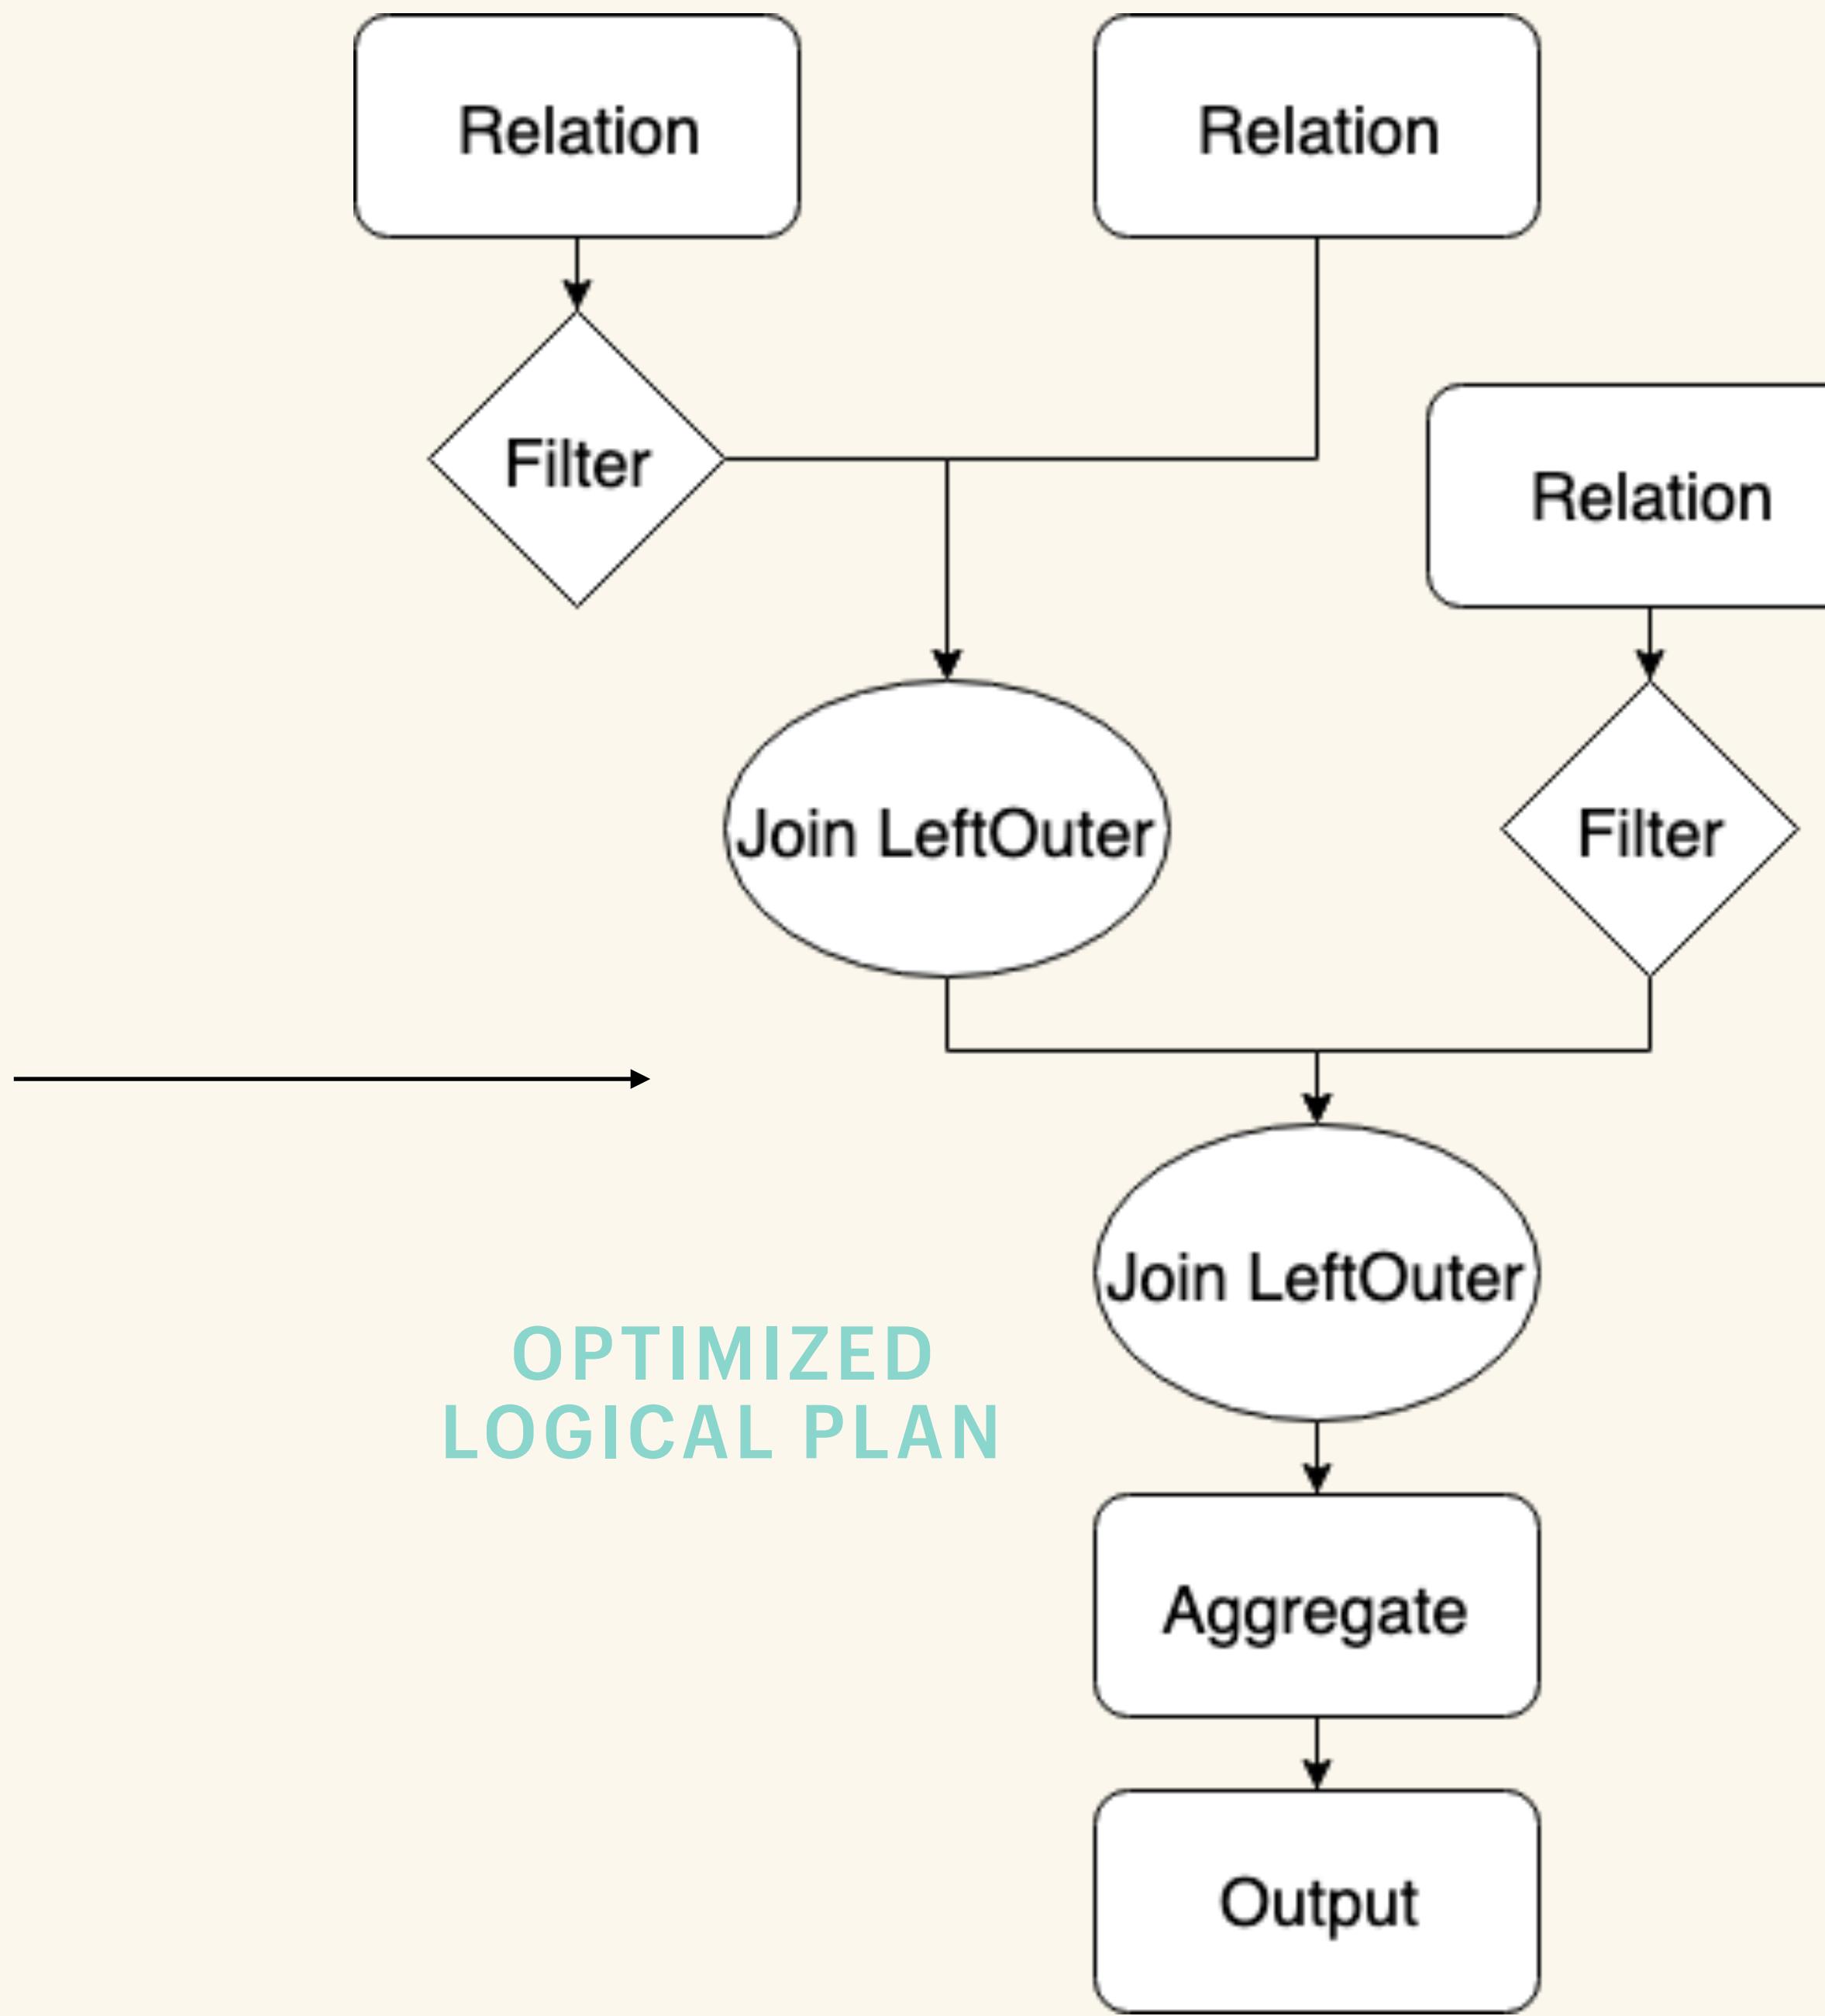
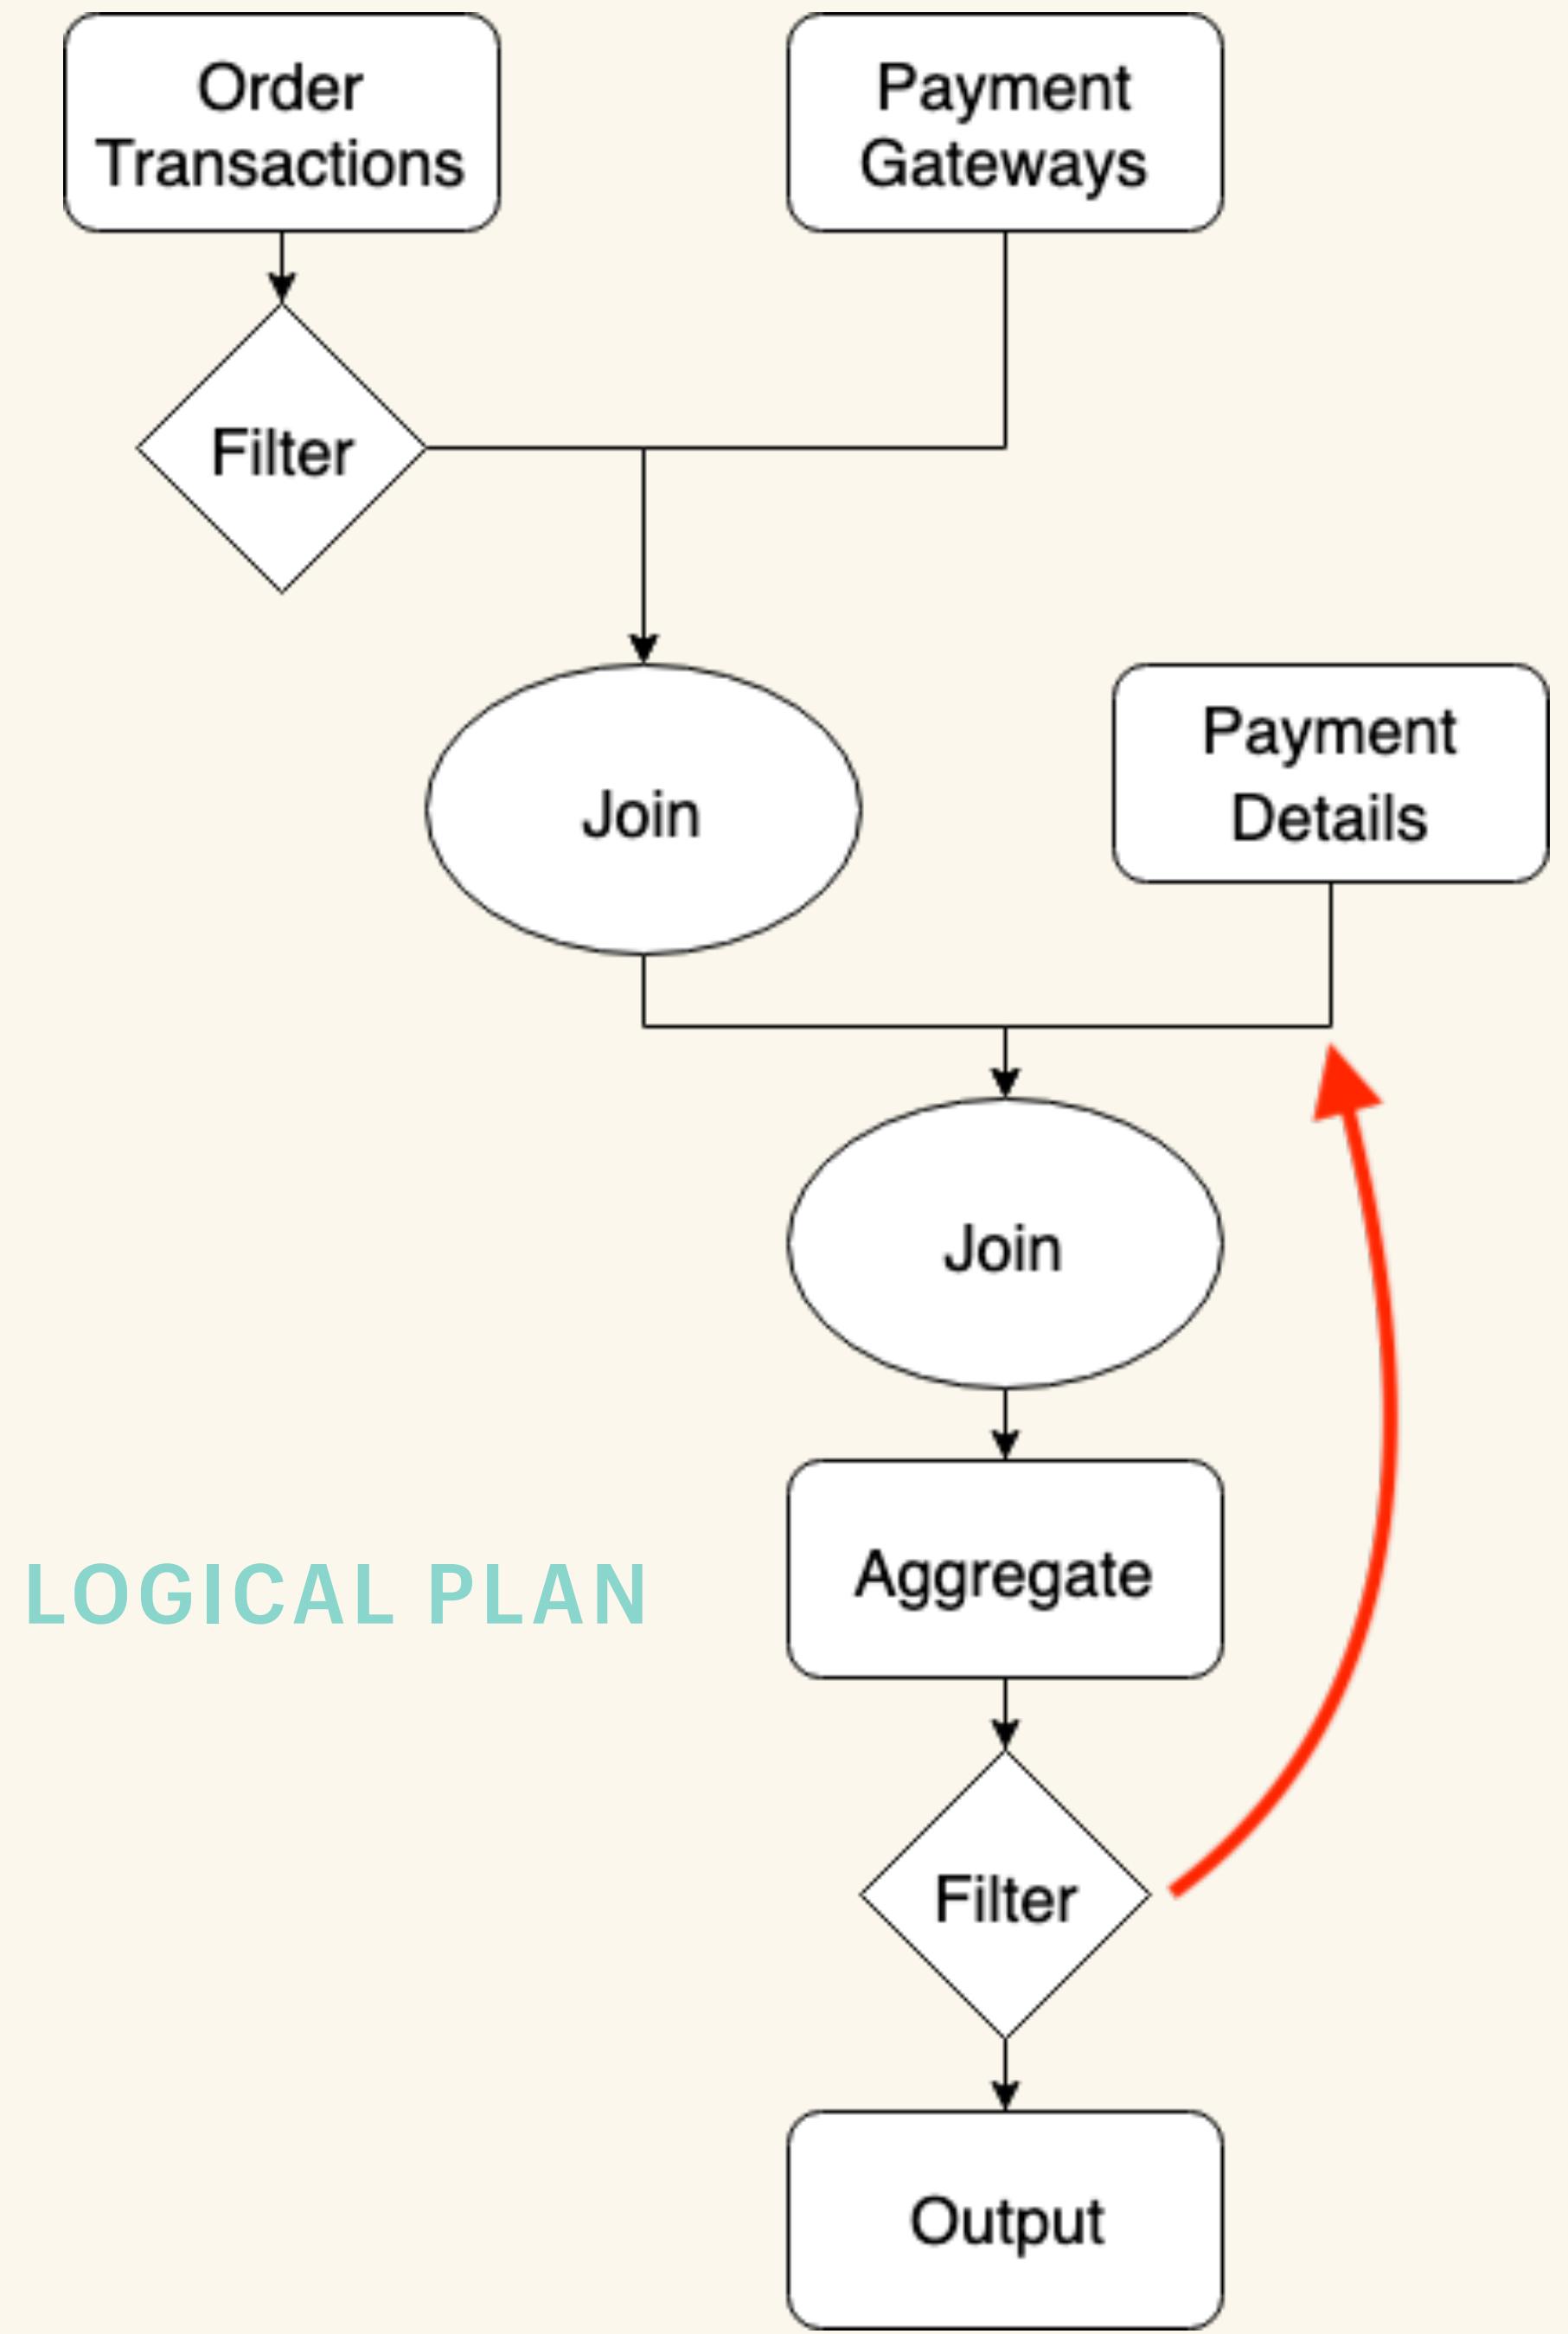
Output

**LOGICAL PLAN**

# Optimized Logical Plan

- Checks if set of expressions are valid.
  - ie. tables and columns exists.
- If valid, plan is passed to the Spark Catalyst Optimizer to be optimized.





# Physical Plan

- Specifies exactly how the logical plan will execute.
- Multiple plans are generated and most optimal is selected.
- Based on the physical attributes of the tables.
  - ie. table size, and partition size.



Relation

Relation

Scan Parquet  
Scan Filter

Scan Parquet  
Scan Filter

Exchange  
Sort

Exchange  
Sort

Filter

Relation

Filter

Join LeftOuter

Filter

SortMergeJoin

Scan Parquet  
Scan Filter

Exchange  
Sort

Exchange  
Sort

Join LeftOuter

PHYSICAL PLAN

OPTIMIZED  
LOGICAL PLAN

Aggregate

SortMergeJoin  
PartialAggregate

Output

HashAggregate

Relation

Relation

Scan Parquet  
Scan Filter

Scan Parquet  
Scan Filter

Exchange  
Sort

Exchange  
Sort

Filter

Relation

SortMergeJoin

Scan Parquet  
Scan Filter

Exchange  
Sort

Join LeftOuter

Filter

Exchange  
Sort

Join LeftOuter

Aggregate

SortMergeJoin  
PartialAggregate

HashAggregate

OPTIMIZED  
LOGICAL PLAN

PHYSICAL PLAN

Output

Relation

Relation

Scan Parquet  
Scan Filter

Scan Parquet  
Scan Filter

Exchange  
Sort

Exchange  
Sort

Filter

Relation

SortMergeJoin

Scan Parquet  
Scan Filter

Exchange  
Sort

Join LeftOuter

Filter

Exchange  
Sort

Join LeftOuter

Aggregate

SortMergeJoin  
PartialAggregate

HashAggregate

OPTIMIZED  
LOGICAL PLAN

PHYSICAL PLAN

Output

Relation

Relation

Scan Parquet  
Scan Filter

Scan Parquet  
Scan Filter

Exchange  
Sort

Exchange  
Sort

Filter

Relation

SortMergeJoin

Scan Parquet  
Scan Filter

Exchange  
Sort

Join LeftOuter

Filter

Exchange  
Sort

Join LeftOuter

PHYSICAL PLAN

OPTIMIZED  
LOGICAL PLAN

Filter

SortMergeJoin  
PartialAggregate

HashAggregate

Aggregate

Output

Relation

Relation

Scan Parquet  
Scan Filter

Scan Parquet  
Scan Filter

Exchange  
Sort

Exchange  
Sort

Filter

Relation

SortMergeJoin

Scan Parquet  
Scan Filter

Exchange  
Sort

Join LeftOuter

Filter

Exchange  
Sort

OPTIMIZED  
LOGICAL PLAN

PHYSICAL PLAN

Join LeftOuter

Aggregate

SortMergeJoin  
PartialAggregate

HashAggregate

Output

Relation

Relation

Scan Parquet  
Scan Filter

Scan Parquet  
Scan Filter

Exchange  
Sort

Exchange  
Sort

Filter

Relation

SortMergeJoin

Scan Parquet  
Scan Filter

Exchange  
Sort

Filter

Join LeftOuter

Exchange  
Sort

Join LeftOuter

PHYSICAL PLAN

OPTIMIZED  
LOGICAL PLAN

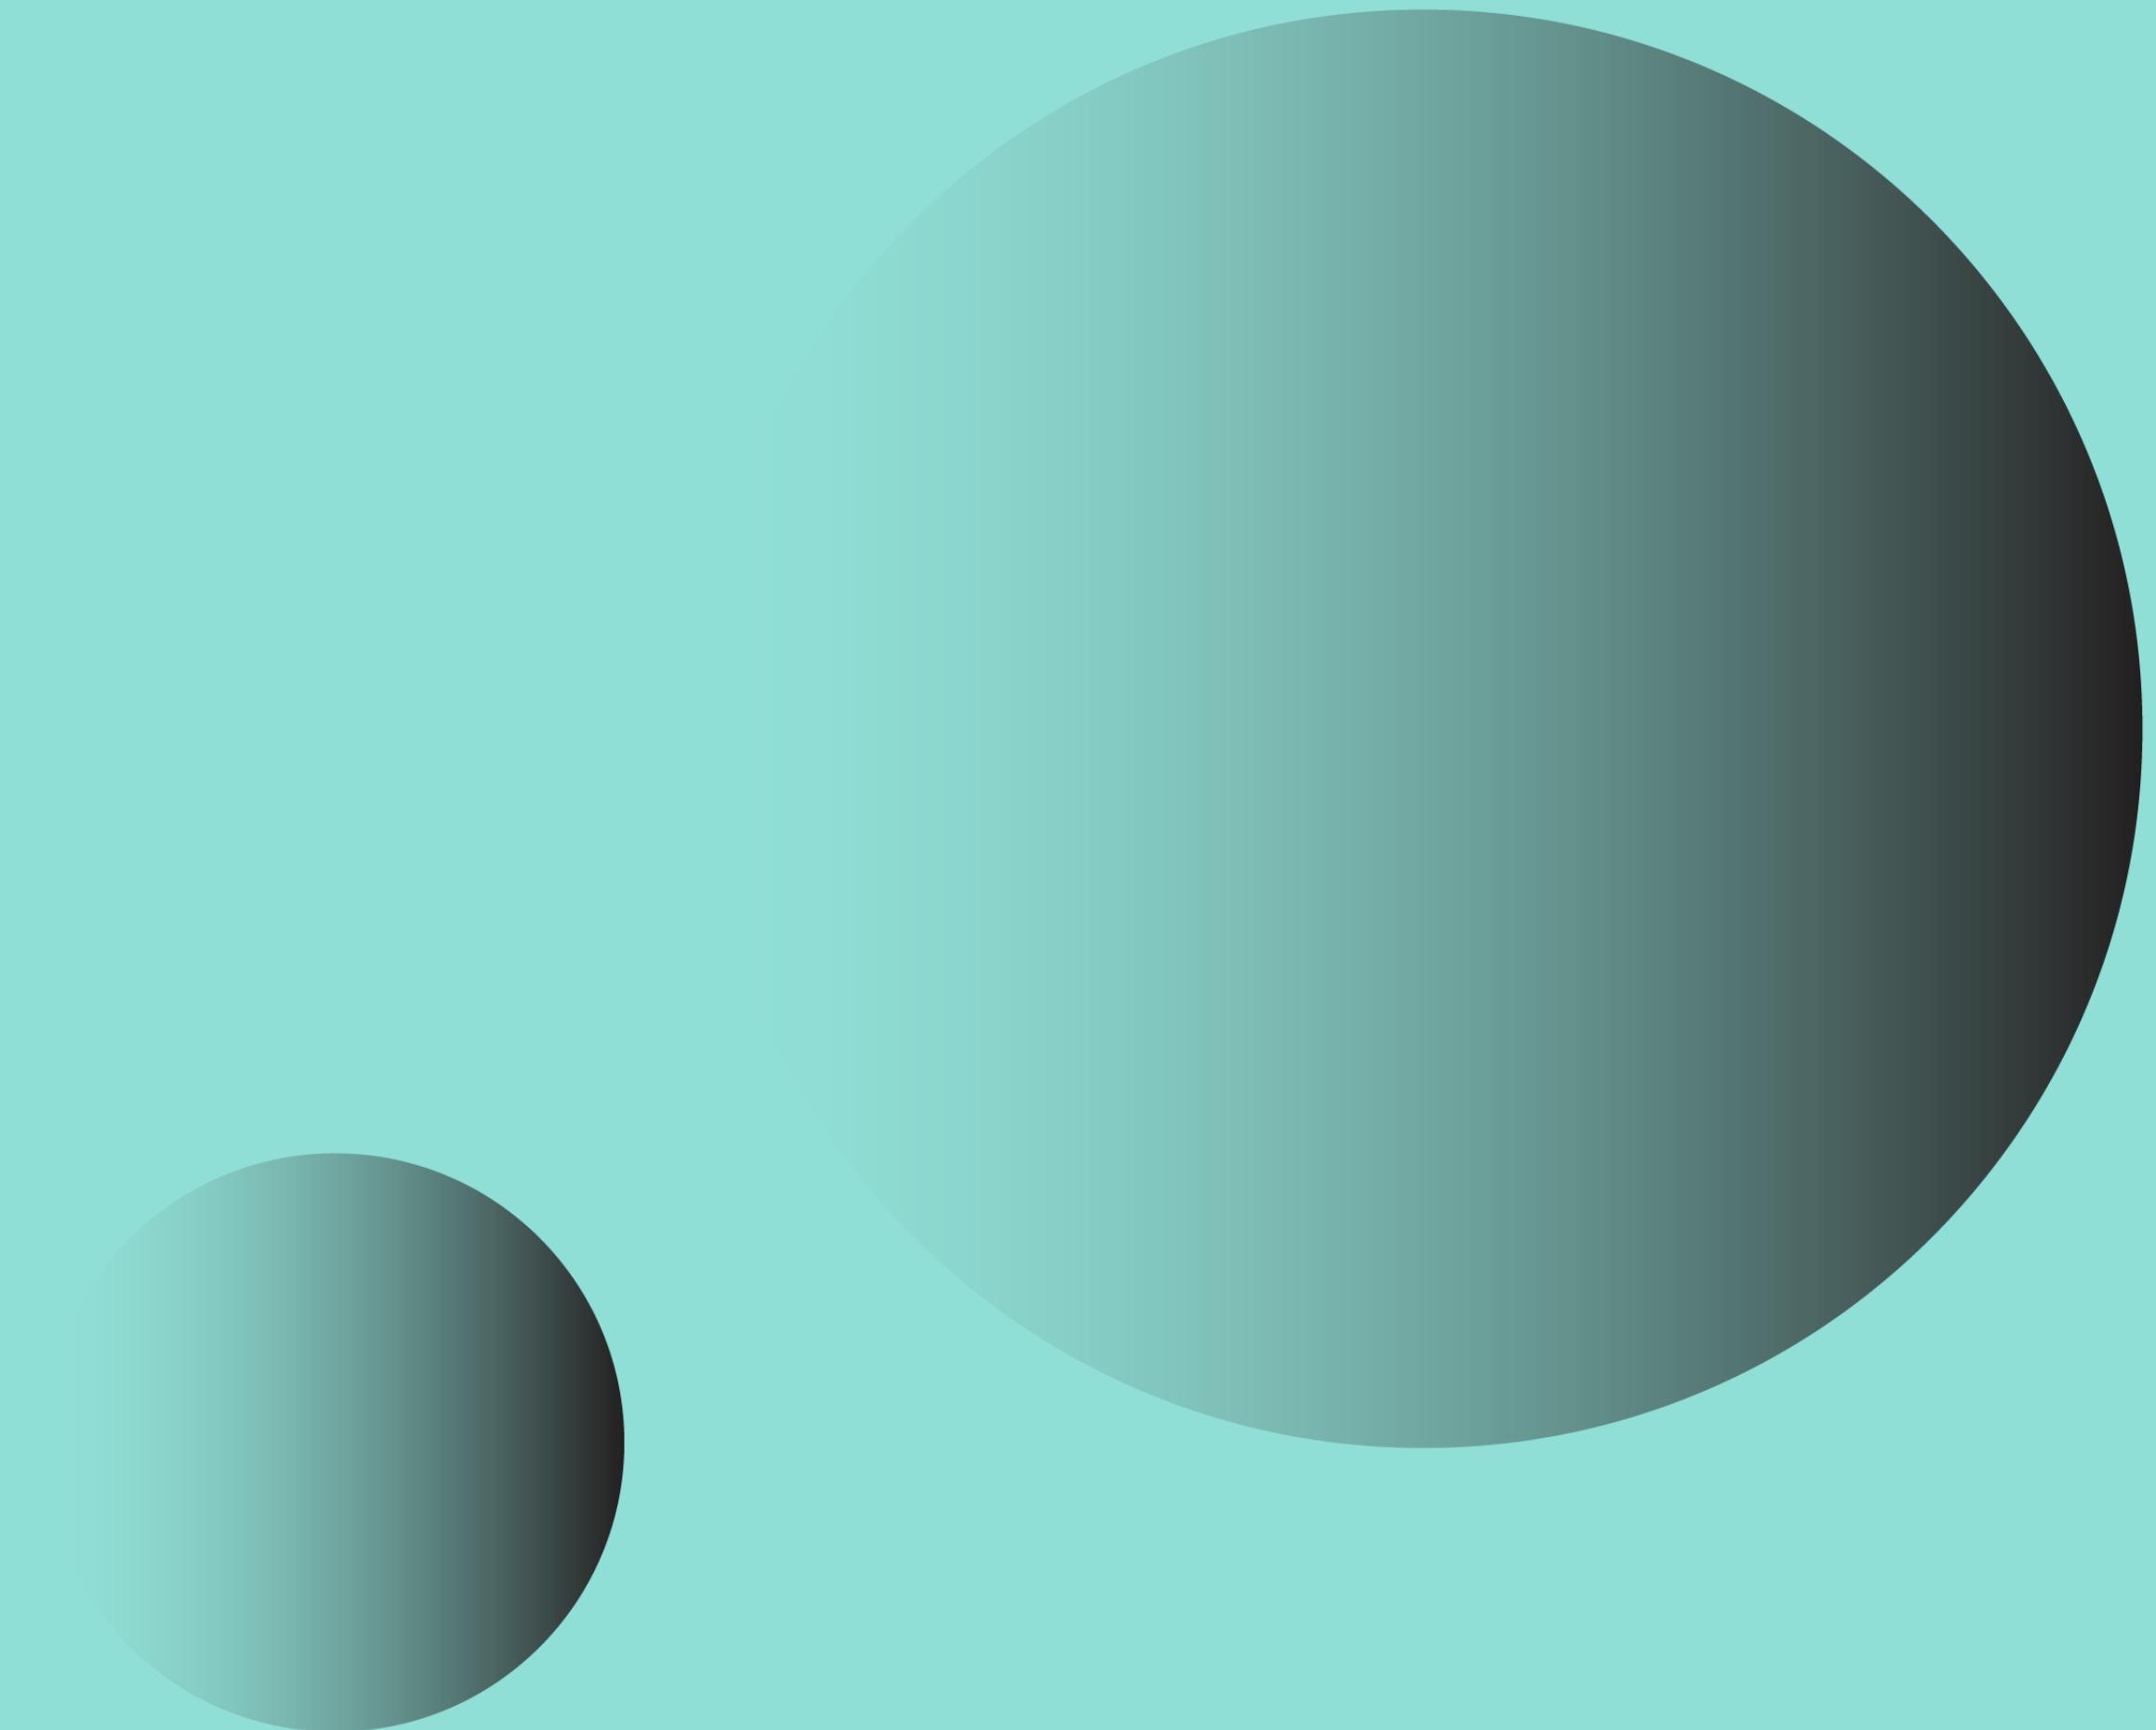
Aggregate

SortMergeJoin  
PartialAggregate

HashAggregate

Output

# **Section 3: Optimization 1 - Improving Joins**



**SparkUI**

**Spark DAG**

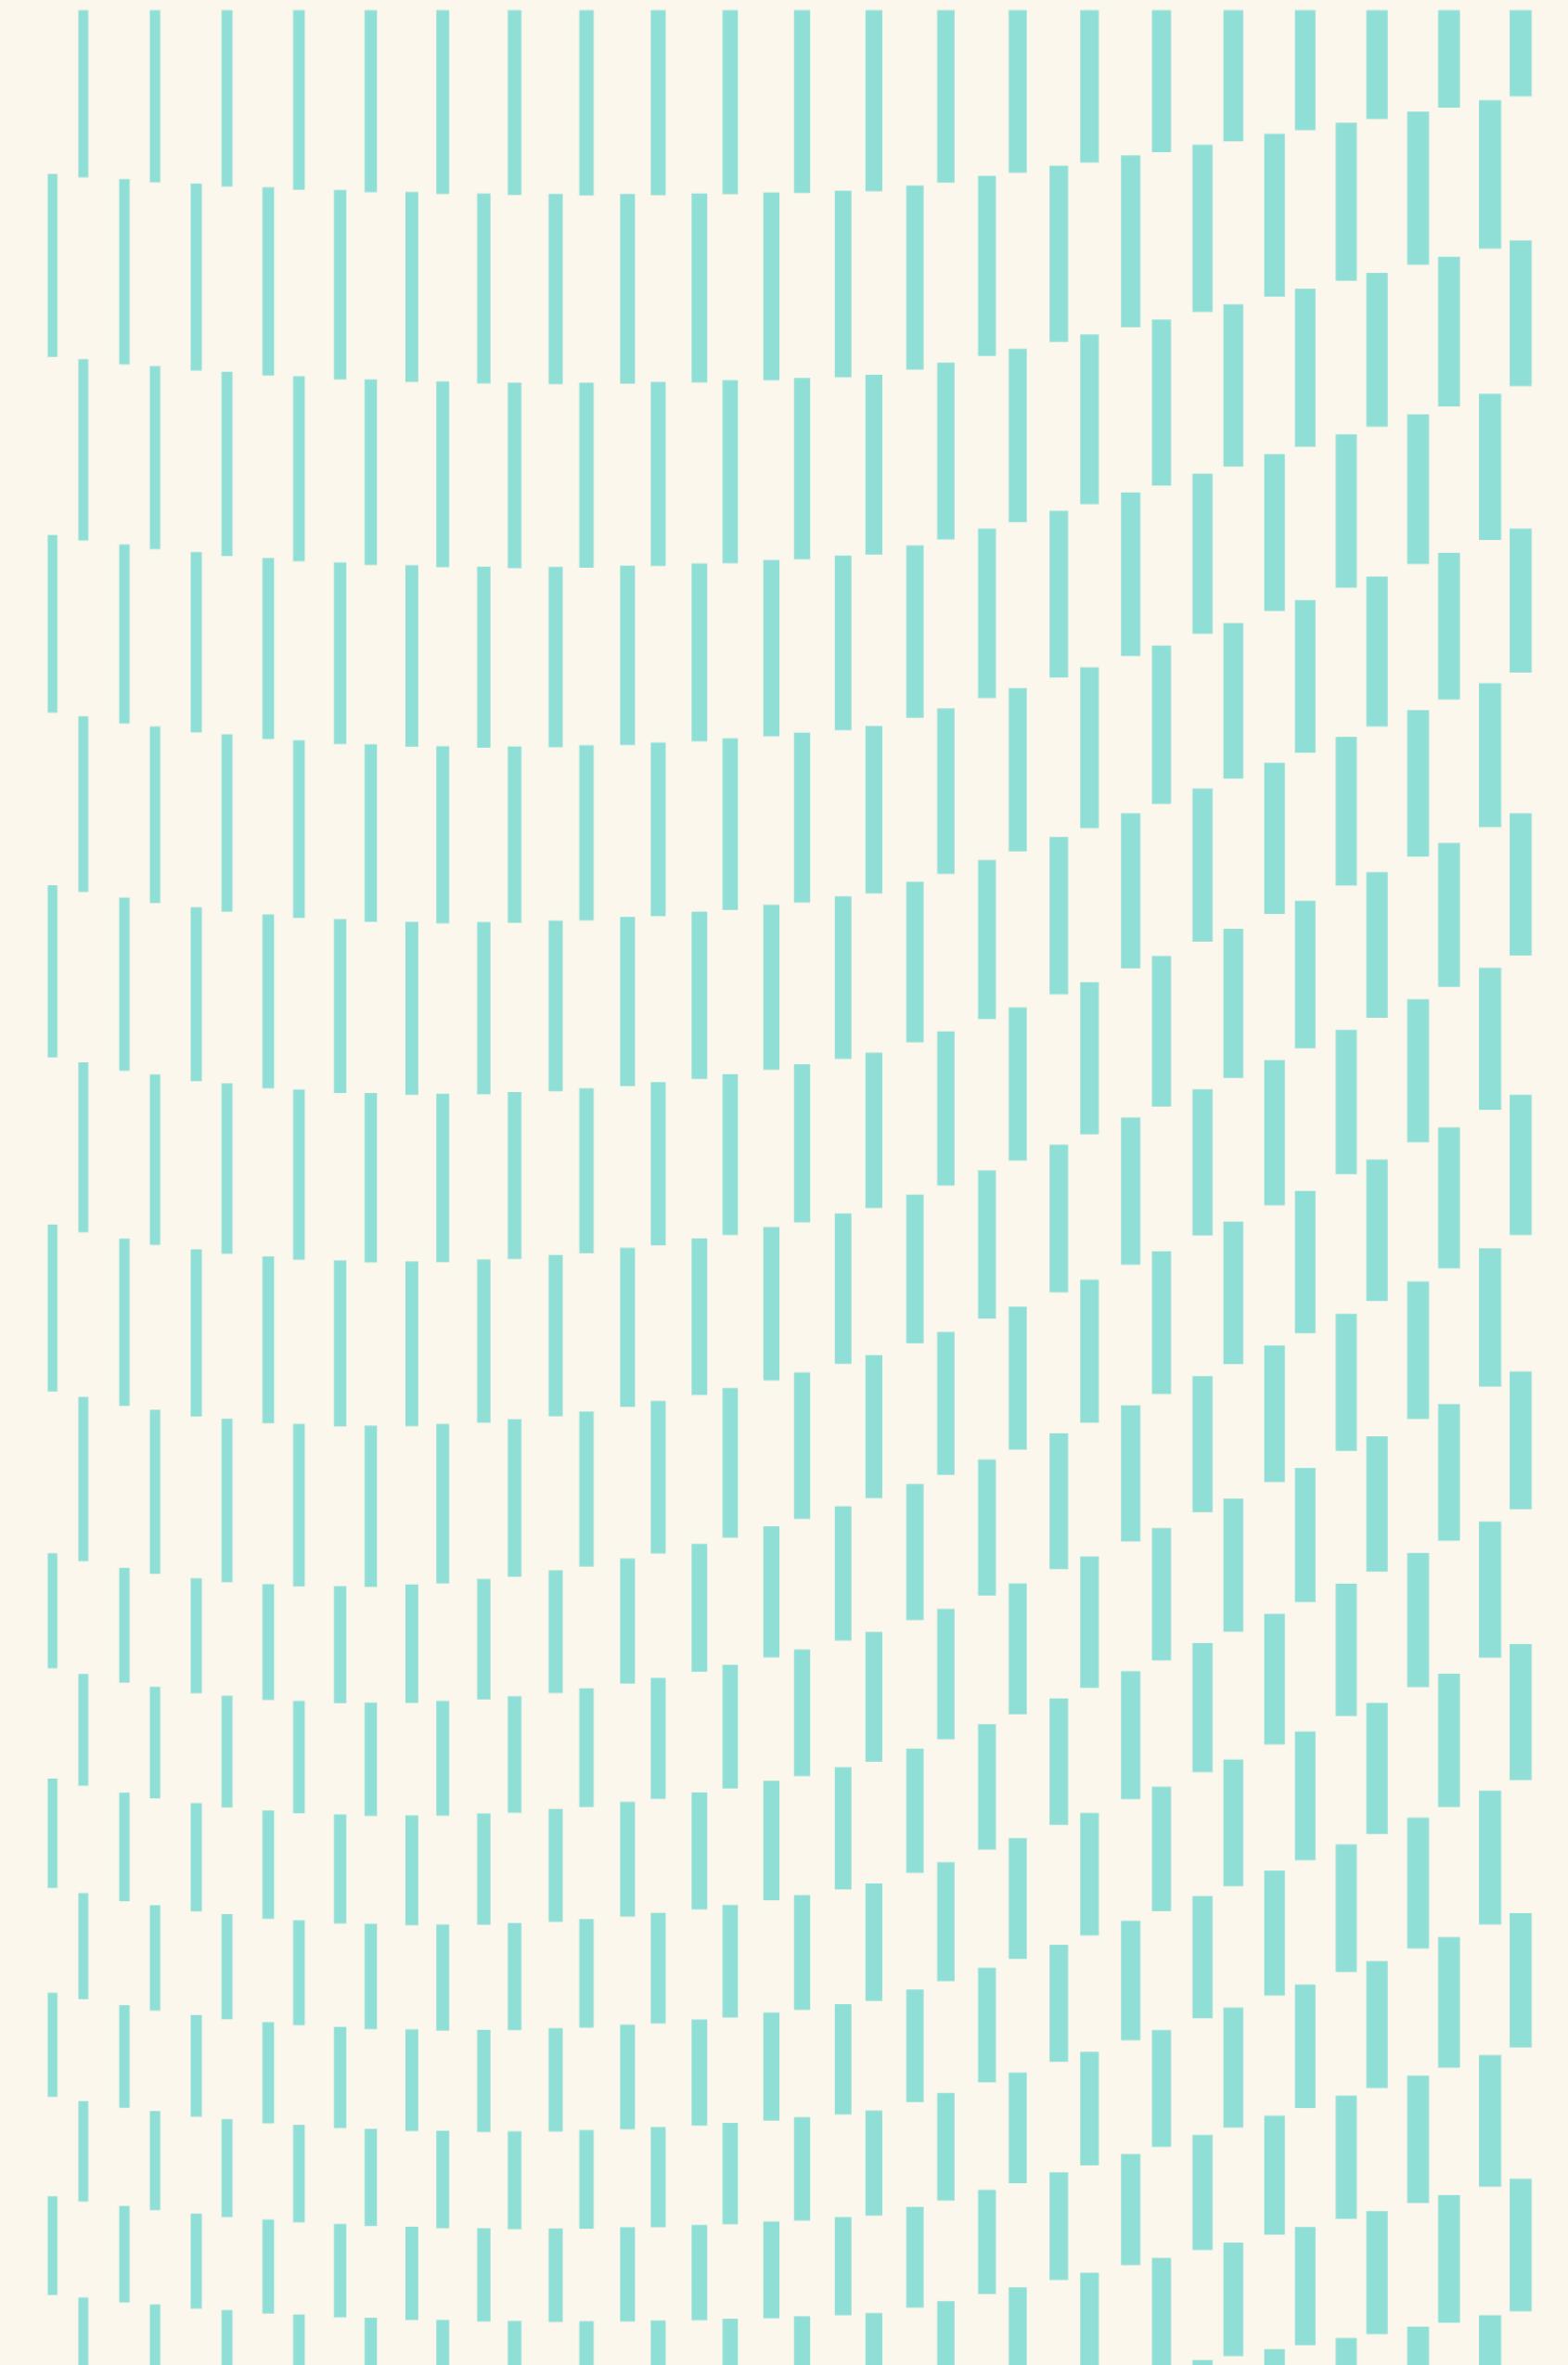
**Shuffle**

**Broadcasting**

**Spark SQL tab**

**Join Skew**

**Join Strategies**



## Spark Jobs [\(?\)](#)

User: ericxiao

Total Uptime: 6.4 h

Scheduling Mode: FIFO

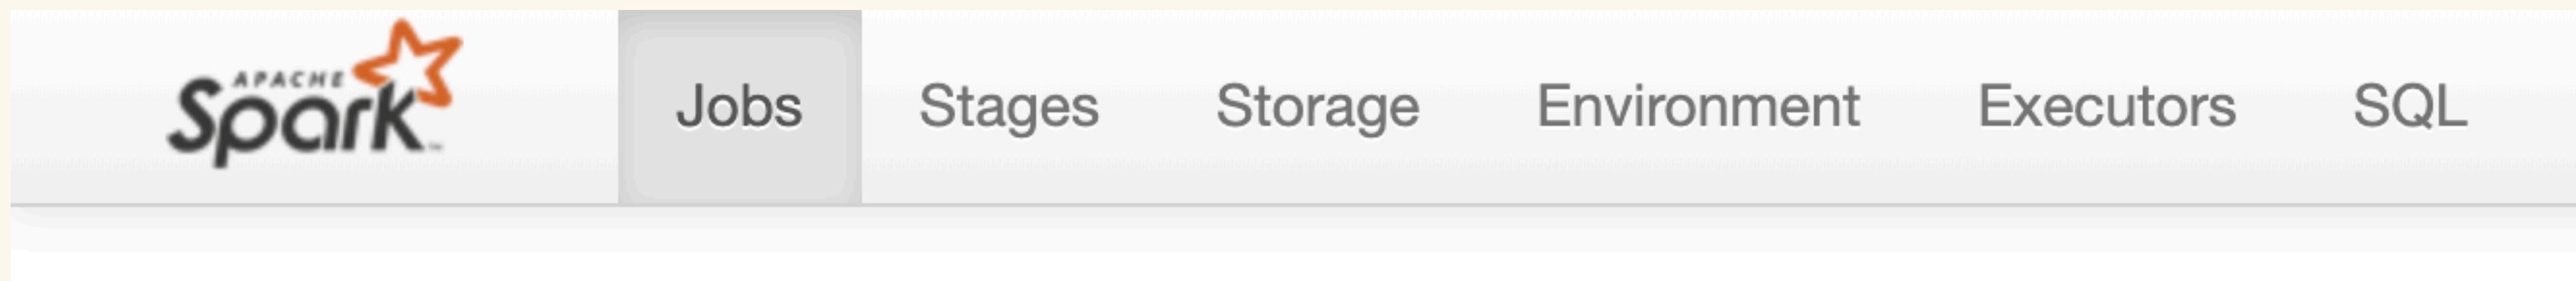
Completed Jobs: 6

▶ Event Timeline

### Completed Jobs (6)

Job Id ▾	Description	Submitted	Duration	Stages: Succeeded/Total	Tasks (for all stages): Succeeded/Total
5	showString at NativeMethodAccessorImpl.java:0 <a href="#">showString at NativeMethodAccessorImpl.java:0</a>	2019/07/25 14:44:53	7.9 min	6/6	31795/31795
4	showString at NativeMethodAccessorImpl.java:0 <a href="#">showString at NativeMethodAccessorImpl.java:0</a>	2019/07/25 14:36:31	8.4 min	6/6	31315/31315
3	showString at NativeMethodAccessorImpl.java:0 <a href="#">showString at NativeMethodAccessorImpl.java:0</a>	2019/07/25 14:19:18	17 min	6/6	16686/16686 (1 killed: another attempt succeeded)
2	parquet at NativeMethodAccessorImpl.java:0 <a href="#">parquet at NativeMethodAccessorImpl.java:0</a>	2019/07/25 14:19:15	0.2 s	1/1	1/1
1	parquet at NativeMethodAccessorImpl.java:0 <a href="#">parquet at NativeMethodAccessorImpl.java:0</a>	2019/07/25 14:19:11	0.1 s	1/1	1/1
0	parquet at NativeMethodAccessorImpl.java:0 <a href="#">parquet at NativeMethodAccessorImpl.java:0</a>	2019/07/25 14:18:47	21 s	1/1	1/1

# SparkUI Tabs

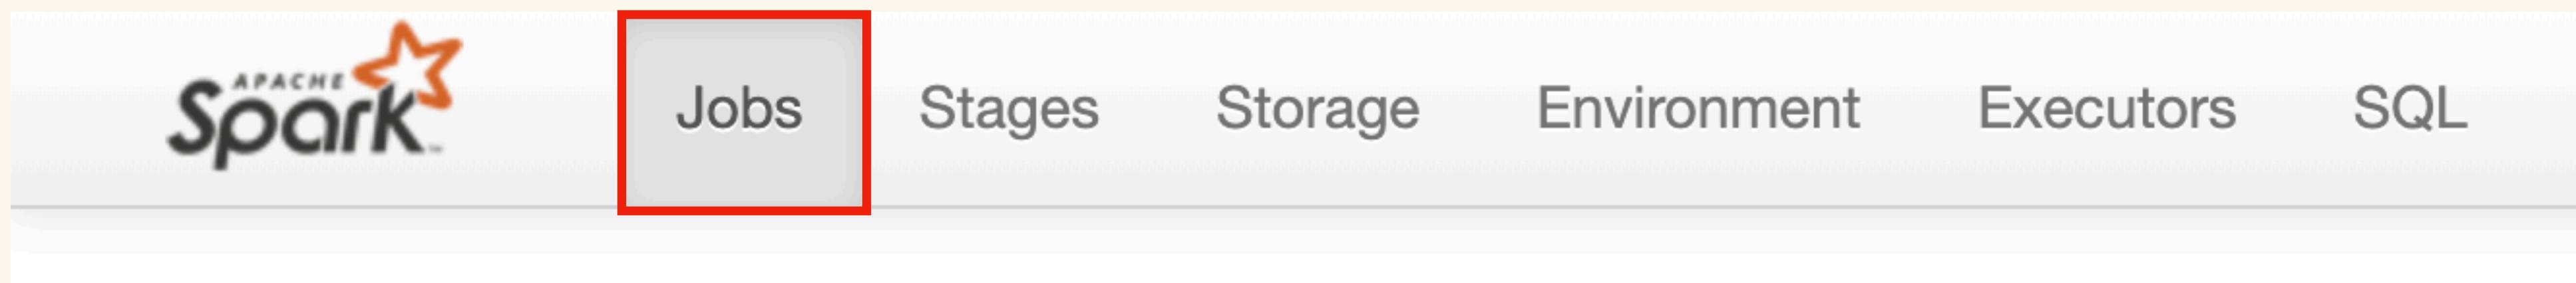


## Main tabs:

- **Jobs**
- **SQL**
- **Storage**



# SparkUI Tabs

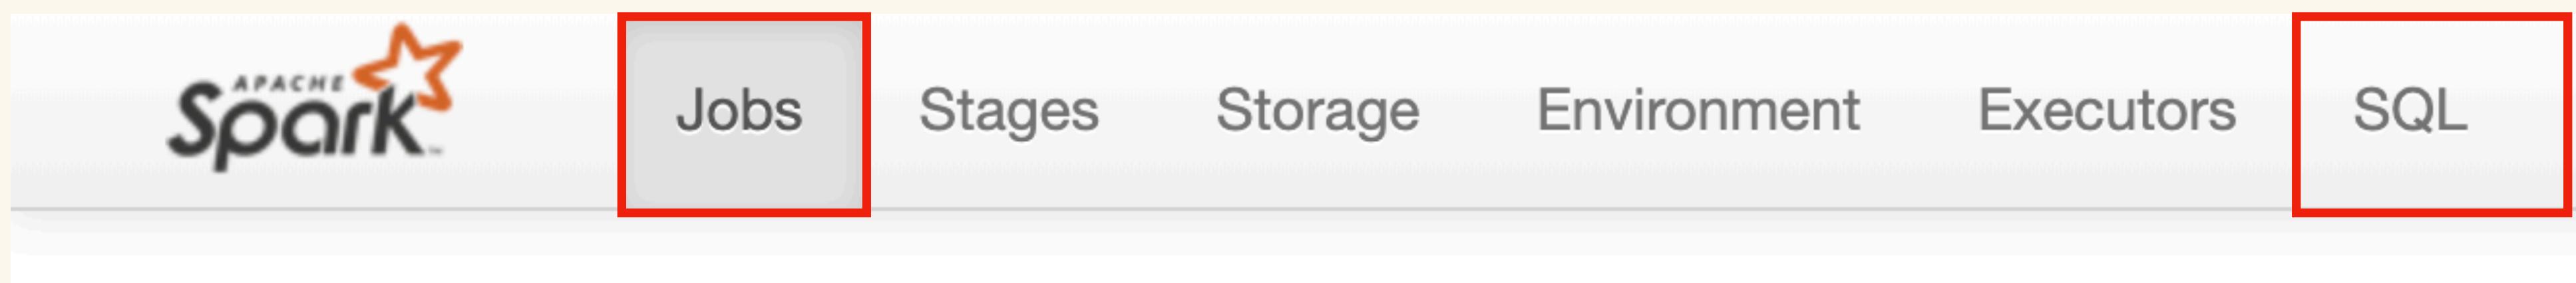


## Main tabs:

- **Jobs**
- **SQL**
- **Storage**



# SparkUI Tabs

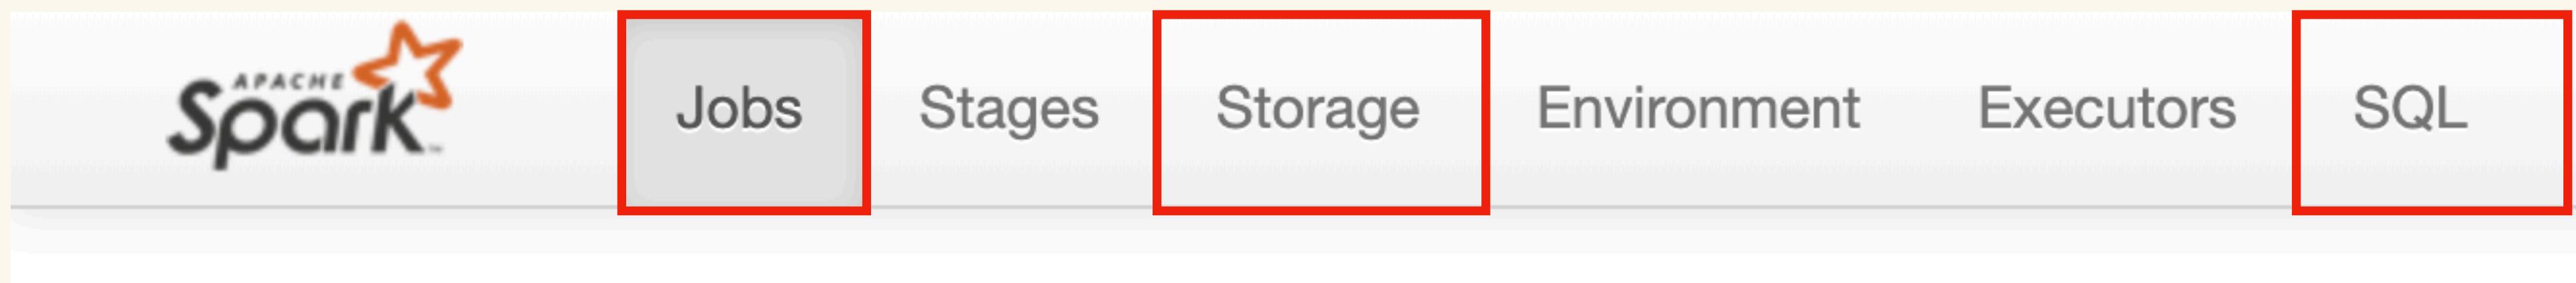


## Main tabs:

- **Jobs**
- **SQL**
- **Storage**



# SparkUI Tabs



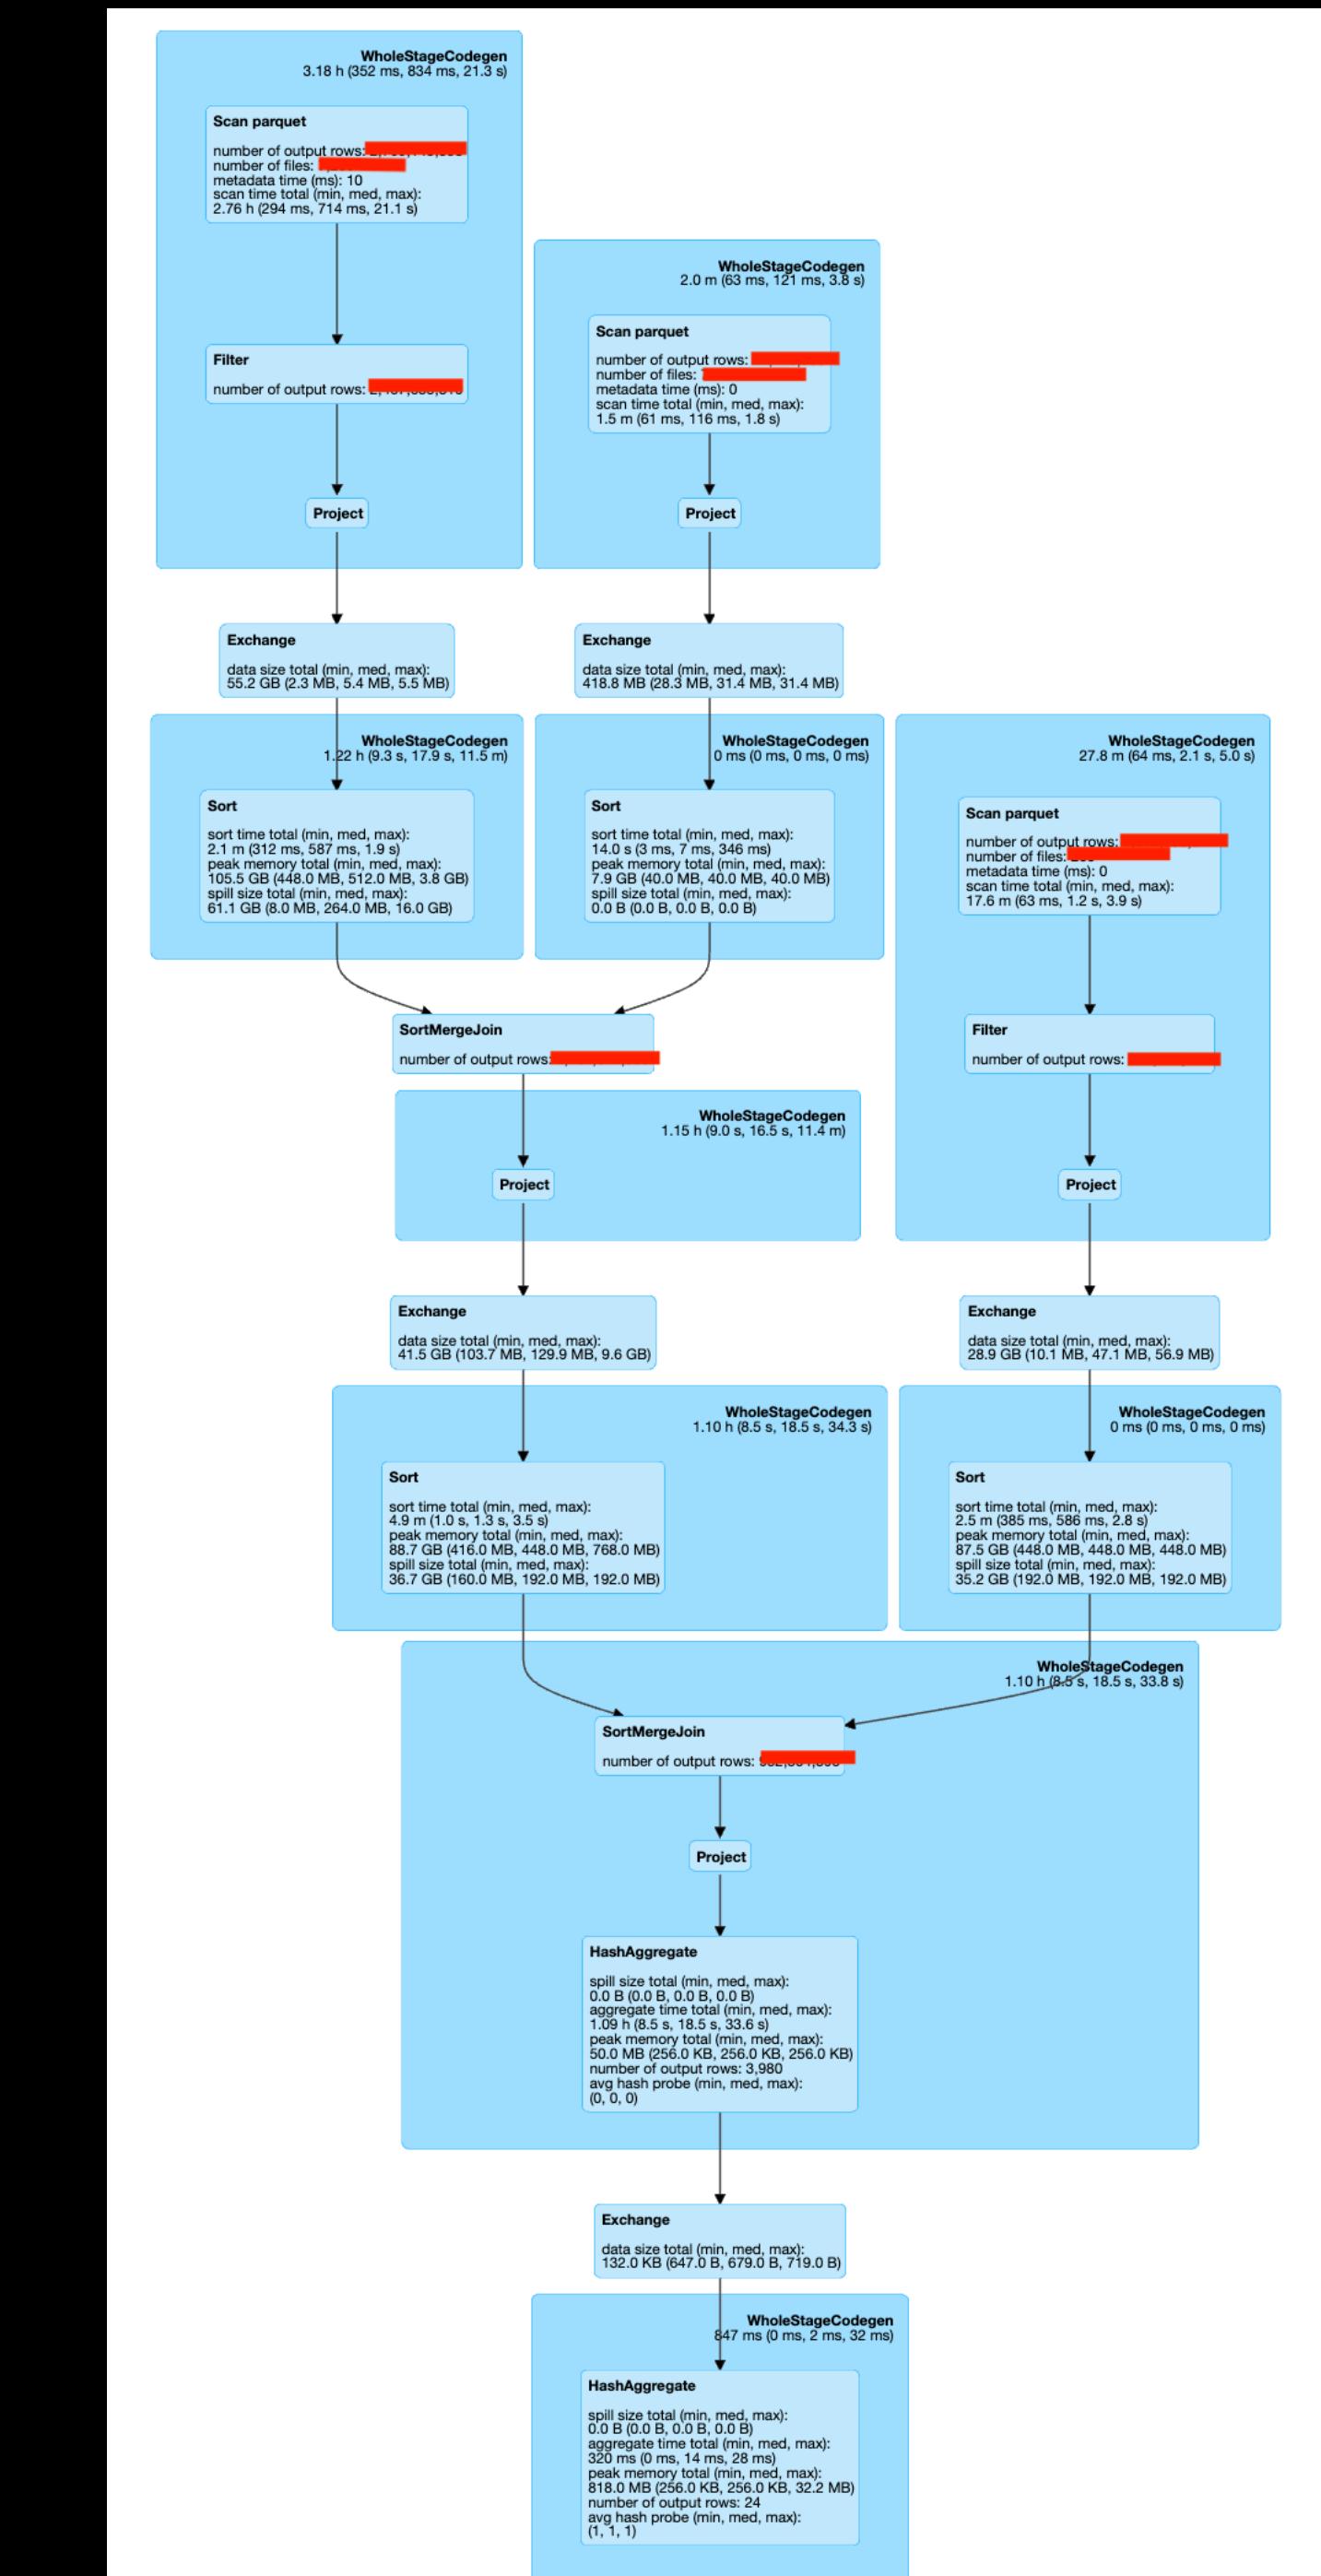
## Main tabs:

- **Jobs**
- **SQL**
- **Storage**



# Spark DAG

- Directed Acyclic Graph.
  - No cycles.
  - Data flows in one direction.
- A Spark DAG represents a Spark Job.
- A spark job consists of multiple stages.
- These stages can run in parallel.



# Shuffle

- Certain transformations in Spark trigger an event known as the Shuffle.
- A shuffle re-distributes data so that it's grouped differently across partitions.
- Why? Not all values needed for a transformation exist on the same partition or machine at the time of the transformation, but they must be co-located to perform the transformation.
- This involves copying data across executors and machines.
- Which is why shuffling data is complex and a costly operation.



reference: <http://people.apache.org/~pwendell/spark-nightly/spark-master-docs/latest/rdd-programming-guide.html#shuffle-operations>

# Shuffle

Operations that *might* cause a shuffle:

- **Joins**
- **GroupBy**
- **Distinct**
- **Repartition**
- **Coalesce**
- **Window**



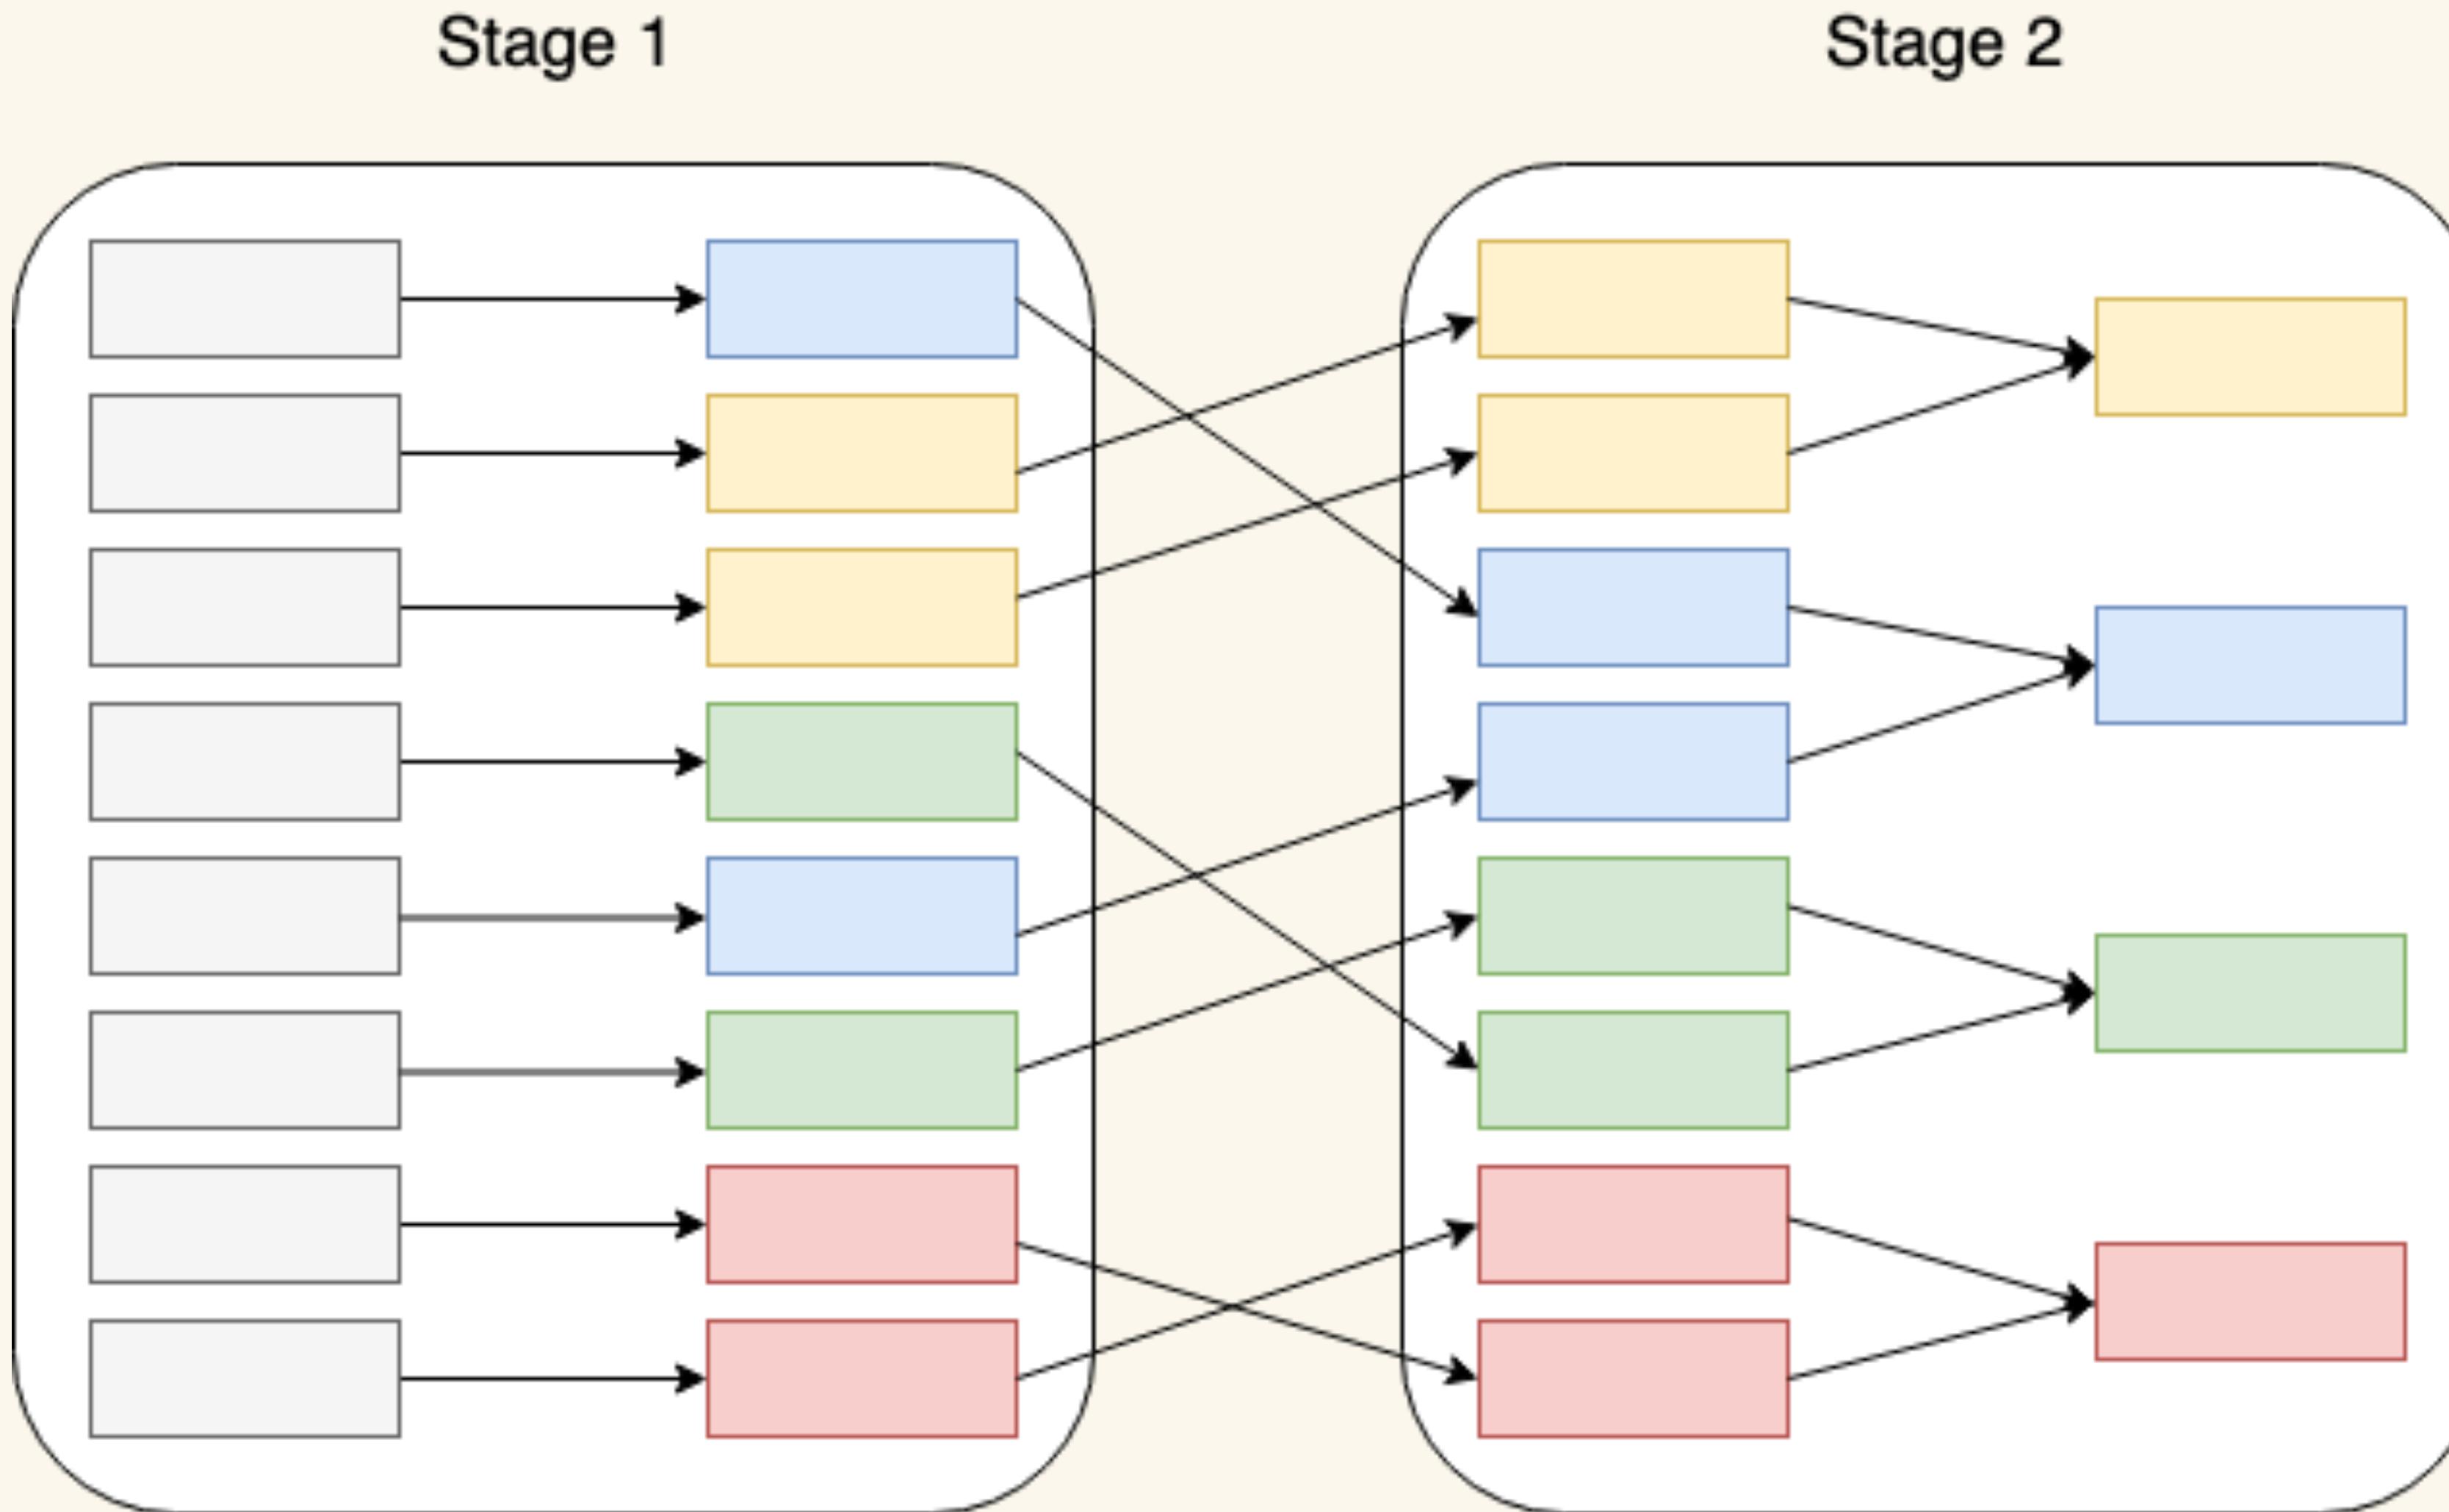
# Shuffle

**When a shuffle isn't performed:**

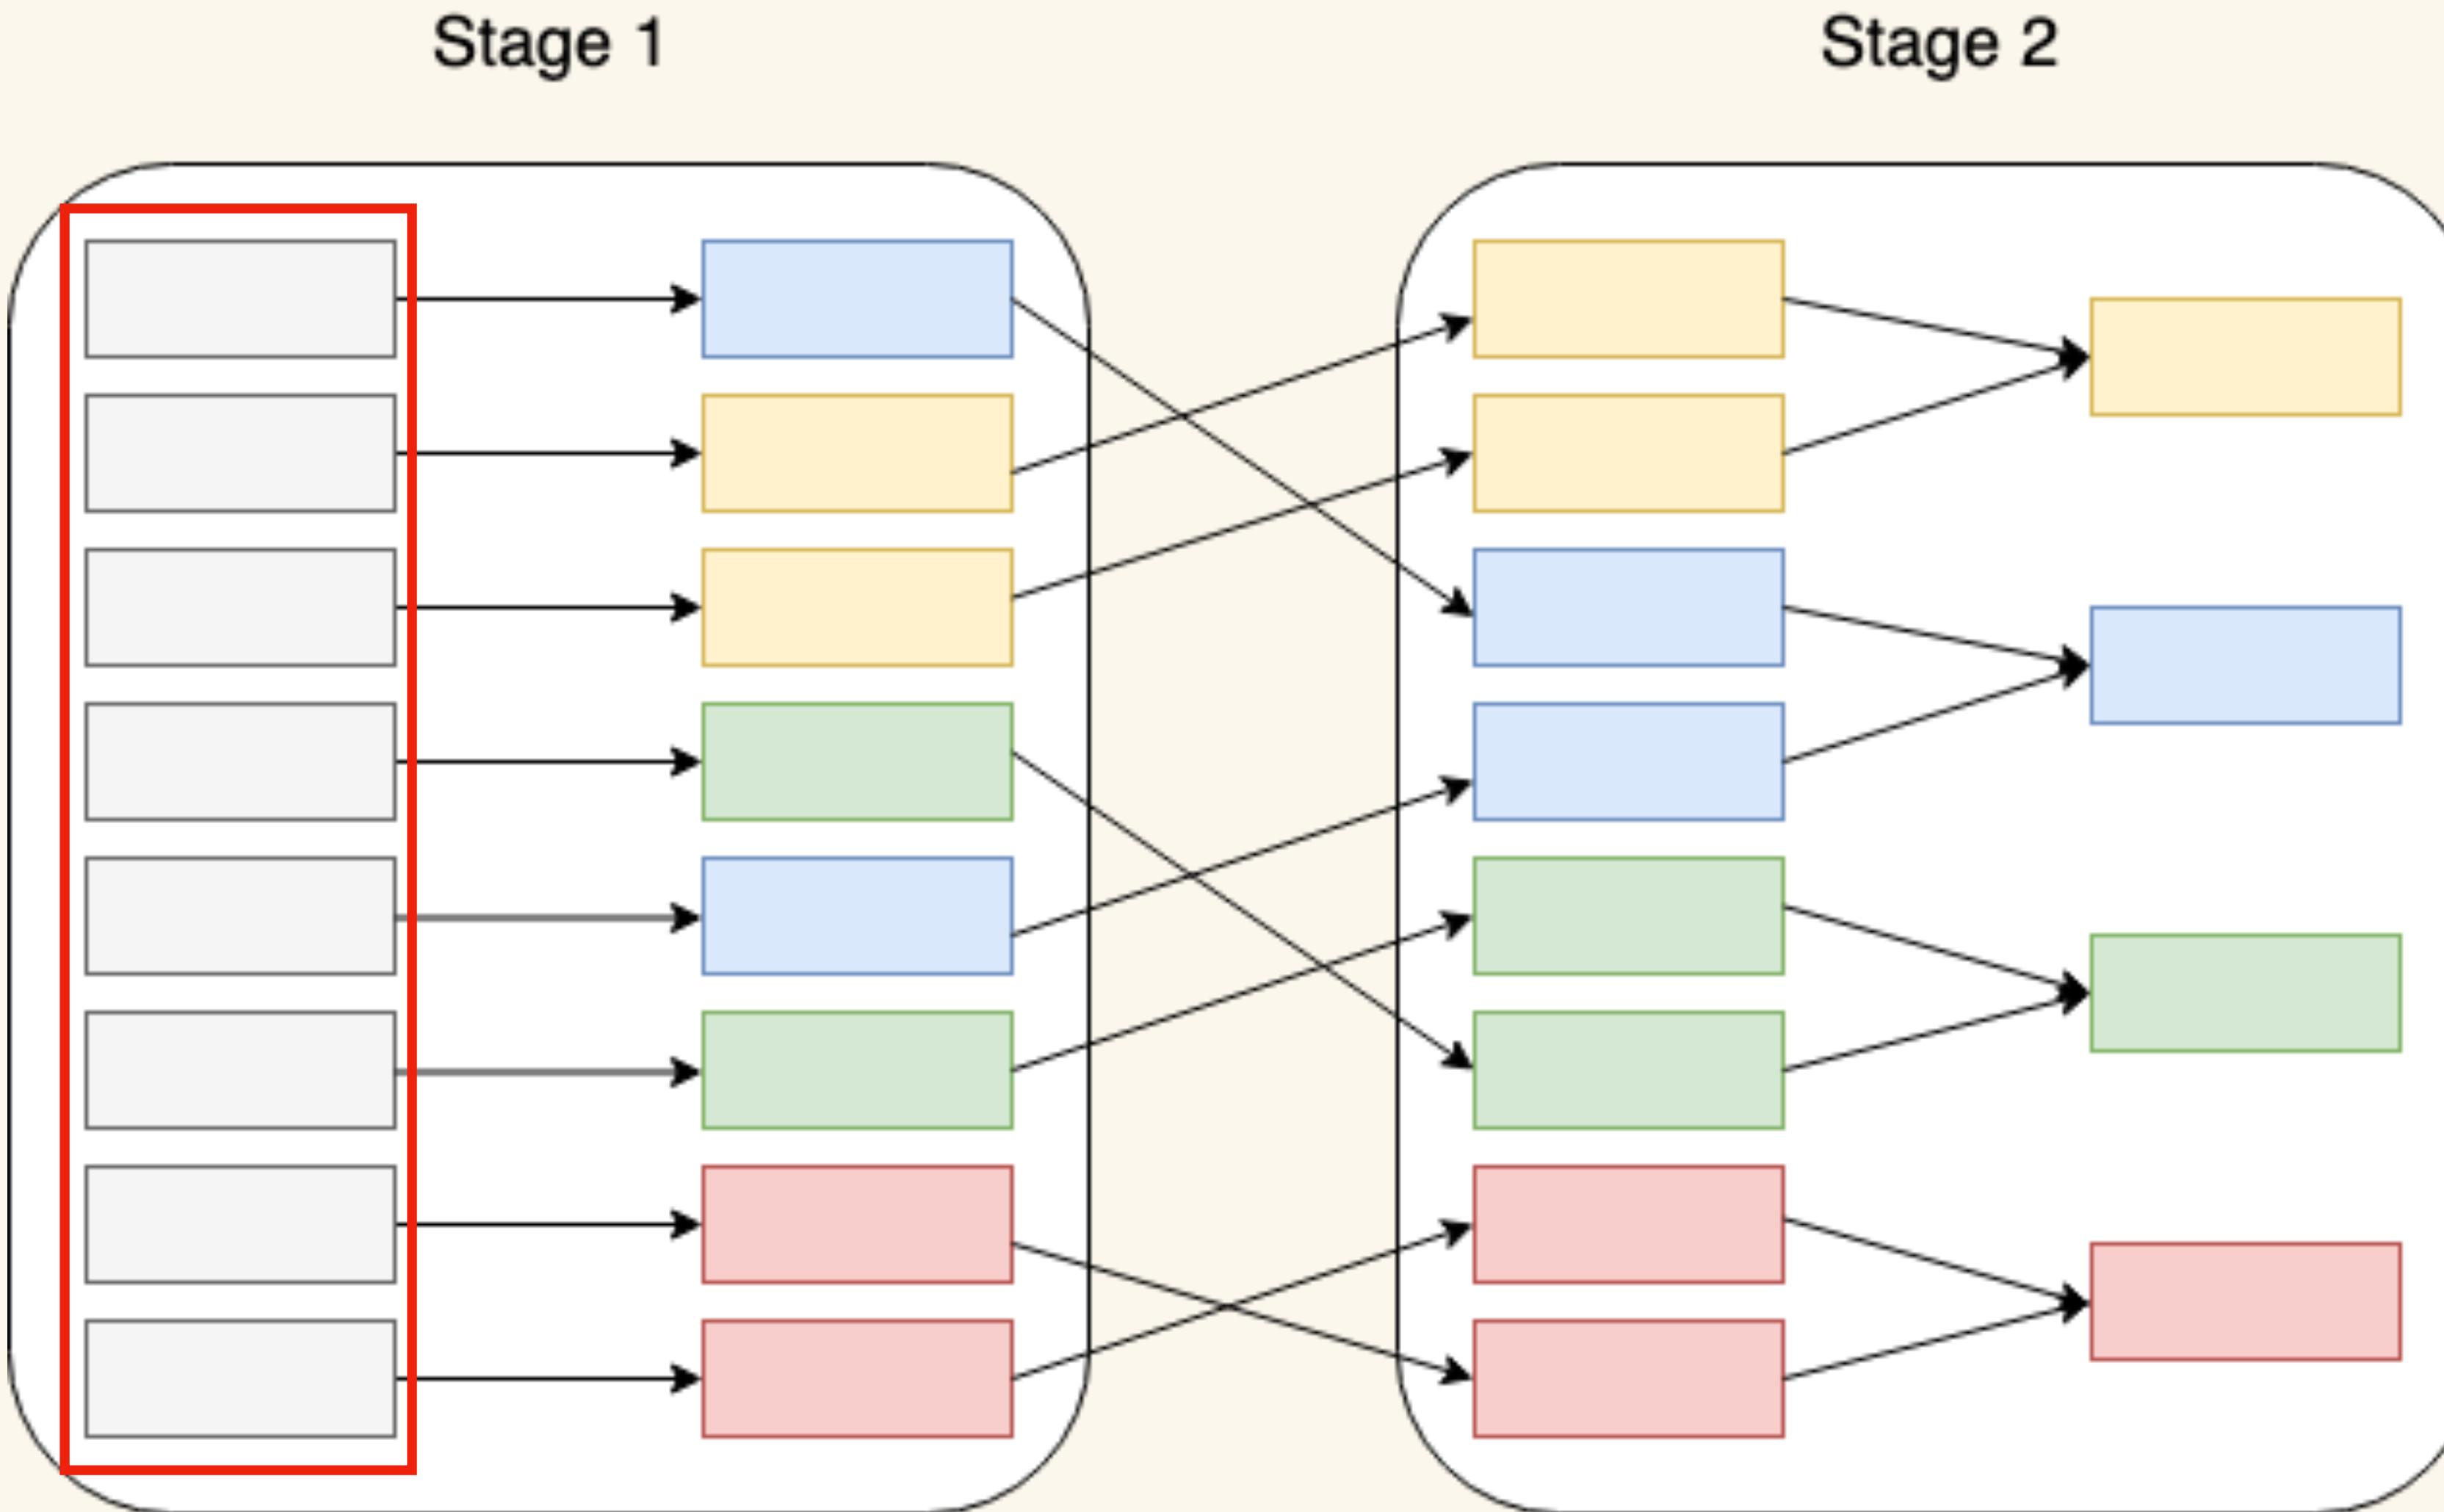
- Data is sorted / bucketed by the partition key on disk.
- Data has already been shuffled from a previous operation.
  - ex. When an aggregate happens after a join but both are performed on the same column.



# Shuffle



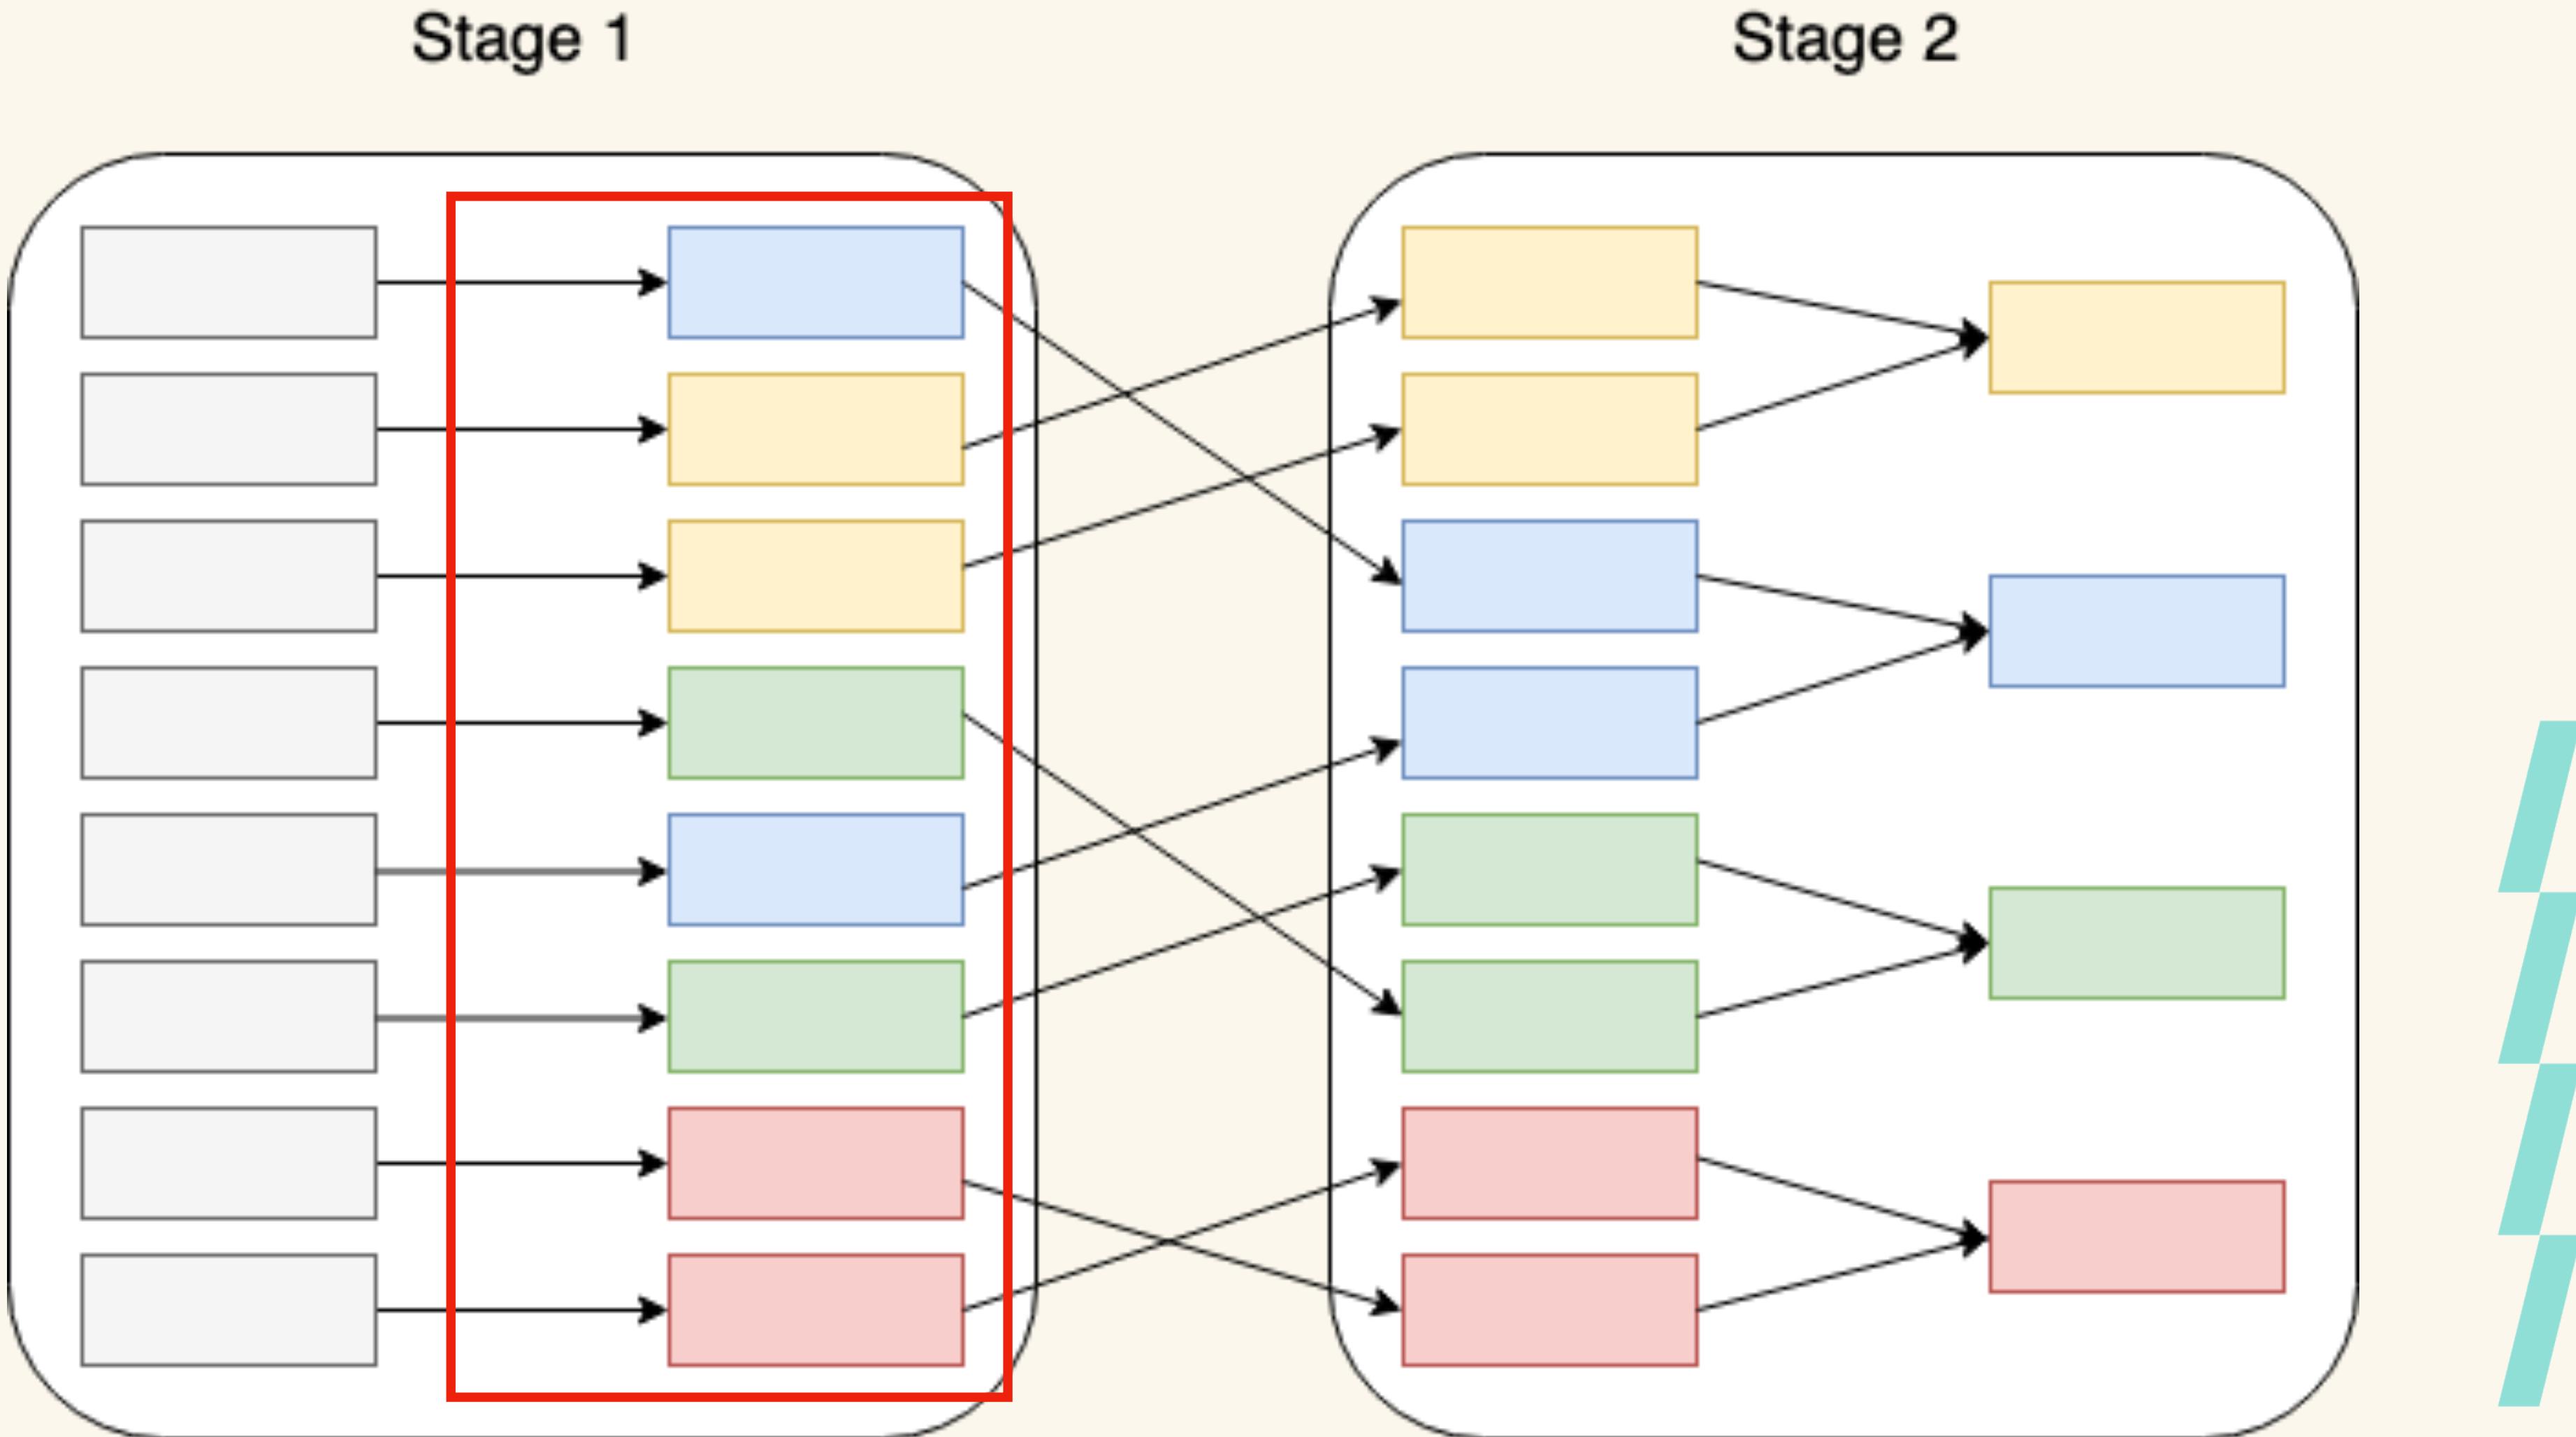
# Shuffle



1. Initial partitions read from disk.

# Shuffle

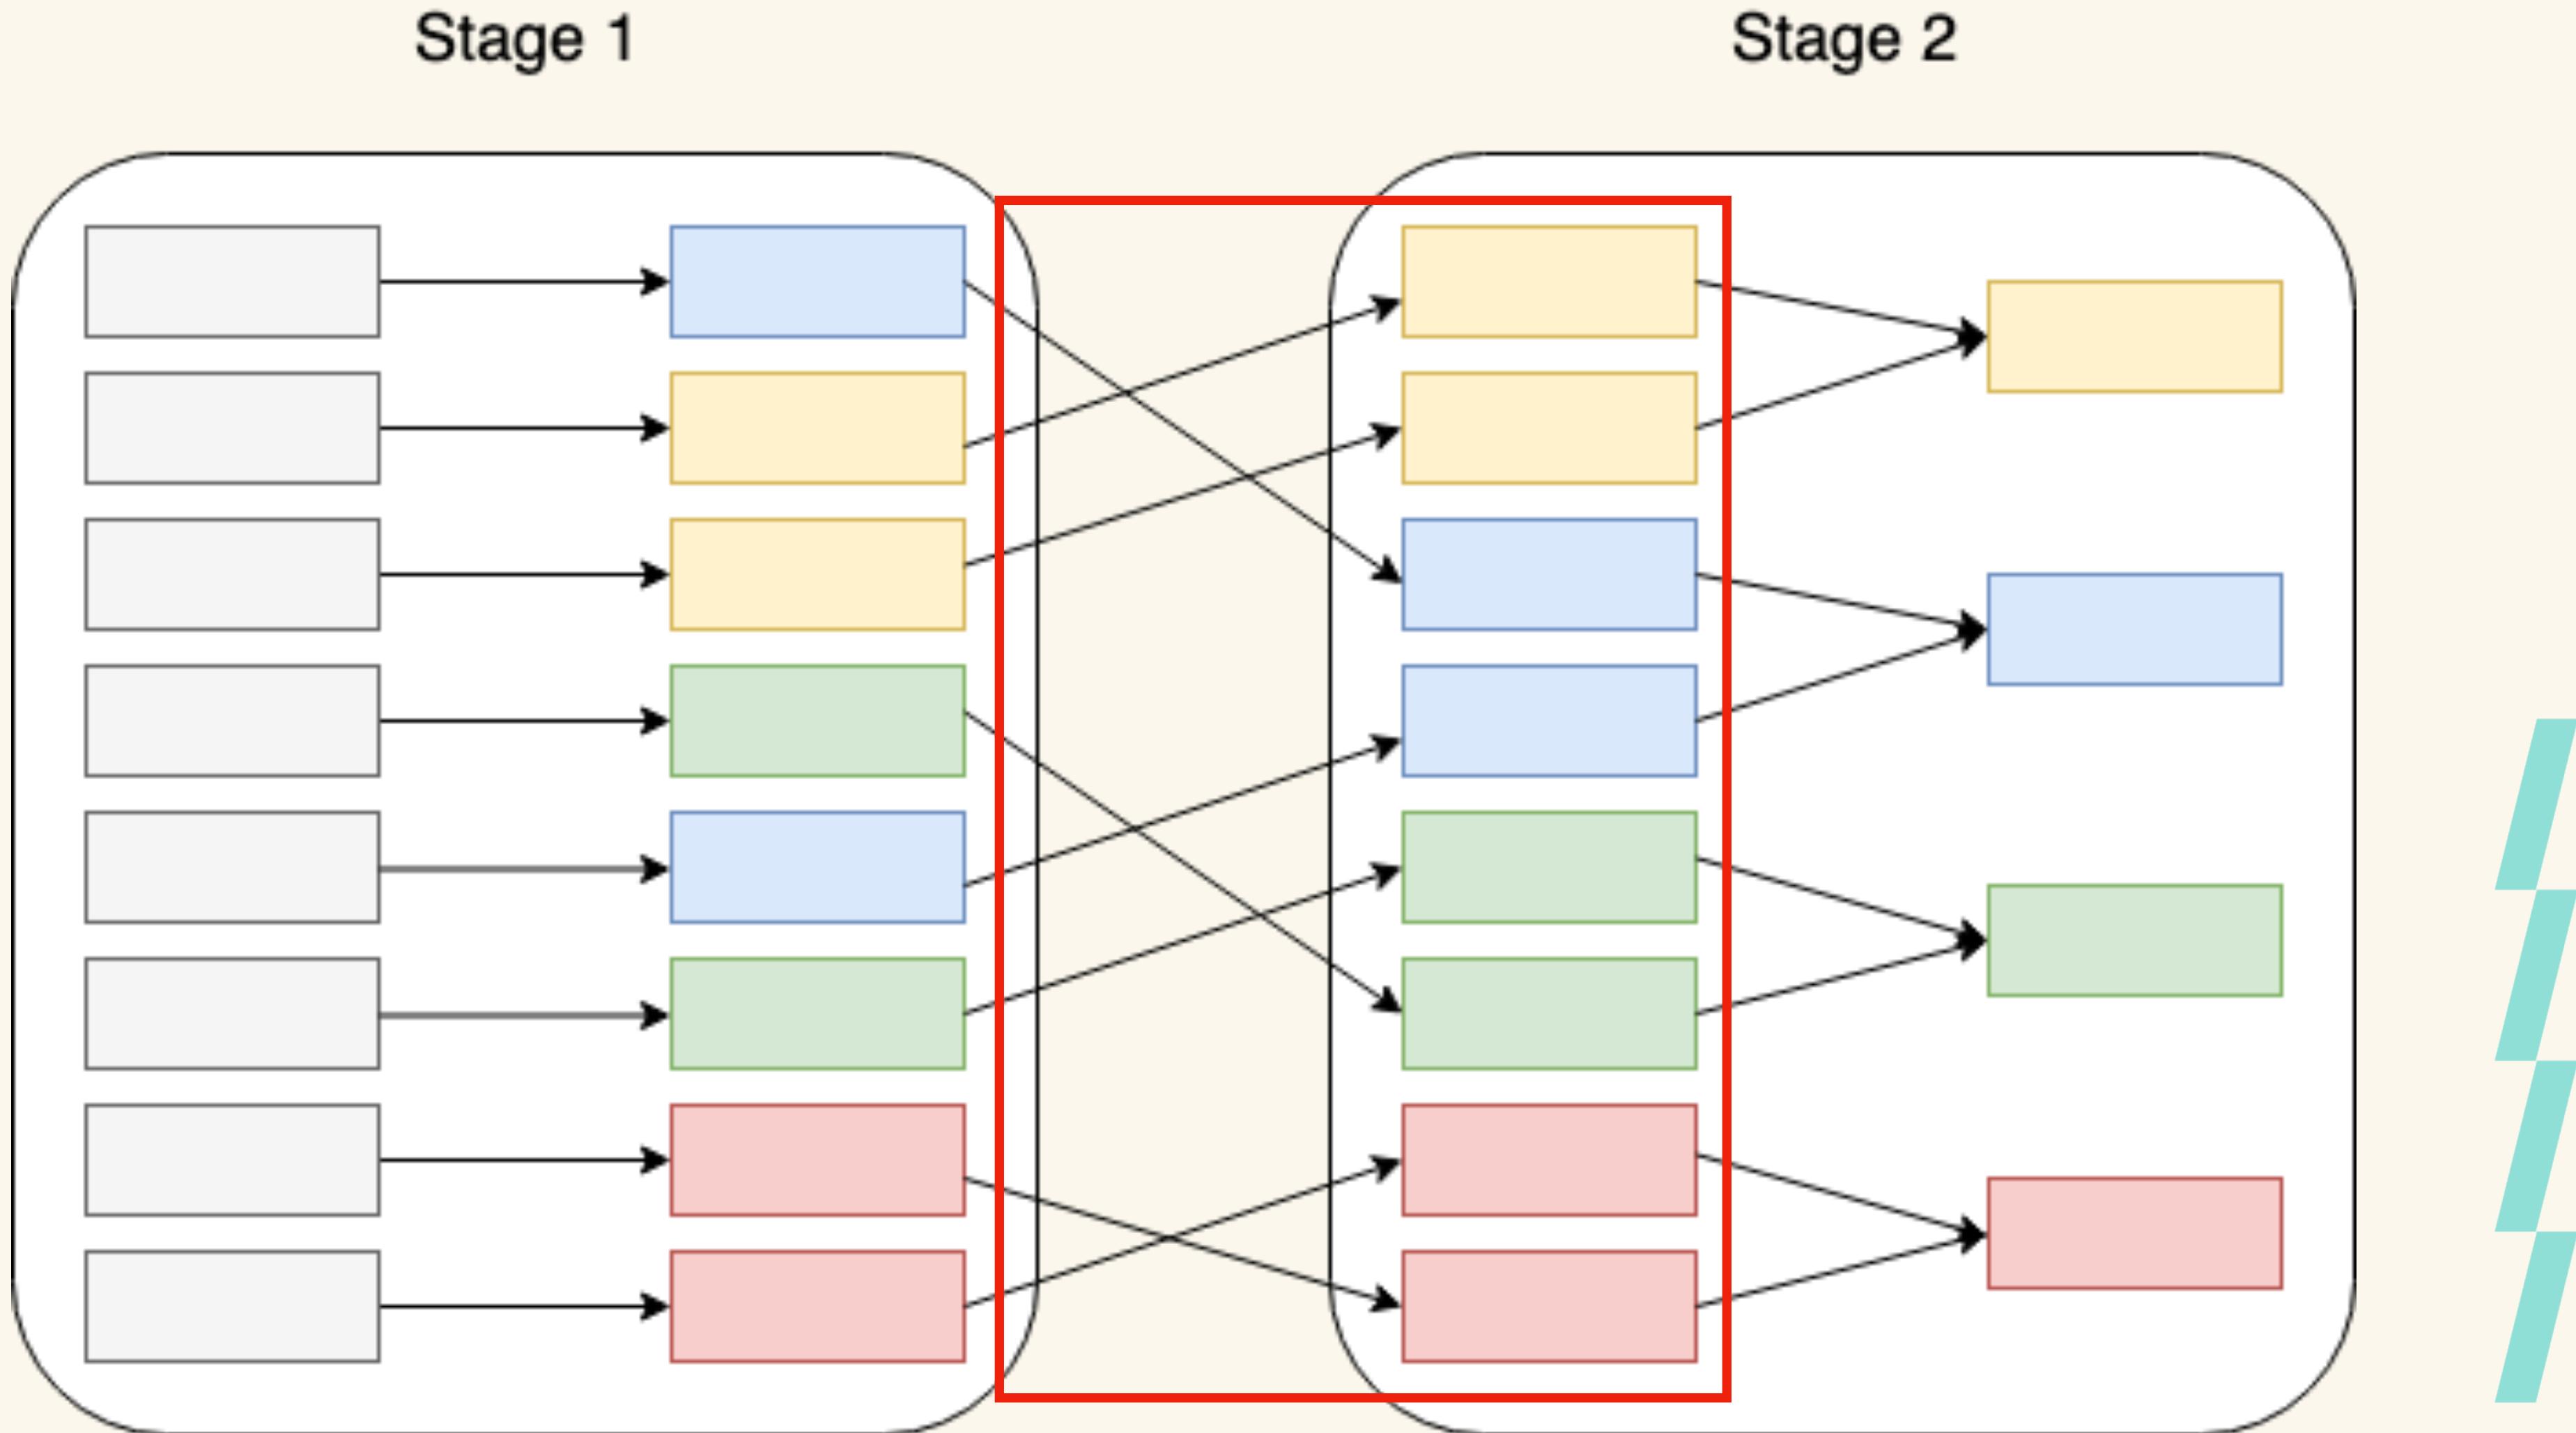
```
ex.  
(  
df  
.withColumn(x, x + 1)  
)
```



2. apply transformations on each partition.

# Shuffle

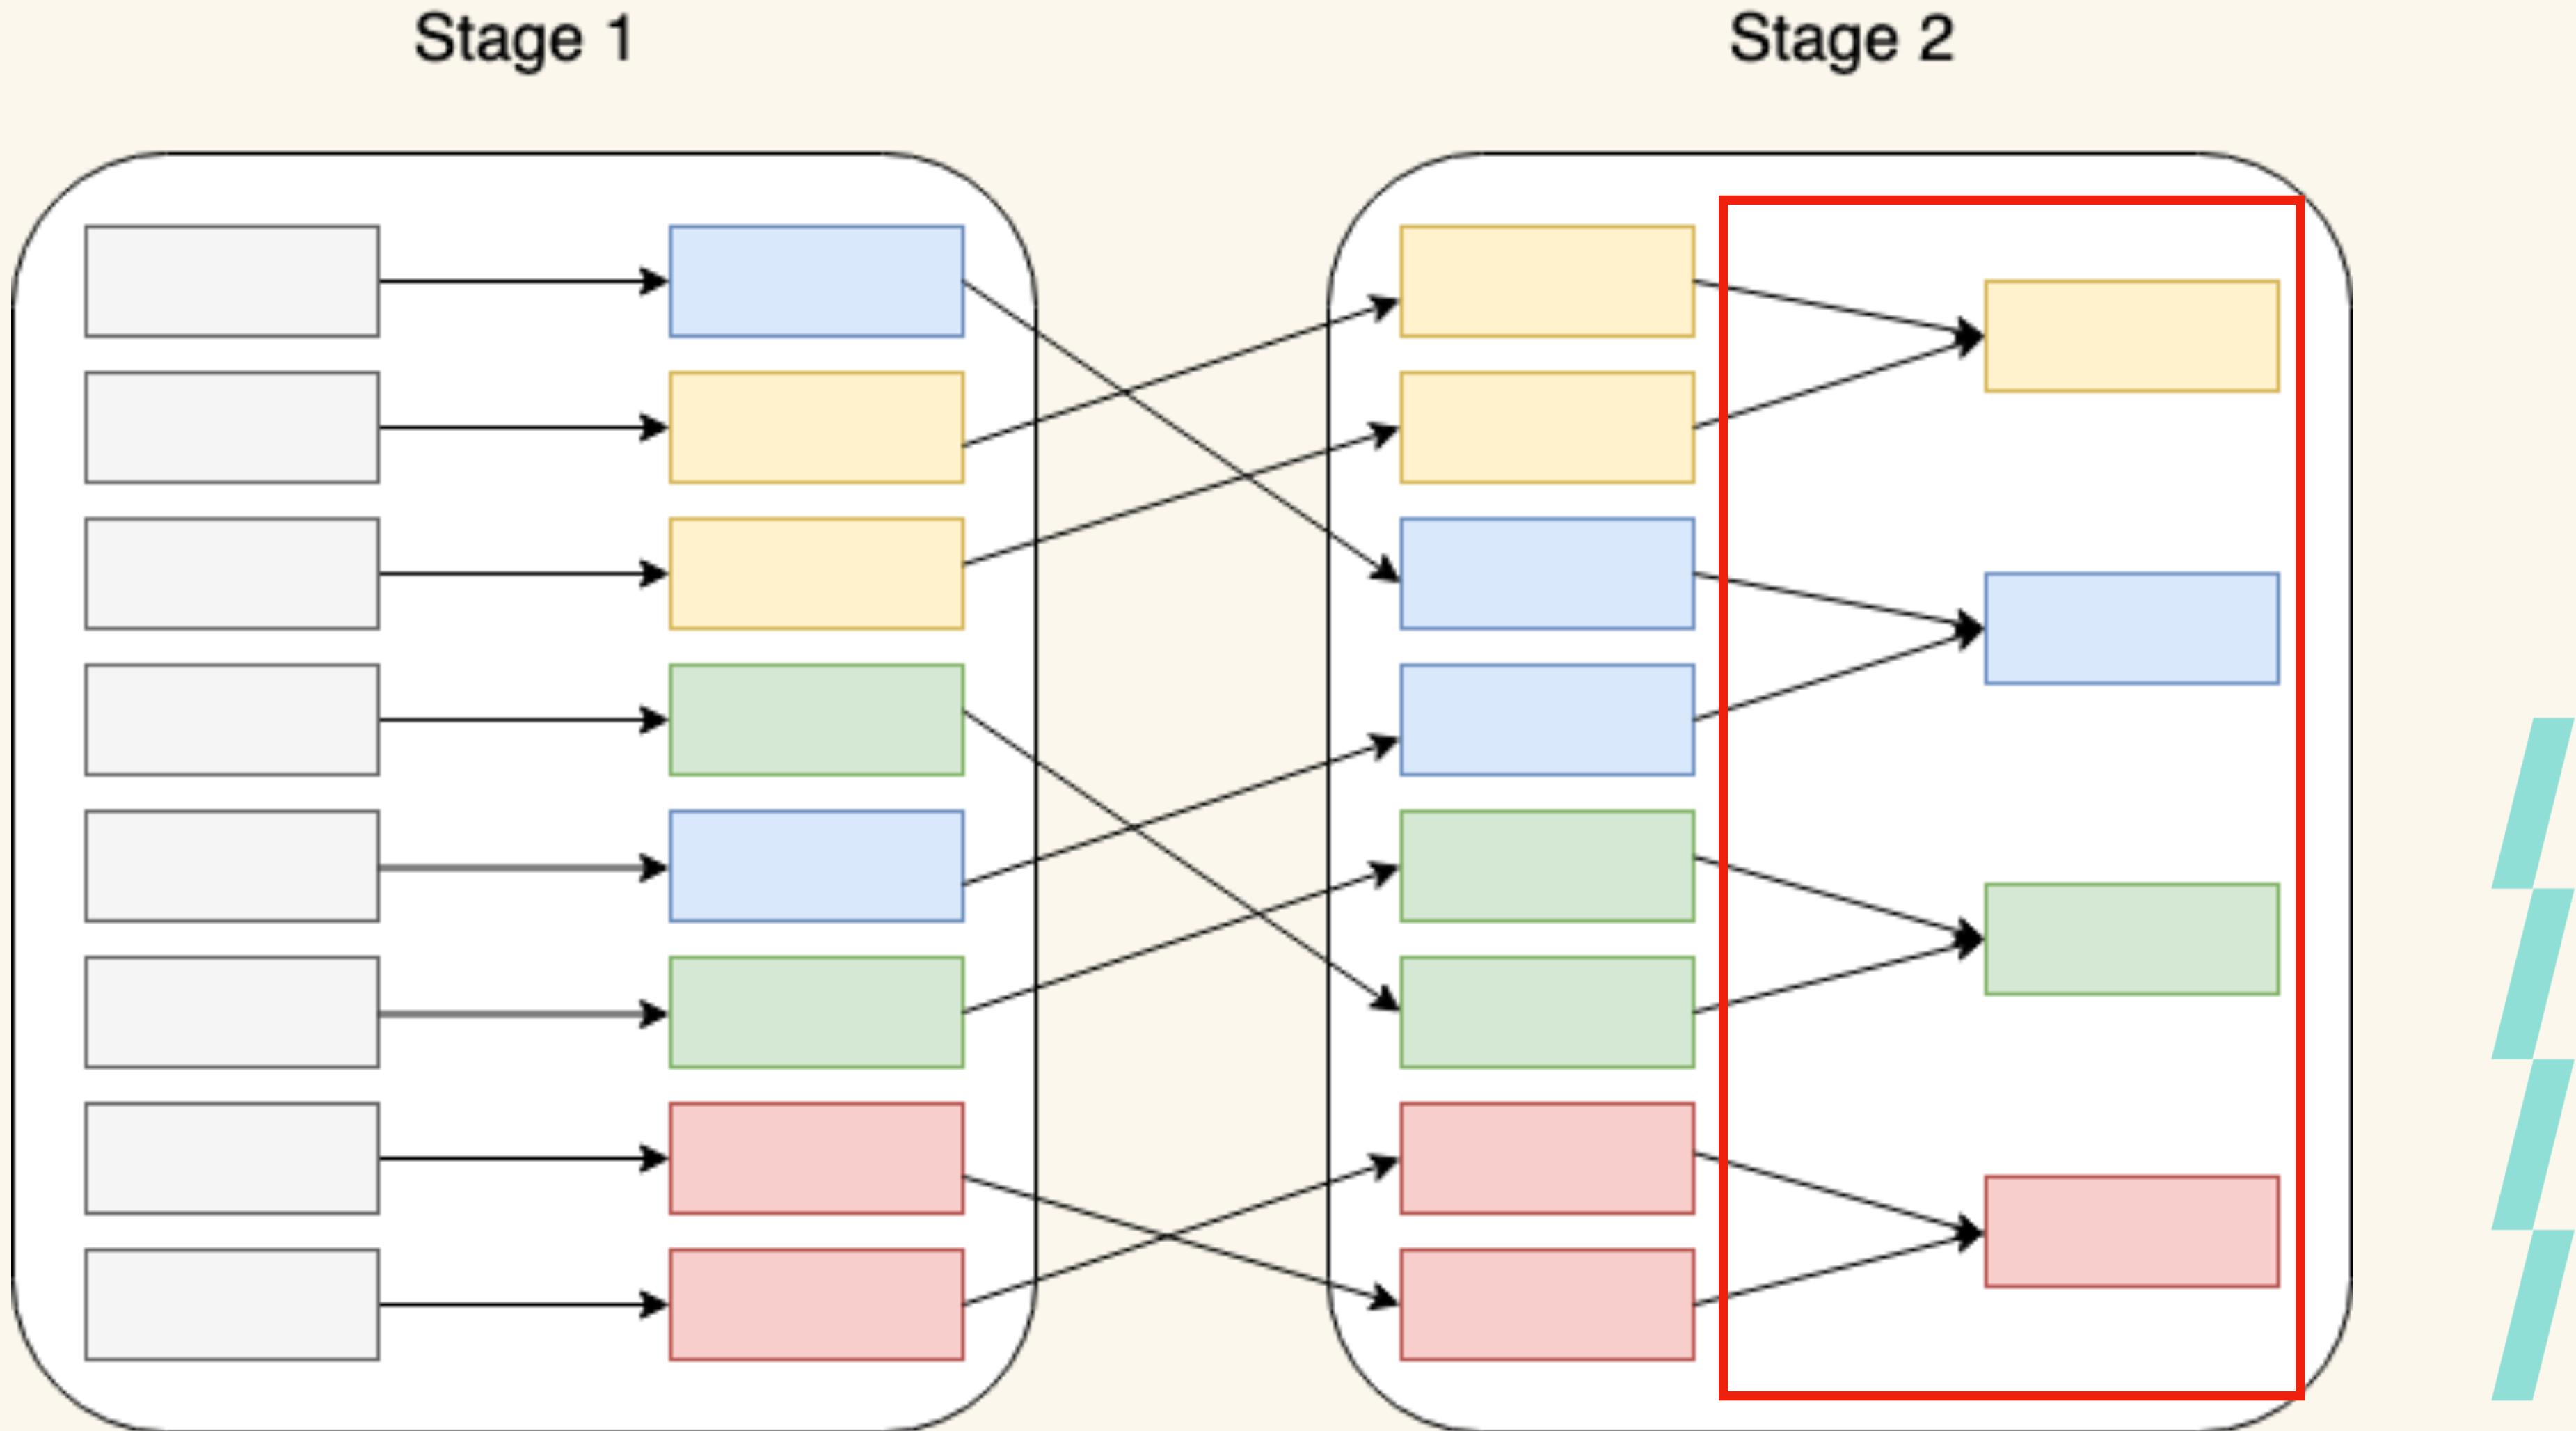
```
ex.  
(  
df  
.groupBy(y)  
.agg(...)  
)
```



3. shuffle data from stage 1 for stage 2.

# Shuffle

```
ex.  
(  
df  
.groupBy(y)  
.agg(...)  
)
```



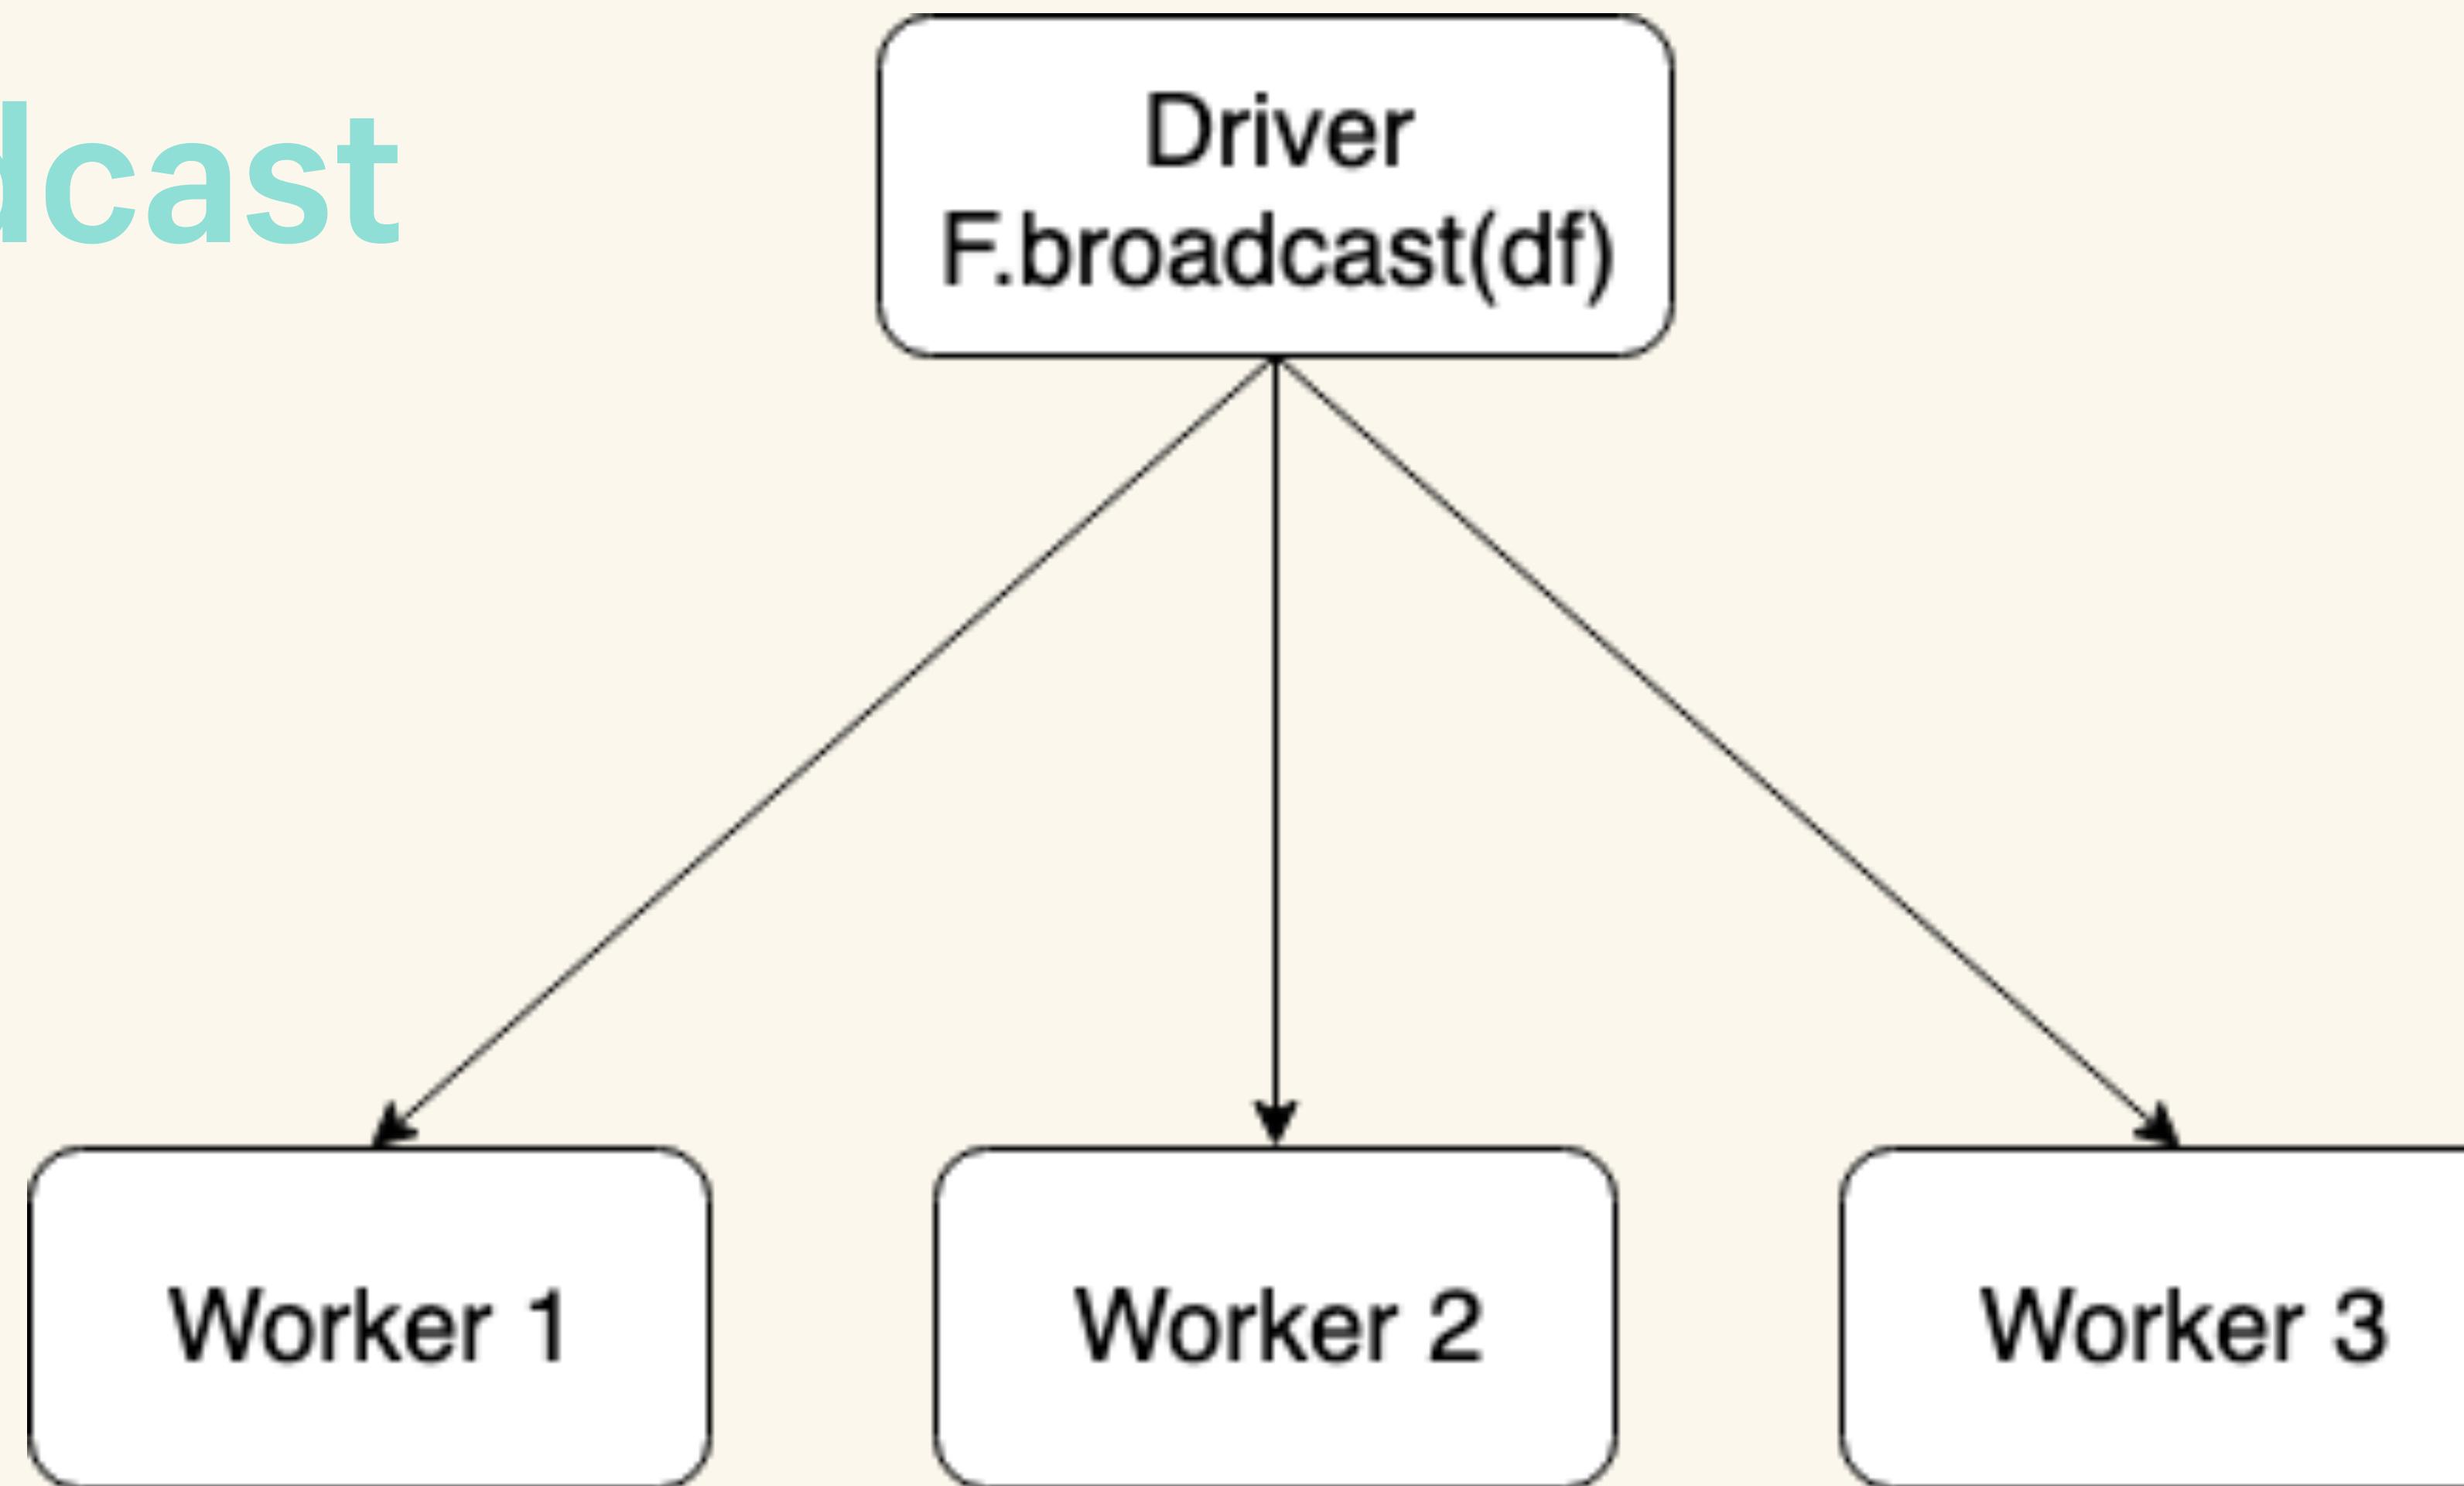
4. perform stage 2 transformations.

# Broadcast

- When data is small enough, Spark can broadcast it.
- DataFrame will be executed.
- Instead of shuffling data, Spark (the driver) will “broadcast” the DataFrame to every executor.

```
1 df = ...
2 broadcasted_df = broadcast(df)
3
4 spark.sql.autoBroadcastJoinThreshold
5
6
7
```

# Broadcast



# Broadcast Gotchas

- As mentioned before the broadcasted `dataFrame` is executed before being broadcasted.
- If the transformations performed are complex or take long enough, this will cause a timeout.
- This wait time is defined by `spark.sql.broadcastTimeout` which is 300s (5 minutes by default).

## RED FLAGS:

- `TimeoutException`: Futures timed out after [300 seconds].

## Solution:

- Increase the timeout, 10 minutes should be the upper limit.



# SQL Tab

- Shows the physical plan of your Spark application.
- Shows an aggregate level of task duration and data volume.
- Dropdown of all the other plans.



# 1. Overall DAG

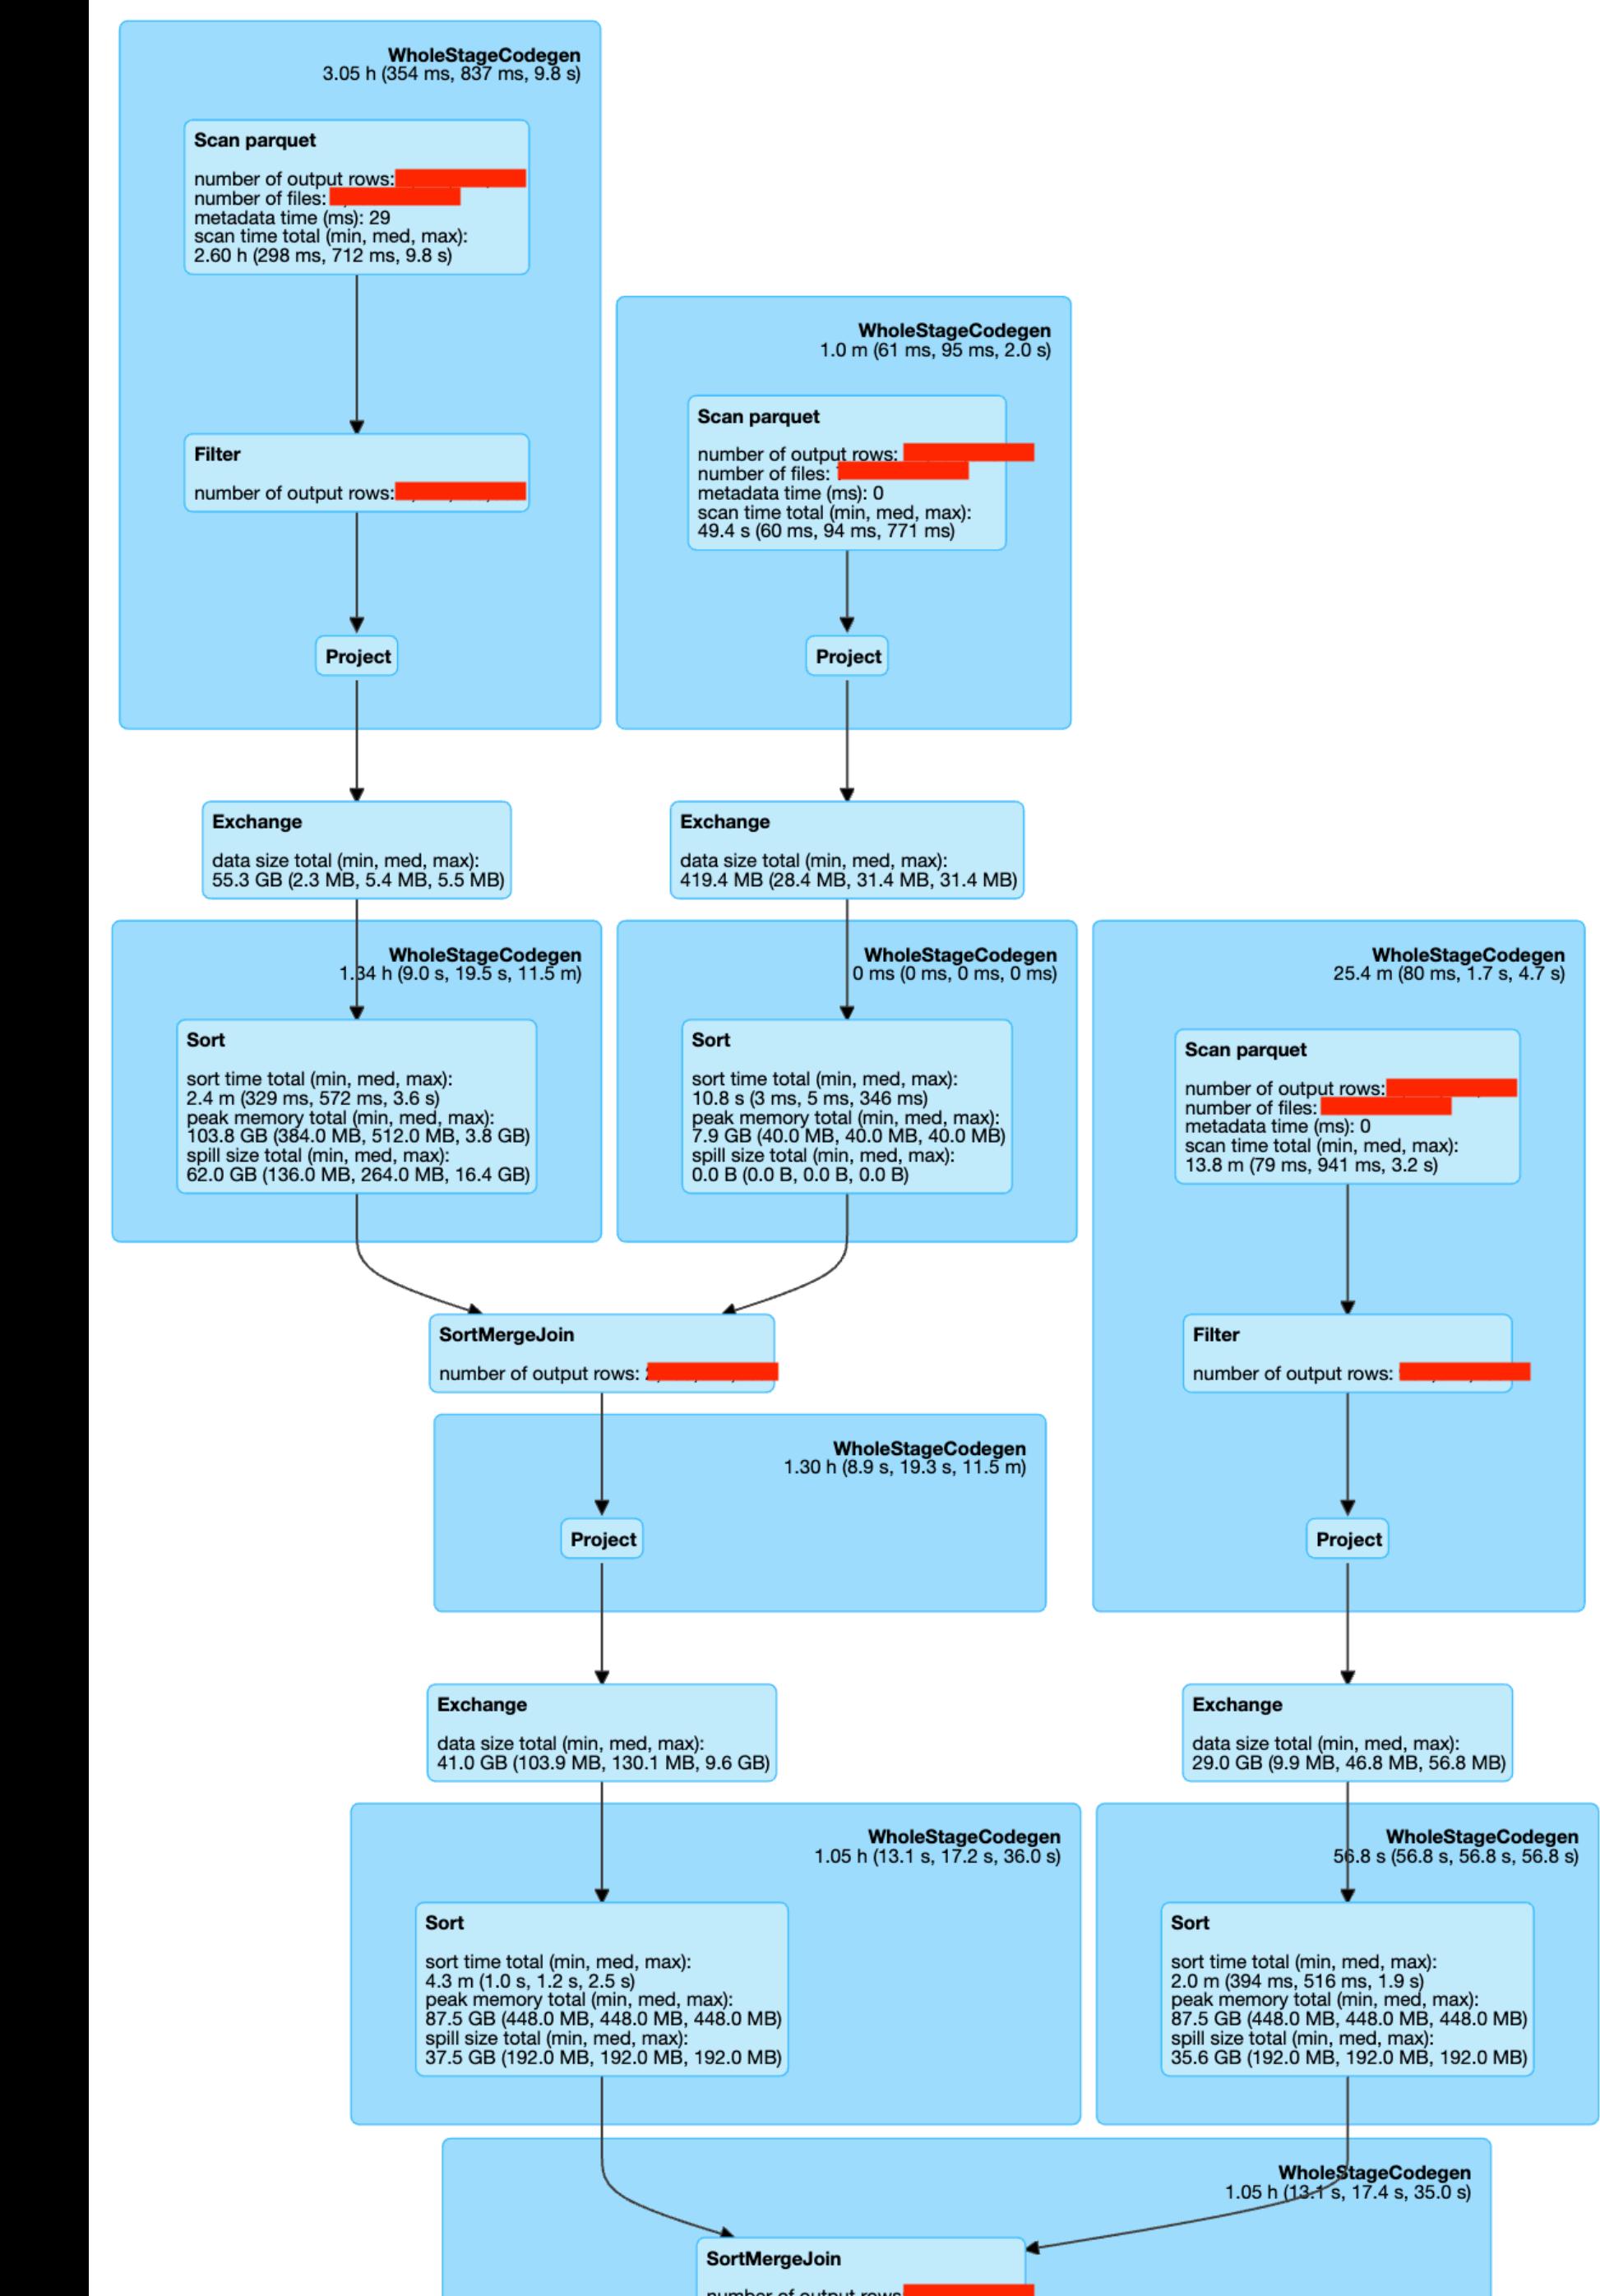
- Visual representation of physical plan.
- It flows from top (reading data) to bottom (writing / outputting data).

## RED FLAGS:

- Too large and wide DAG.

## Solution:

- Reduce application complexity.



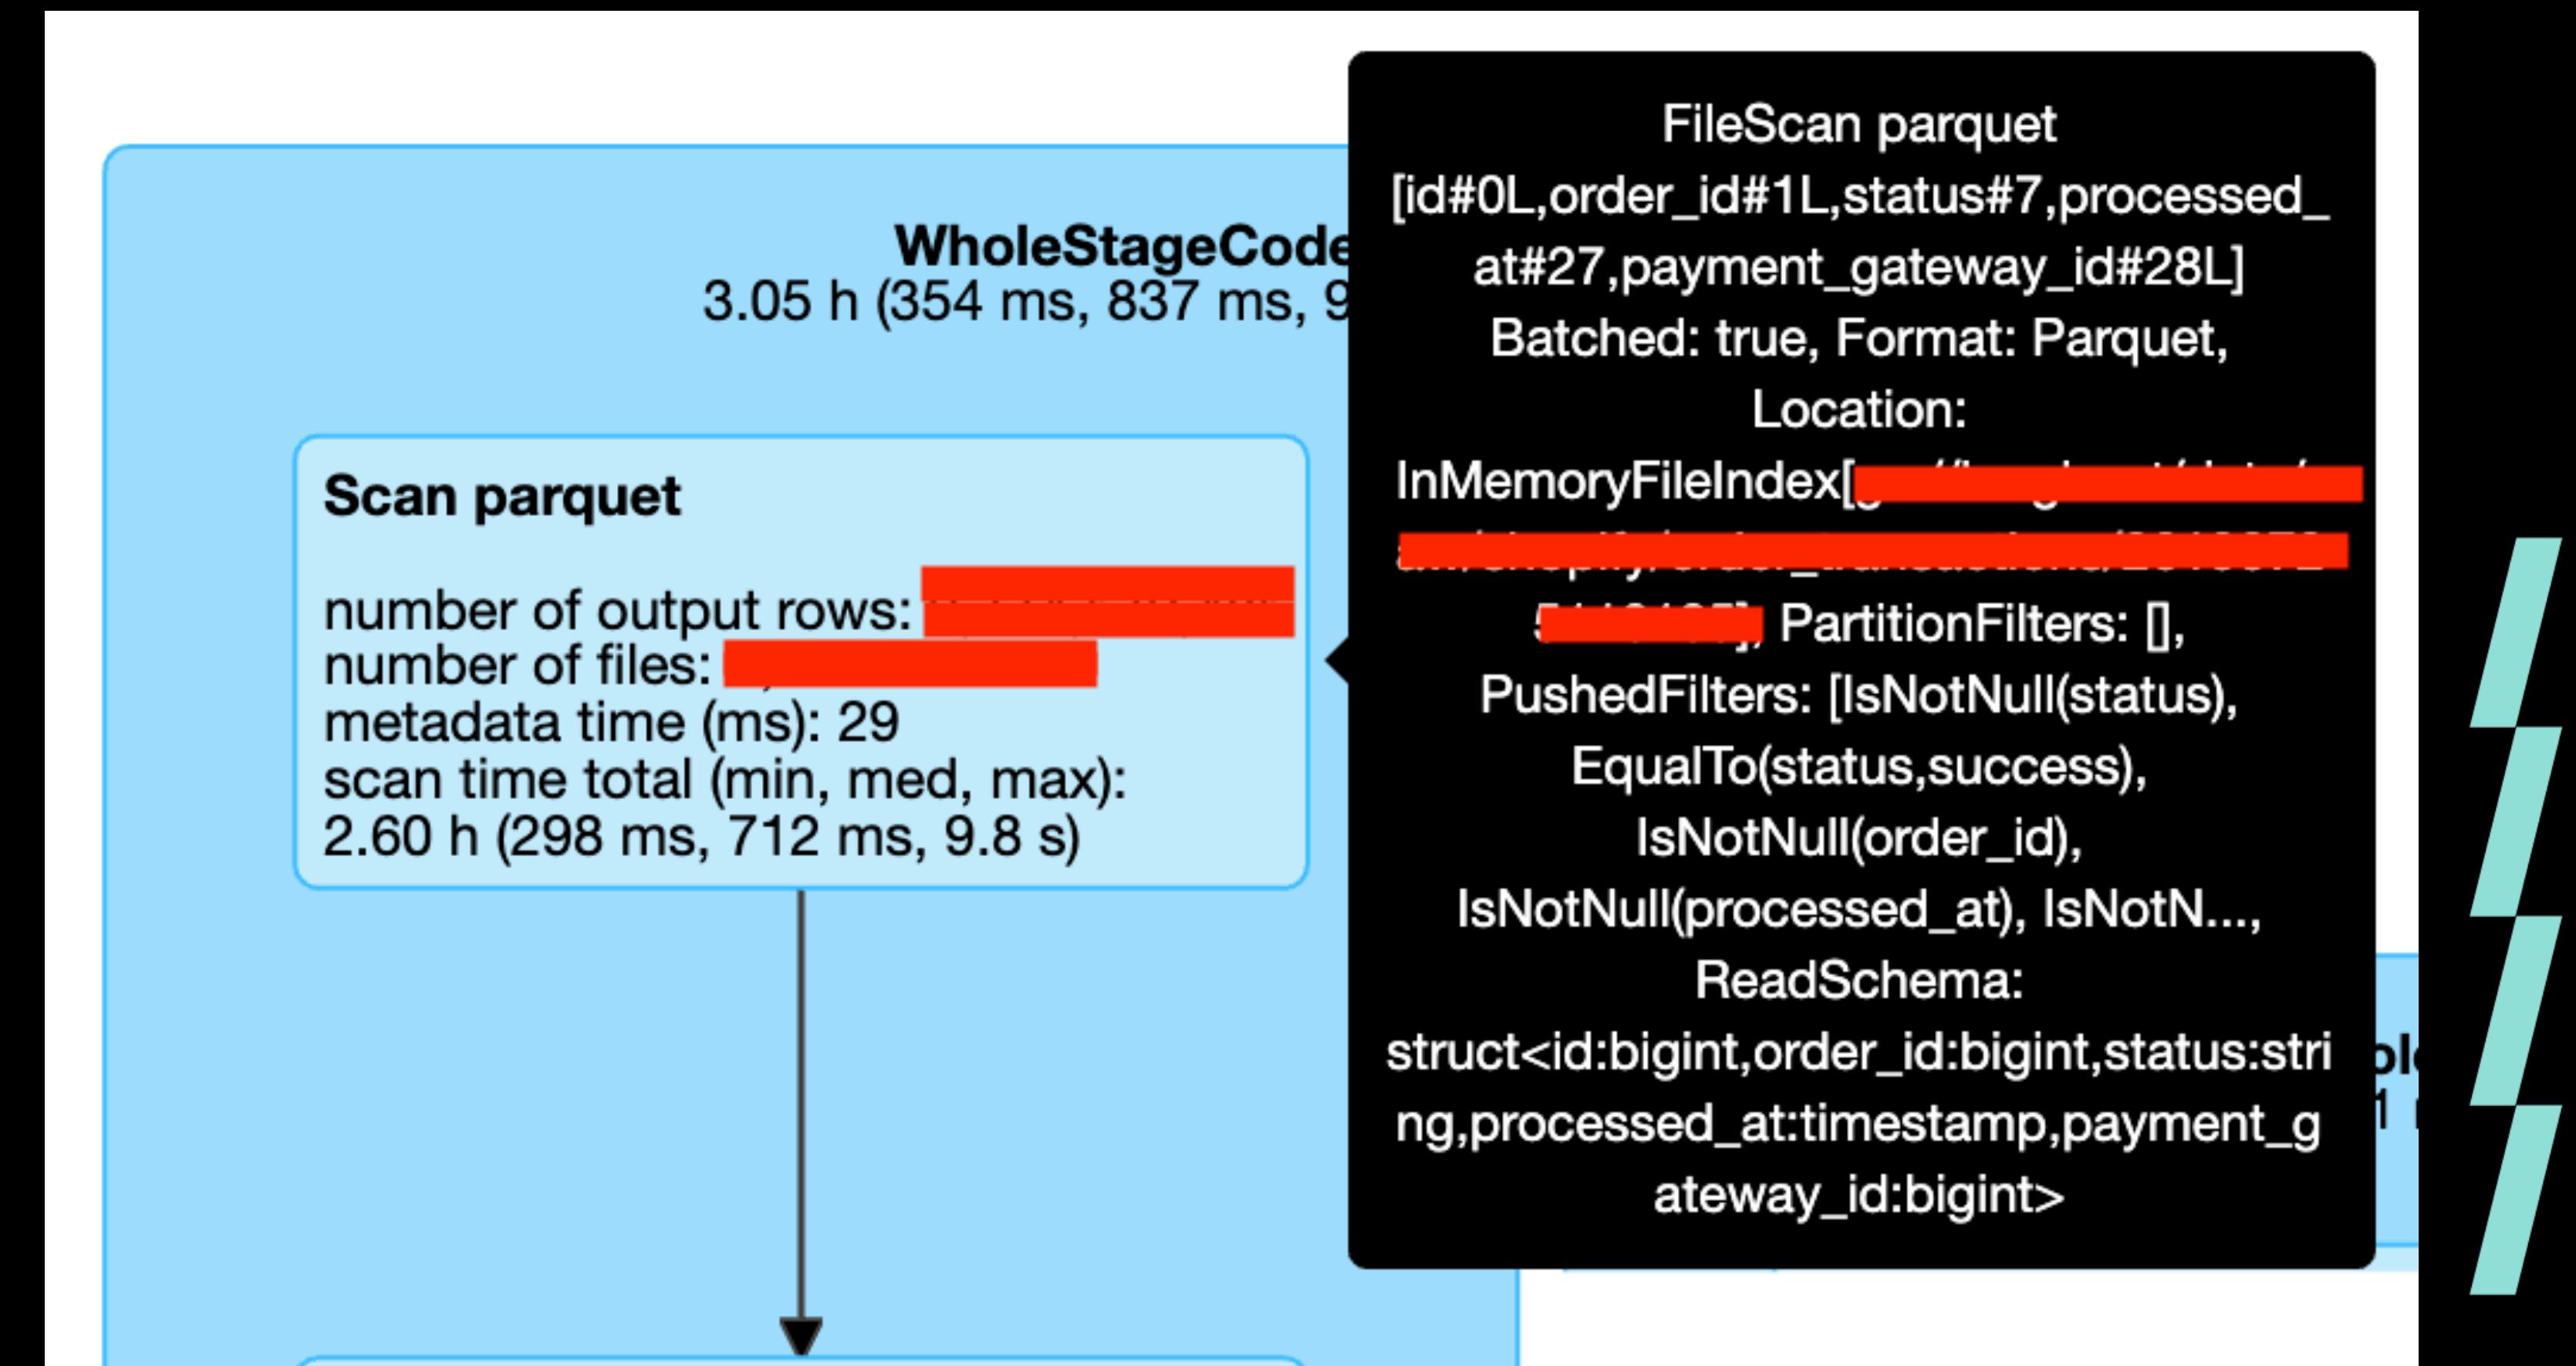
# 2. Spark Plans

▼ Details

```
== Parsed Logical Plan ==
GlobalLimit 6
+- LocalLimit 6
  +- AnalysisBarrier
    +- Project [cast(card_type#229 as string) AS card_type#300, cast(count#289L as string) AS count#301]
      +- Sort [count#289L DESC NULLS LAST], true
      +- Filter isnotnull(card_type#229)
        +- Aggregate [card_type#229], [card_type#229, sum(cast(1 as bigint)) AS count#289L]
          +- Project [order_transaction_id#202L, payment_gateway_id#28L, created_at#9, processed_at#27,
payment_gateway_request_id#39L, provider_id#101L, name#96, card_type#229]
            +- Join LeftOuter, (order_transaction_id#202L = order_transaction_id#160L)
              :- Project [payment_gateway_id#28L, order_transaction_id#202L, created_at#9, processed_
payment_gateway_request_id#39L, provider_id#101L, name#96]
                :  +- Join LeftOuter, (payment_gateway_id#28L = payment_gateway_id#221L)
                :    :- Project [id#0L AS order_transaction_id#202L, created_at#9, processed_at#27, sh
payment_gateway_request_id#39L]
                  :     :  +- Filter (((status#7 = success) && isnotnull(order_id#1L)) && isnotnull(proce
                  :     :    +- Project [id#0L, created_at#9, processed_at#27, shop_id#21L, order_id#1L,
payment_gateway_request_id#39L]
```

# 3. Transformation Details

- Hover over light blue blobs to see more code details.
- Gives hints to line of code in Spark application.

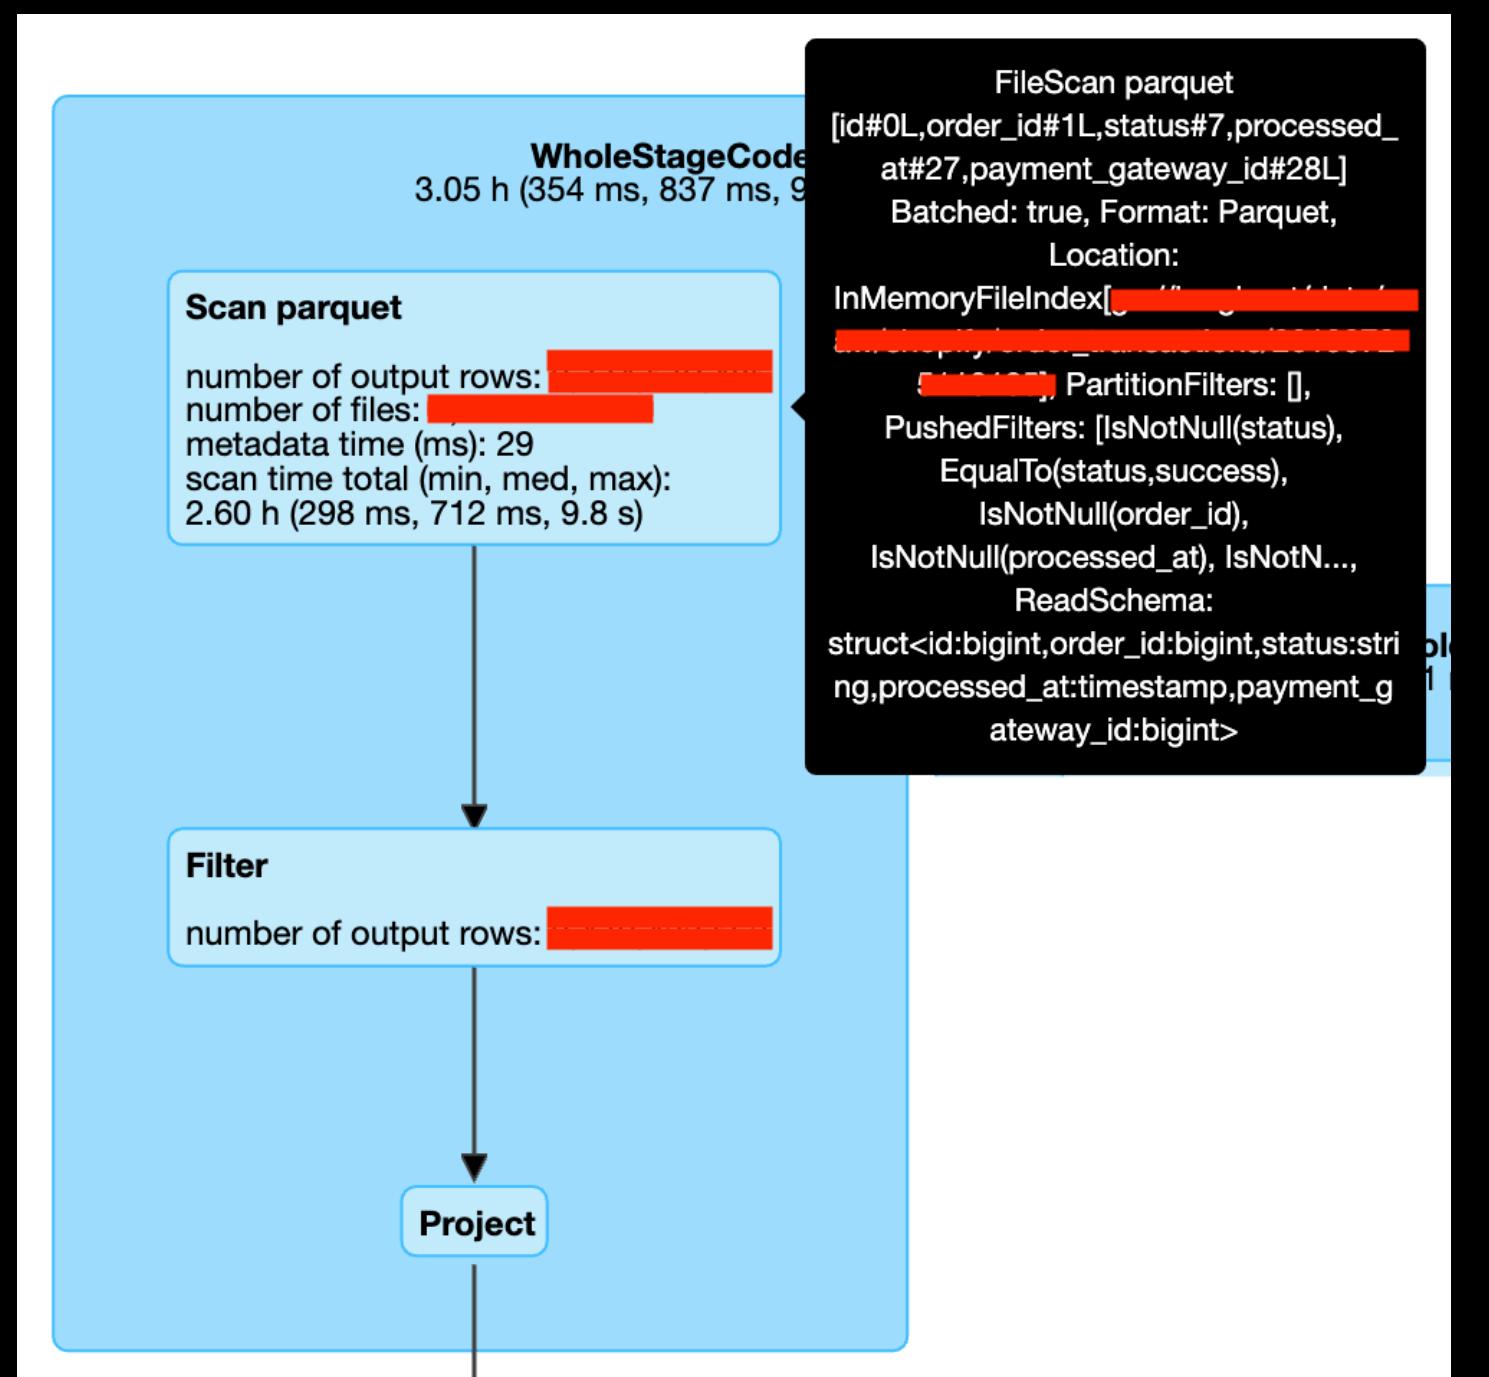


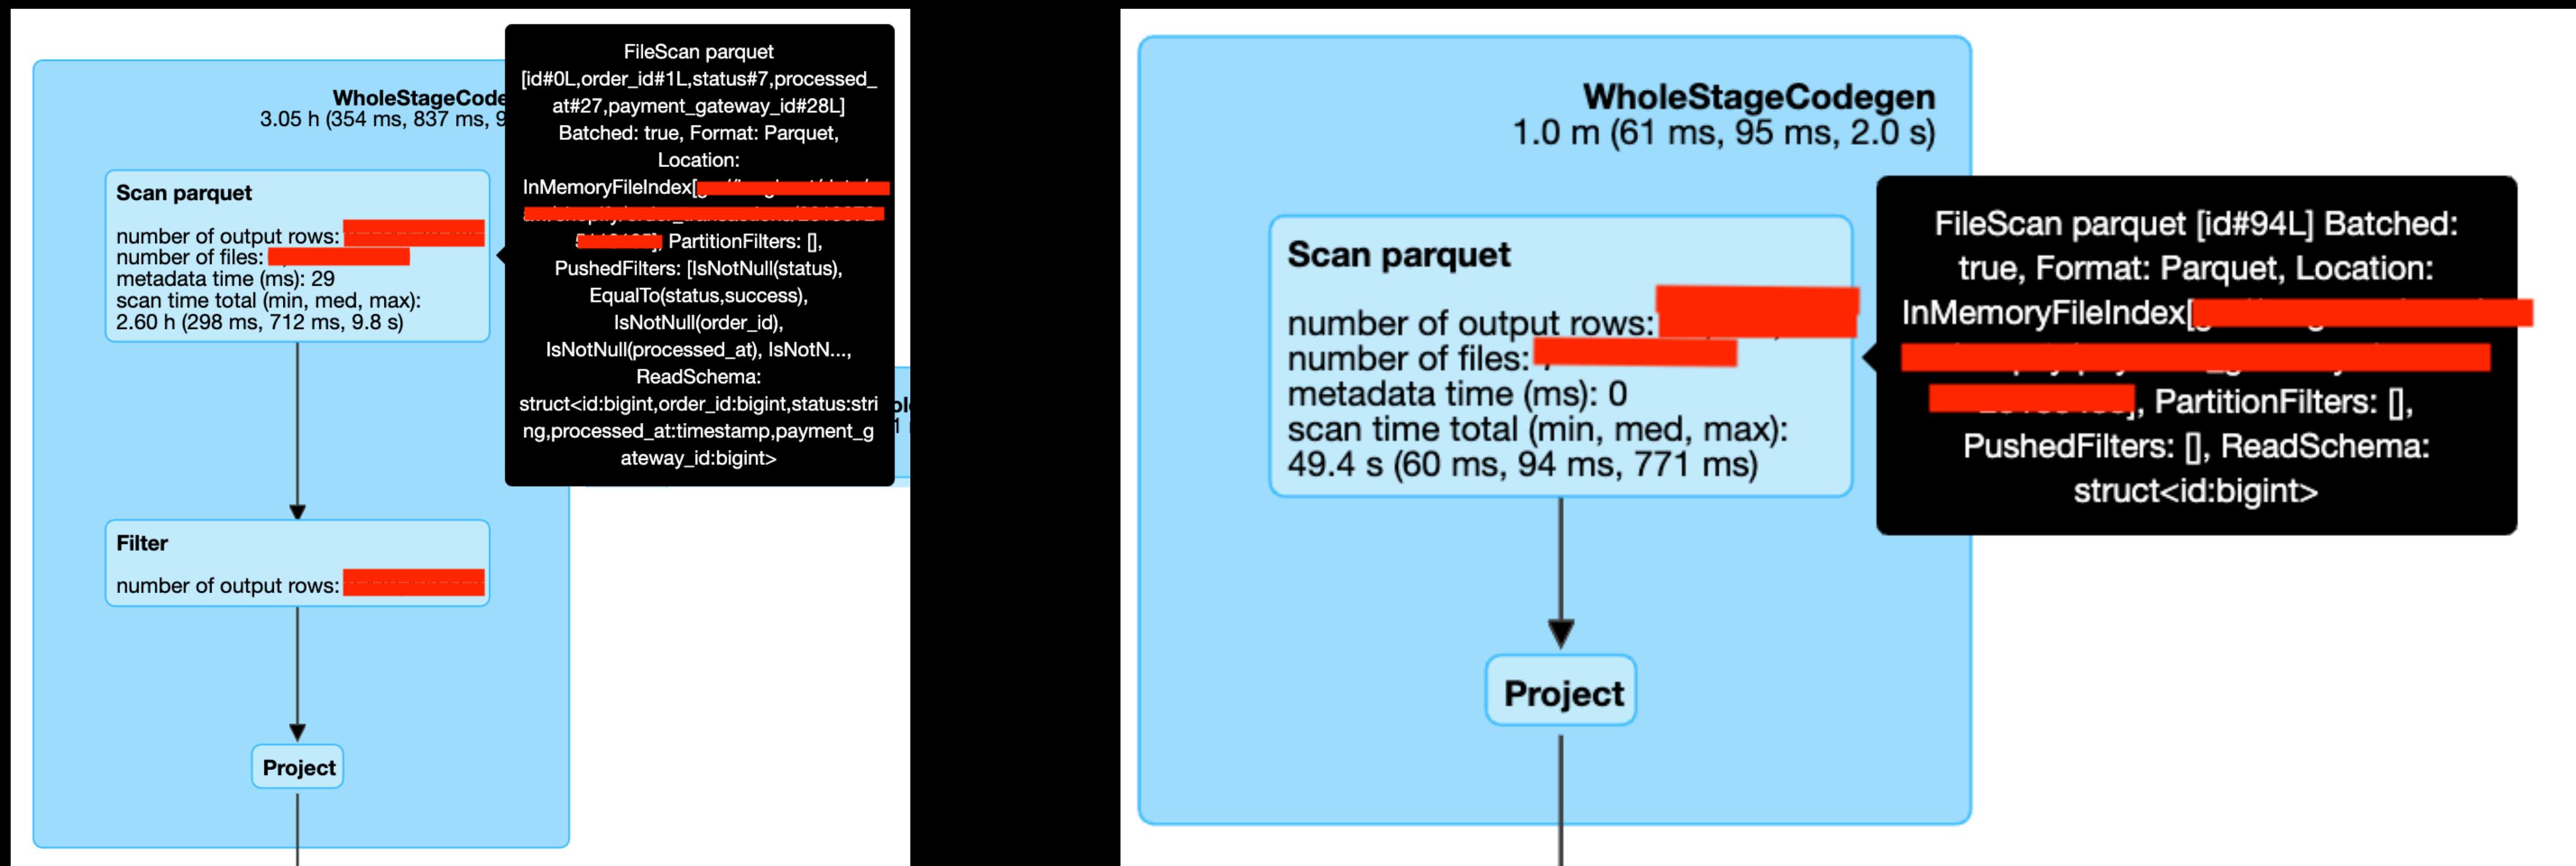
# SQL Query (Again)

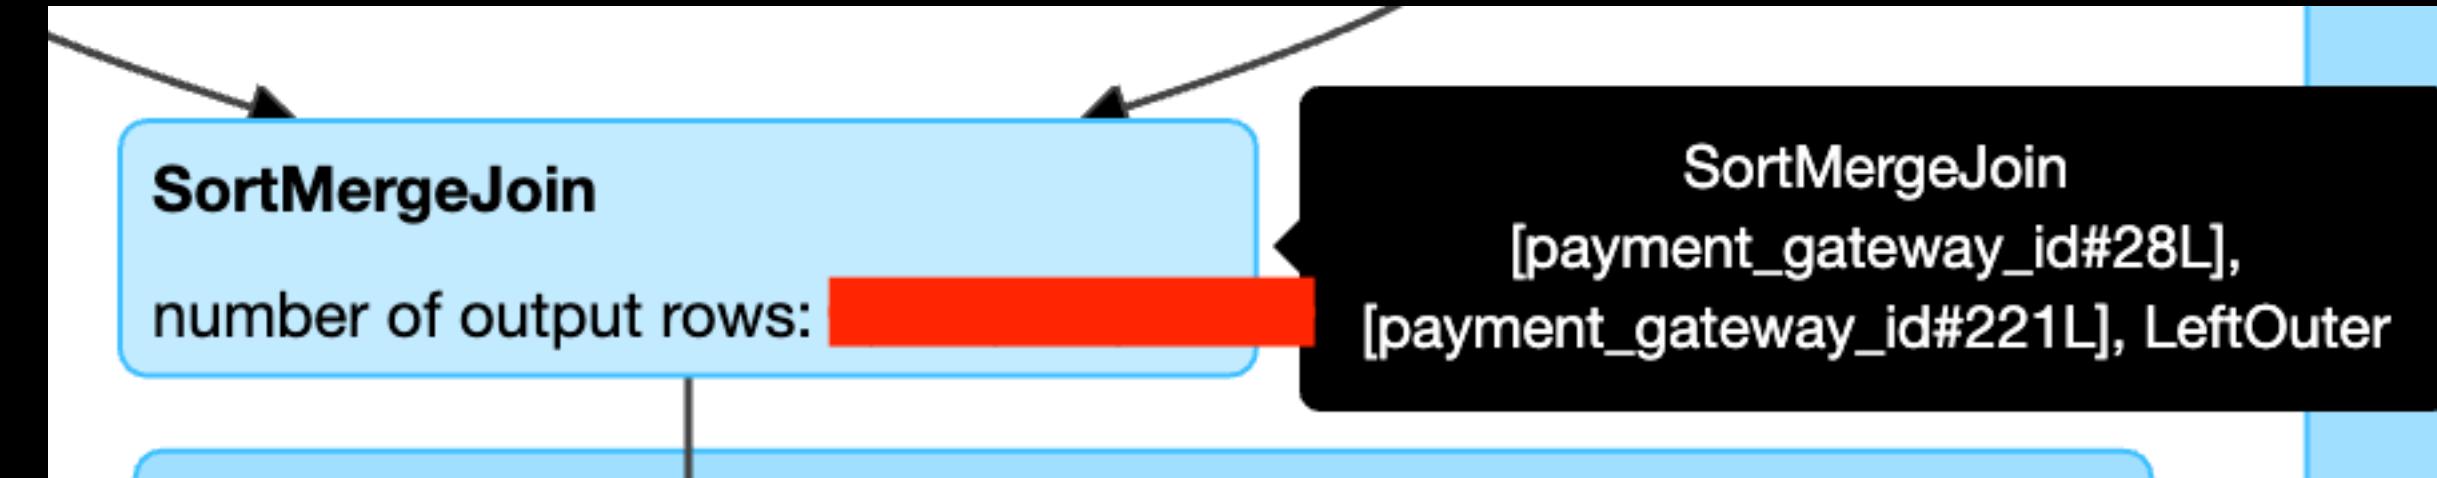
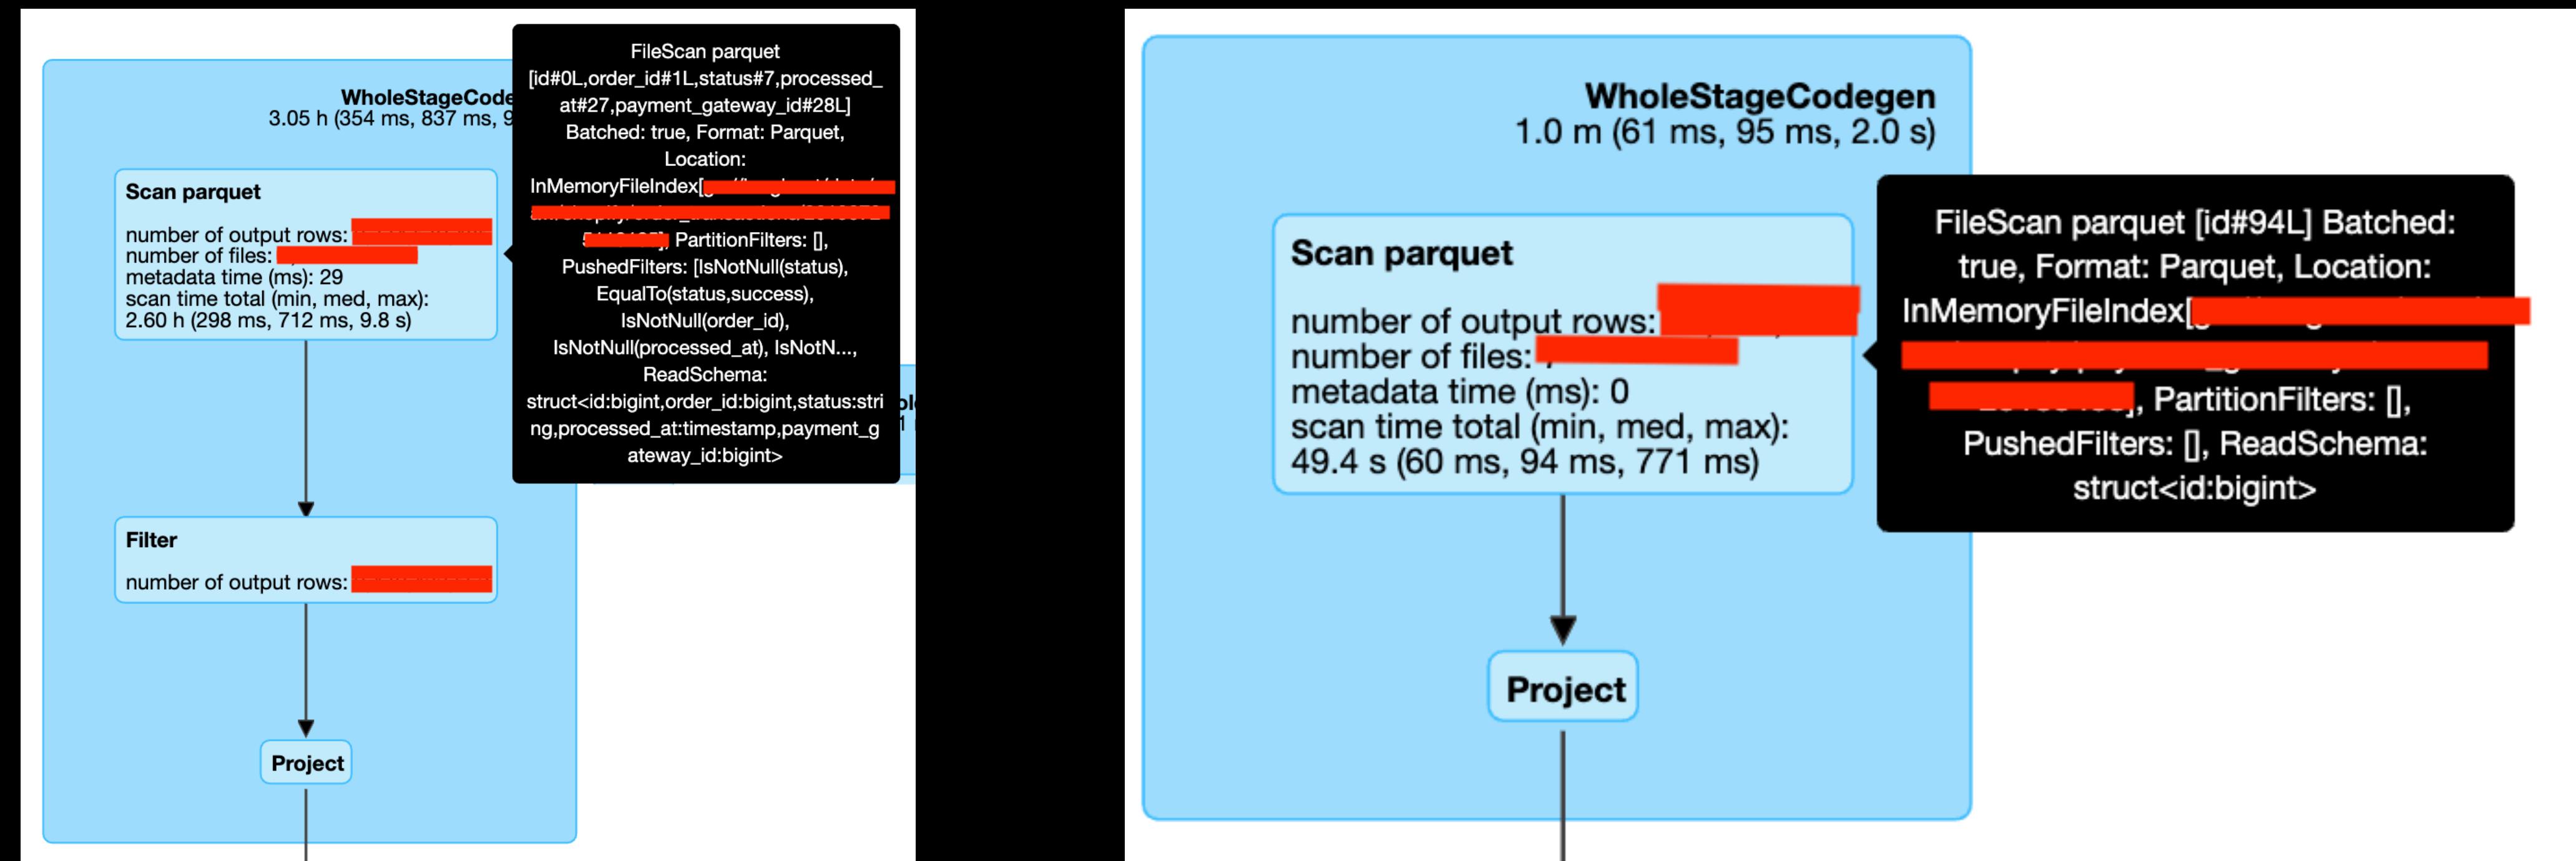
```
SELECT  
    card_type,  
    count(*)  
FROM order_transactions  
LEFT JOIN payment_gateways USING (payment_gateway_id)  
LEFT JOIN payment_details USING (payment_detail_id)  
GROUP BY 1  
WHERE ....
```





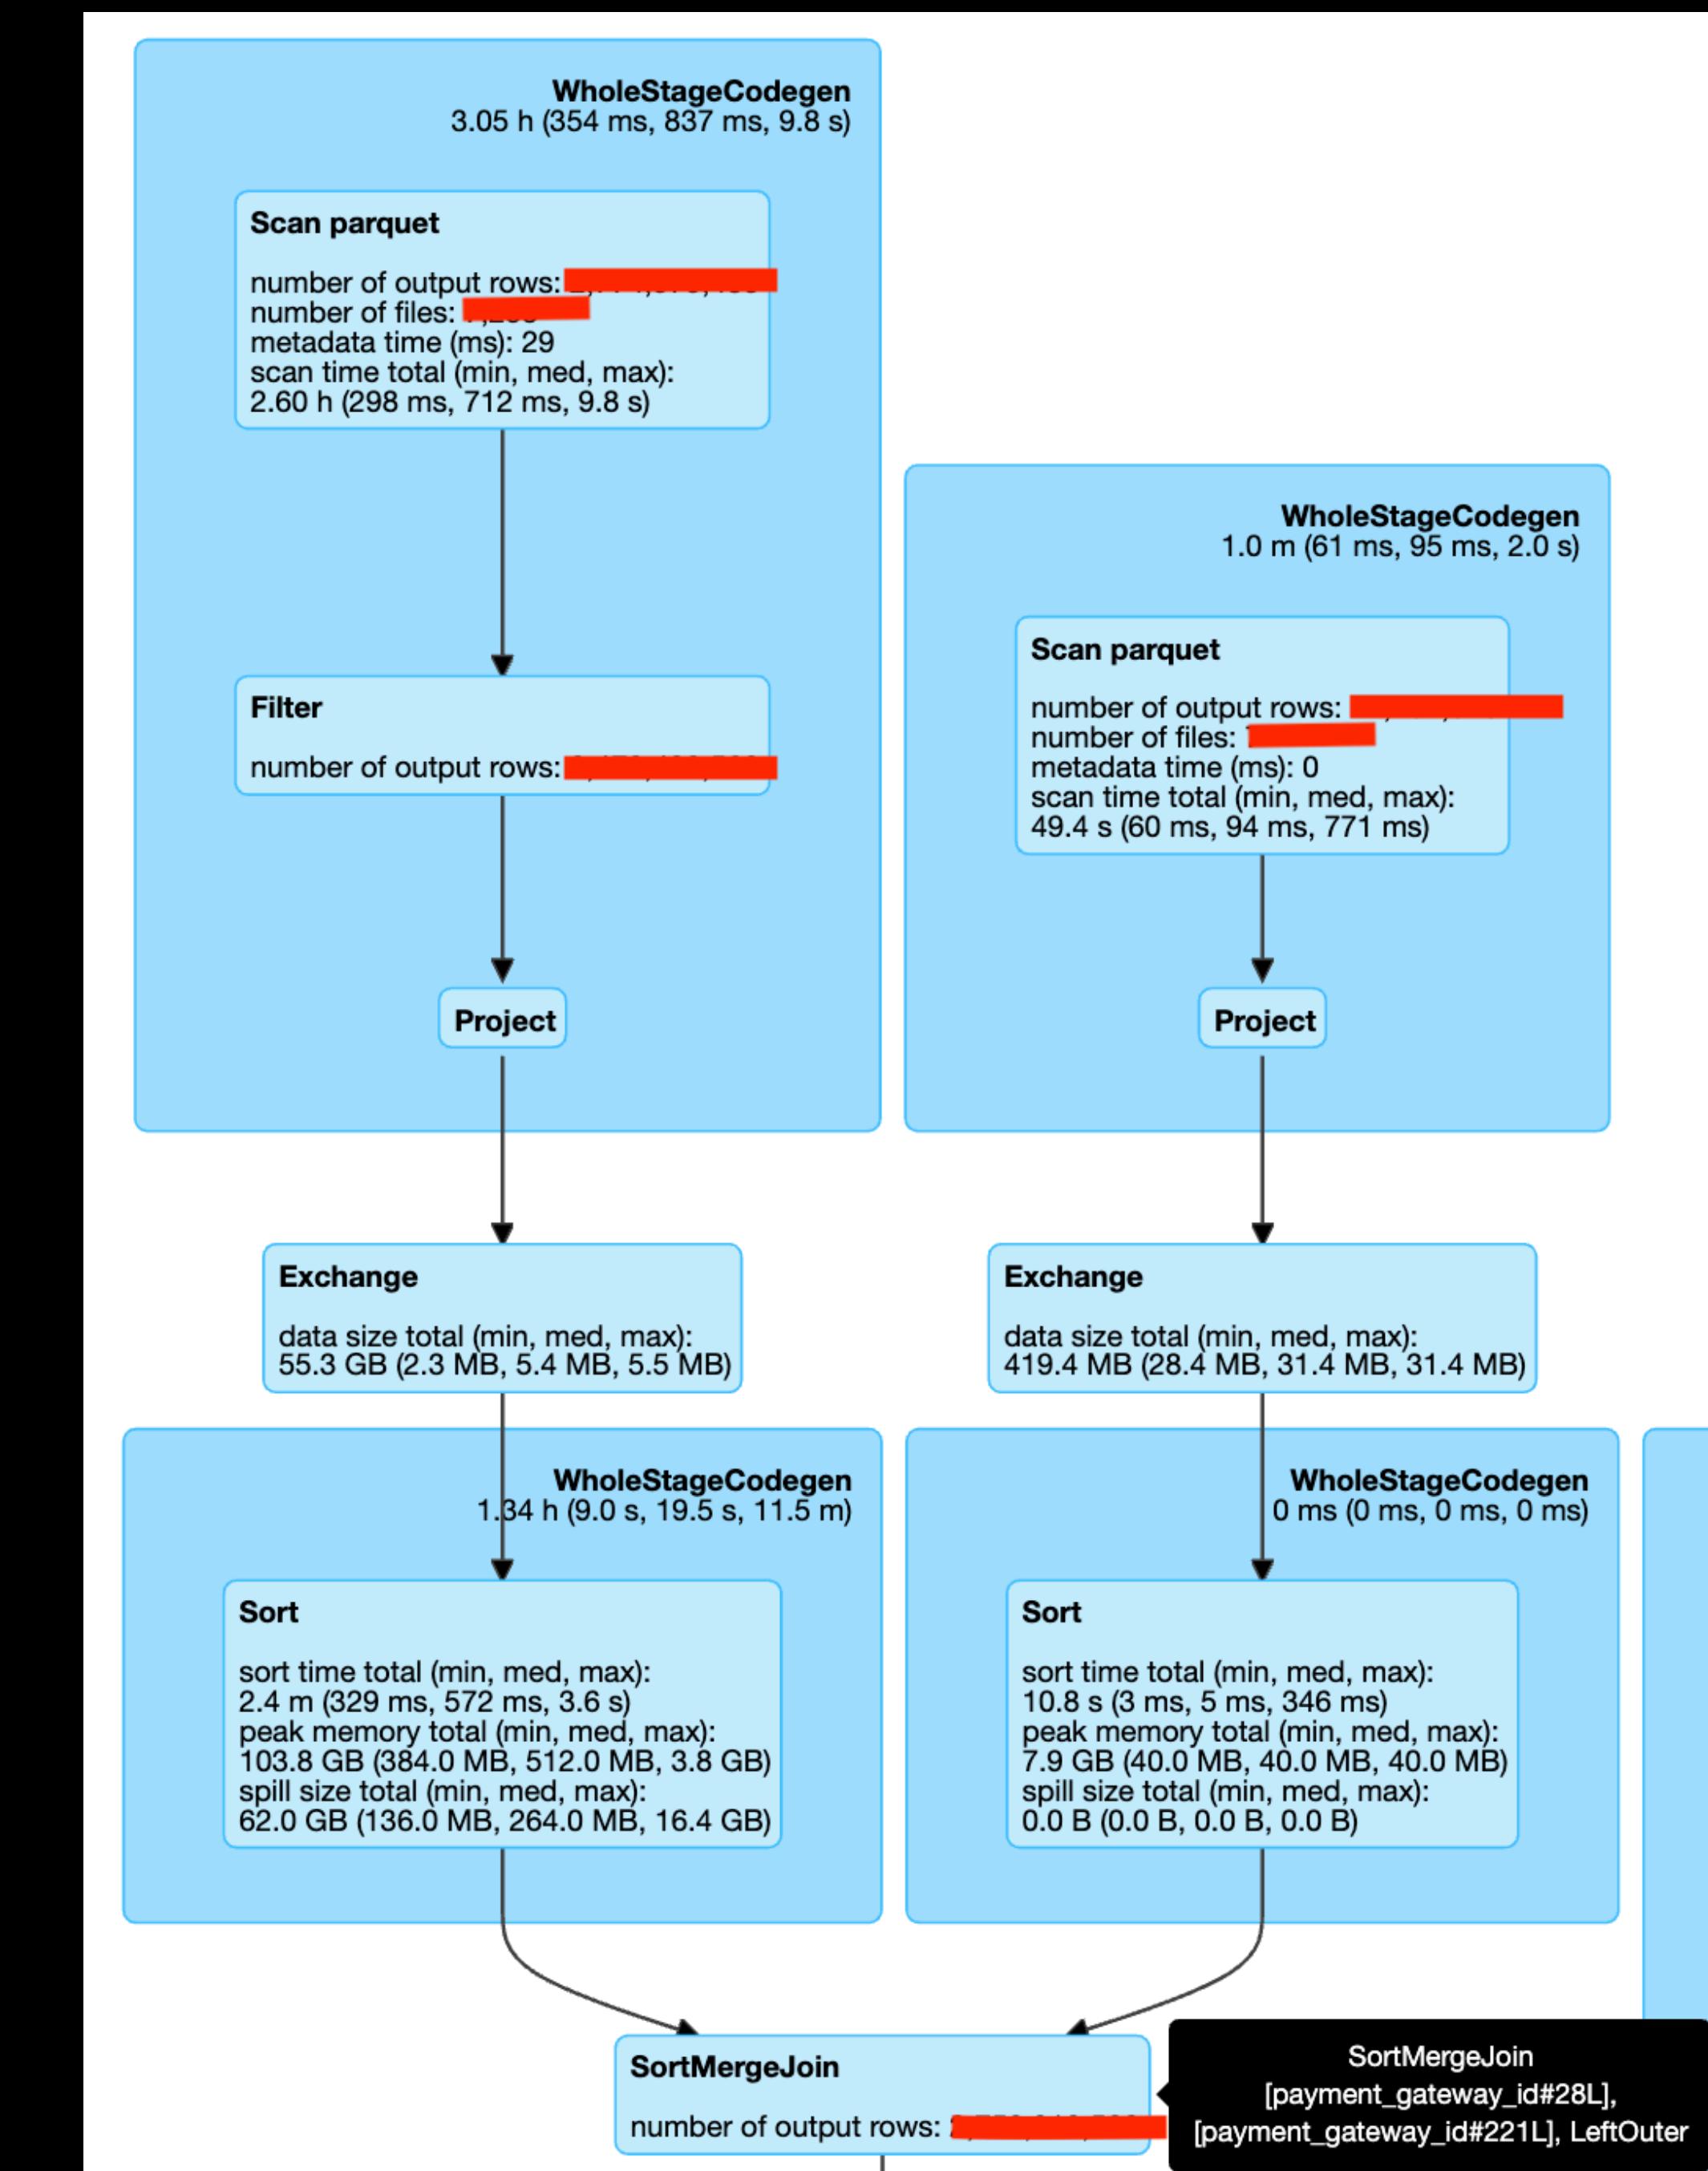






```

SELECT
  card_type,
  count(*)
FROM order_transactions
JOIN payment_gateways ON payment_gateway_id
JOIN payment_details ON payment_detail_id
GROUP BY 1
WHERE ....
  
```



# 4. Aggregate Level Stats

- Spark will provide some overall values for the min / median / max values.
- These values are for the task duration and data volume .
- They are provided on a per stage / transformation level.
- Focus on the blocks around a **join** block.

## RED FLAGS:

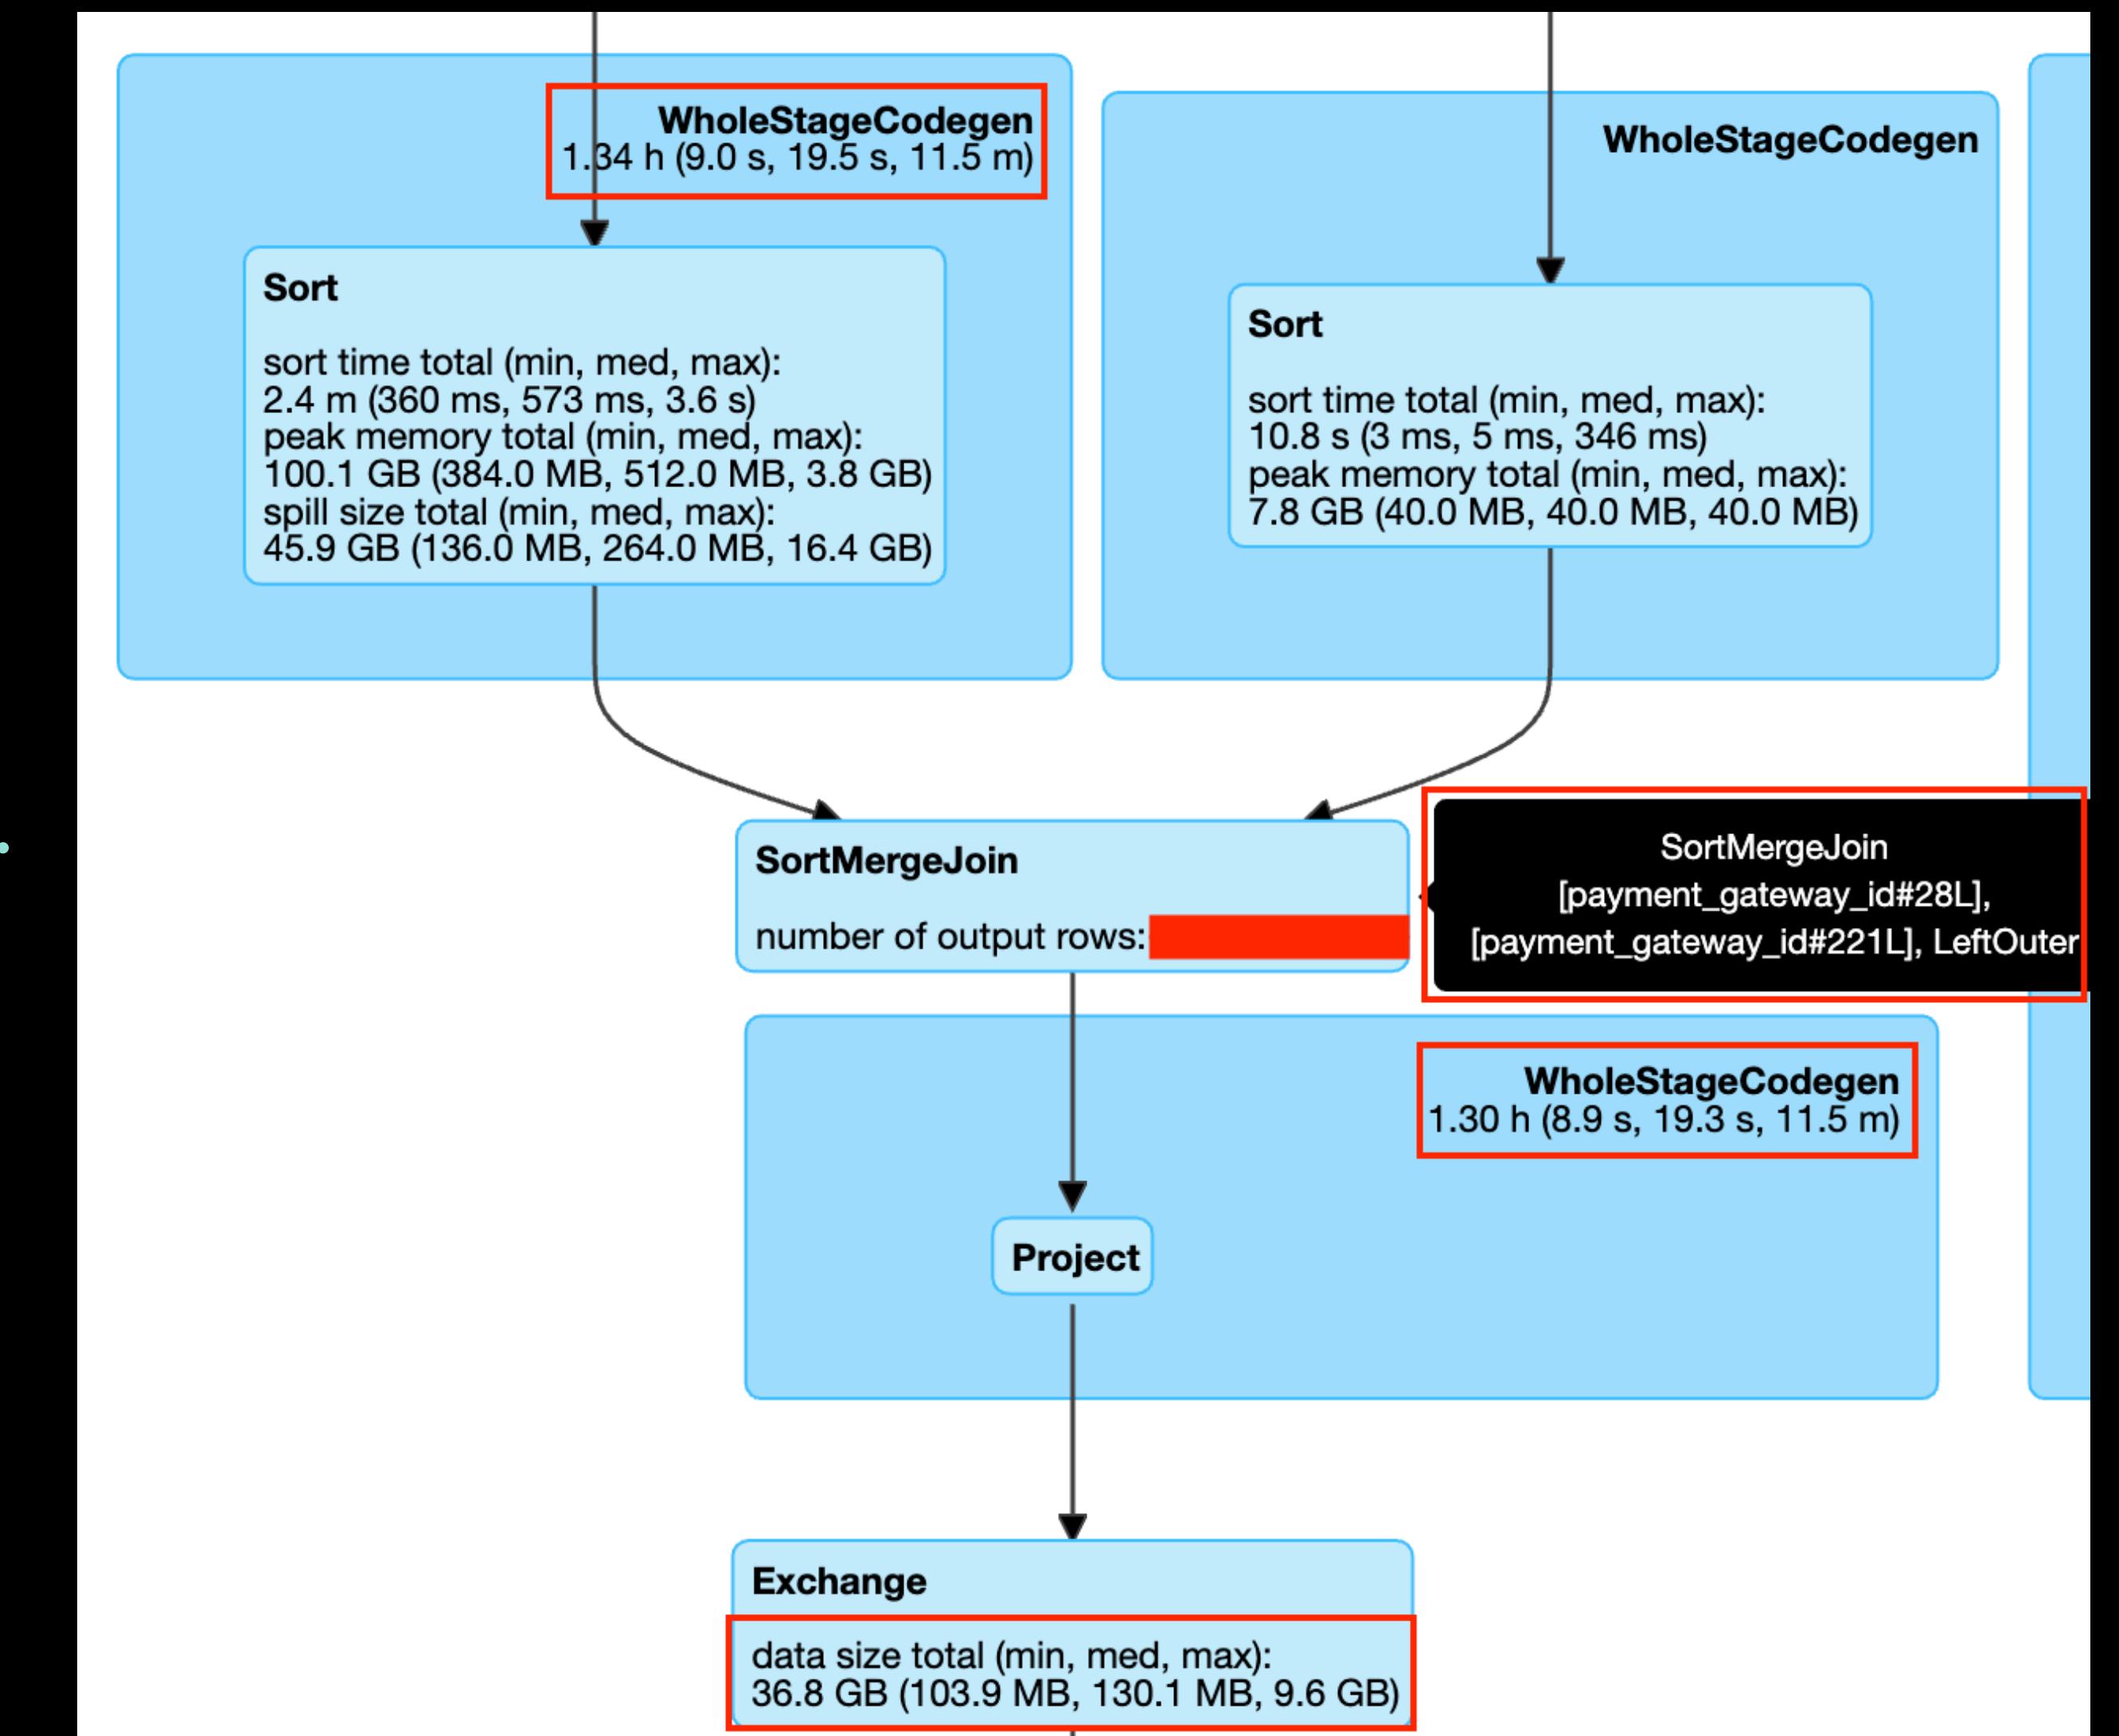
- Highly varying times or data volume.

## Solution:

- Skewed join, use a join optimizer.



- 9 seconds vs 11.5m
- Around the join onto payment\_gateway\_id.



# Goal 1

**Increase Join Performance. Addressing slow running tasks in a (skewed) join.**

**Motivation:**

**When Spark performs a join, spark distributed the data based on the join key to the tasks. If the number of rows per task is not evenly distributed the time it takes per task can vary, thus causing the slower stages to bottleneck the whole application.**

**Sometimes the skew is so large that the data struggles or can't fit in memory.**



# Joins (skew)

2 types of skewed joins:

1. NULL Skewed (only for left joins).
2. Key(s) Skewed (only left and inner joins).

These problems can also appear in our resolvers as well, but they might be implemented differently.



# Determining the Skew Type

## Steps:

1. Run an analysis (dev skew ...) on the join key.
2. If **NULL values** make up a majority of the data, data is **NULL Skewed**.
  1. Replace the join with our **NullValueSkewHelper** join class.
3. If a **single non-null key** makes up a large portion of the dataset, **Key Skewed**.
  - Replace the join with our **SkewPartitioner** join class.
4. If **multiple non-null keys** make up a large portion of the dataset, **Keys Skewed**.
  - Replace the join with our **FrequentValuesSkewHelper** join class.



# Dev Skew

- Calculates the distribution of a provided column in a given dataset.
- Finds the minimal and maximum partition size.

```
1 dev skew analyze --field 'column' --path  
2 gcs_path  
3  
4  
5  
6  
7
```

\*If you have to include legal or source

# Dev Skew

- Shows top 5 most frequent values.
- The count and percentage.
- Shows min and max partition data size.

High Frequency Values (Top 5)

Value	Frequency	Percentage
NULL		23.61
4312925		1.18
6993804		0.74
57217927		0.63
94425479		0.58

# Partitions	Minimum Row Count (Percentage, Size)	Maximum Row Count (Percentage, Size)
200	7826449 (0.28, 5.1 GiB)	663687829 (23.91, 433.6 GiB)

# NullValueSkewHelper Under the Hood

## Motivation:

- Eliminate the shuffle of the null rows.
- Too many rows to fit into memory.

## Solution:

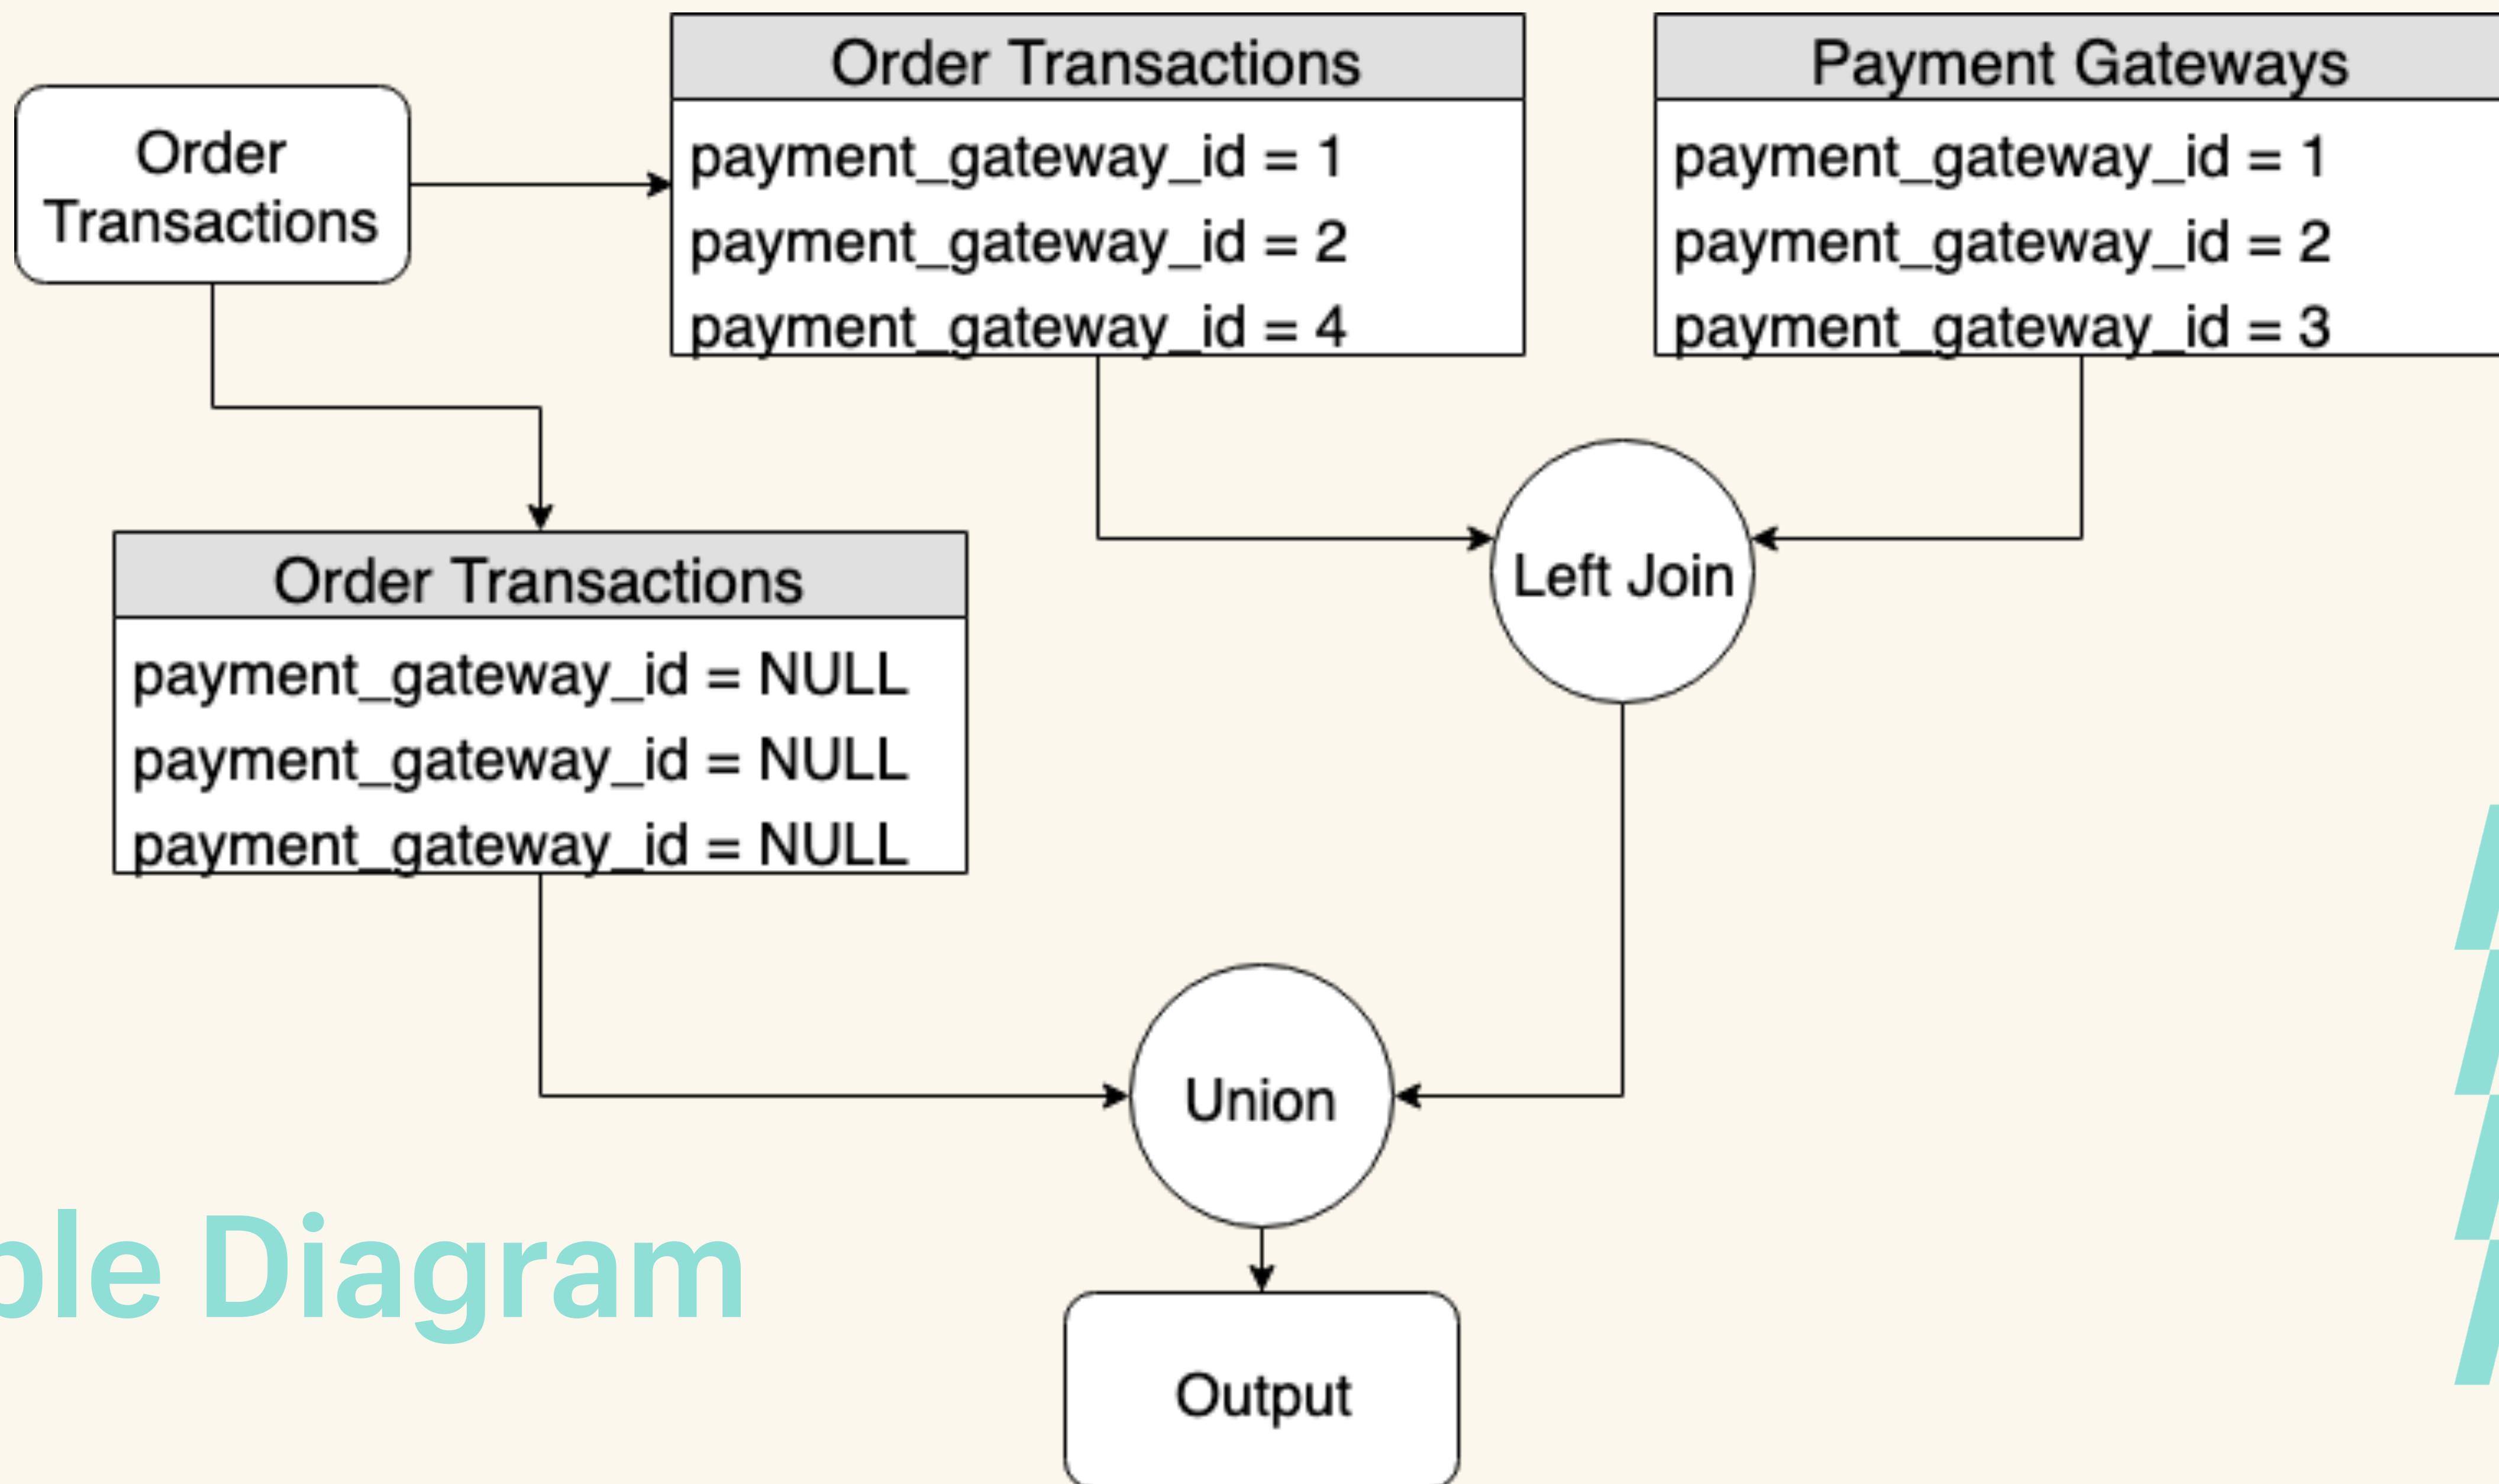
- Split the dataset into null and not null.
- Perform the left join on the not null subset.
- Union result with null subset.
- Unions don't cause a shuffle.



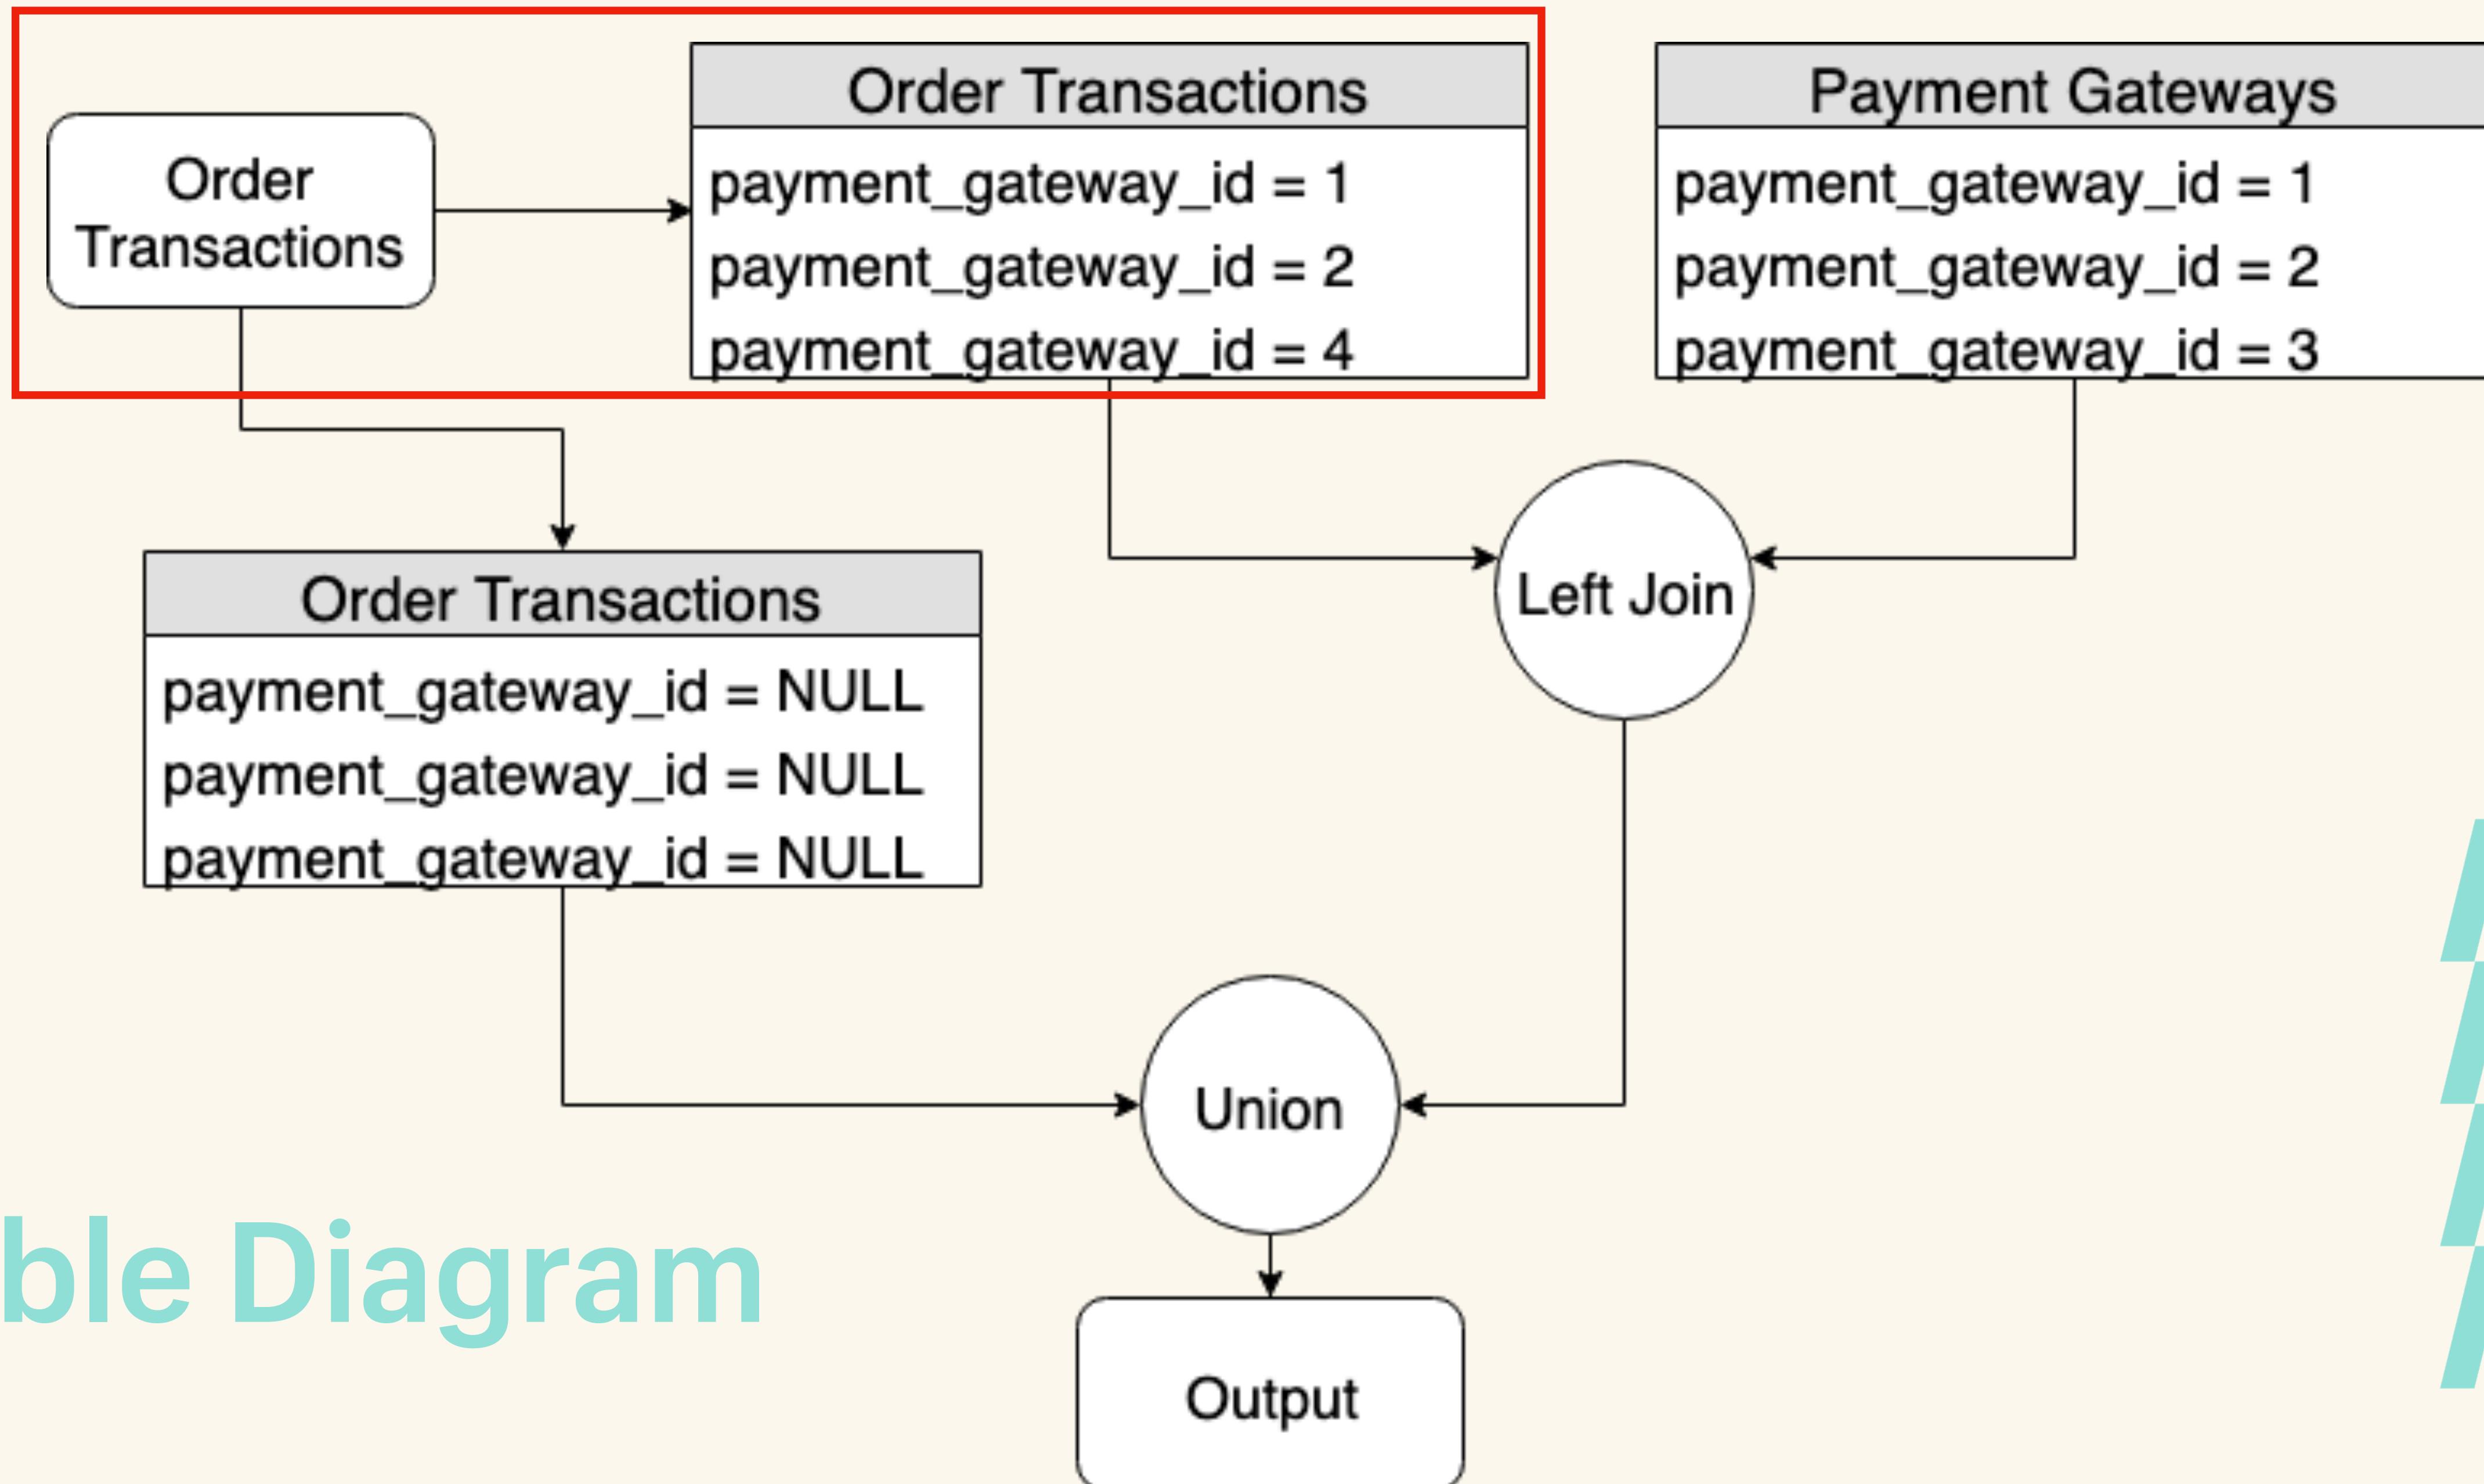
# NullValueSkewHelper Under the Hood

- Filter the null skewed dataset to not null values.
- Left joins the to the right table.
- Filter the null skewed dataset to null values.
- Create the right table columns with nulls.
- Unions the joined results with the null values.

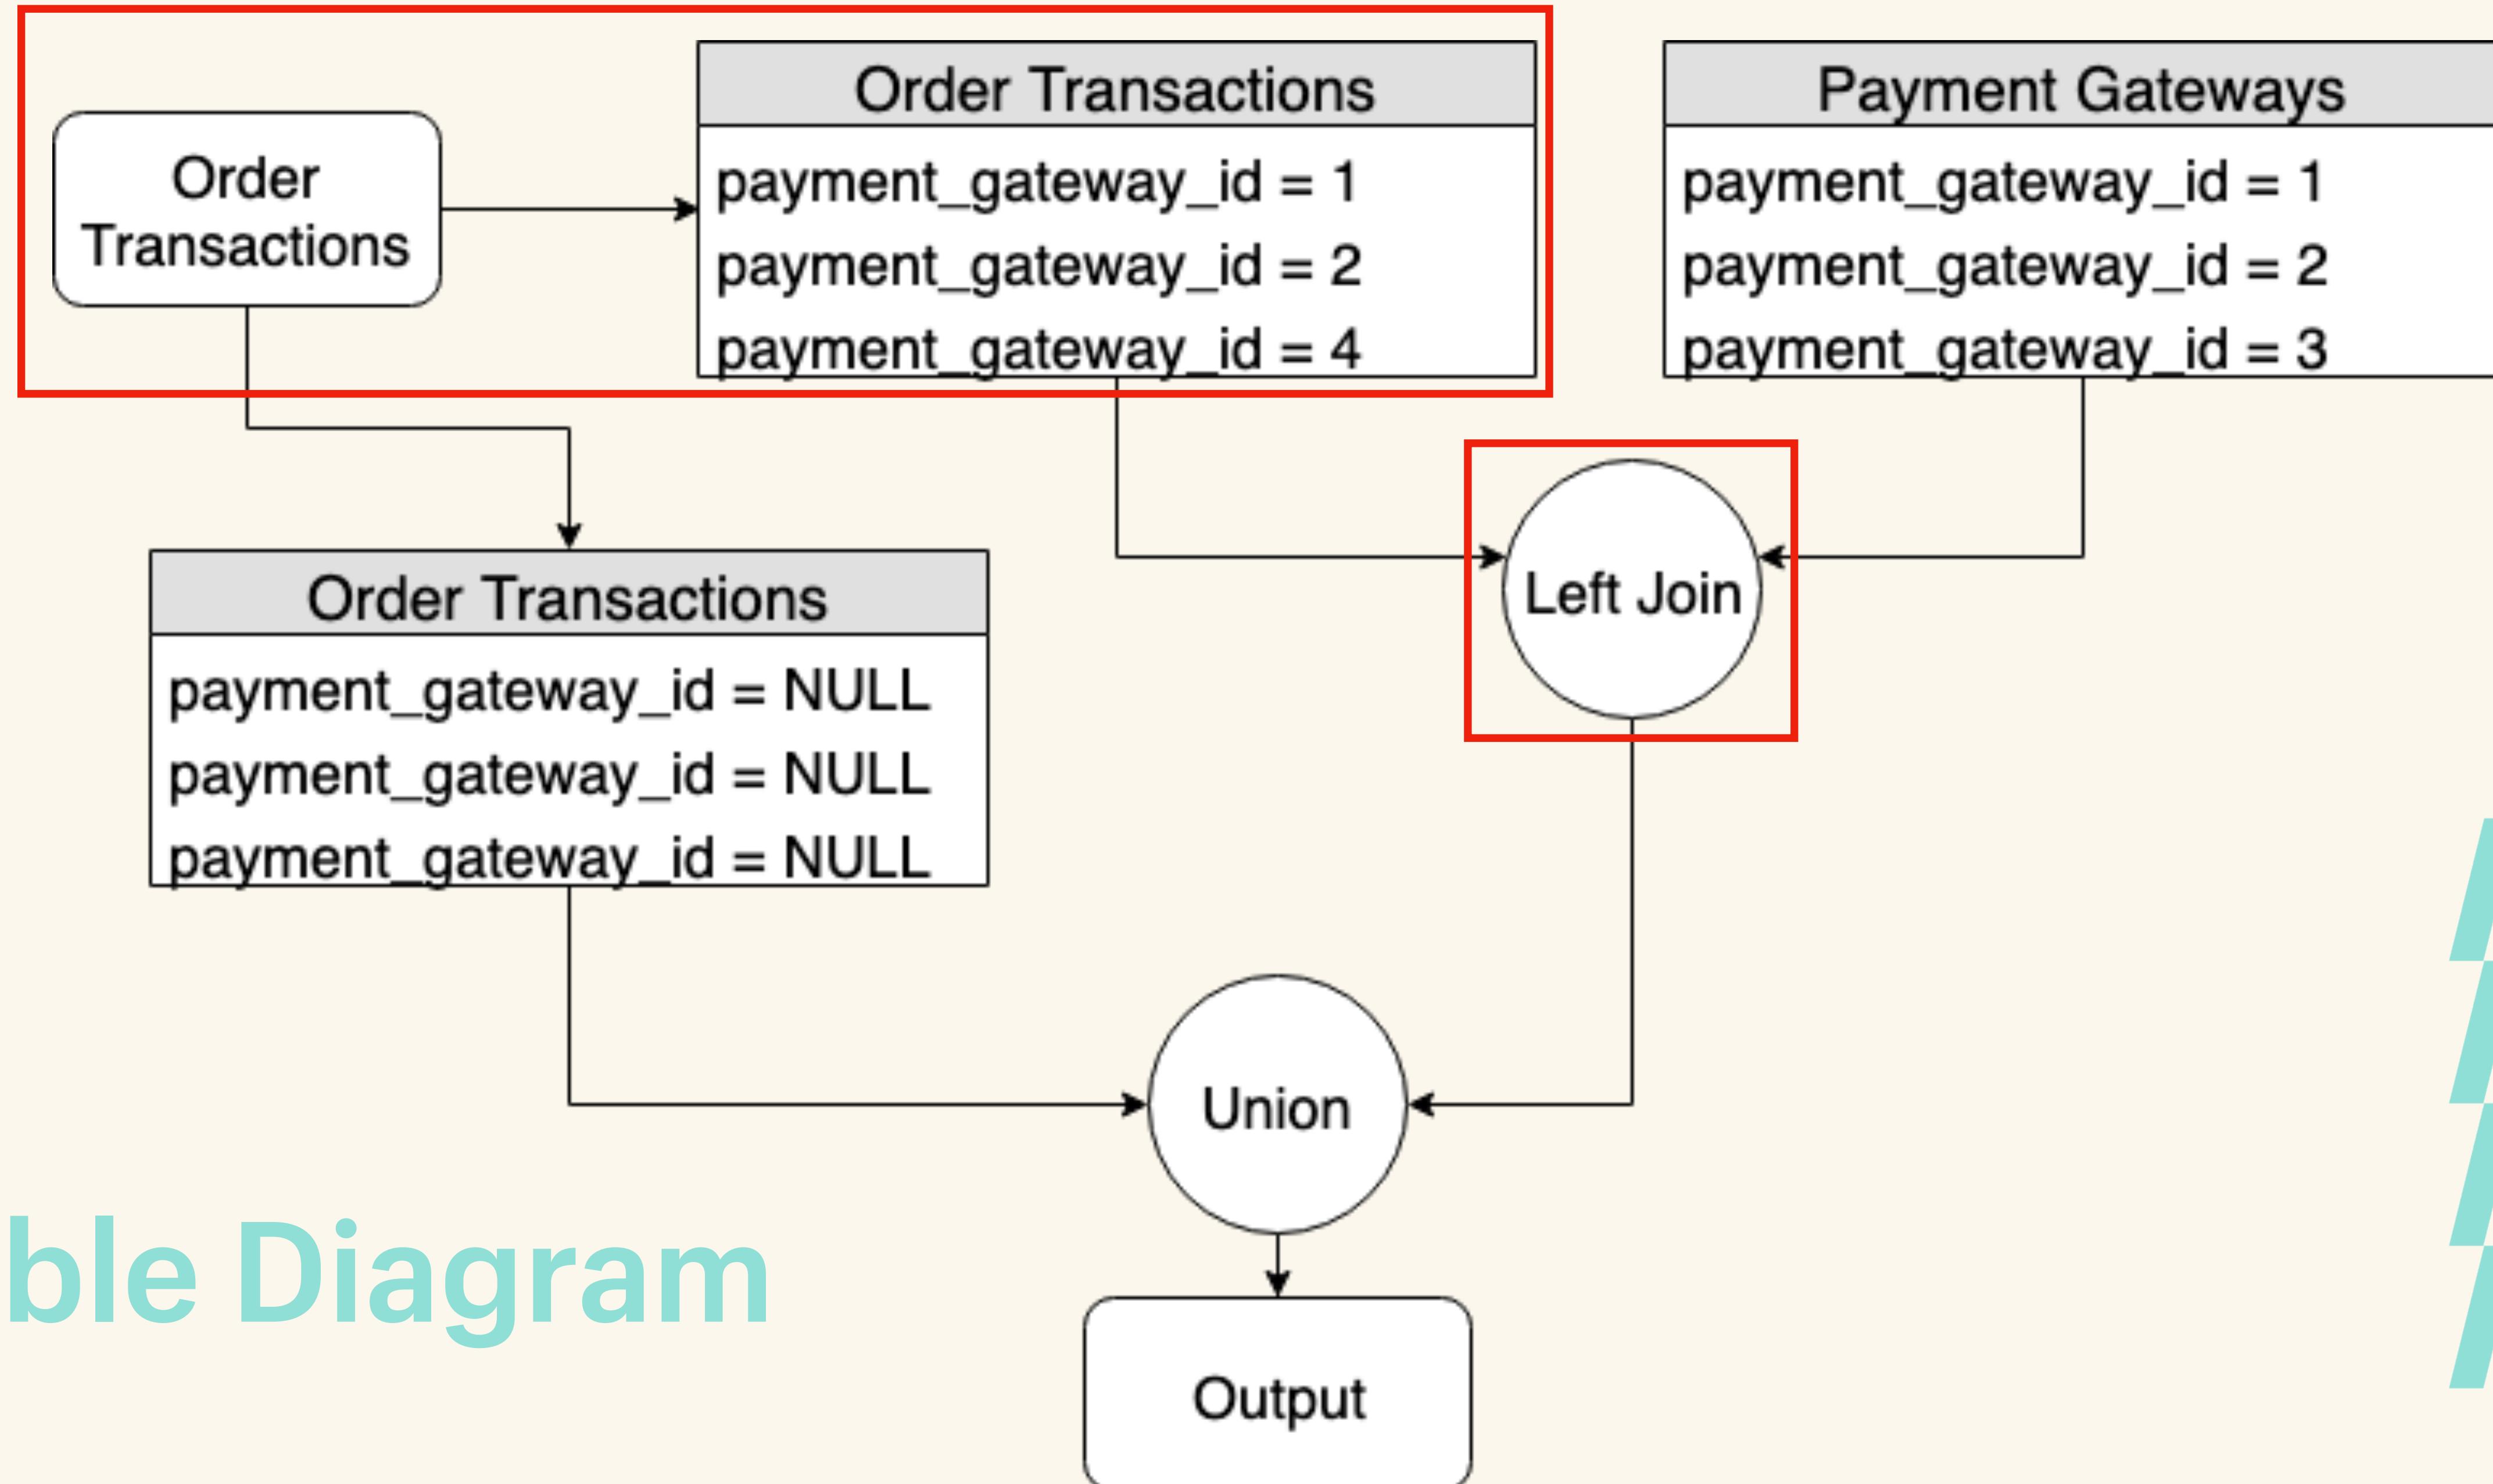




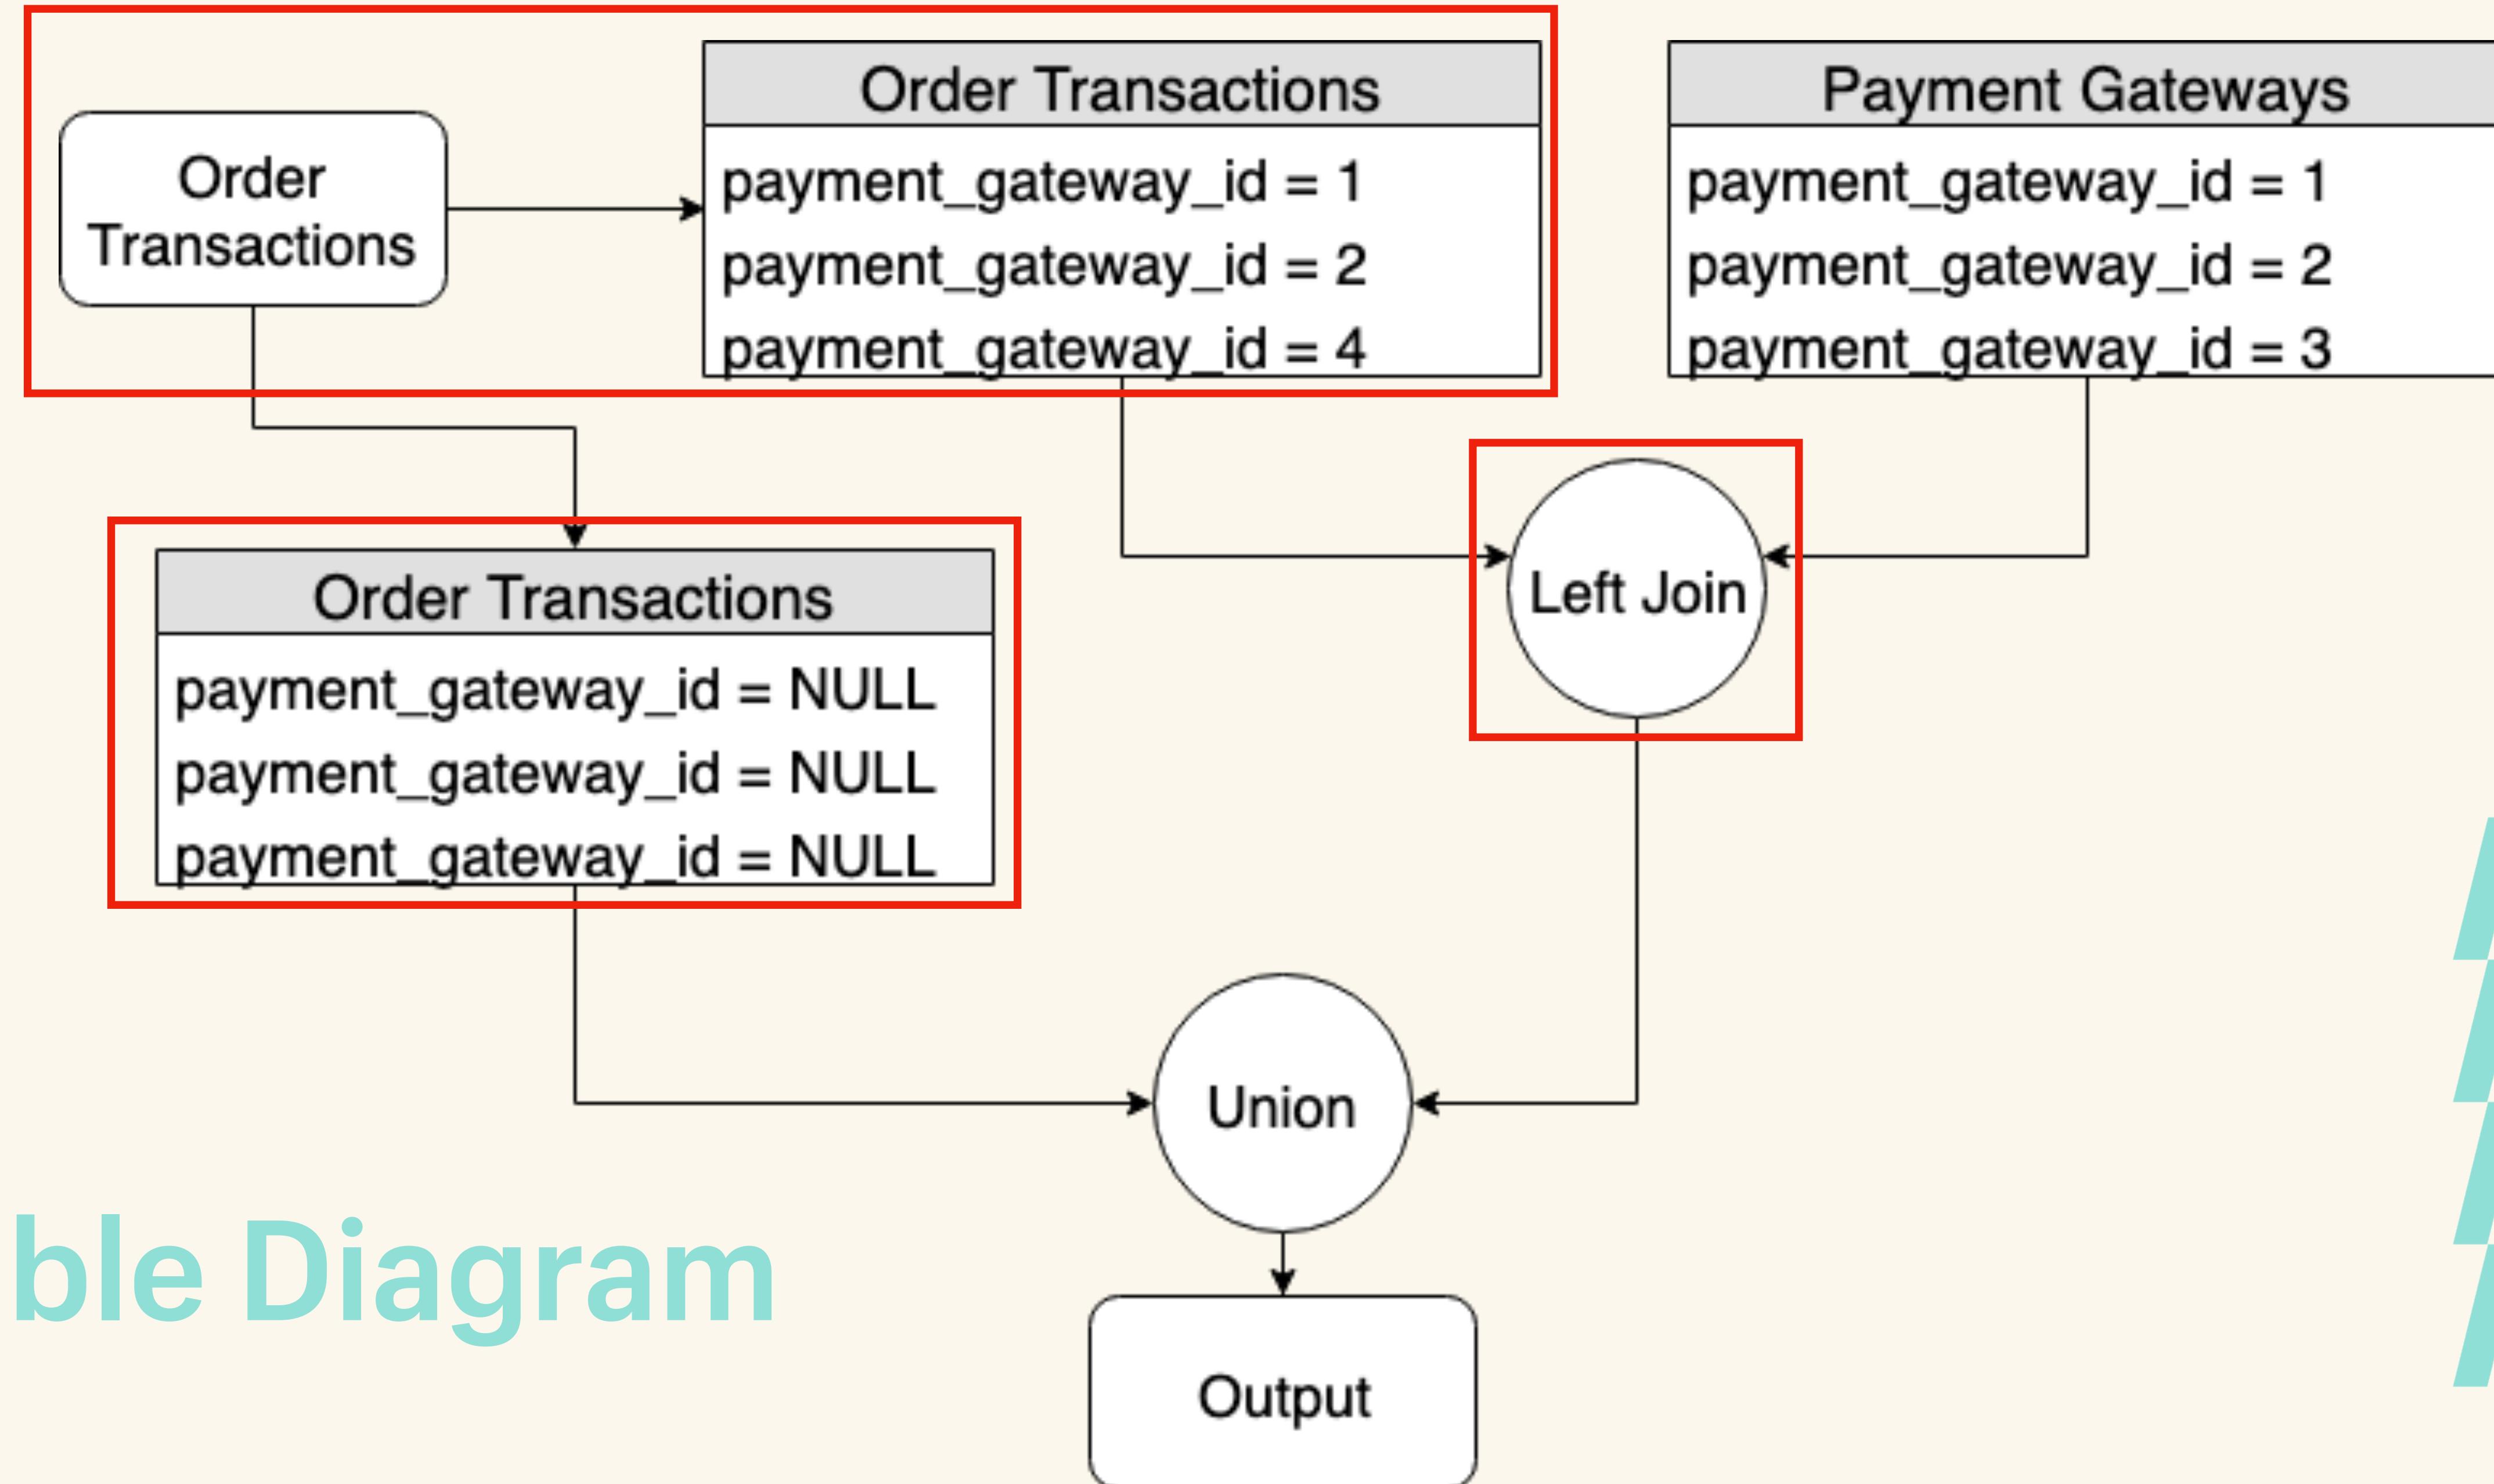
# Table Diagram



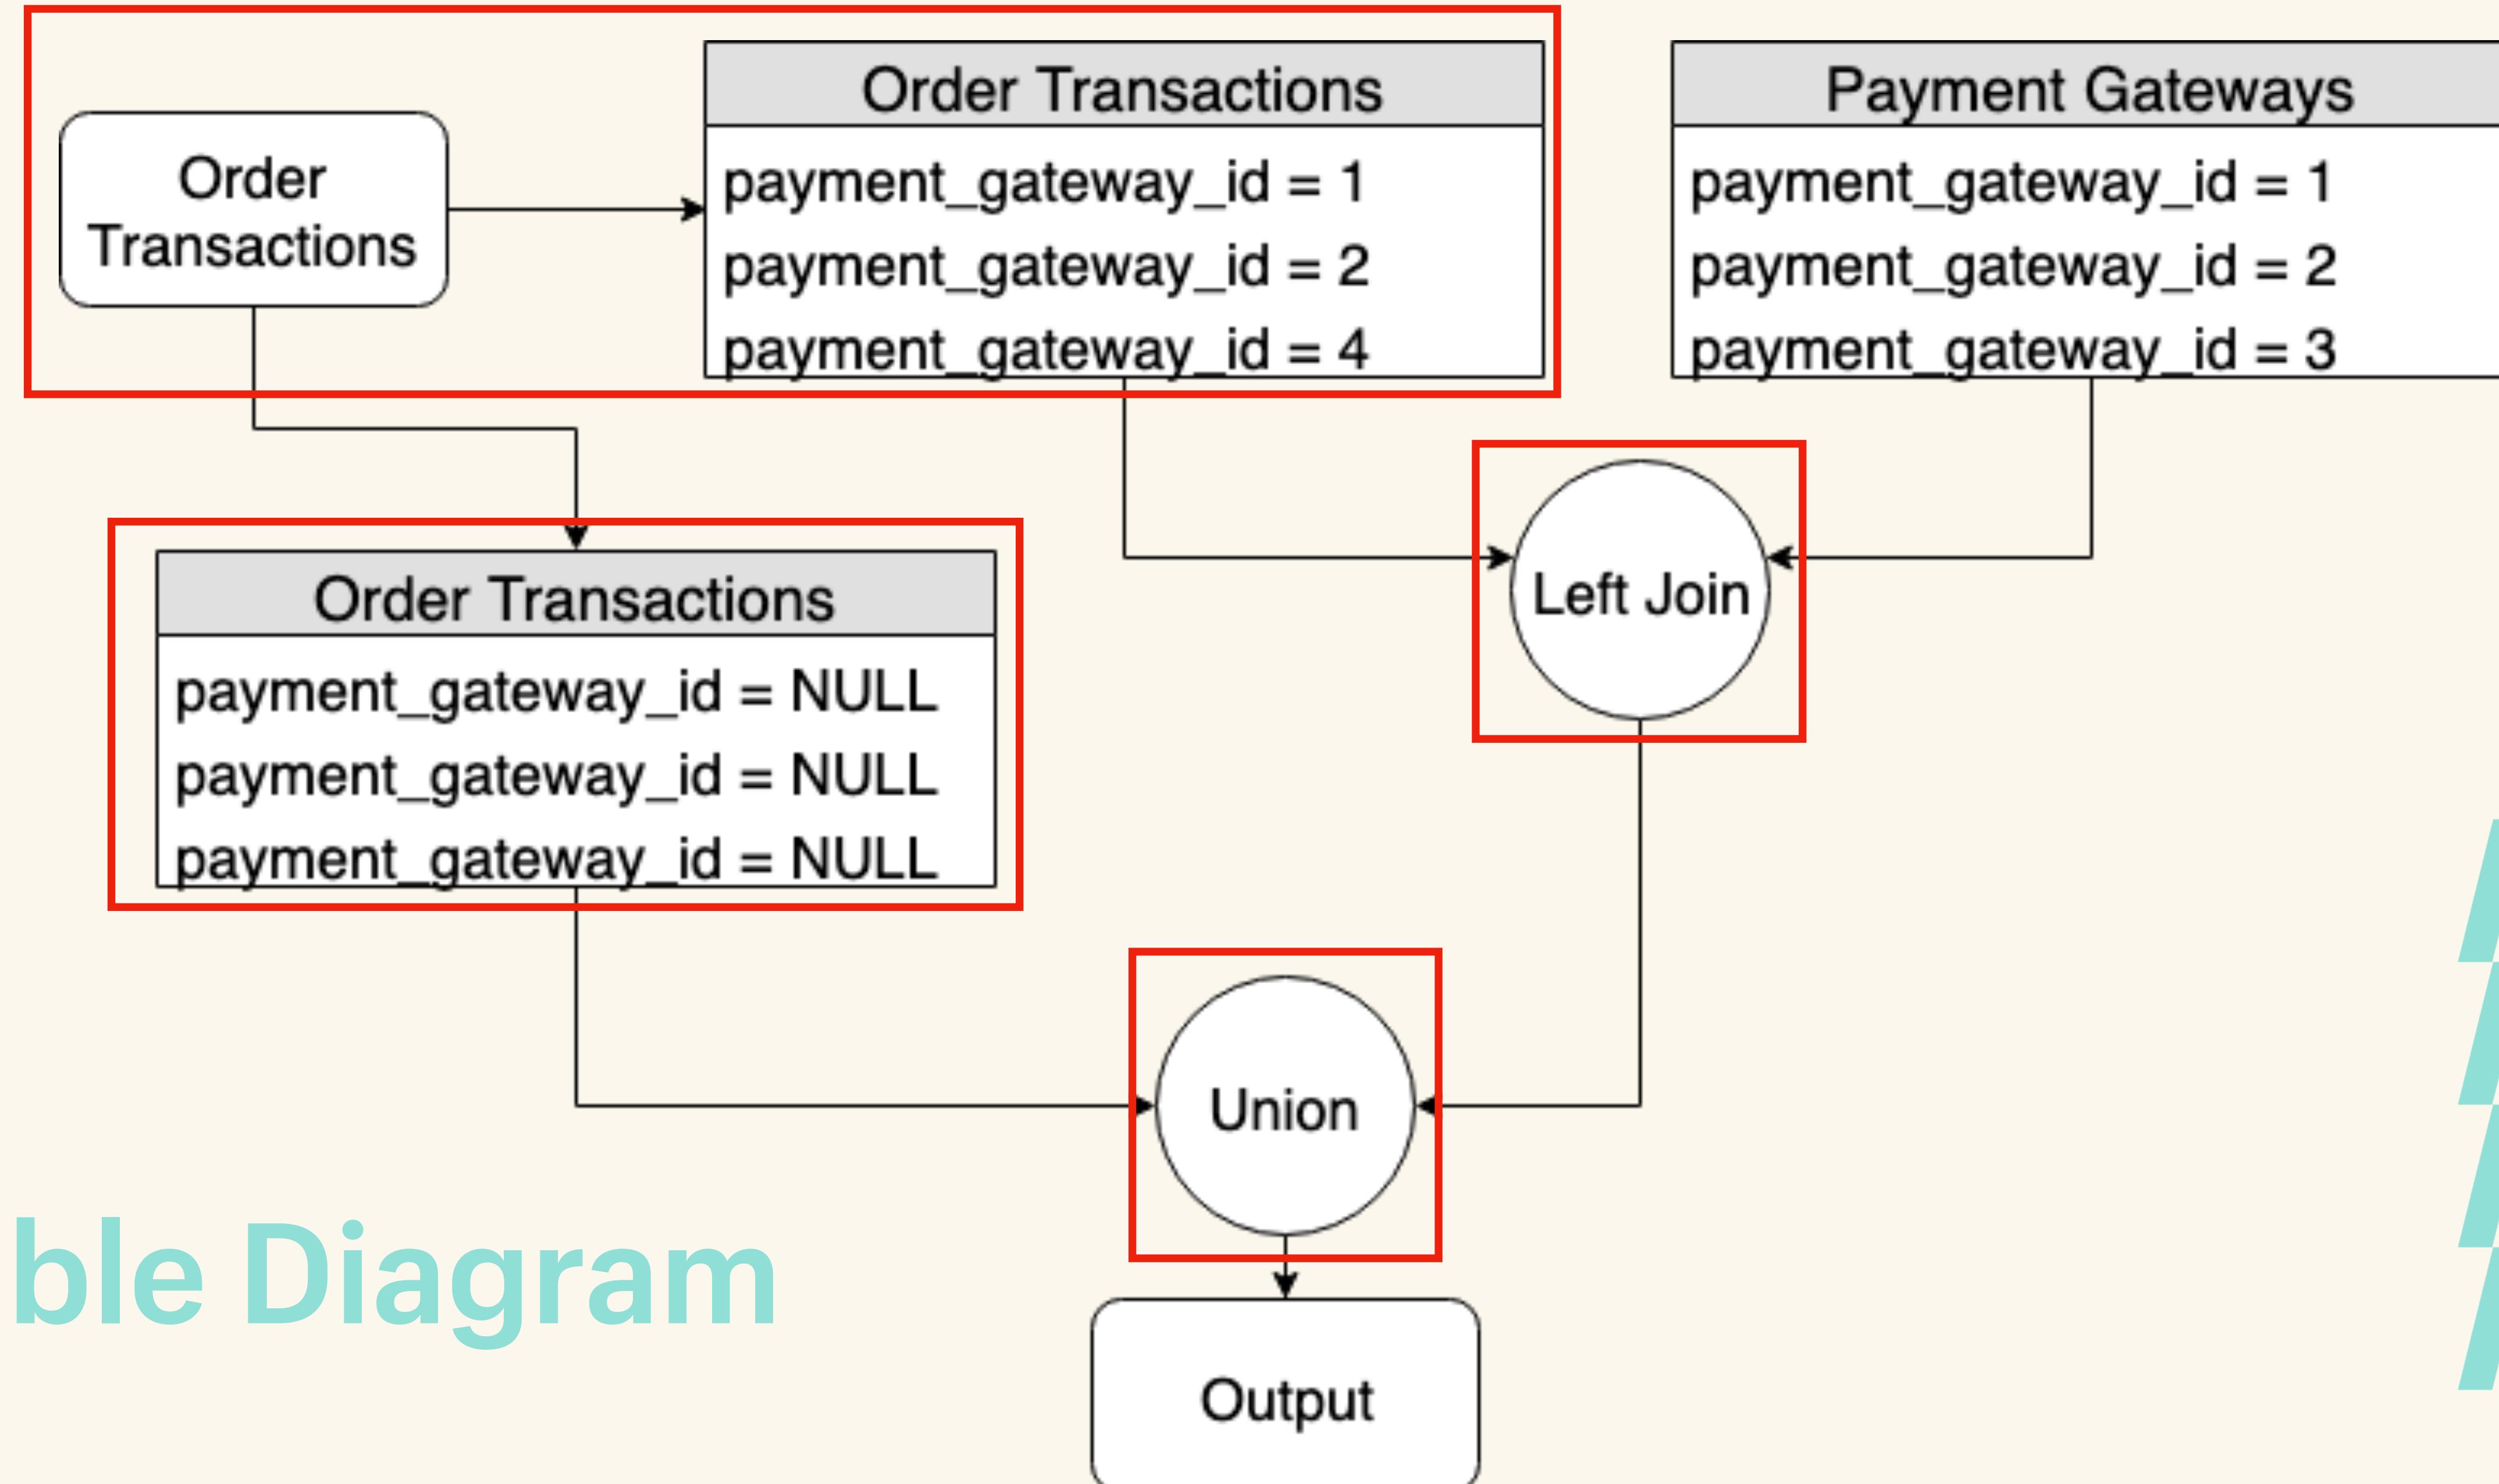
# Table Diagram



# Table Diagram

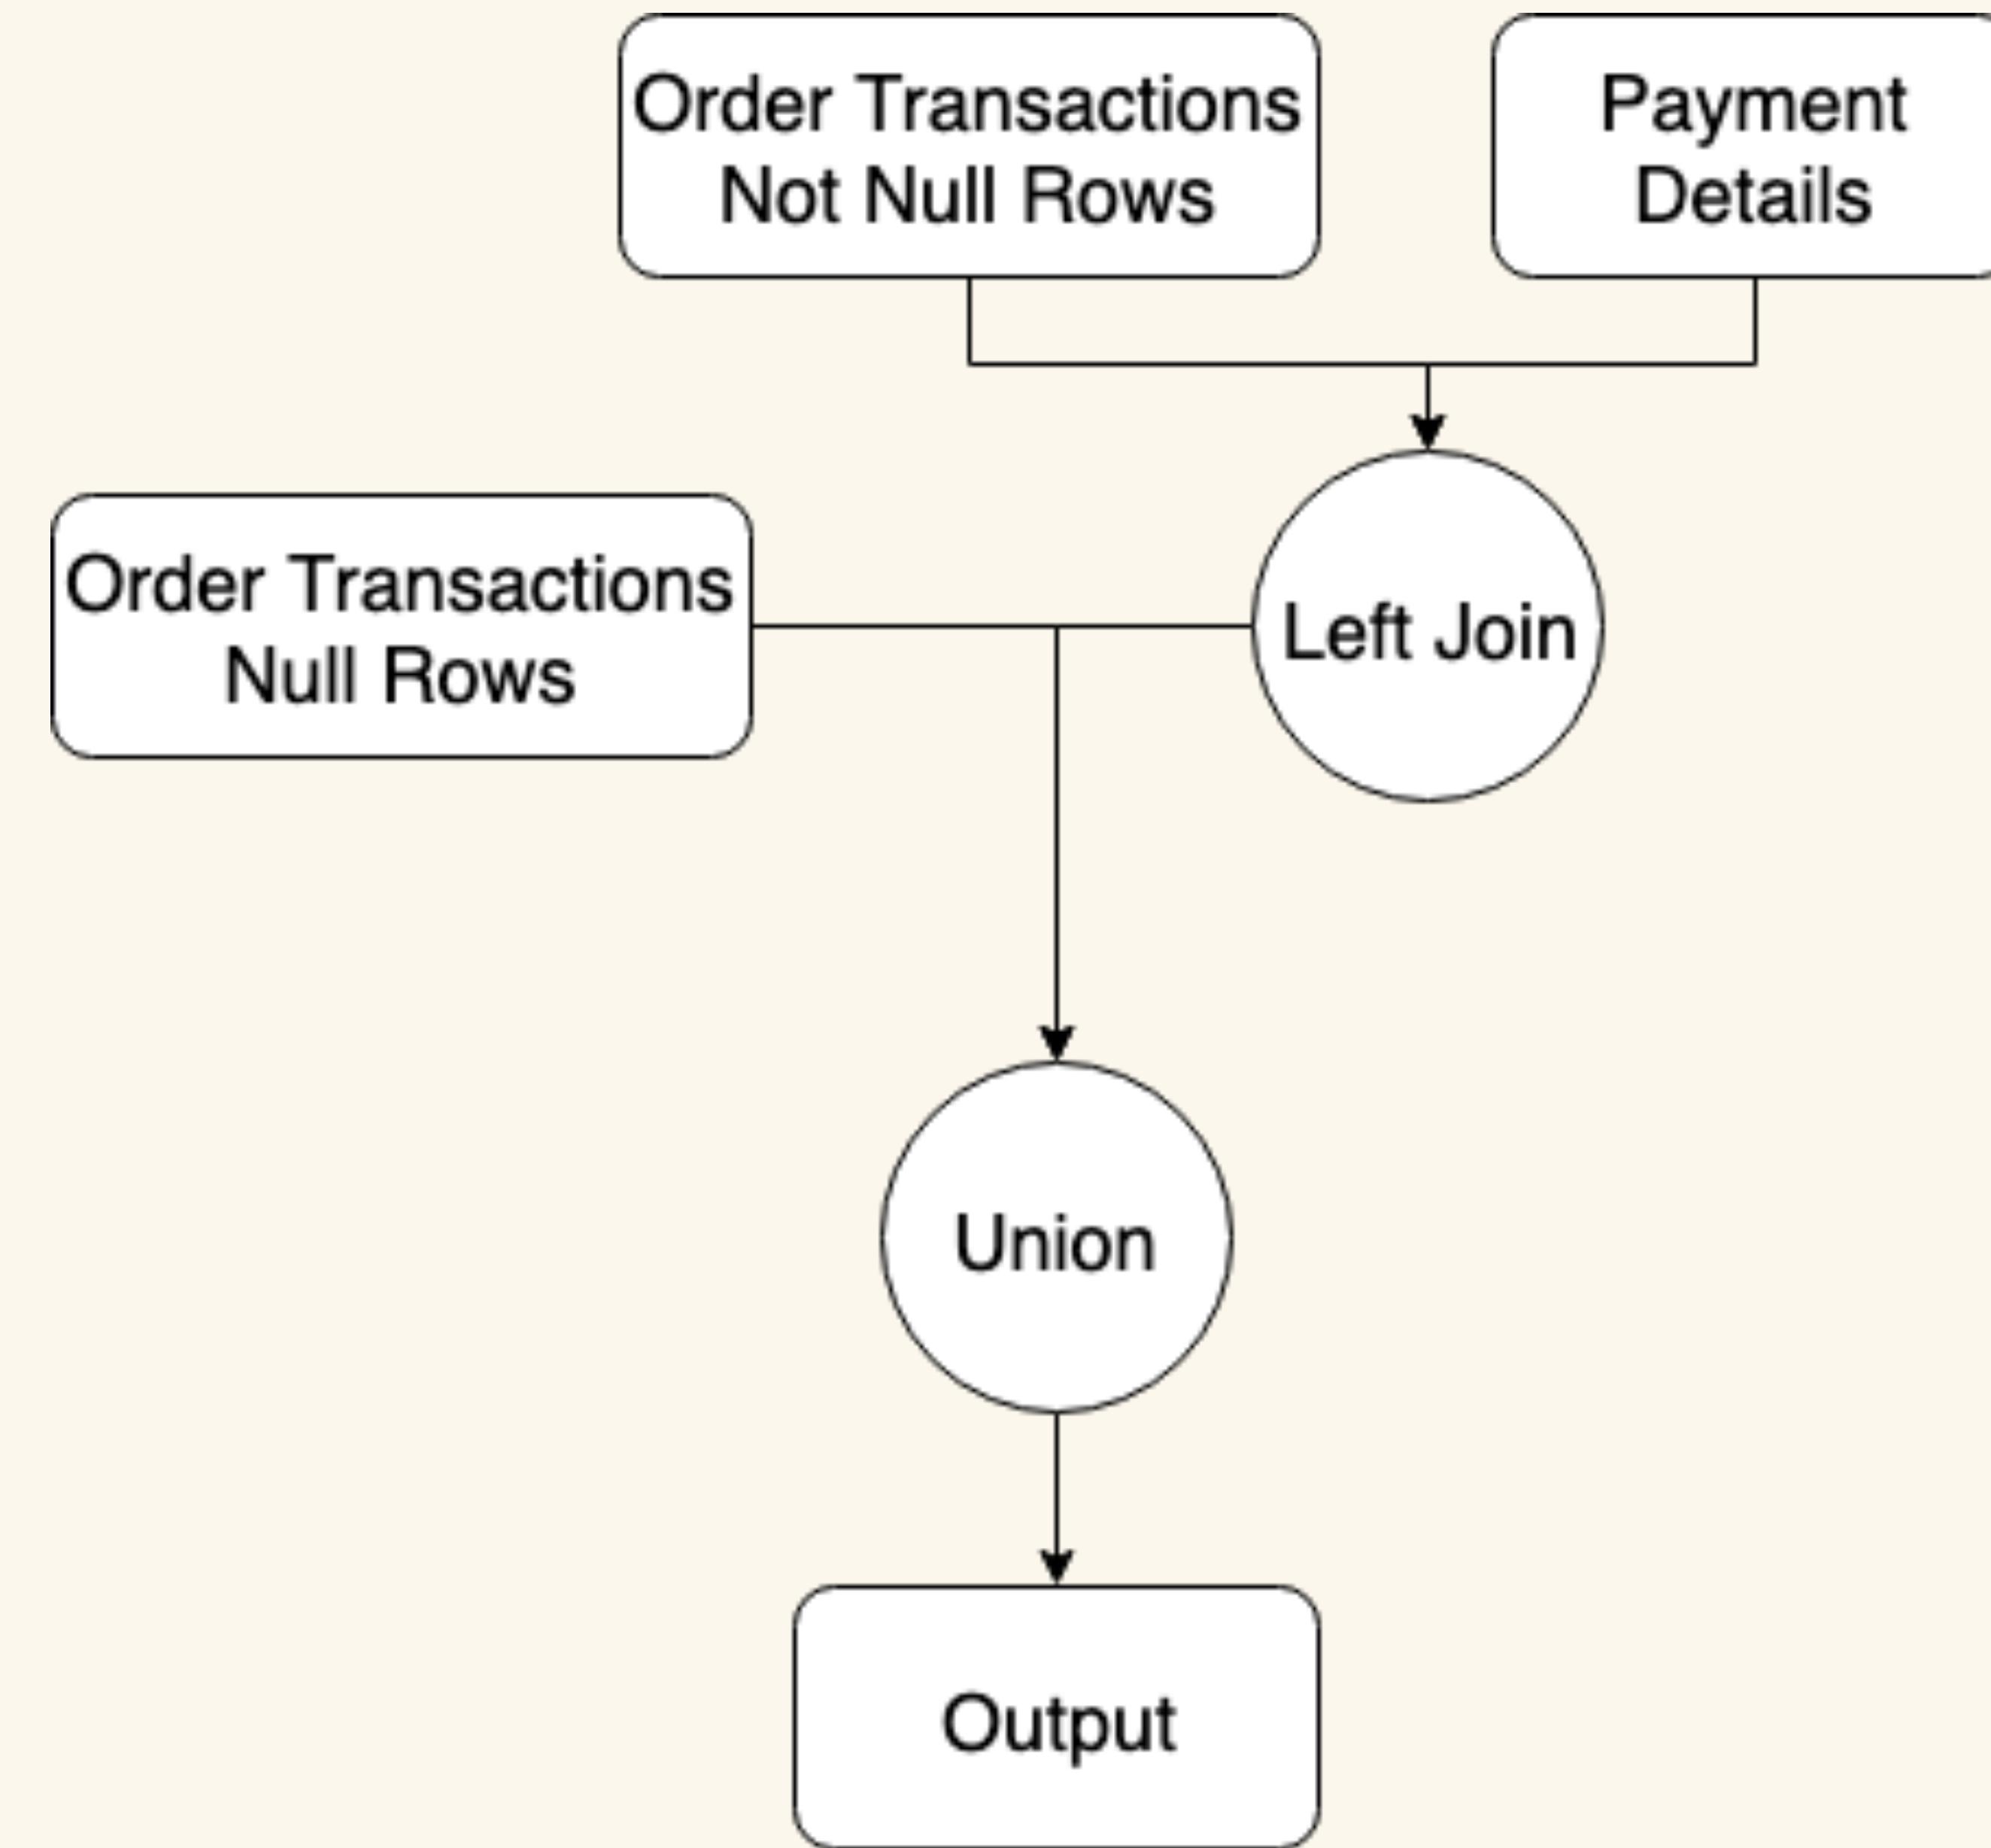


# Table Diagram

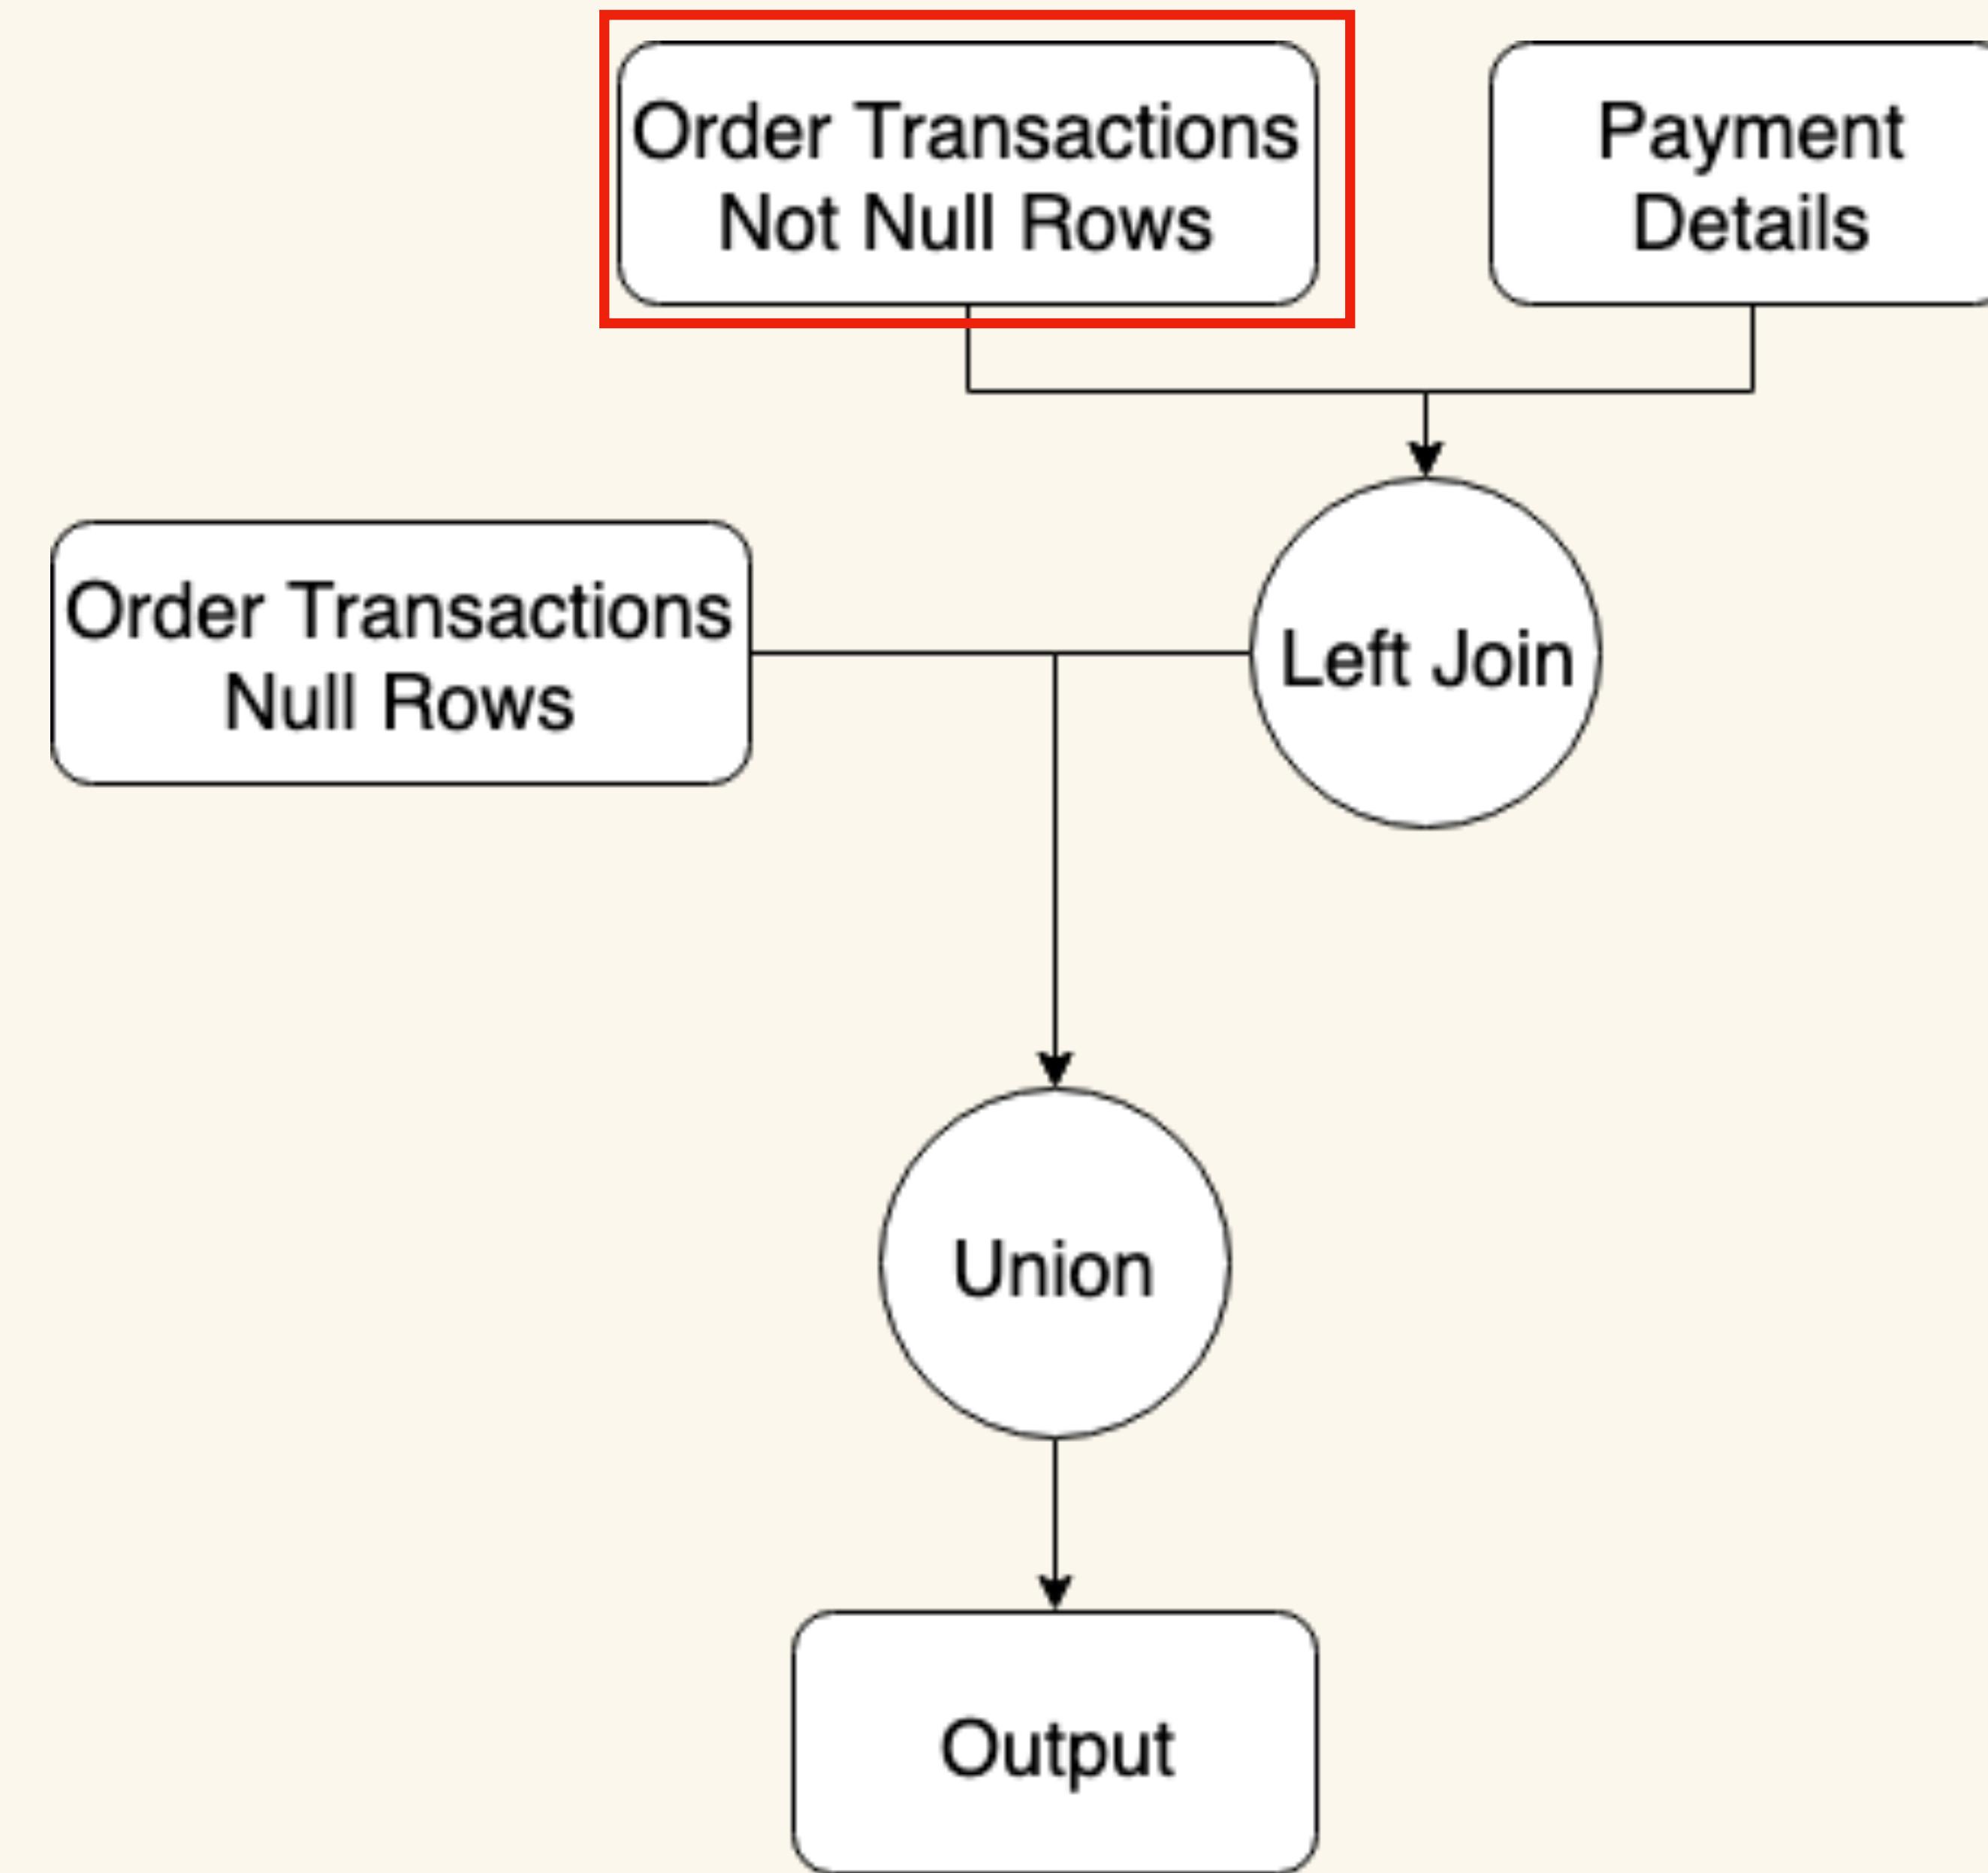


# Table Diagram

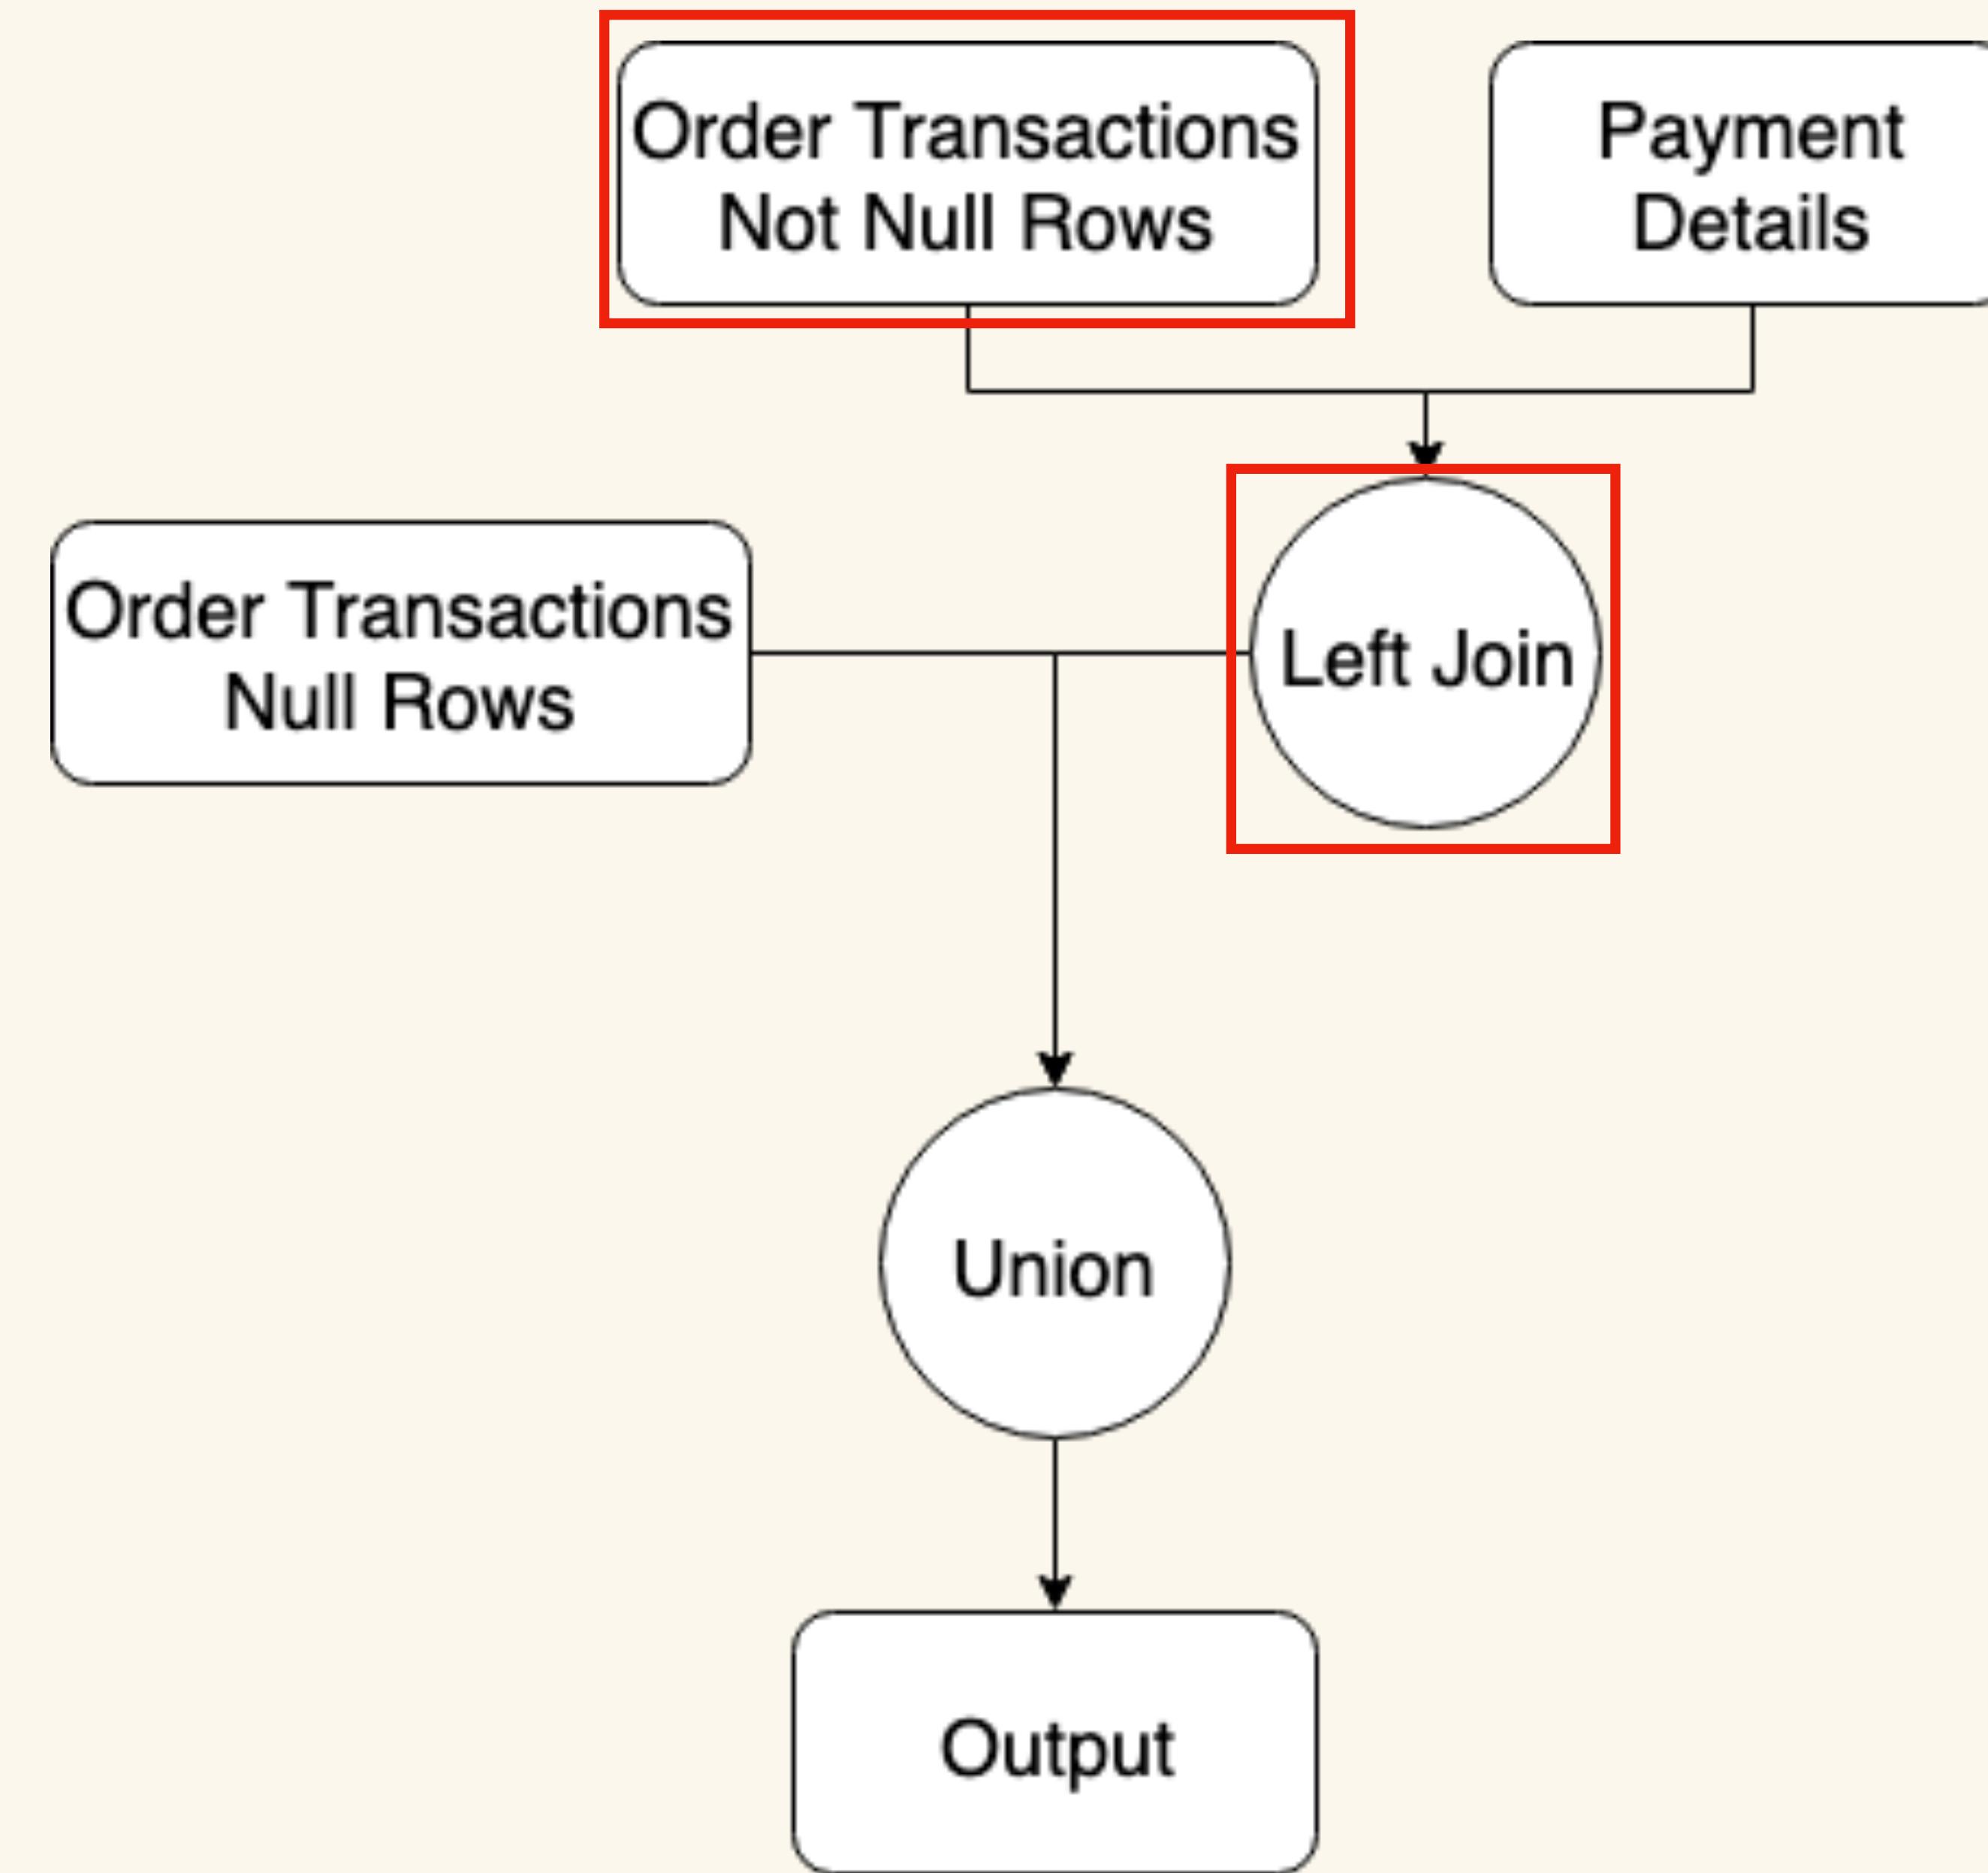
# SQL Diagram



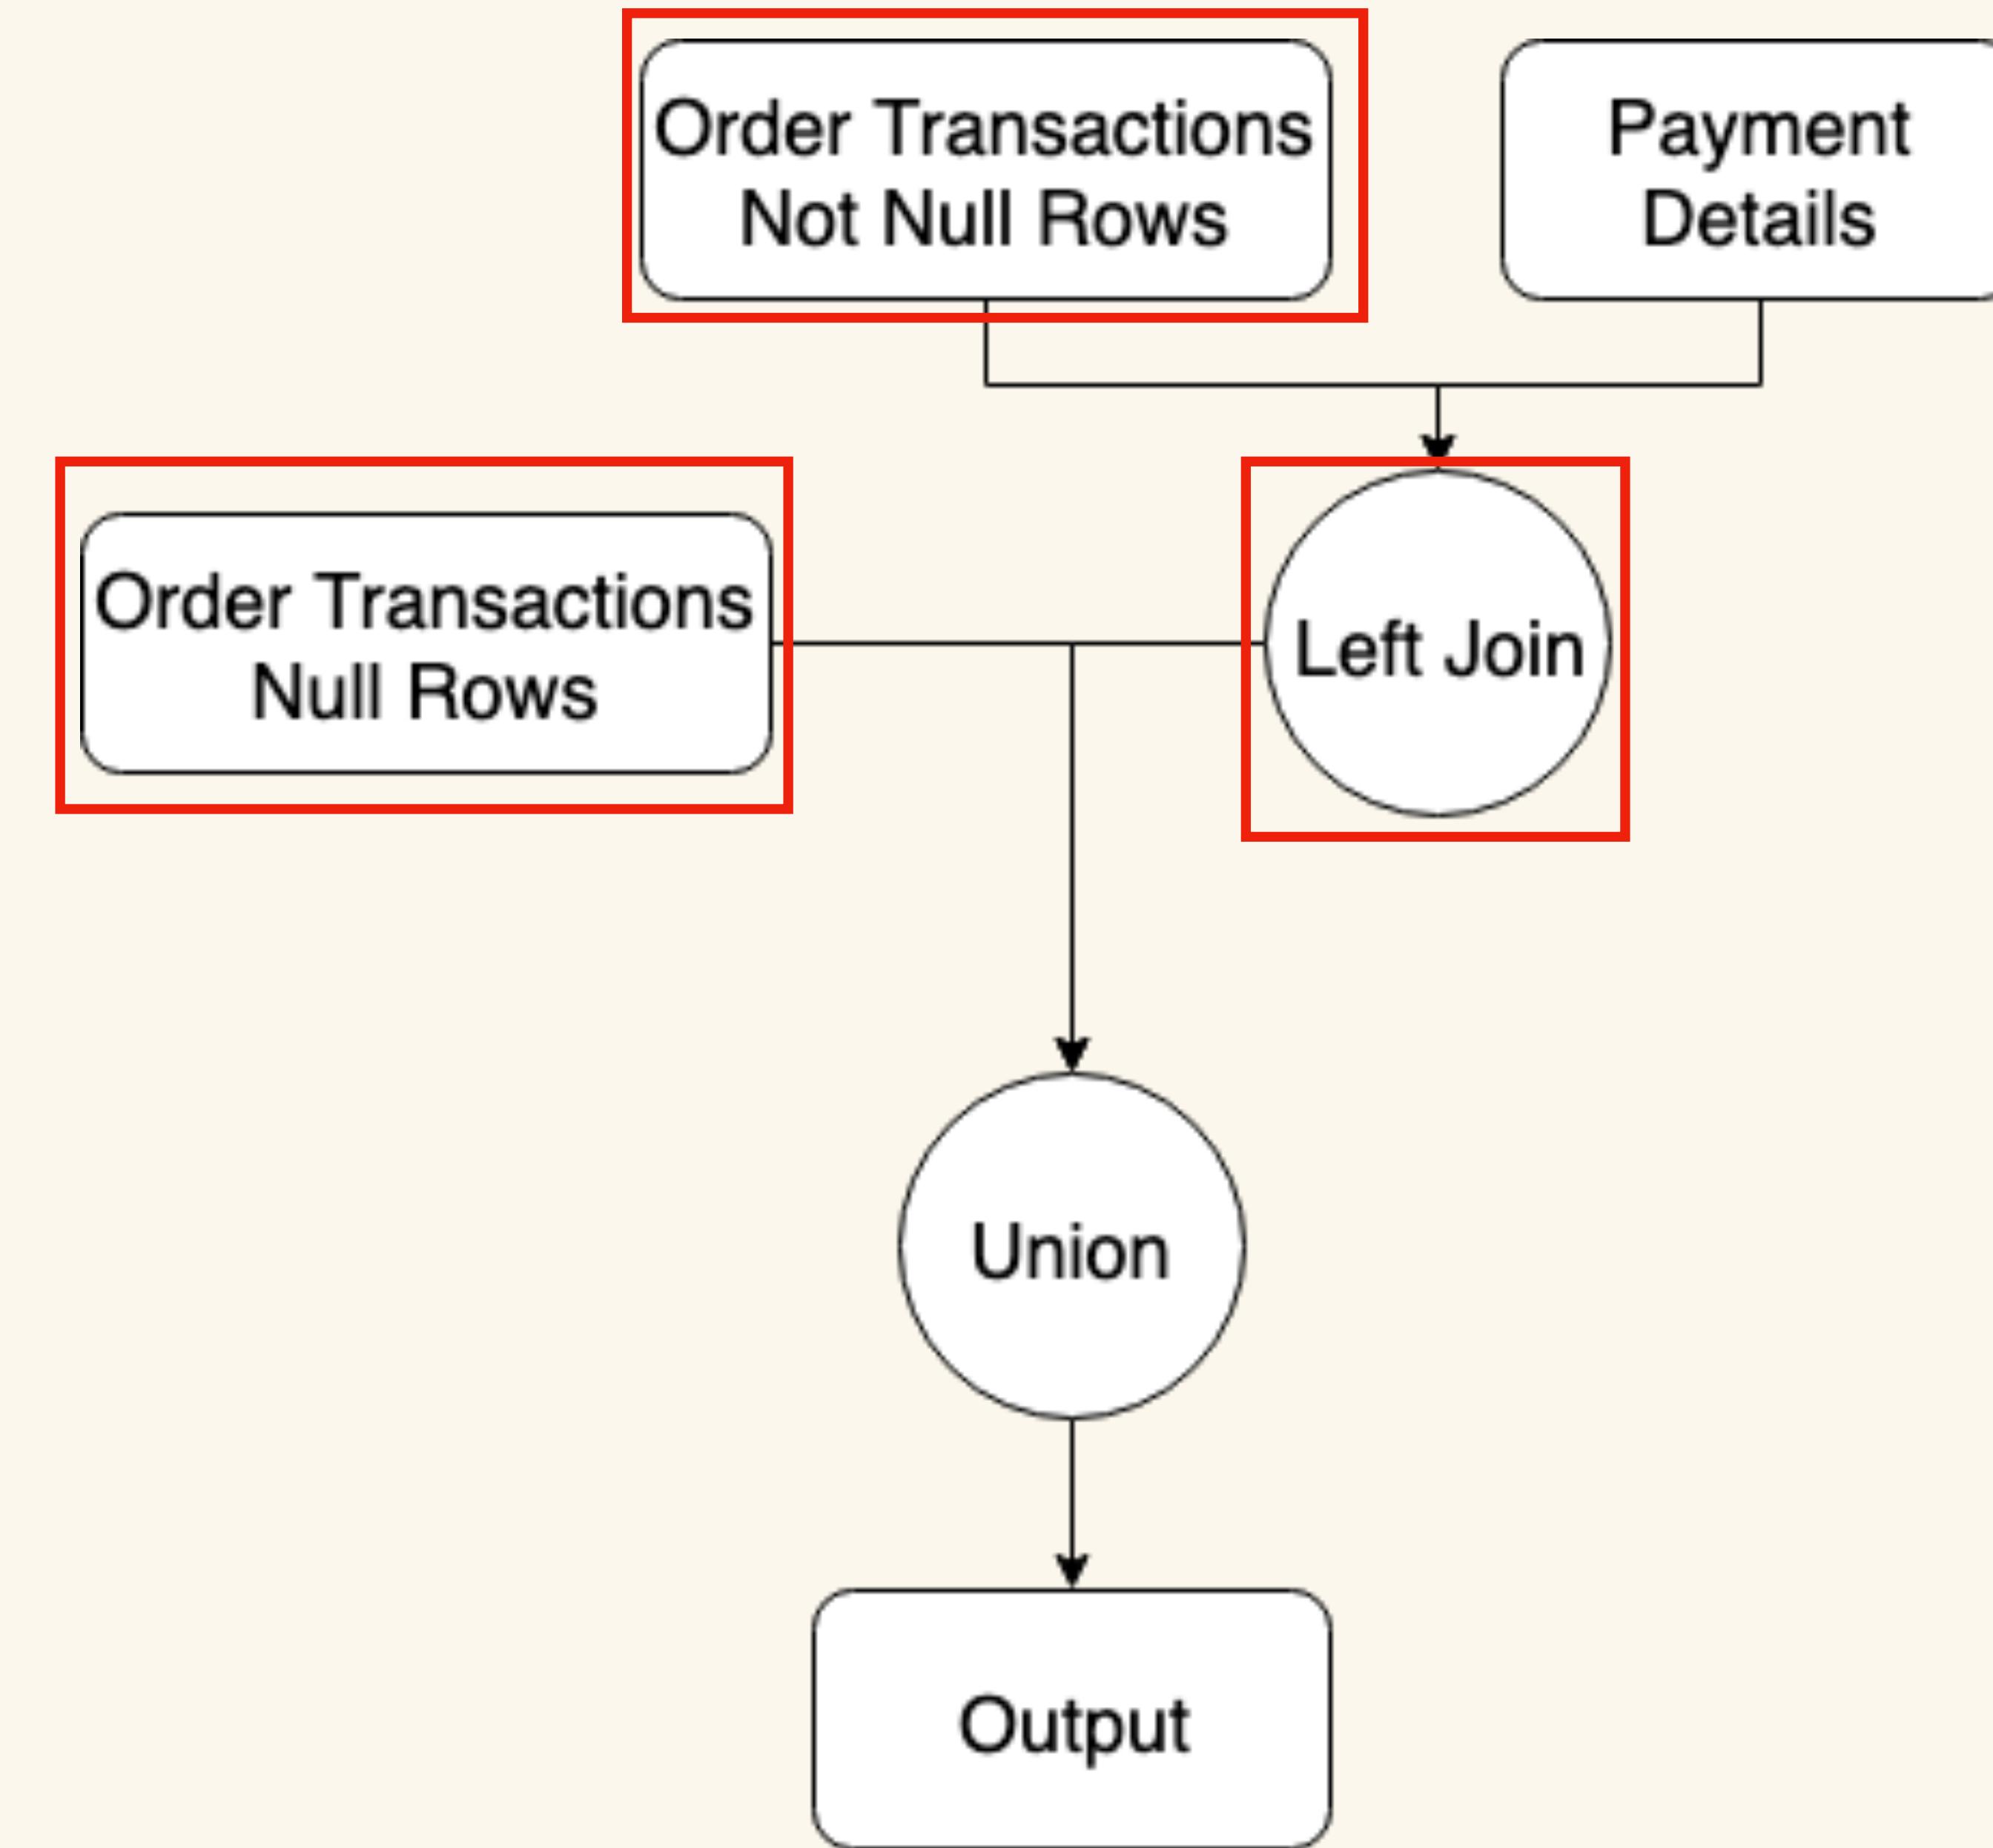
# SQL Diagram



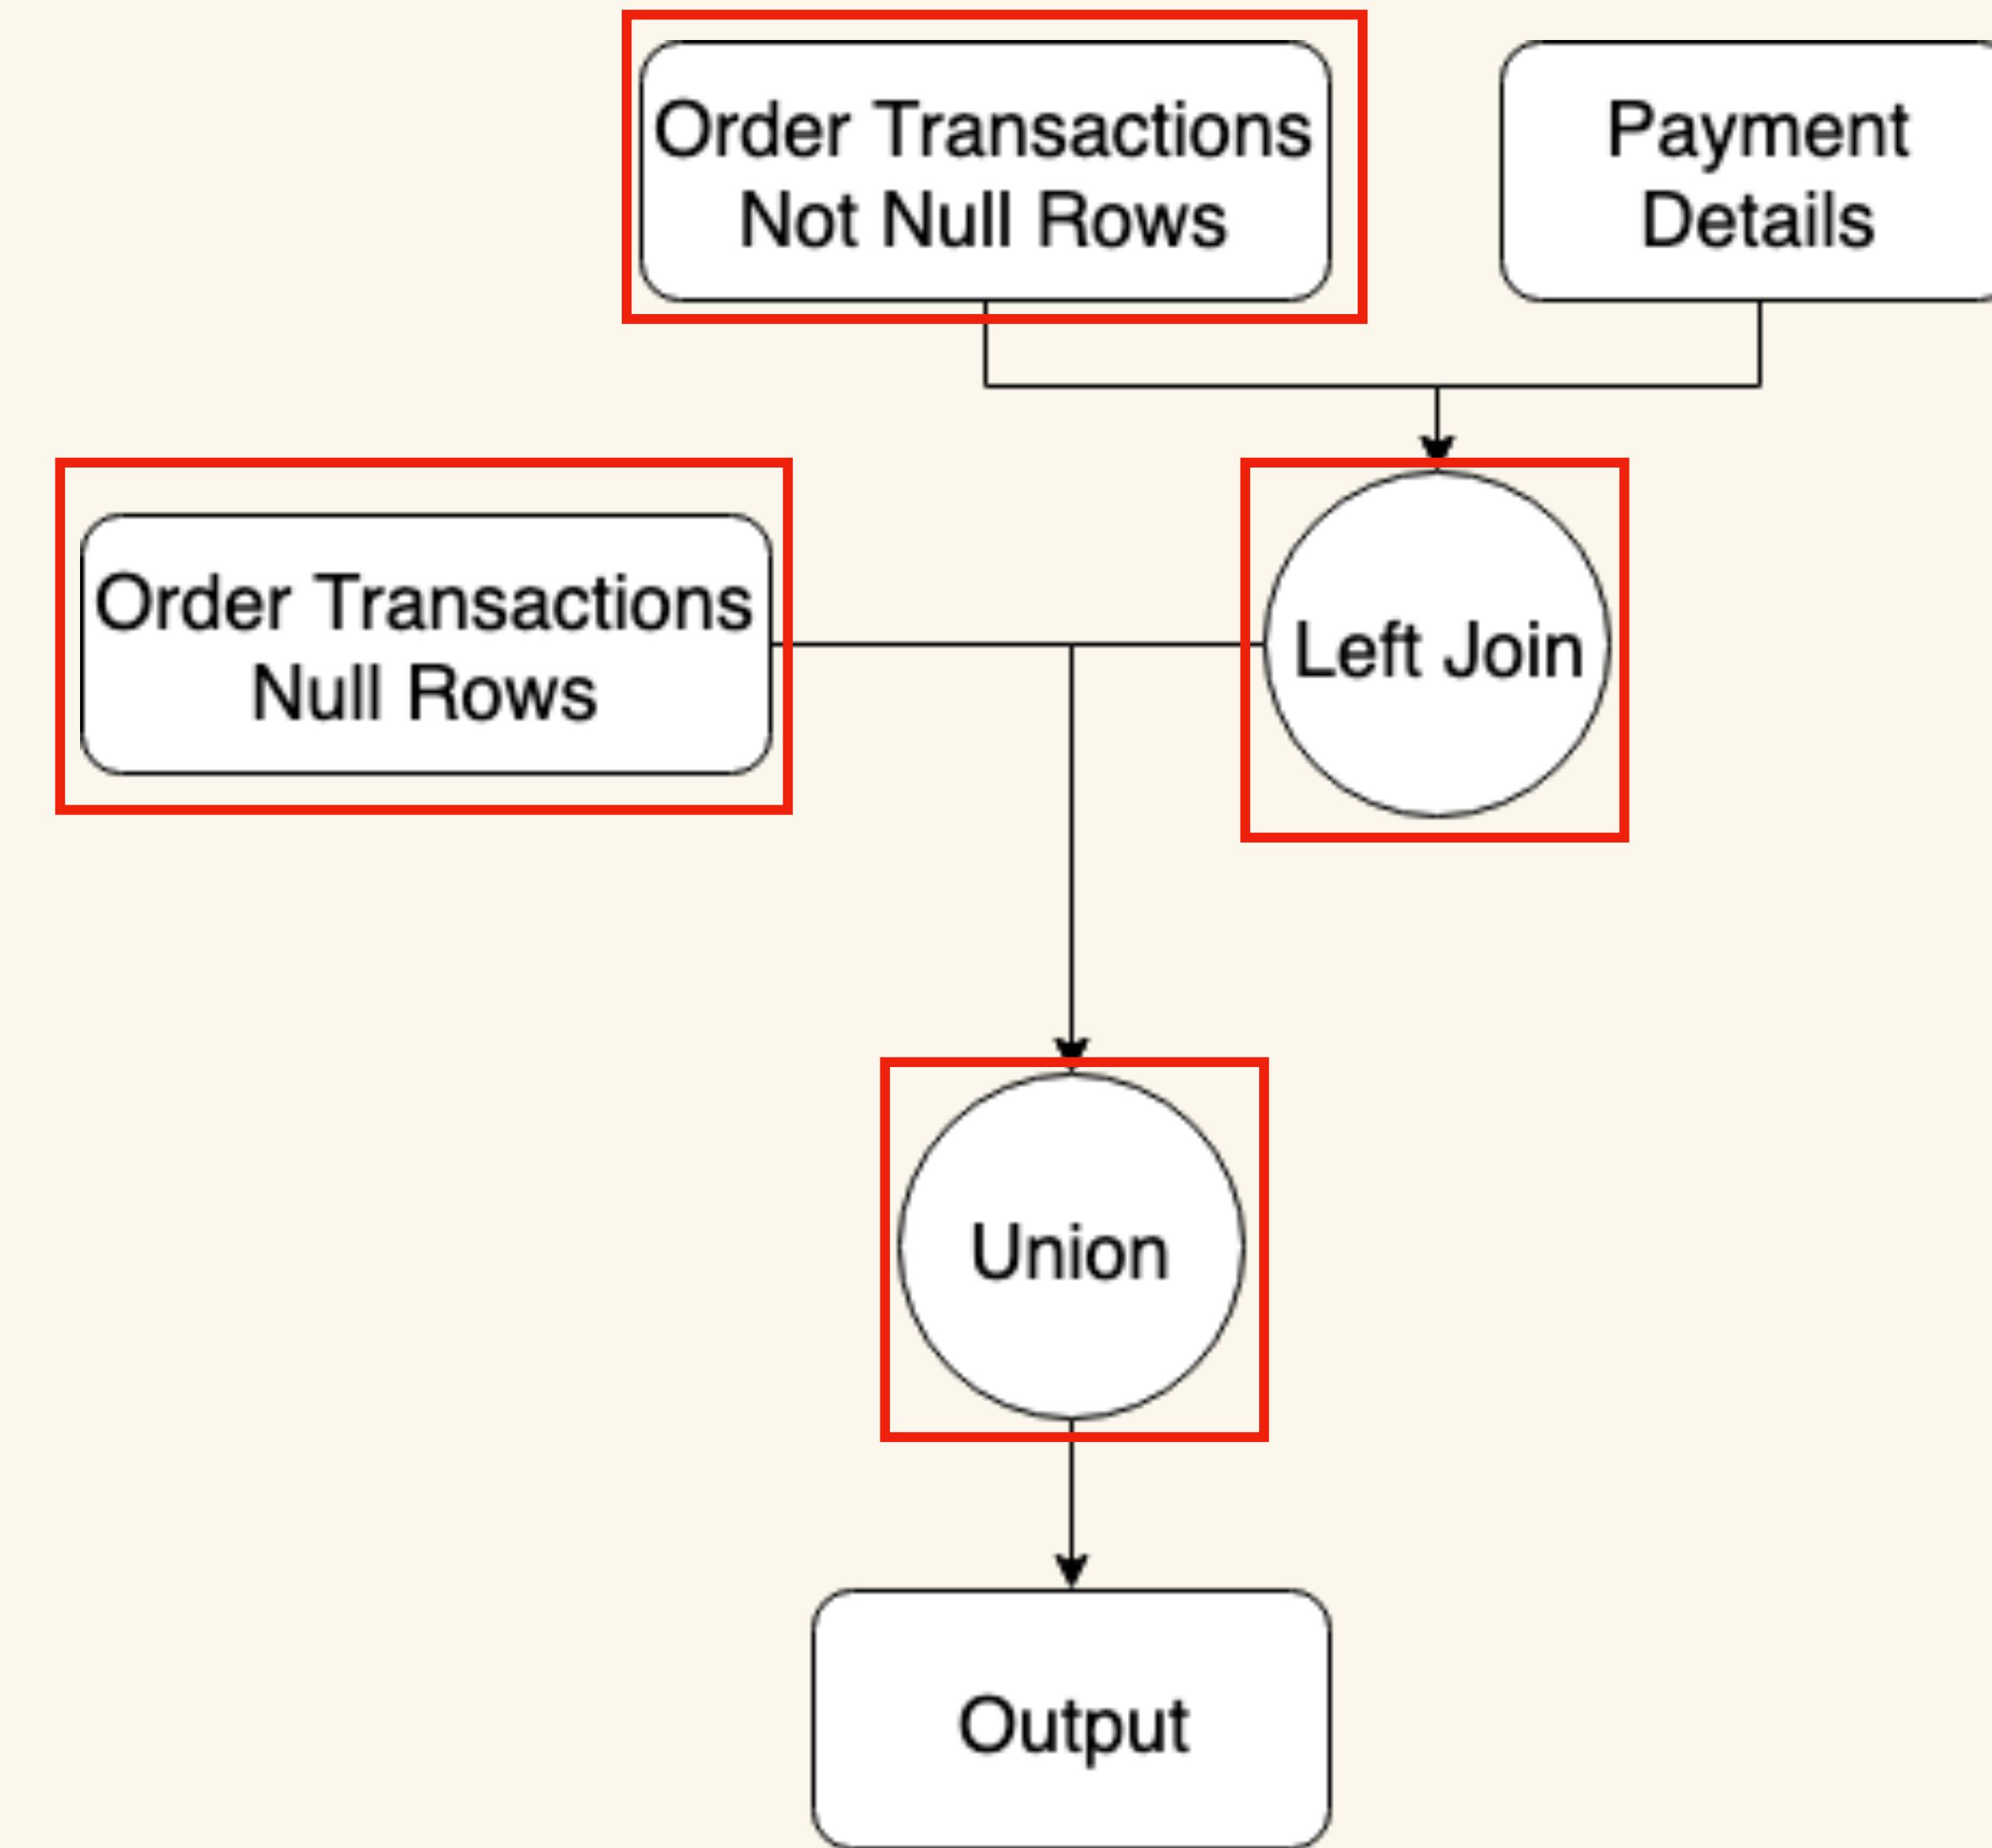
# SQL Diagram



# SQL Diagram



# SQL Diagram



# Downsides

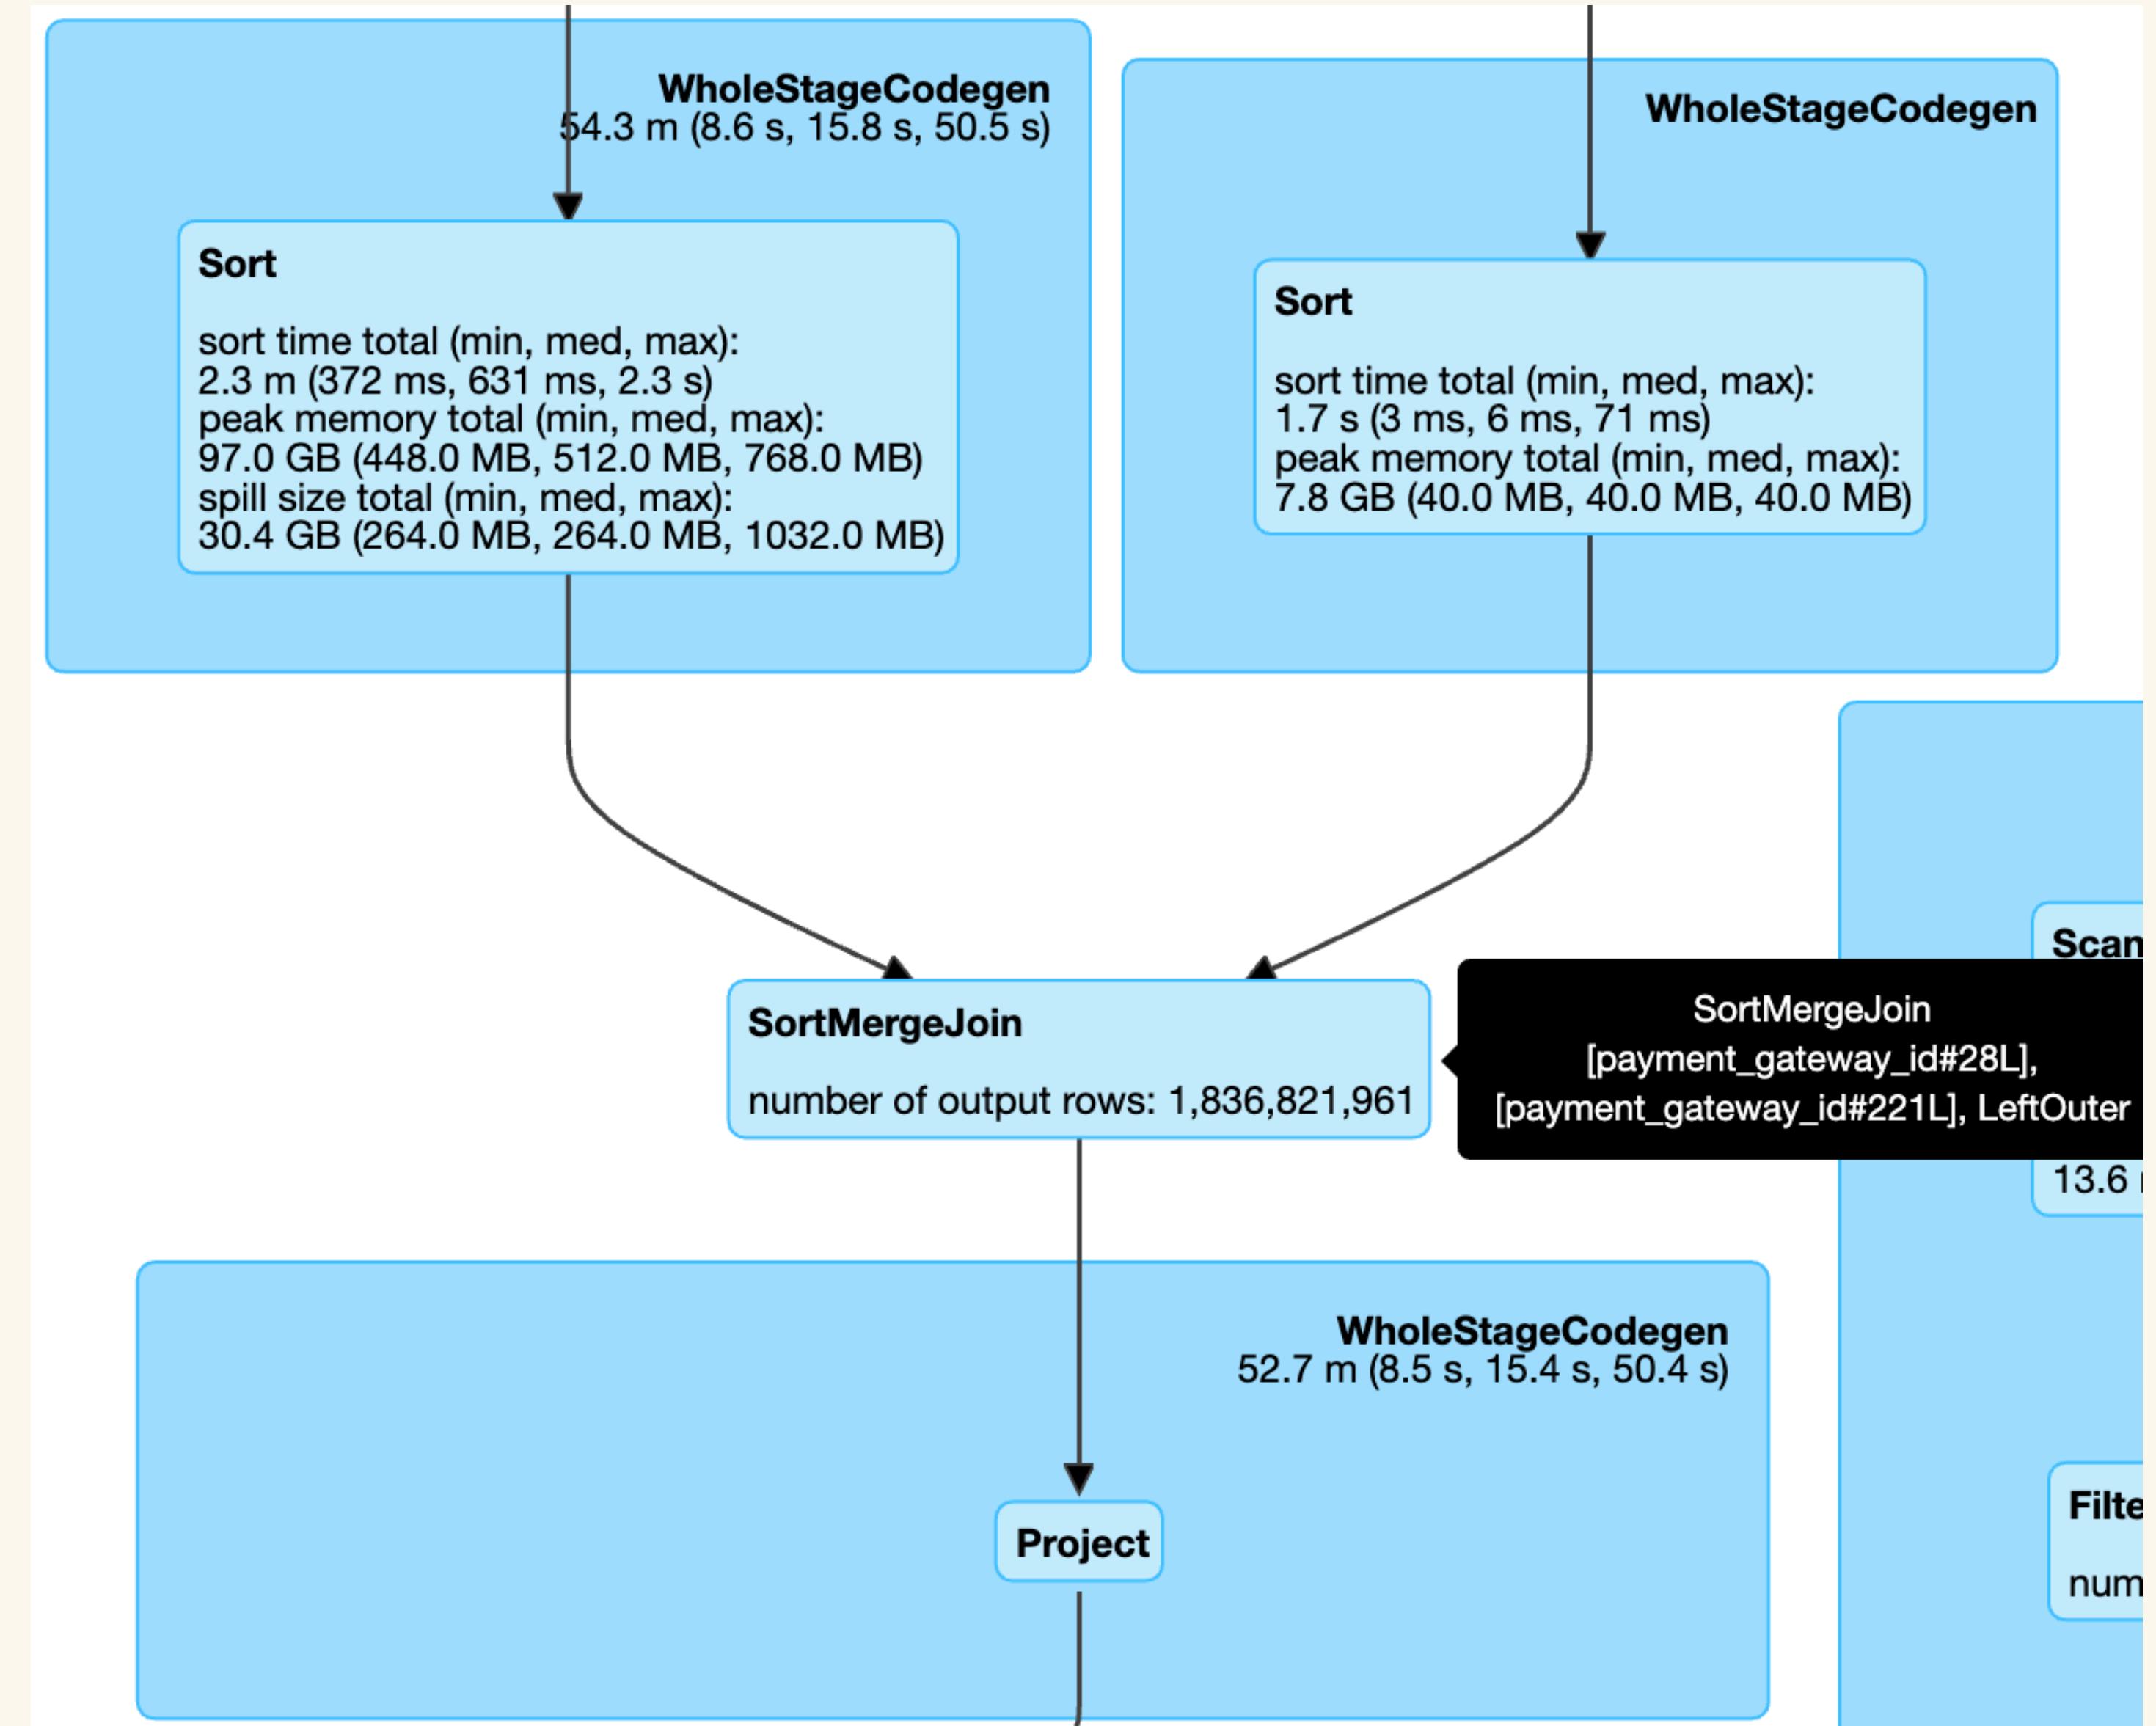
- Spark will recognize that you want to do 2 separate transformations on the skewed dataset:
  - Filter for null values.
  - Filter for not null values.
- Spark will do this read in parallel, so you'll read in twice as much data.
- We recognize this, and “cache”d the skewed dataset, which reduces the read, but now we storing the dataset in memory.
- Thus lowering the amount of total memory for our spark application.
- This can cause memory issues.

# Results

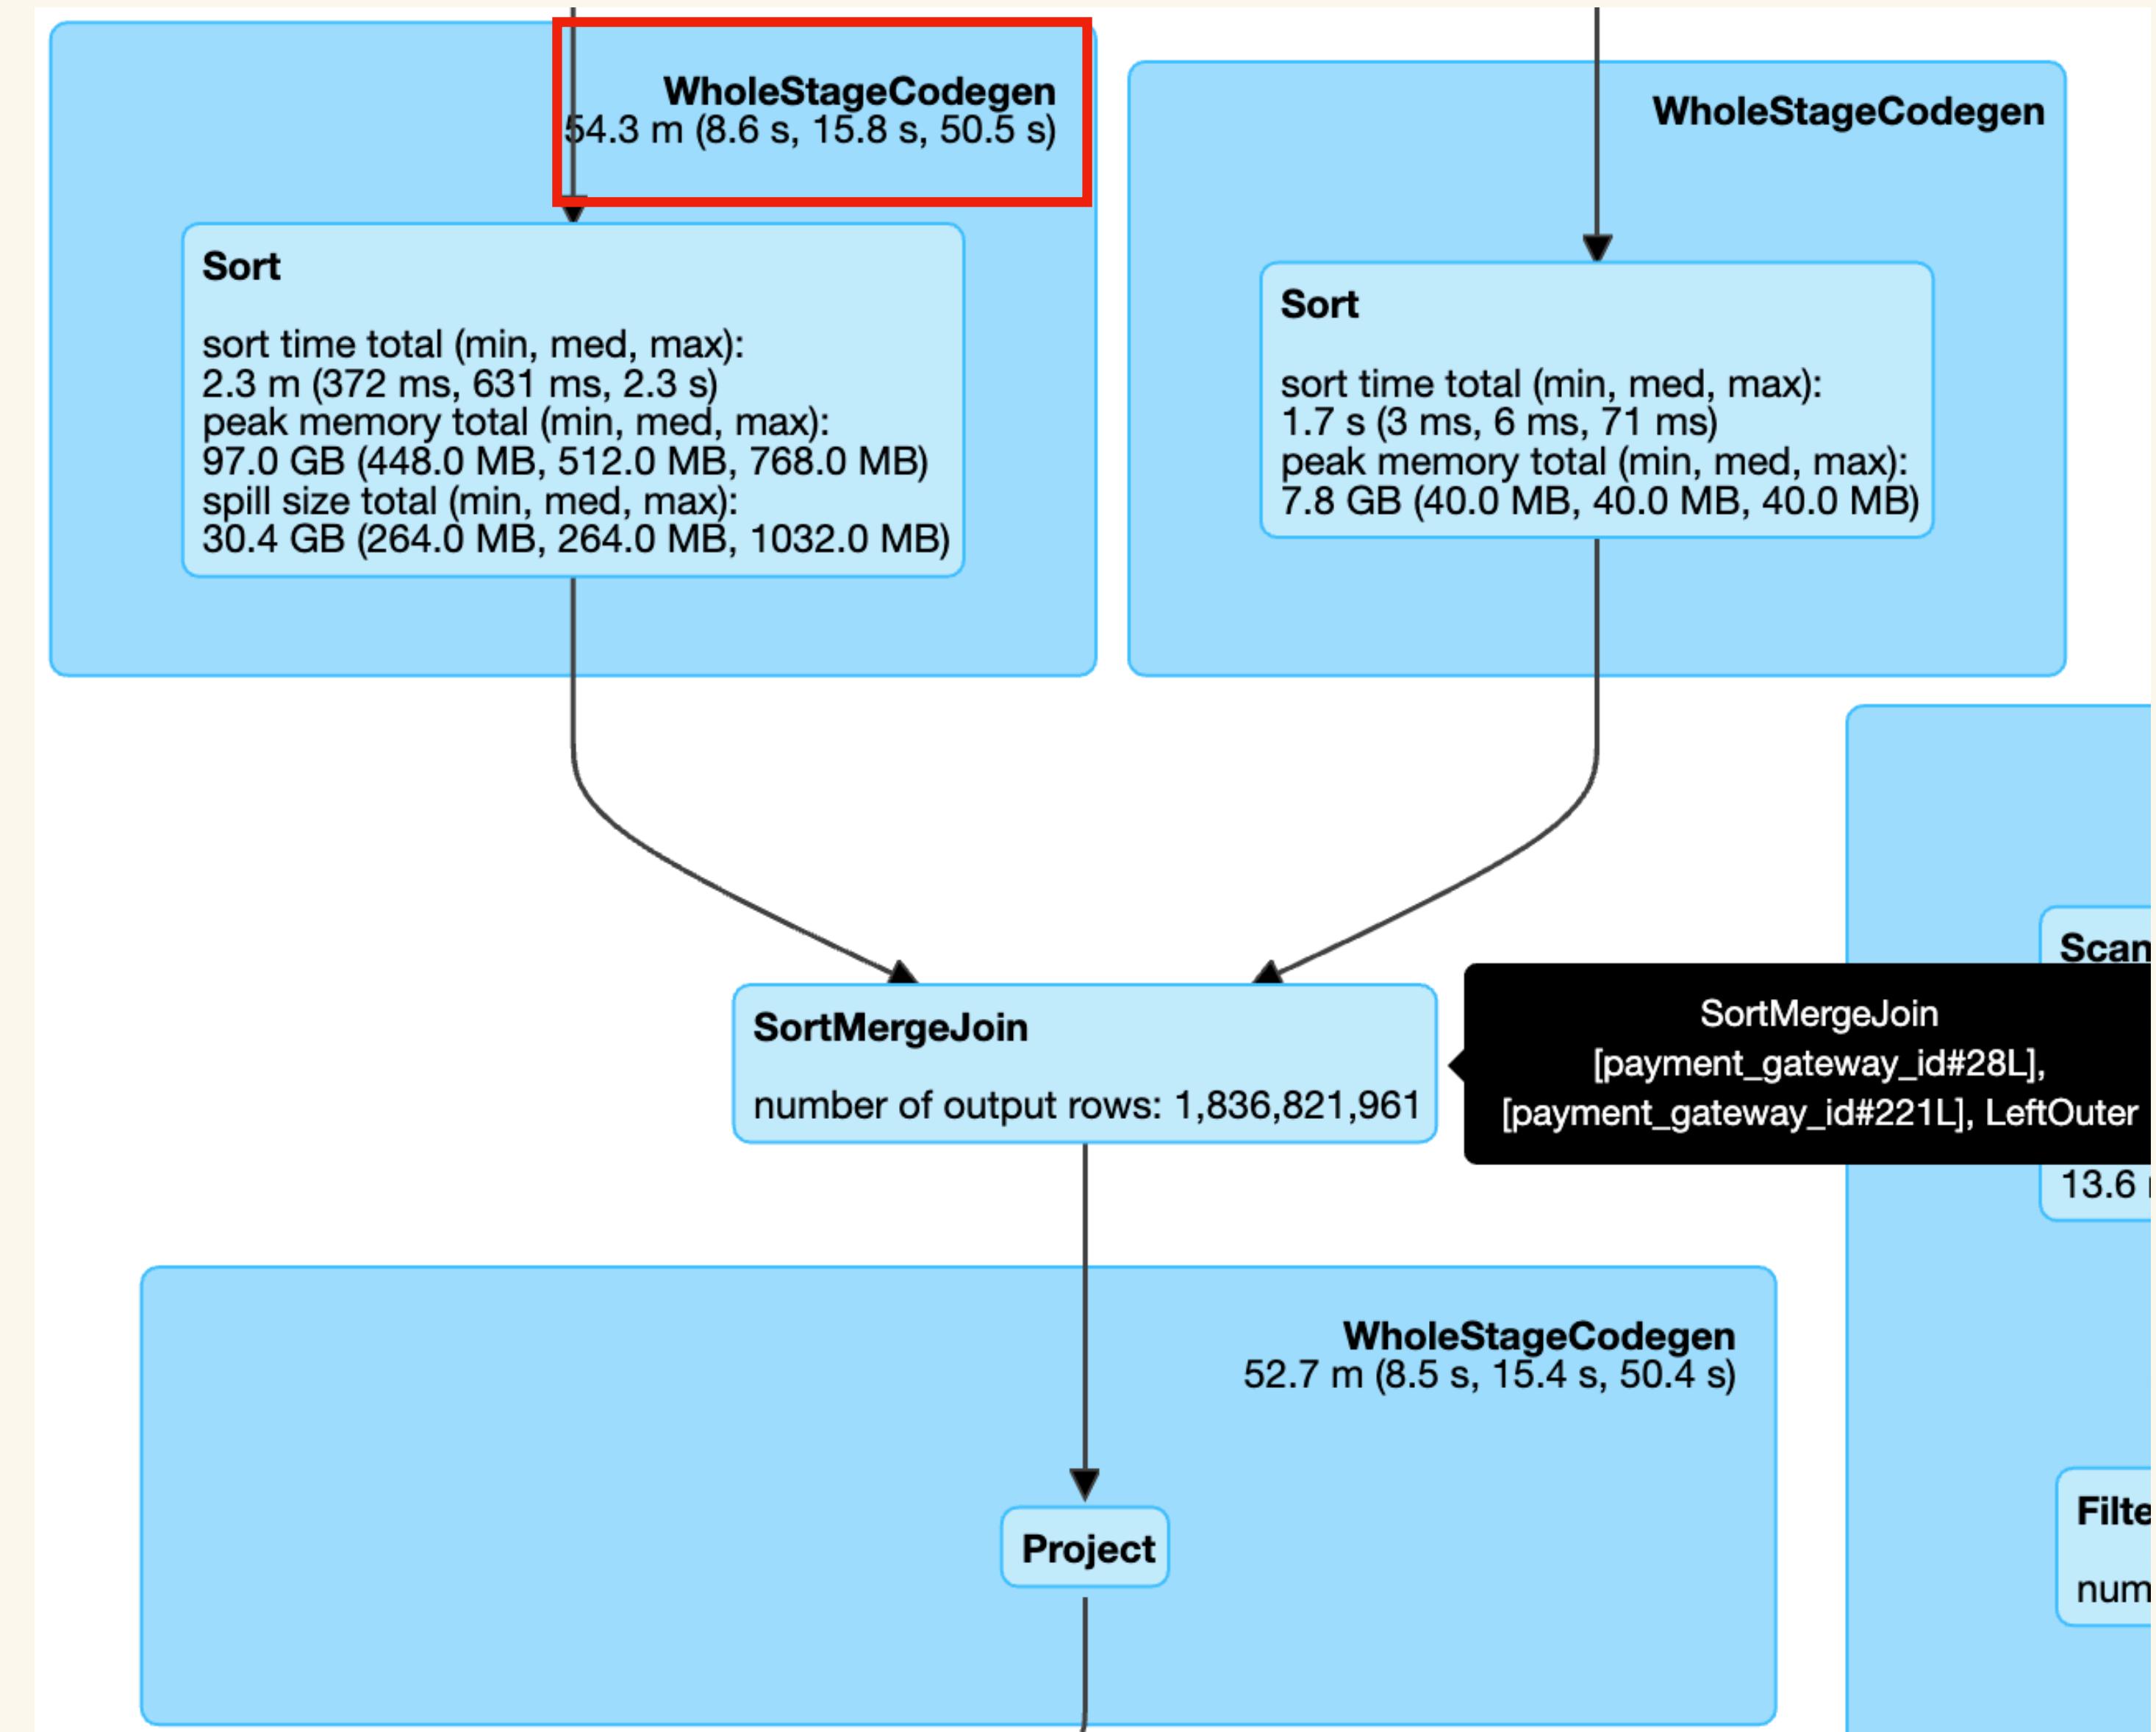
## Completed Queries (3)

ID	Description	Submitted ▲	Duration
0	<a href="#">showString at NativeMethodAccessorImpl.java:0</a>	2019/07/24 21:10:15	22 min
1	<a href="#">showString at NativeMethodAccessorImpl.java:0</a>	2019/07/24 21:31:52	10 min

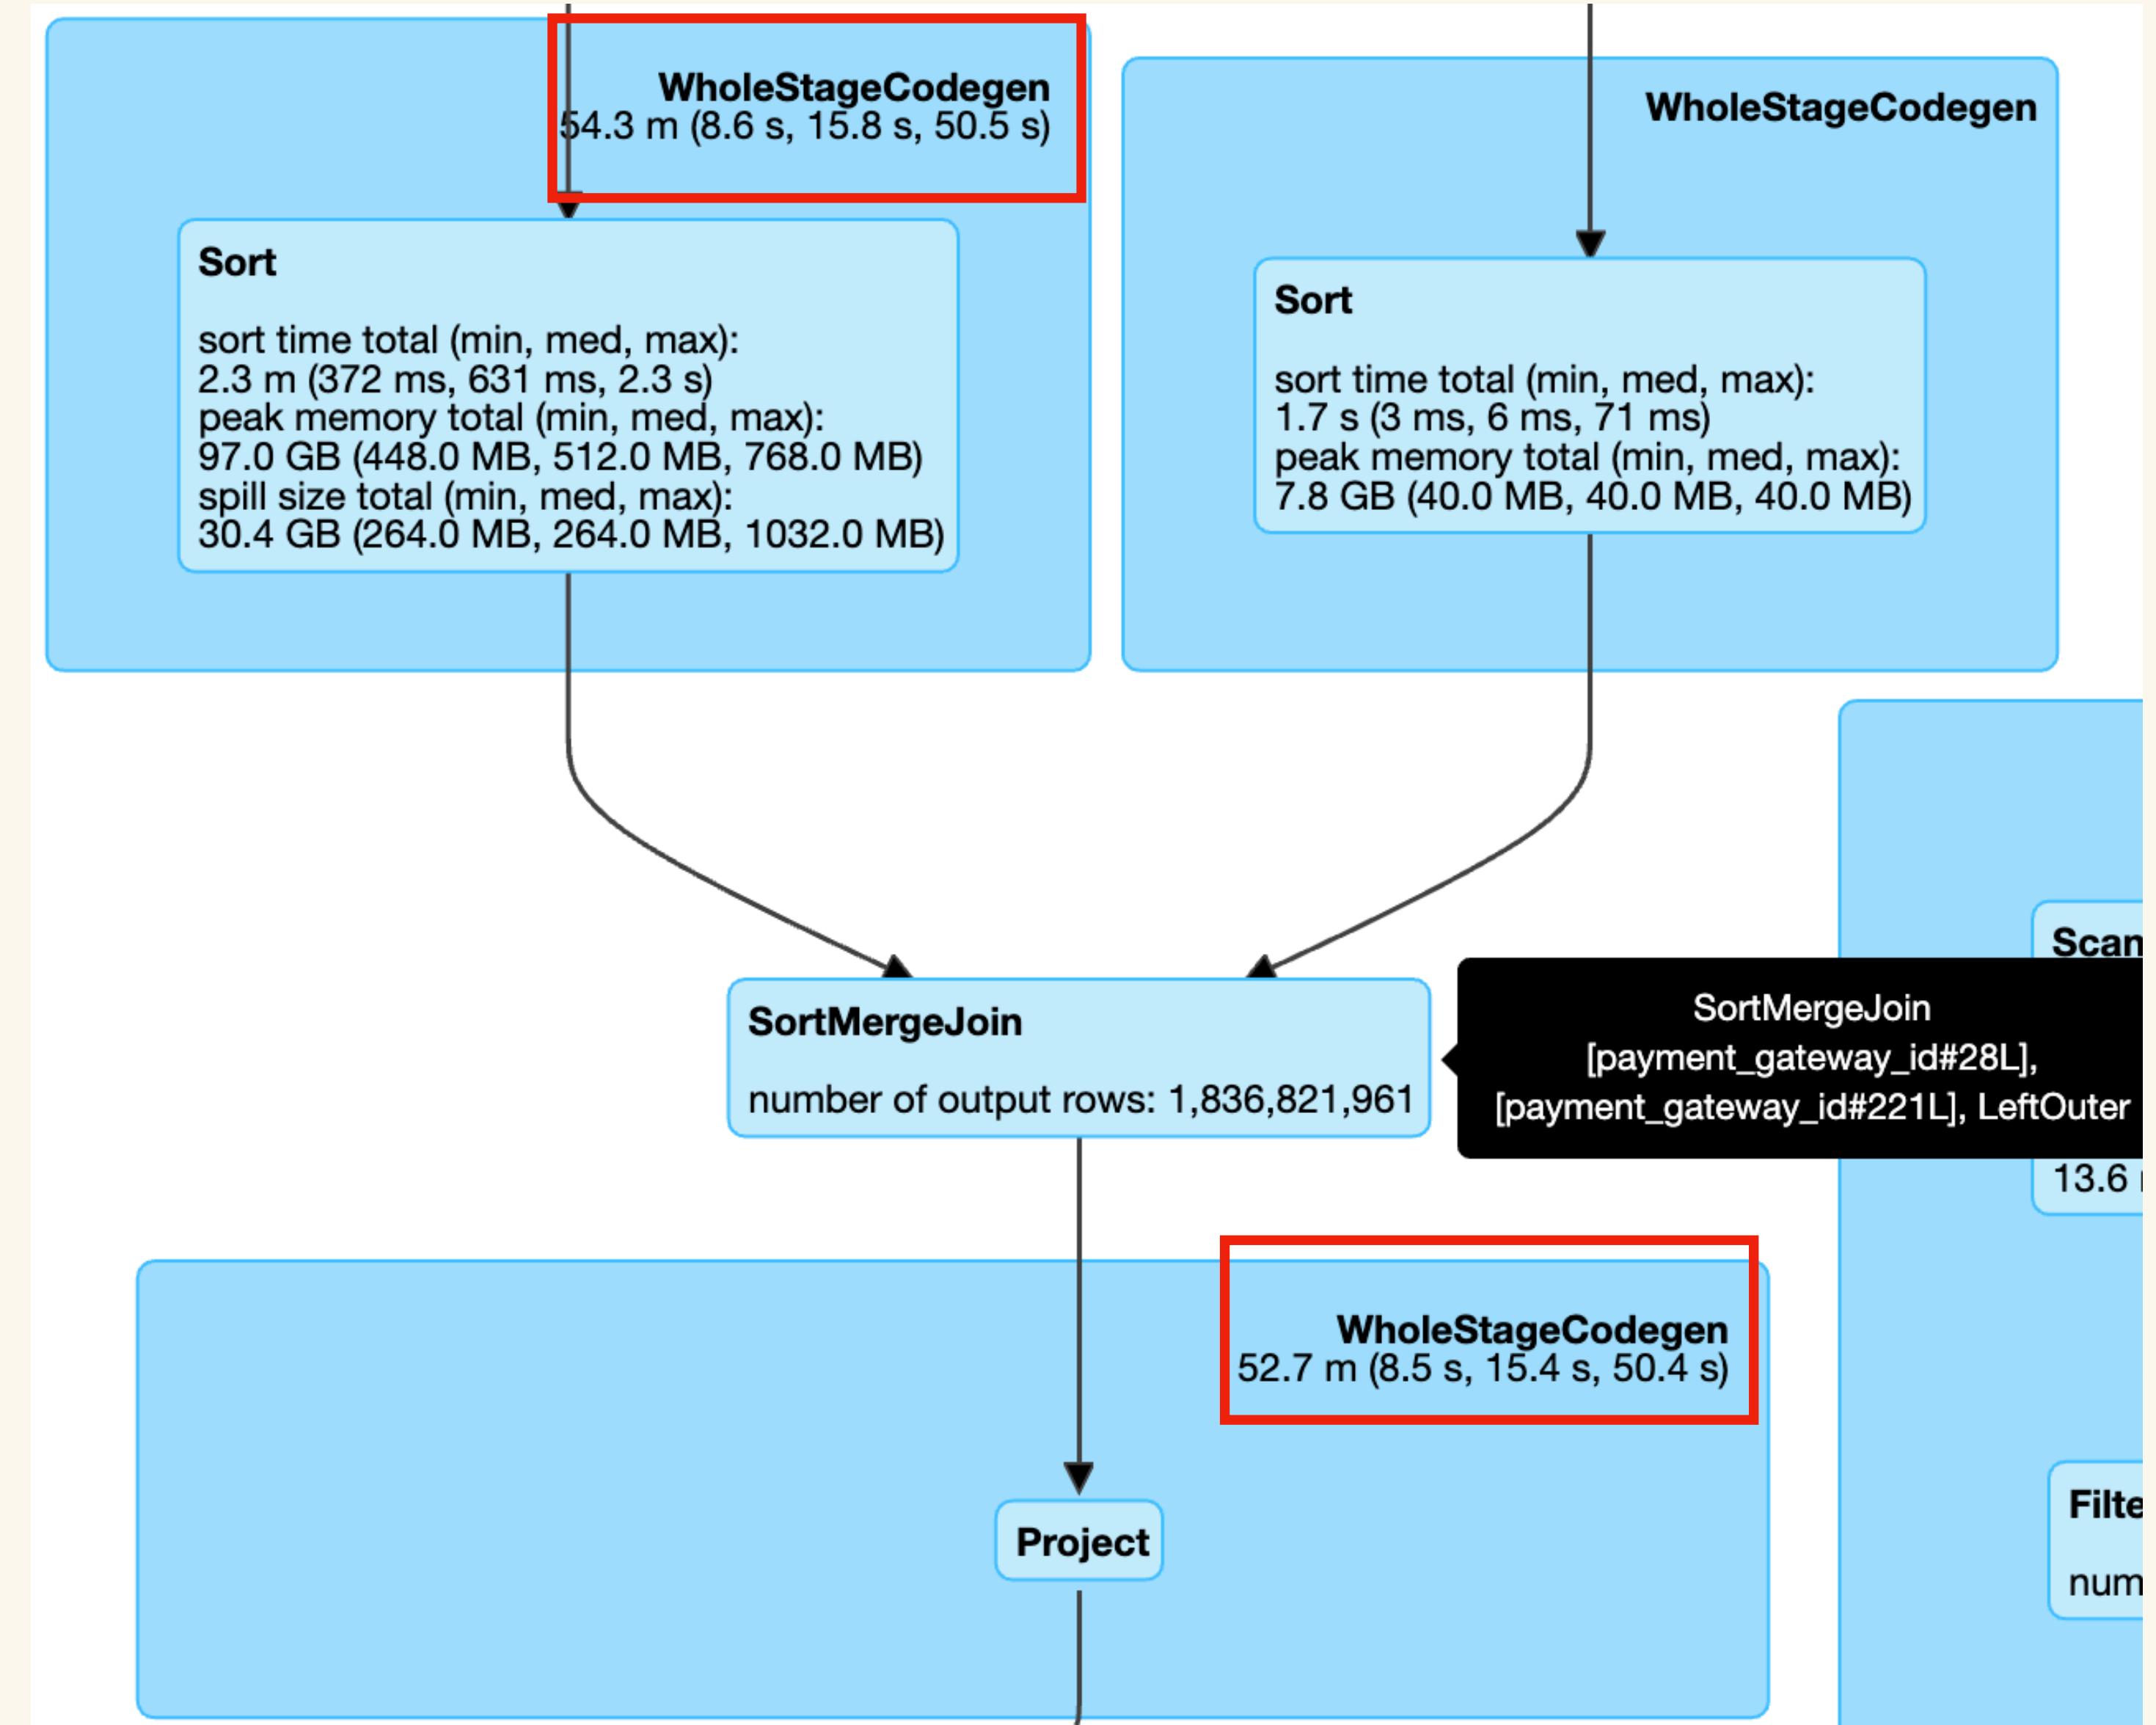
# Results



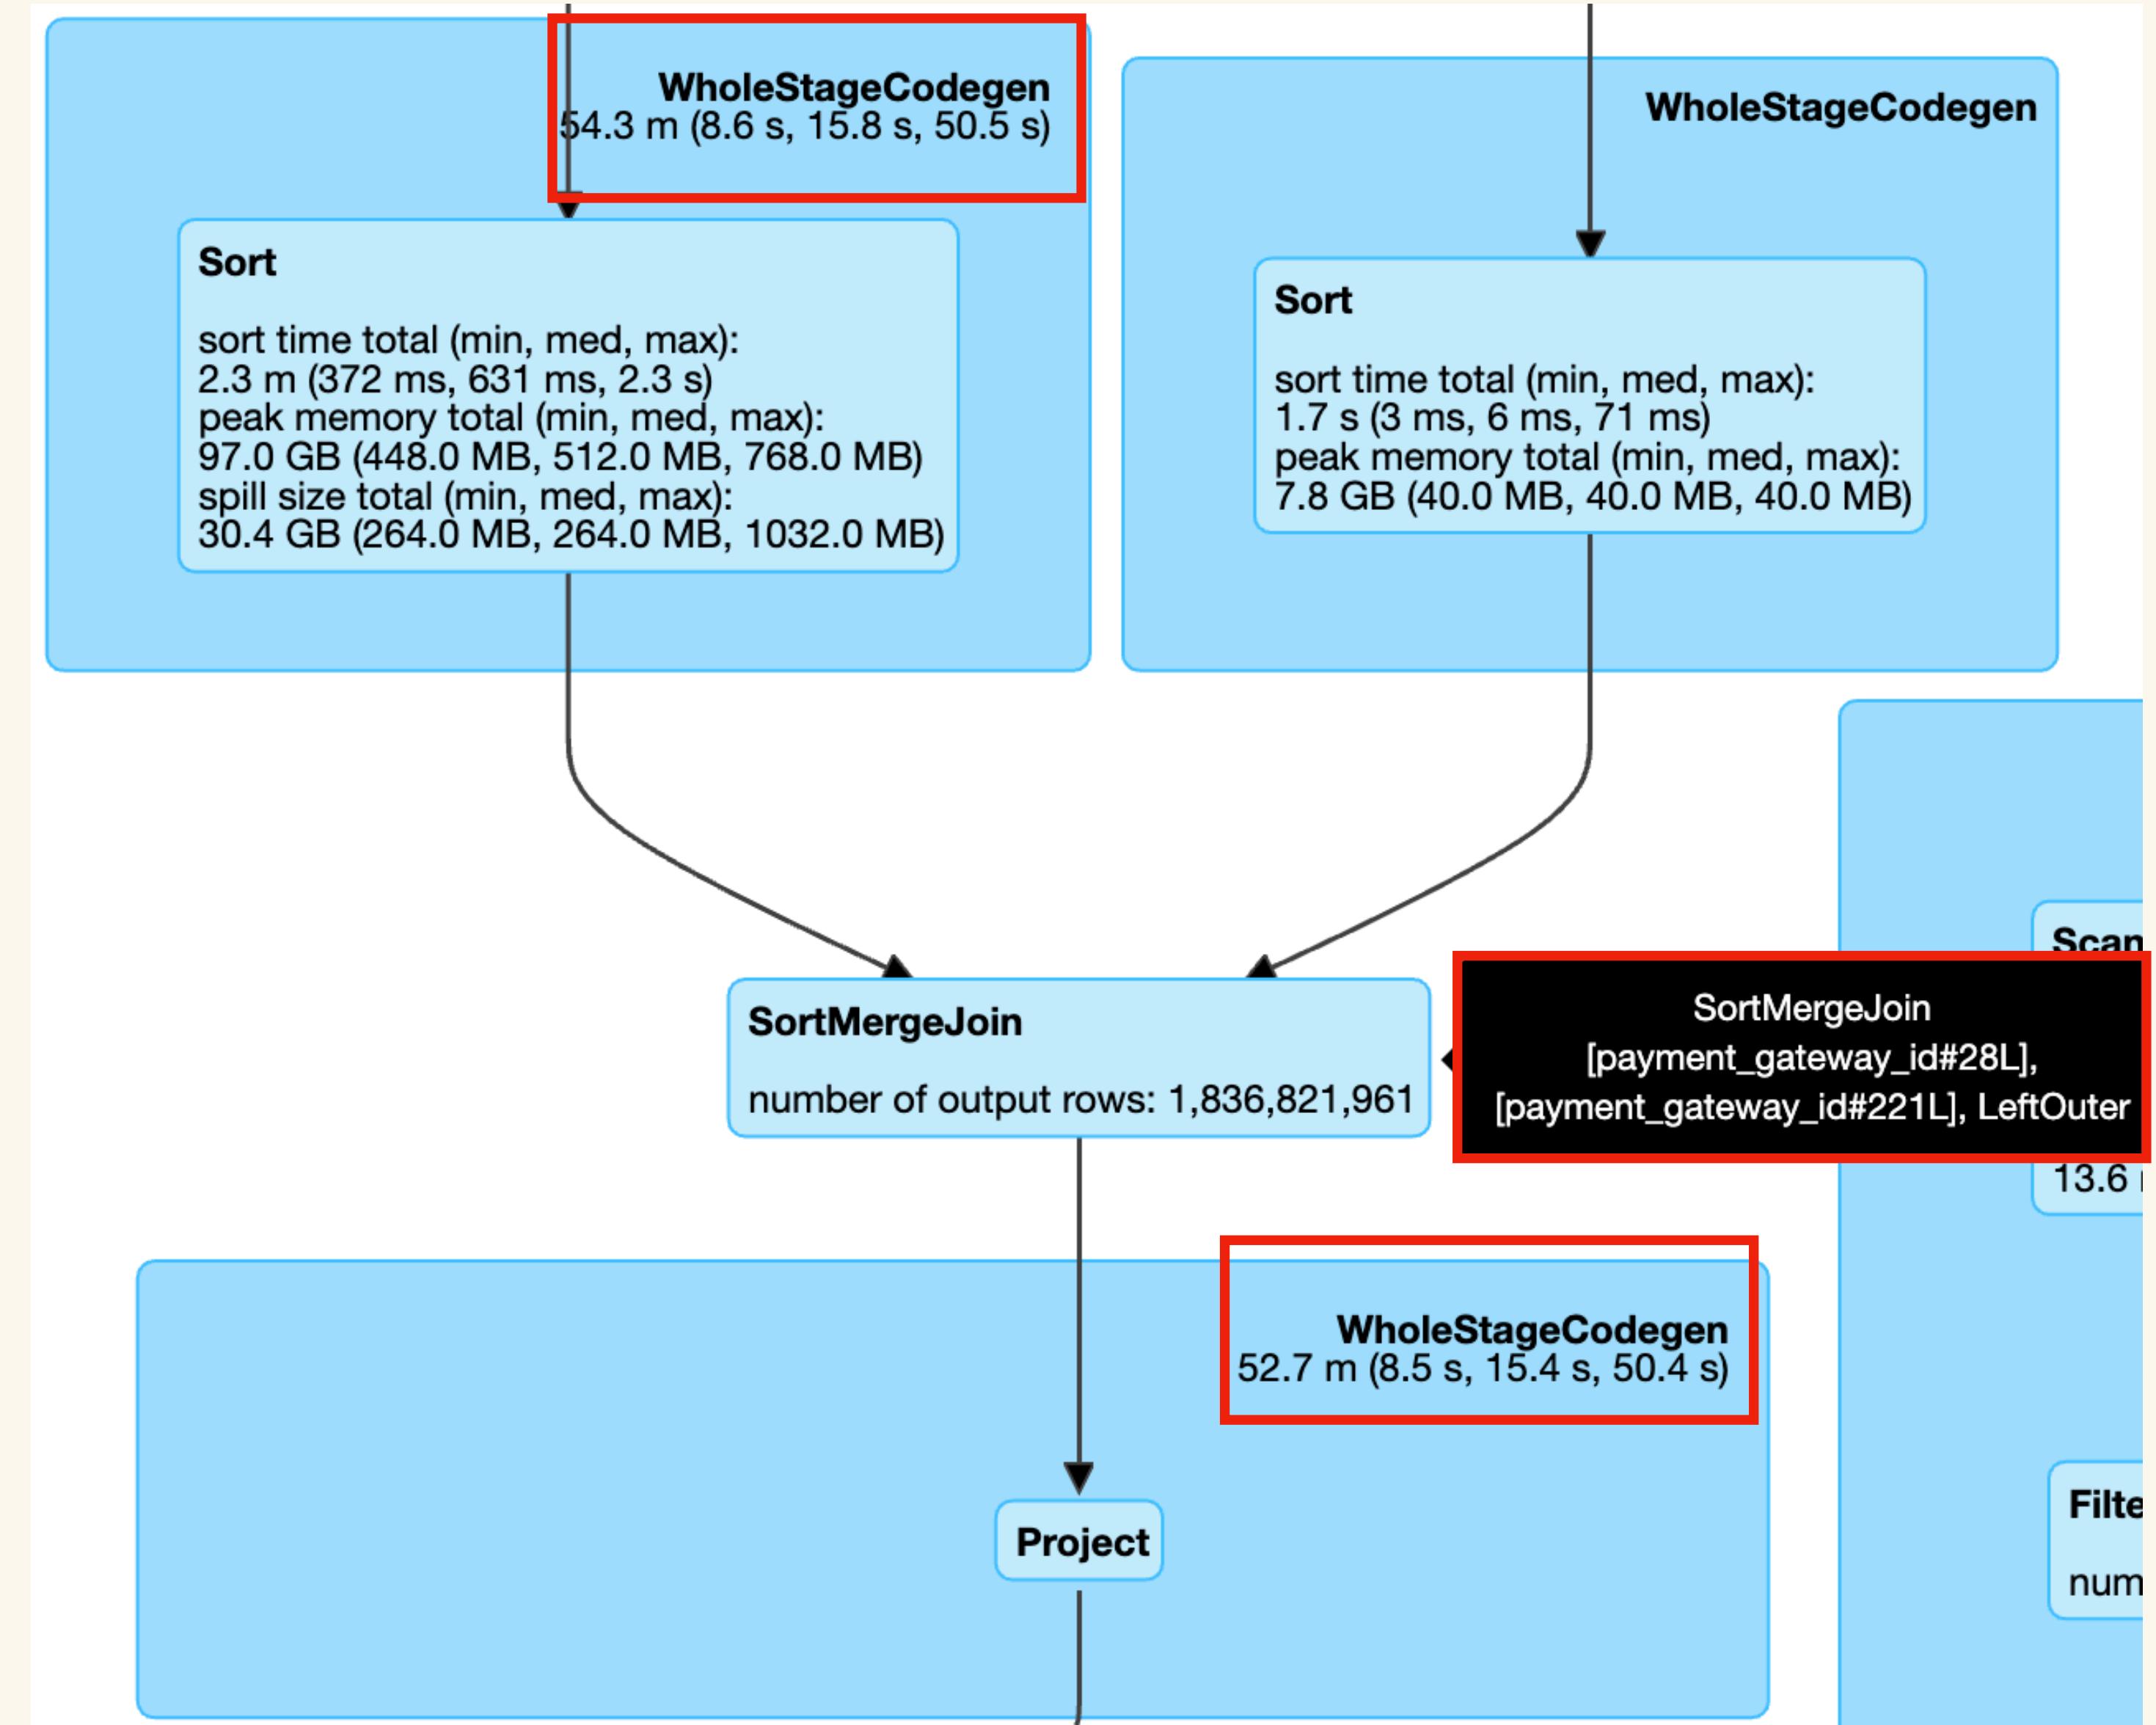
# Results



# Results



# Results



# SkewPartitioner

## Motivation:

- Reduce the amount of rows that contain the skewed key.

## Solution:

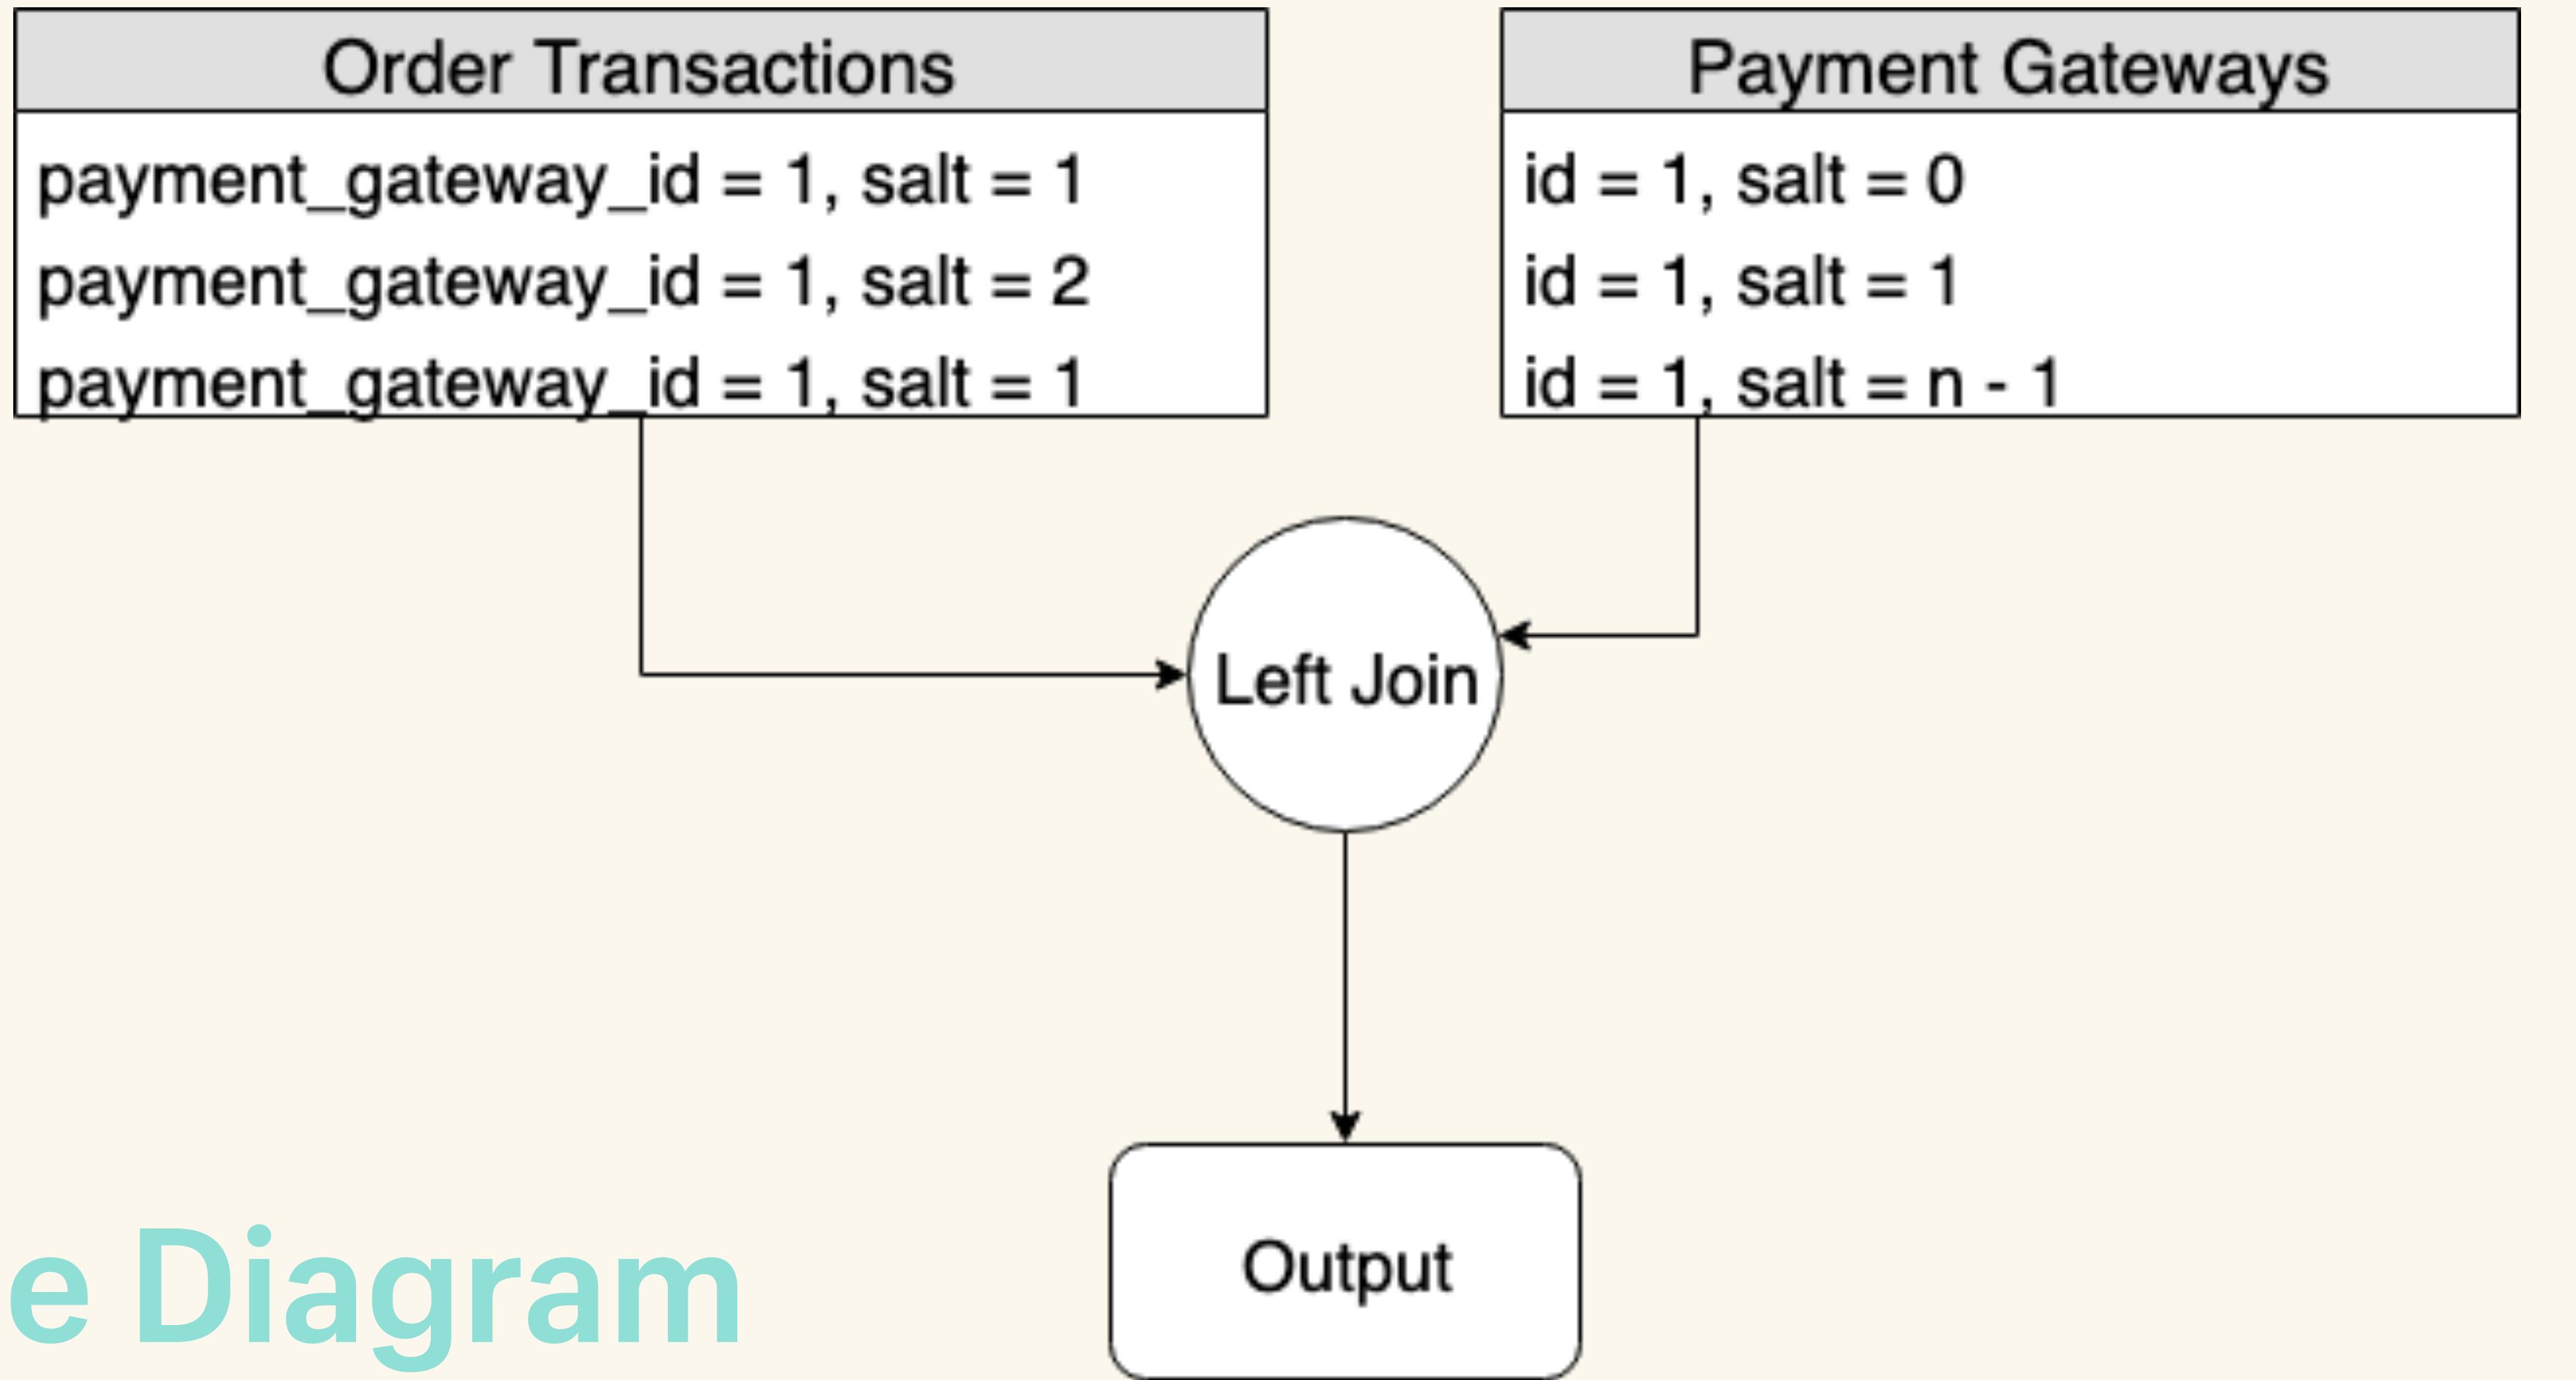
- We randomly assign a number from 0 to  $n - 1$ , called a “salt”, to each row.
- This splits the number of rows by a factor  $n$ .



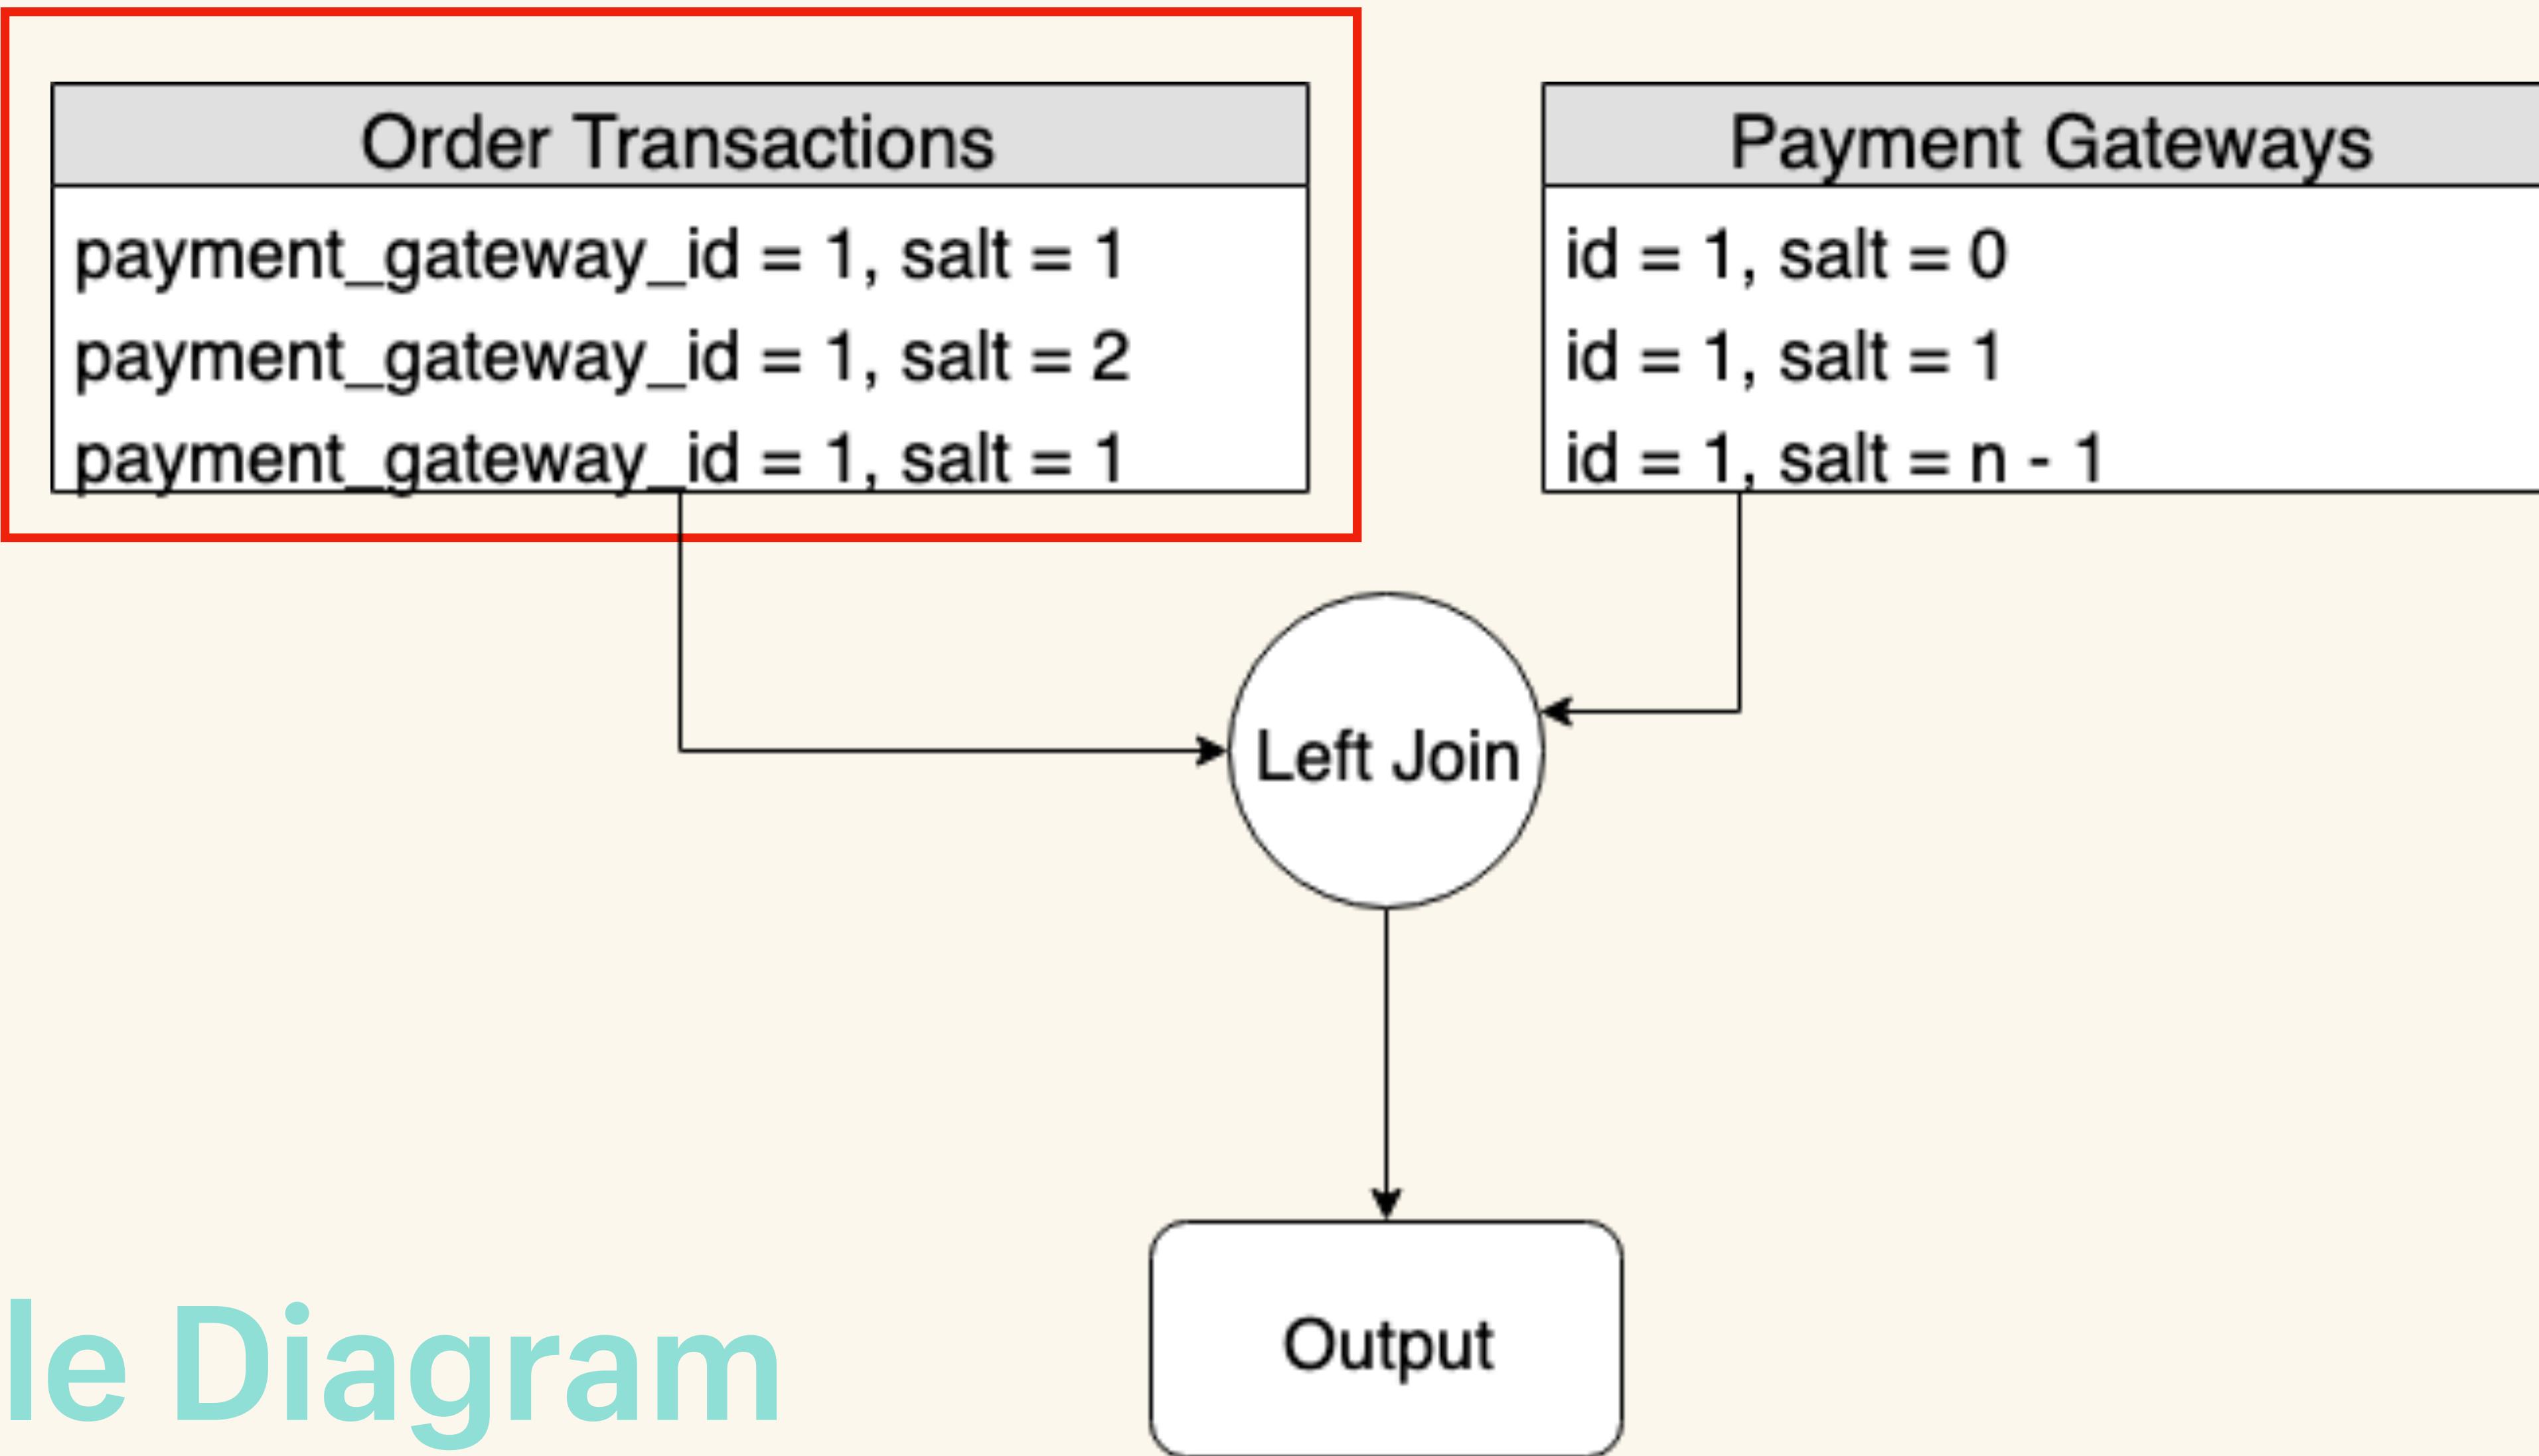
# SkewPartitioner Under the Hood

- Assigns a “salt” number (`random % num_partitions`) to each row of the left `dataFrame`.
- Duplicates the right `dataFrame` `num_partitions` times.
- Assign a different salt value to the every duplicate.
- Joins the two datasets on the (join key + salt number).



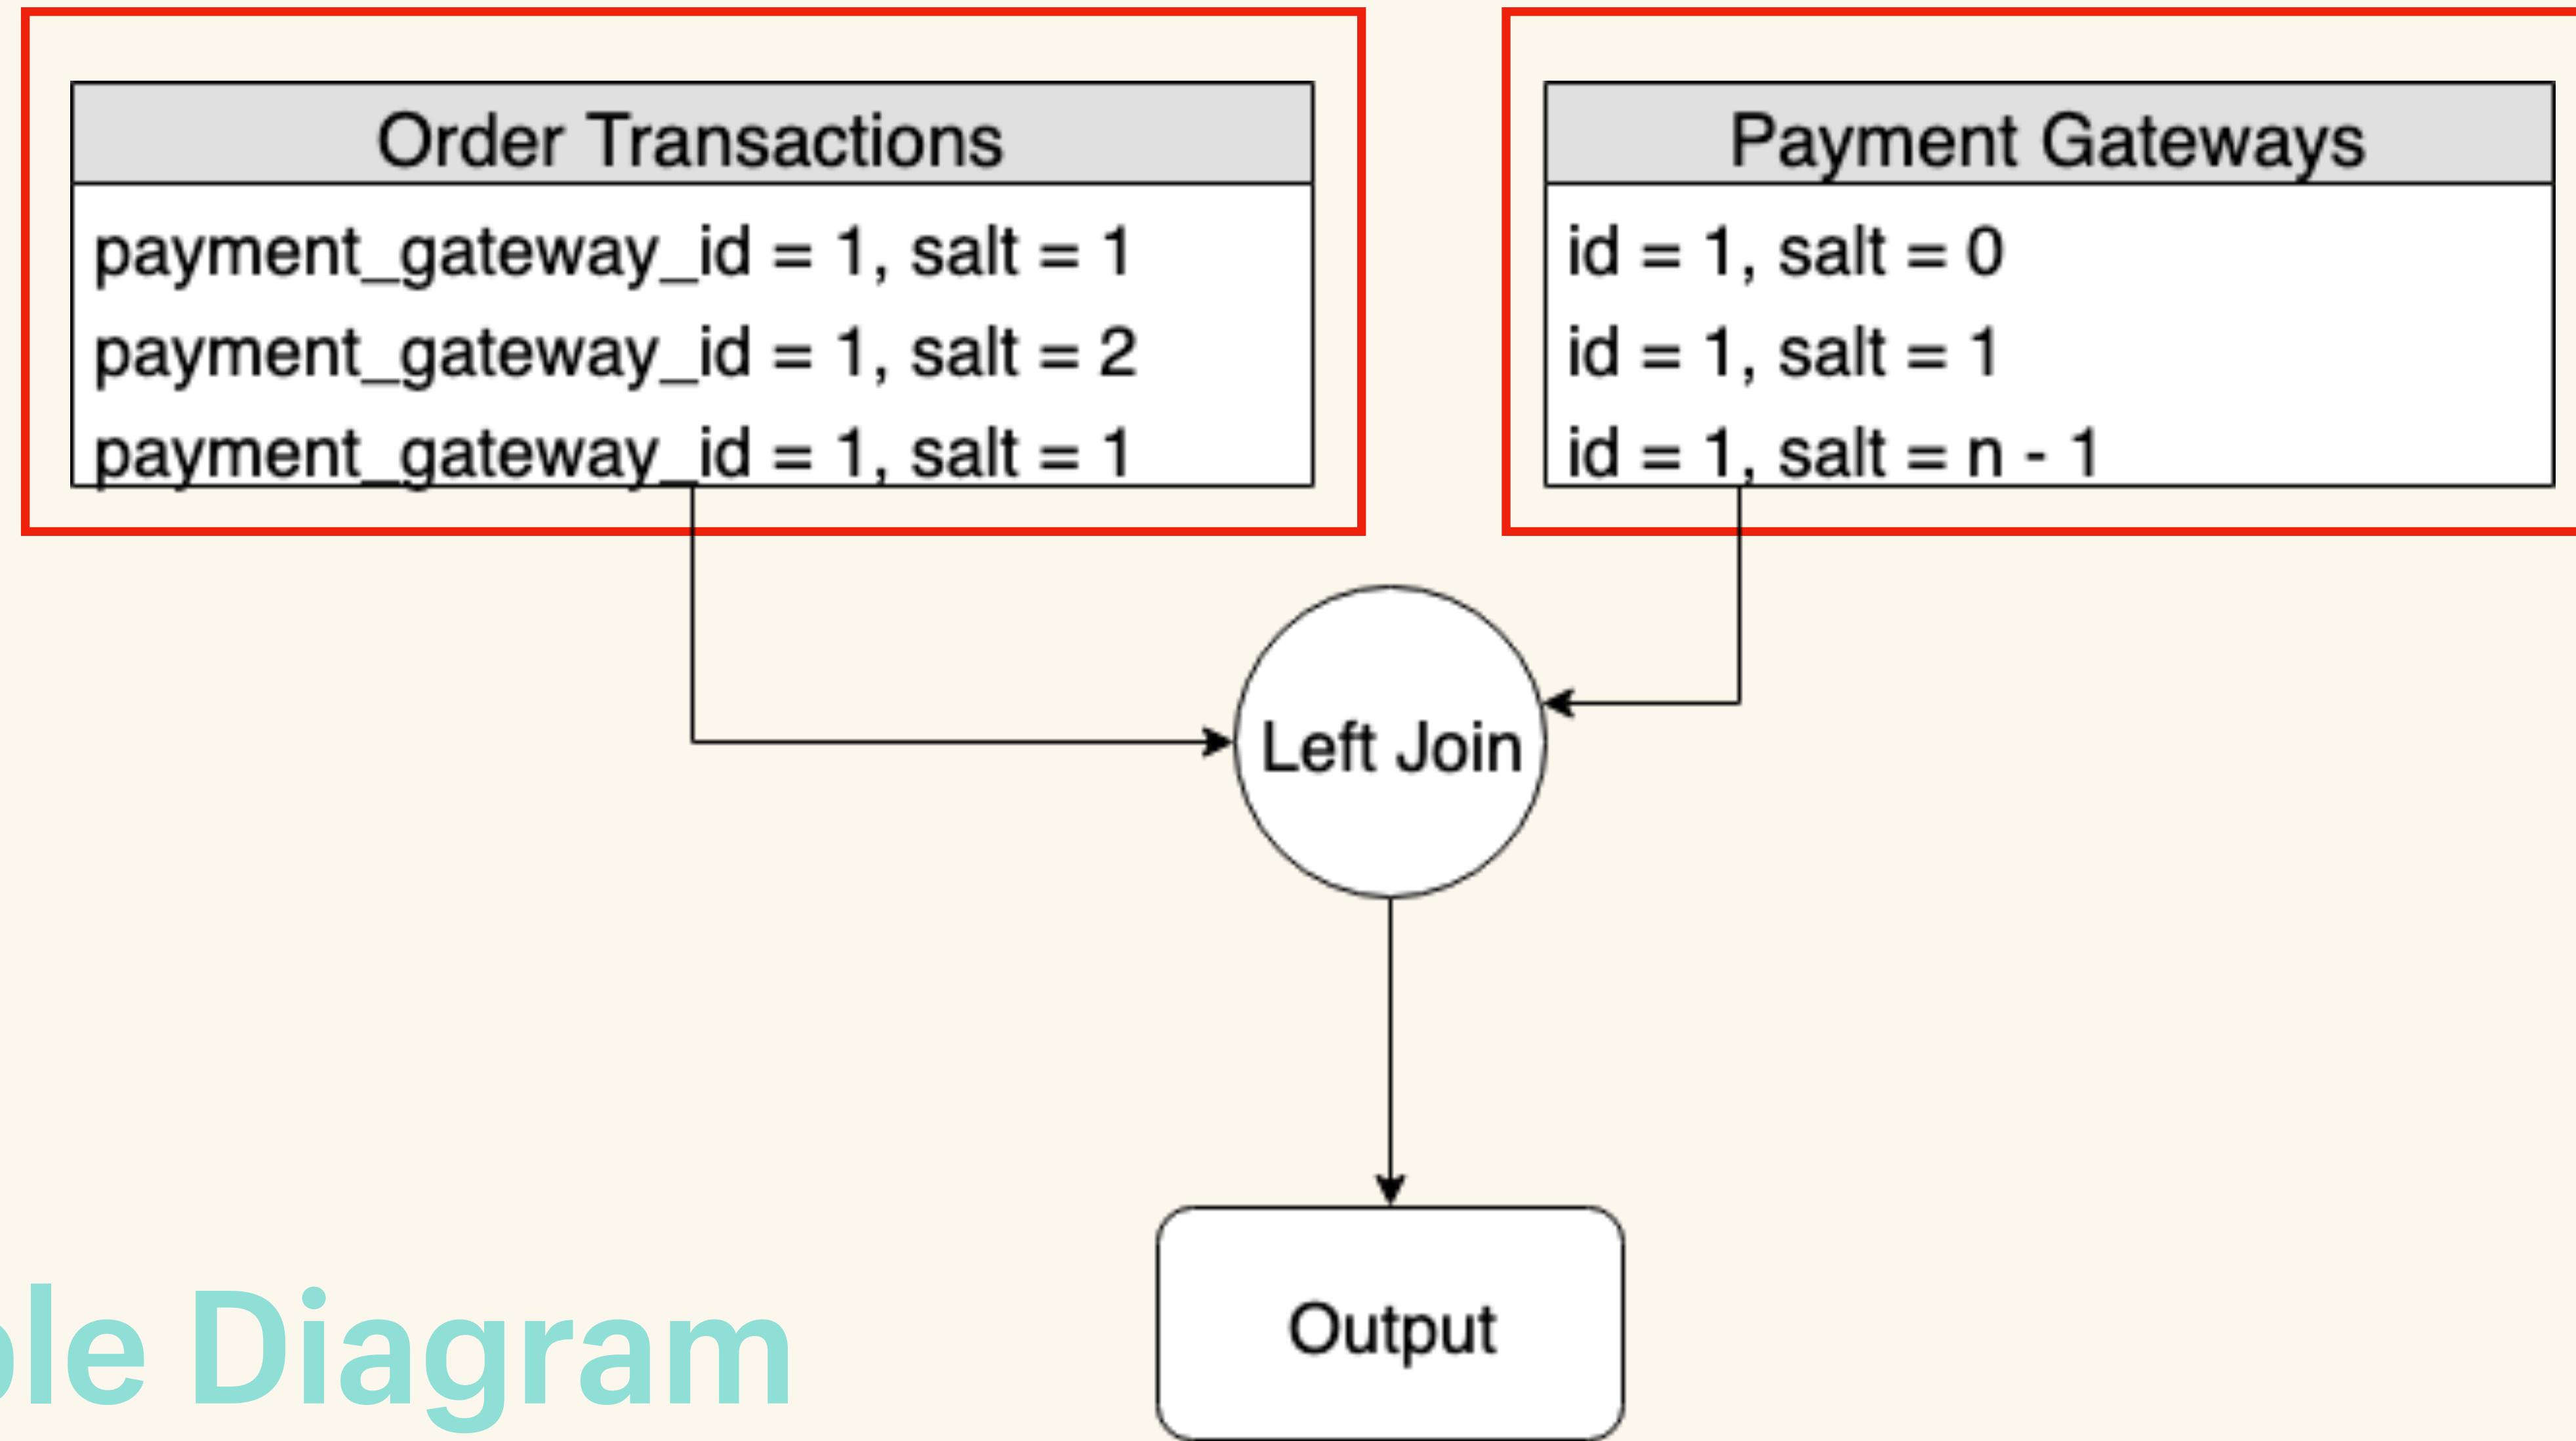


# Table Diagram

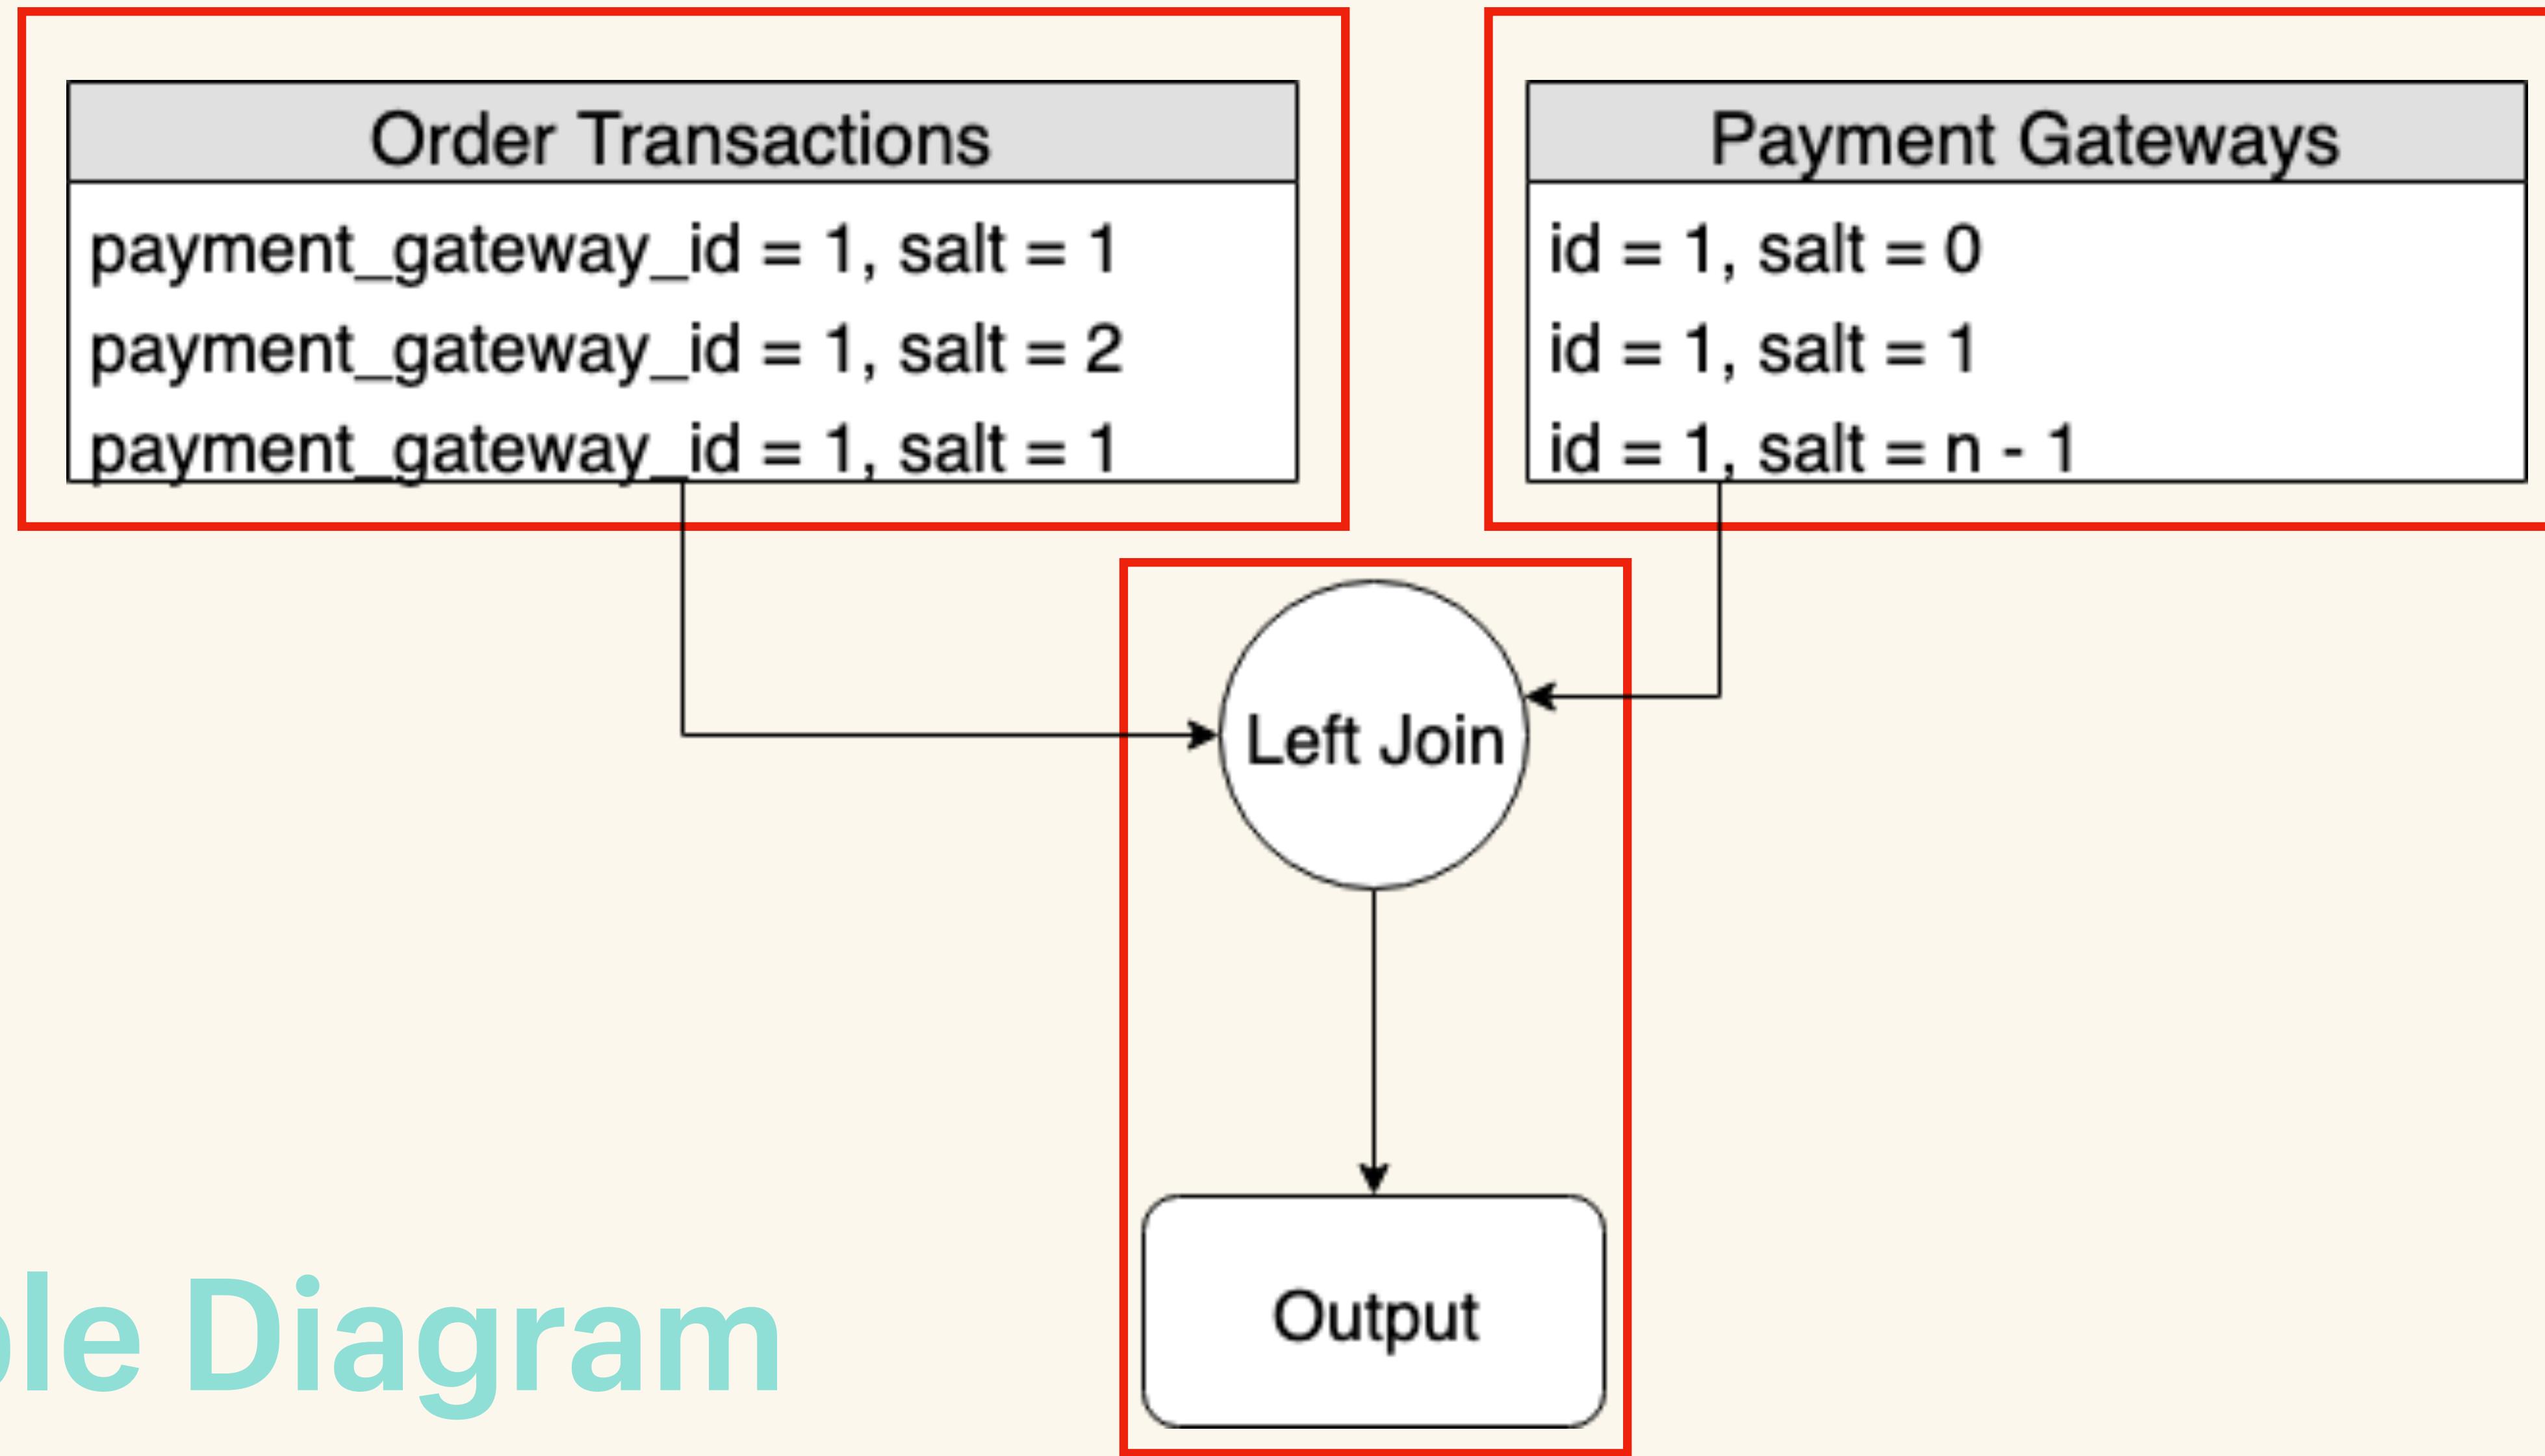


# Table Diagram

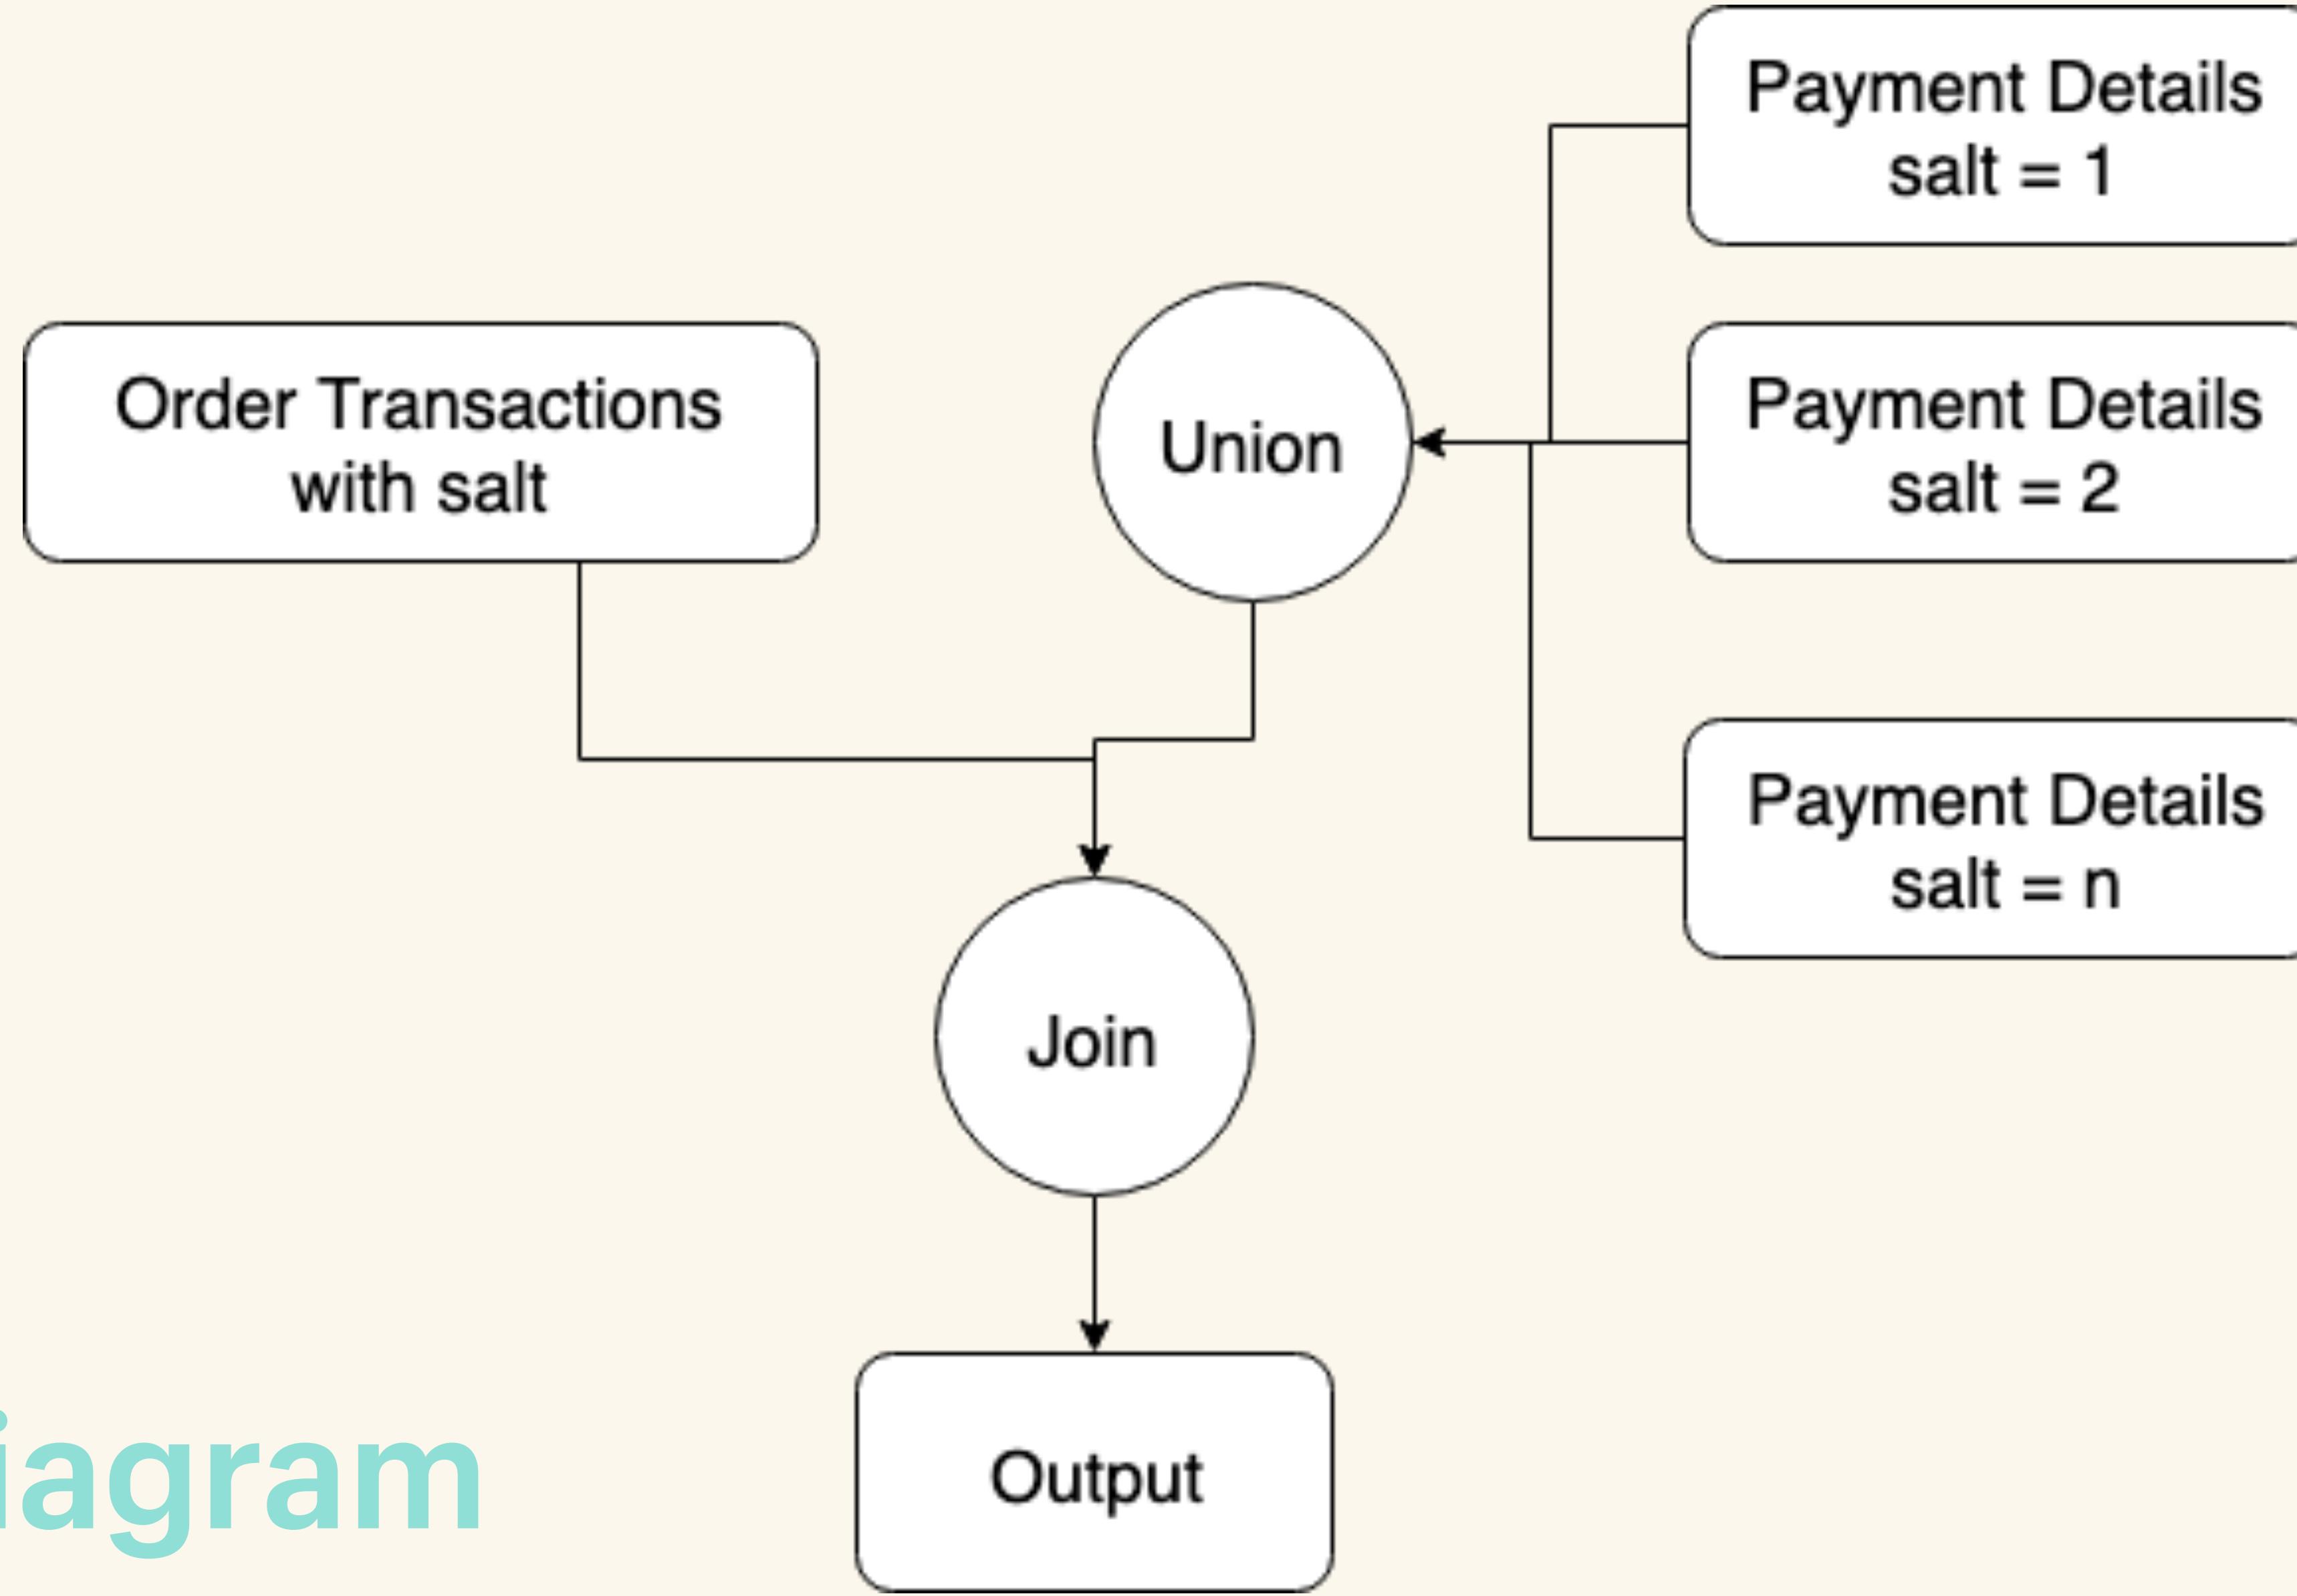
# Table Diagram



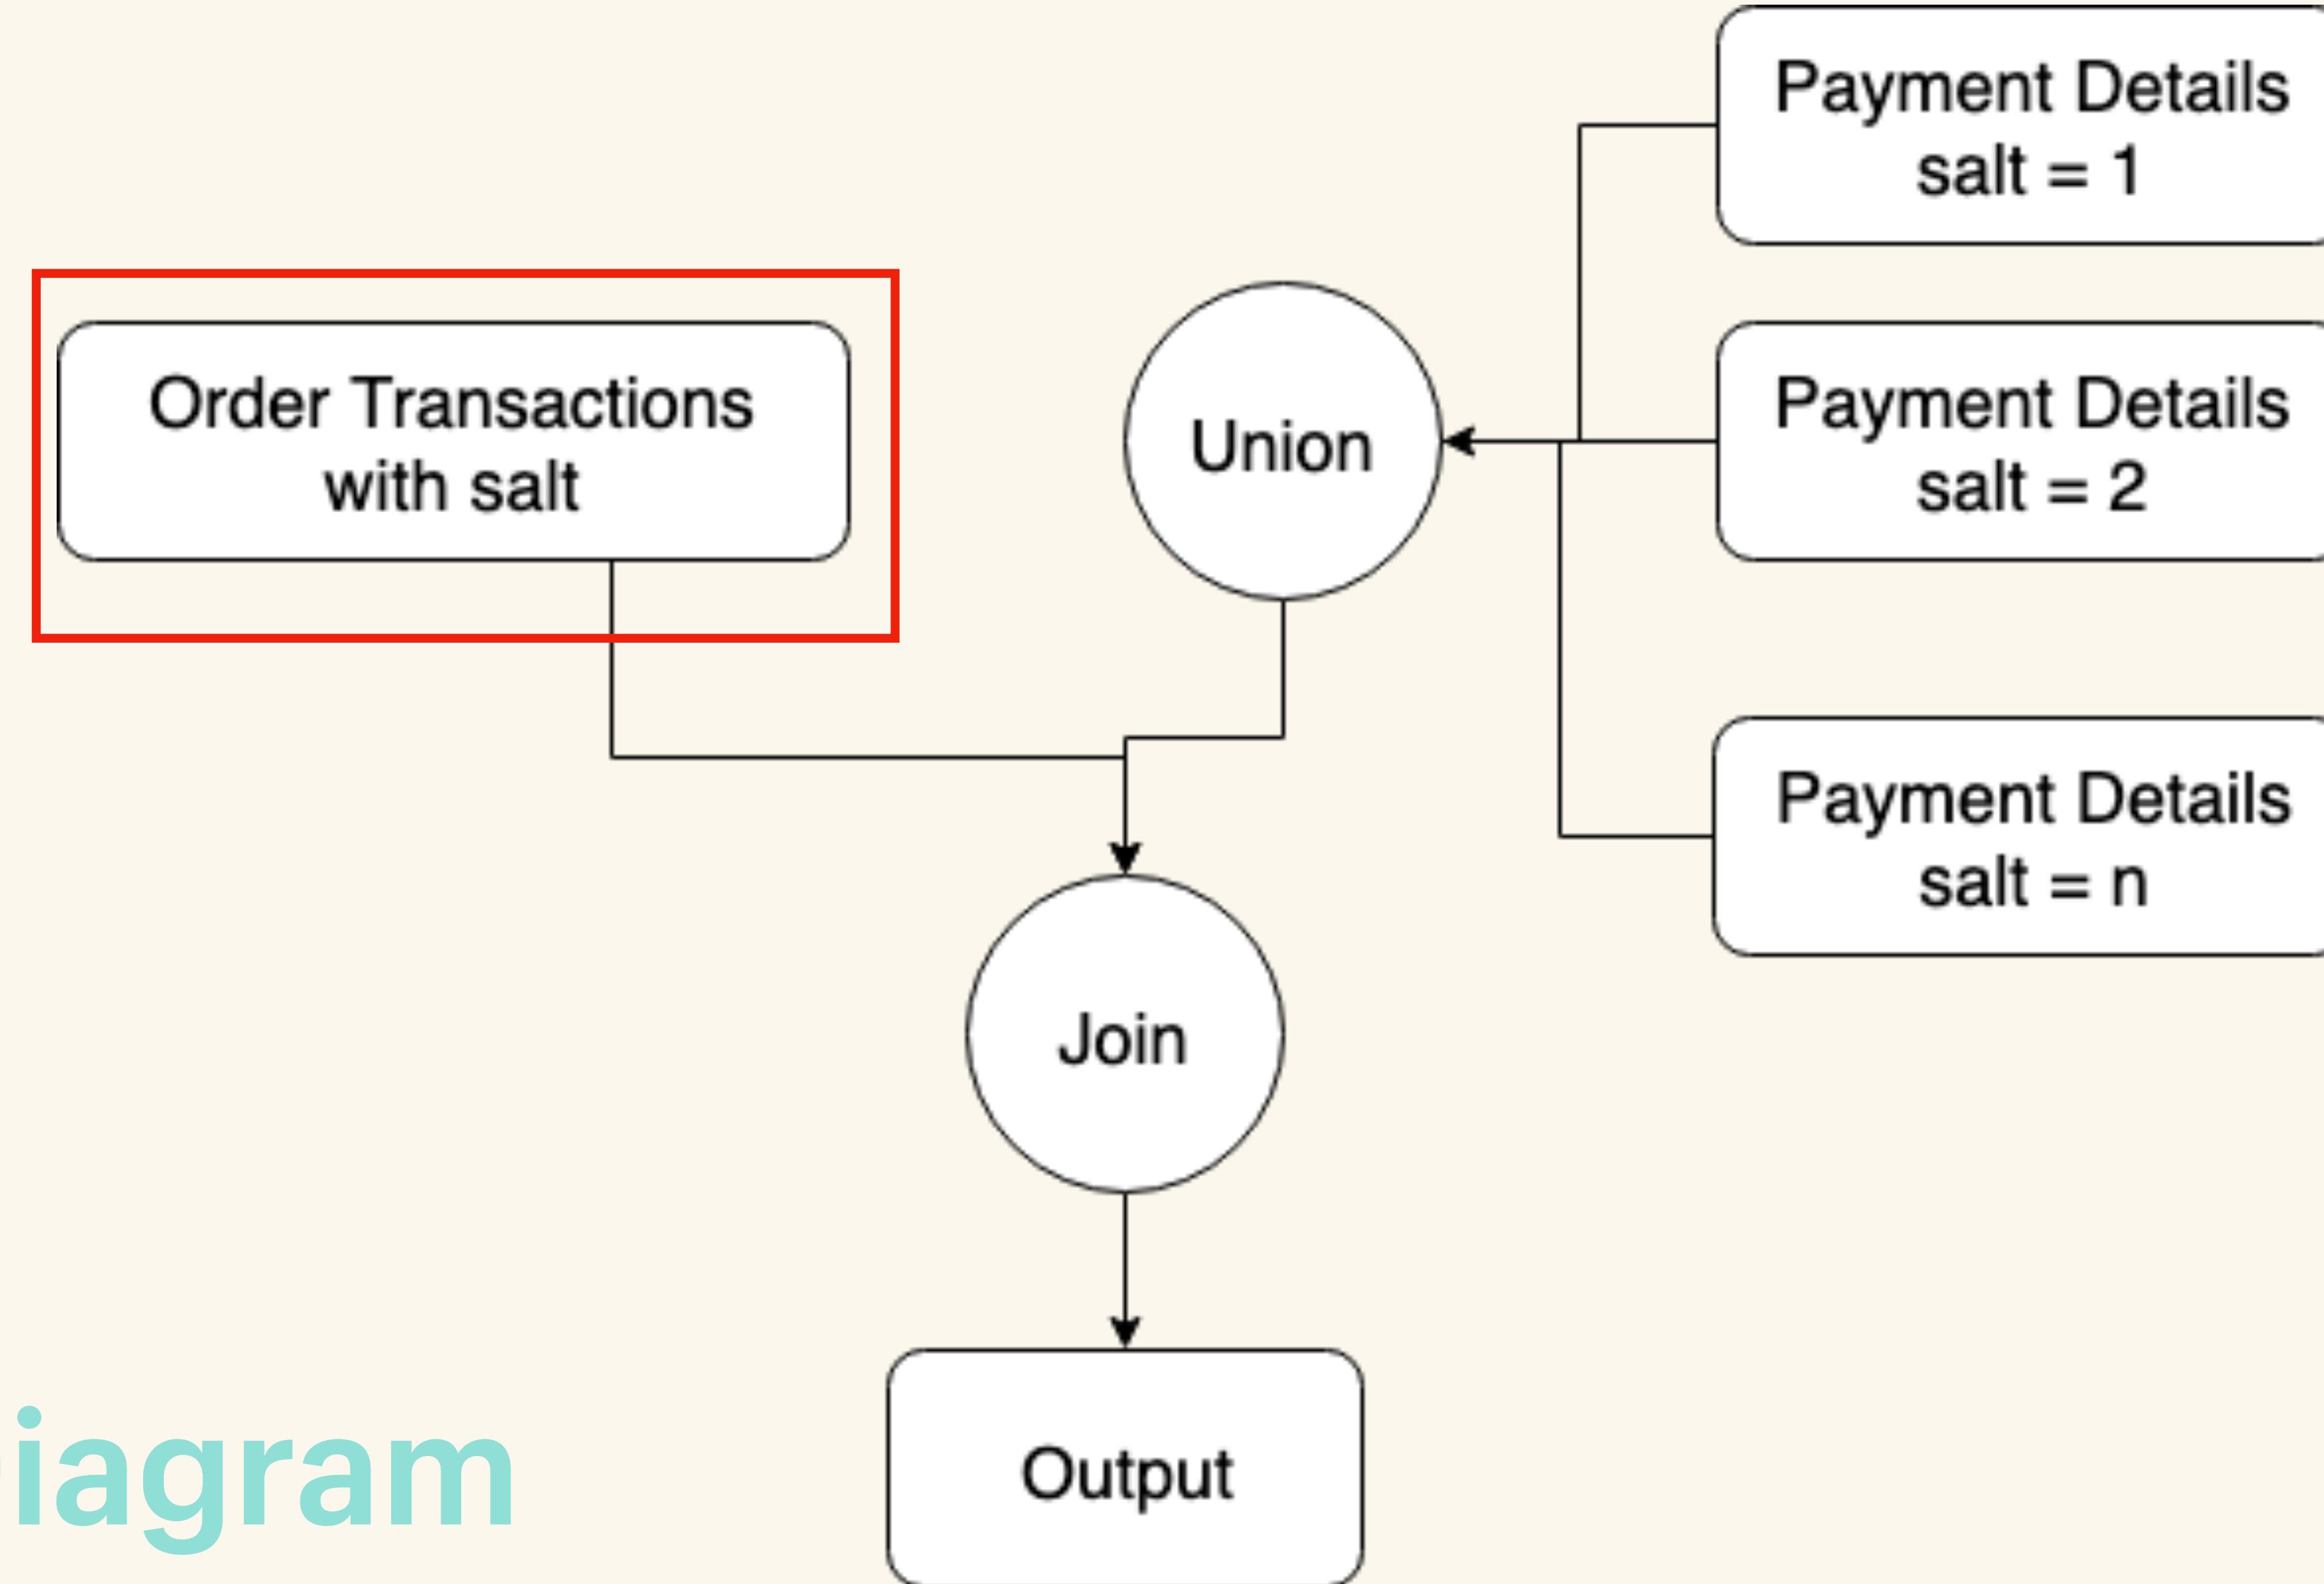
# Table Diagram

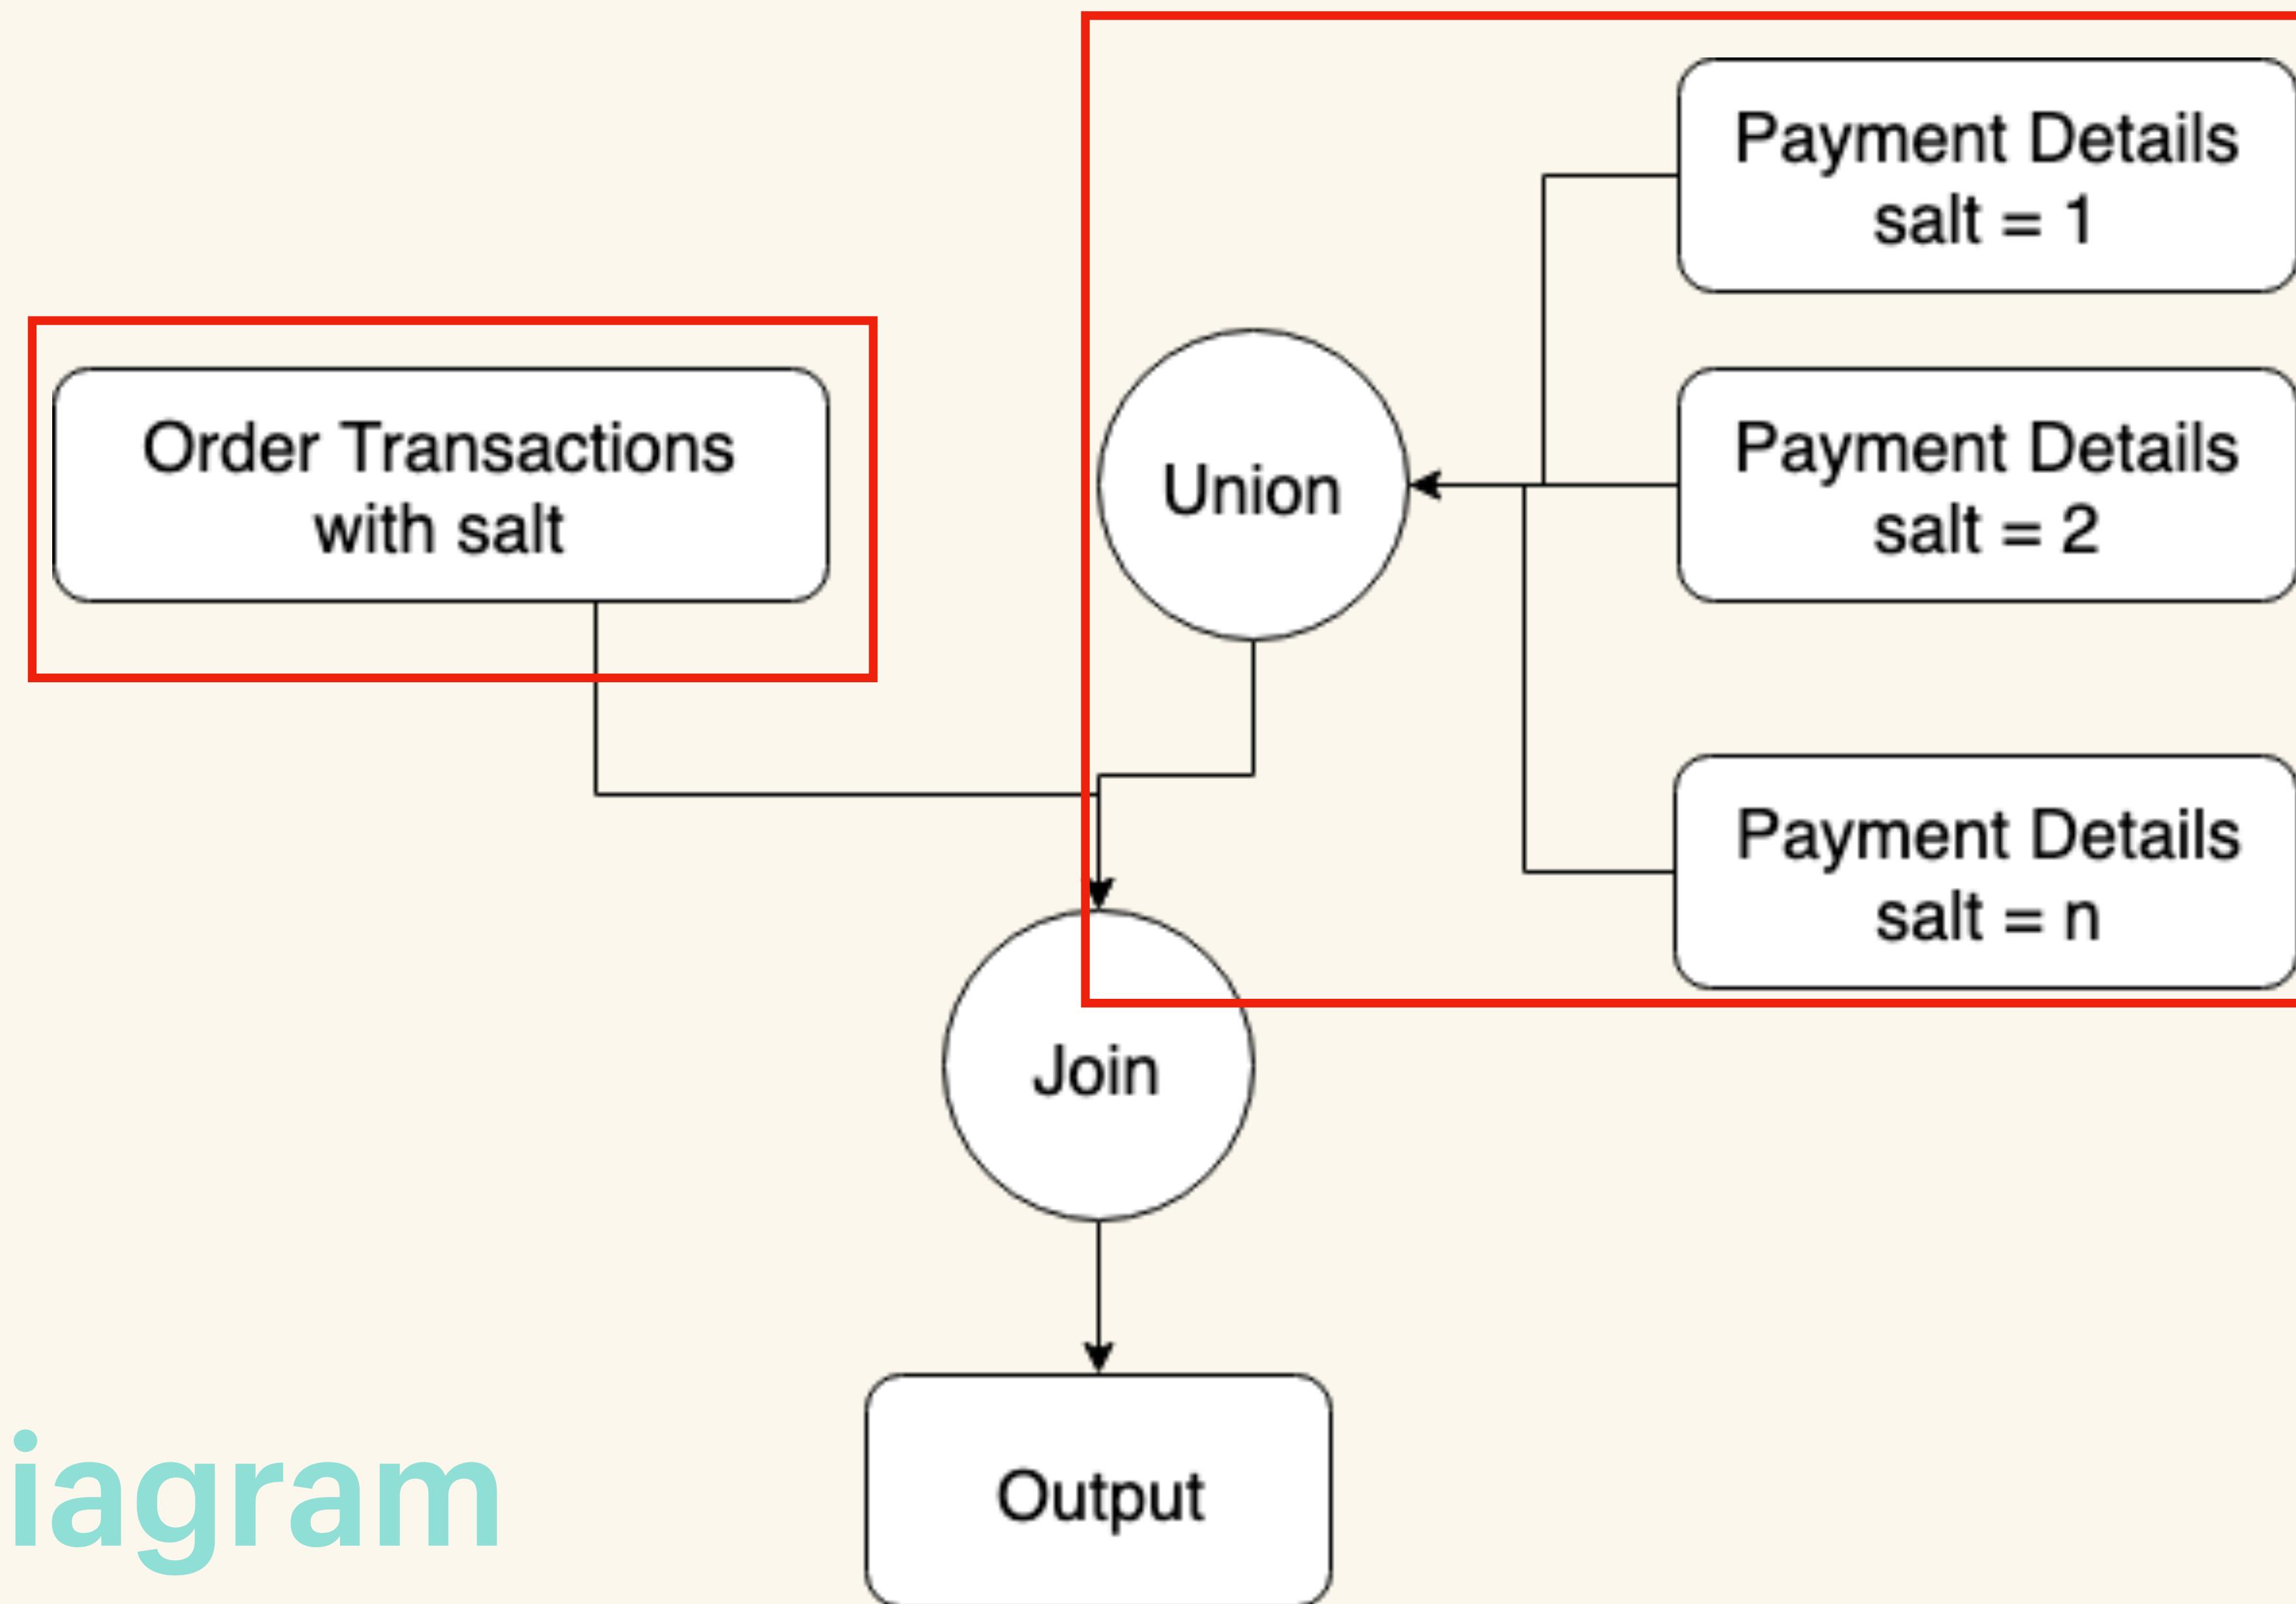


# SQL Diagram



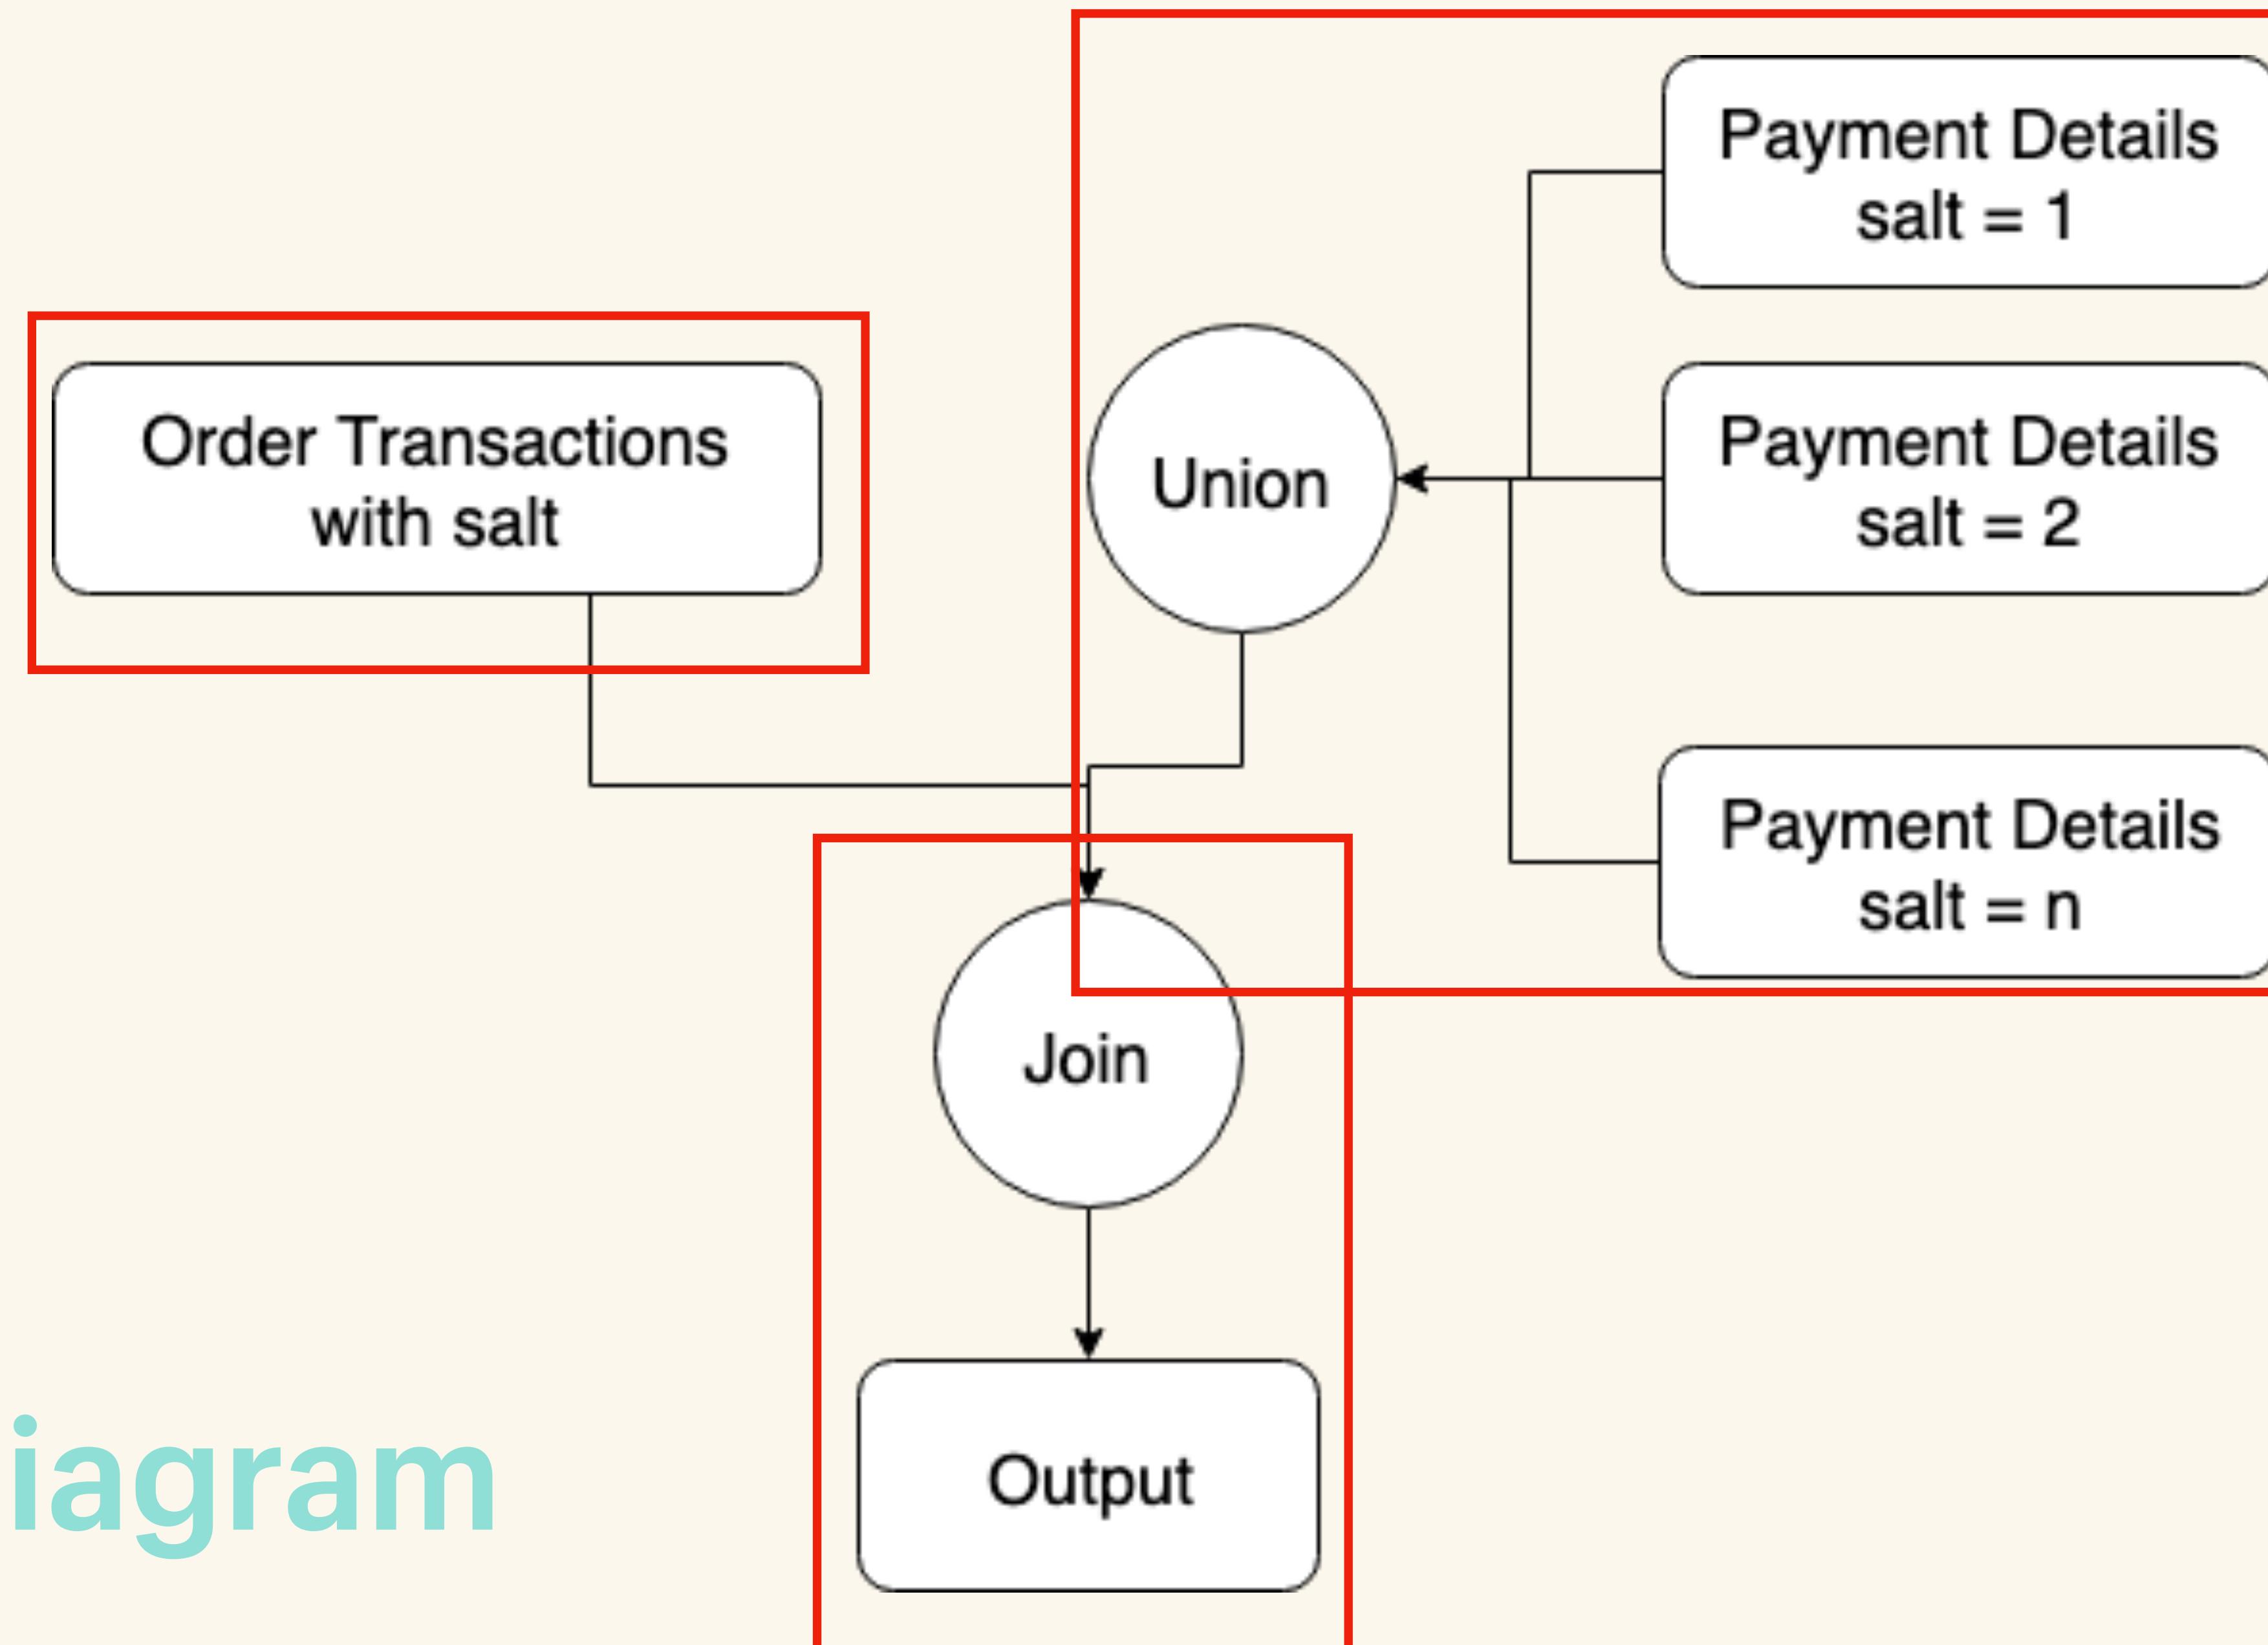
# SQL Diagram





# SQL Diagram

# SQL Diagram



# Downsides

- The right side of the join (the dimension) grows linearly with  $n$ .
- Depending how large the right side is, this can potentially be bad.



# FrequentValuesSkewHelper

## Motivation:

- Eliminate the shuffle of rows containing the subset of skewed keys.
- These rows are too big to fit onto tasks.



# FrequentValuesSkewHelper

## Solution:

- Split the skewed dataset into 2 subsets.
- First subset contains only skewed keys.
- Second contains non skewed keys.
- Do same with right dataset.
- Perform a regular join with non skewed datasets.
- Do a broadcast join on skewed dataset.
- Then union that with the joined result.
- Broadcast join and union do not produce a shuffle. \*\*



We can broadcast cause only a small subset (under broadcast range) of keys are considered skewed.

# FrequentValuesSkewHelper

## Under the Hood

- Calculates the keys that make up a large portion of the left dataset.
- Splits left and right dataset into 2 parts:
  1. All the rows that make up a large portion of the dataset, call this the “high frequency” DataFrame.
  2. Everything else, call this the “low frequency” DataFrame.
- Broadcast the high frequency right side, this eliminates shuffle of left high frequency dataframe.
- Performs the join on both the “high frequency” Dataframes.
- Performs the join on both the “low frequency” Dataframes.
- Unions the two joined datasets together.

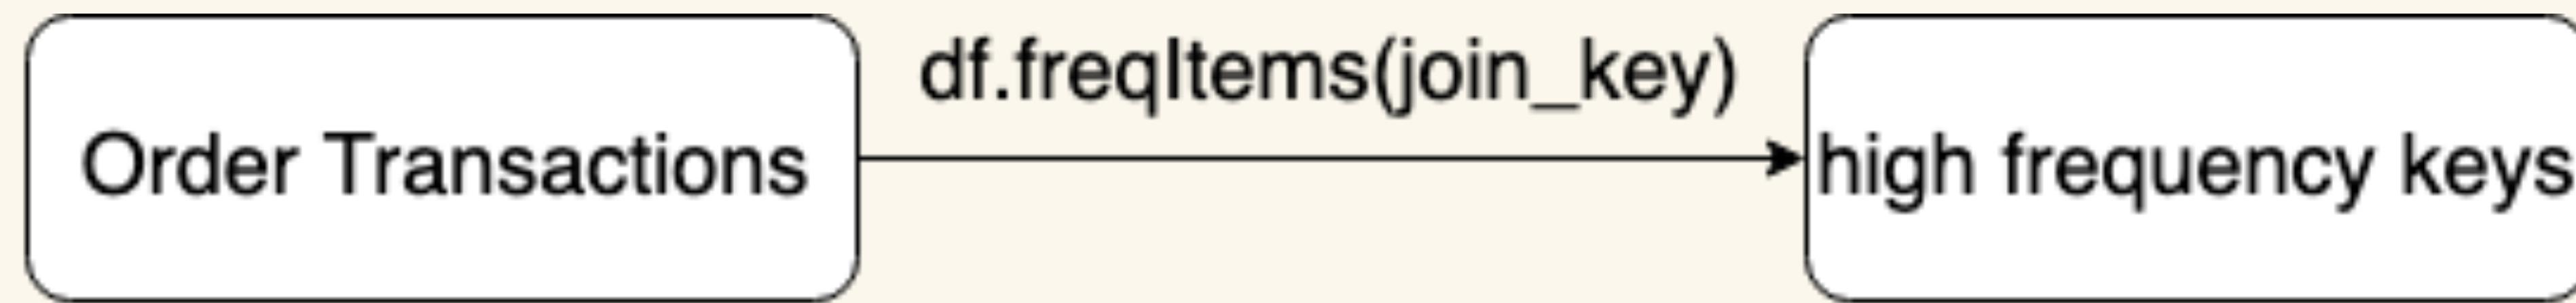




# Table Diagram

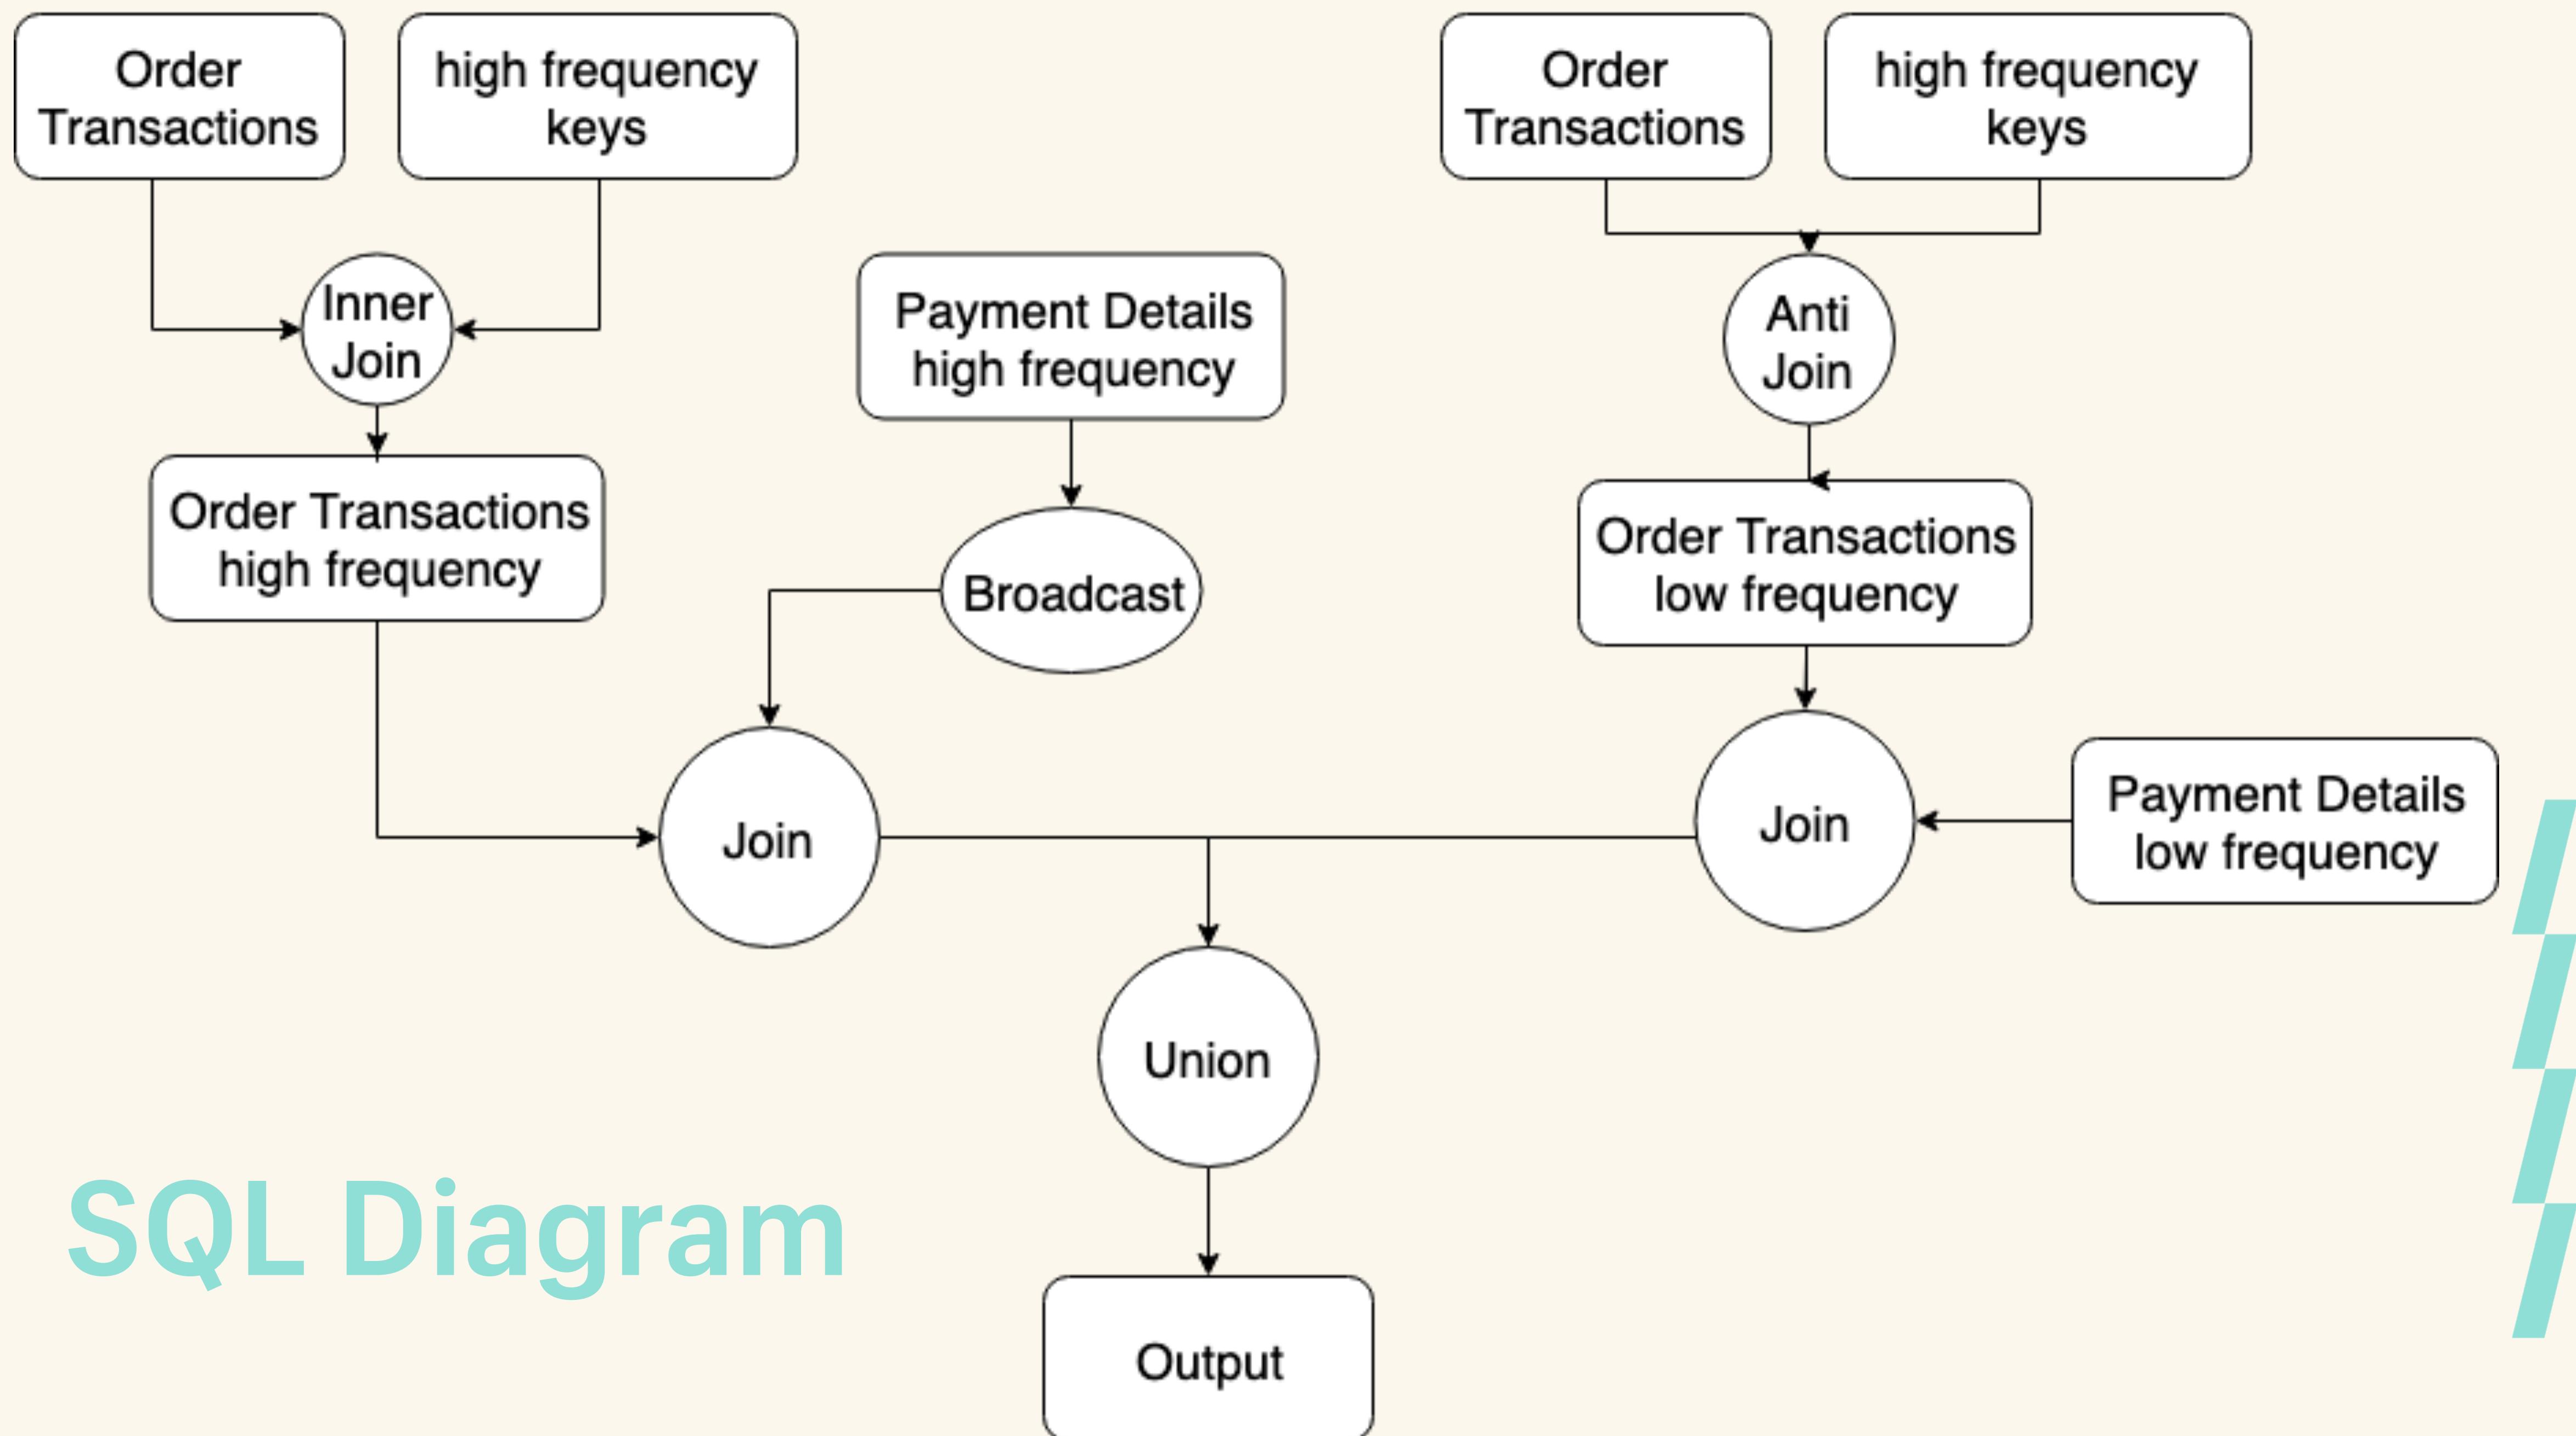


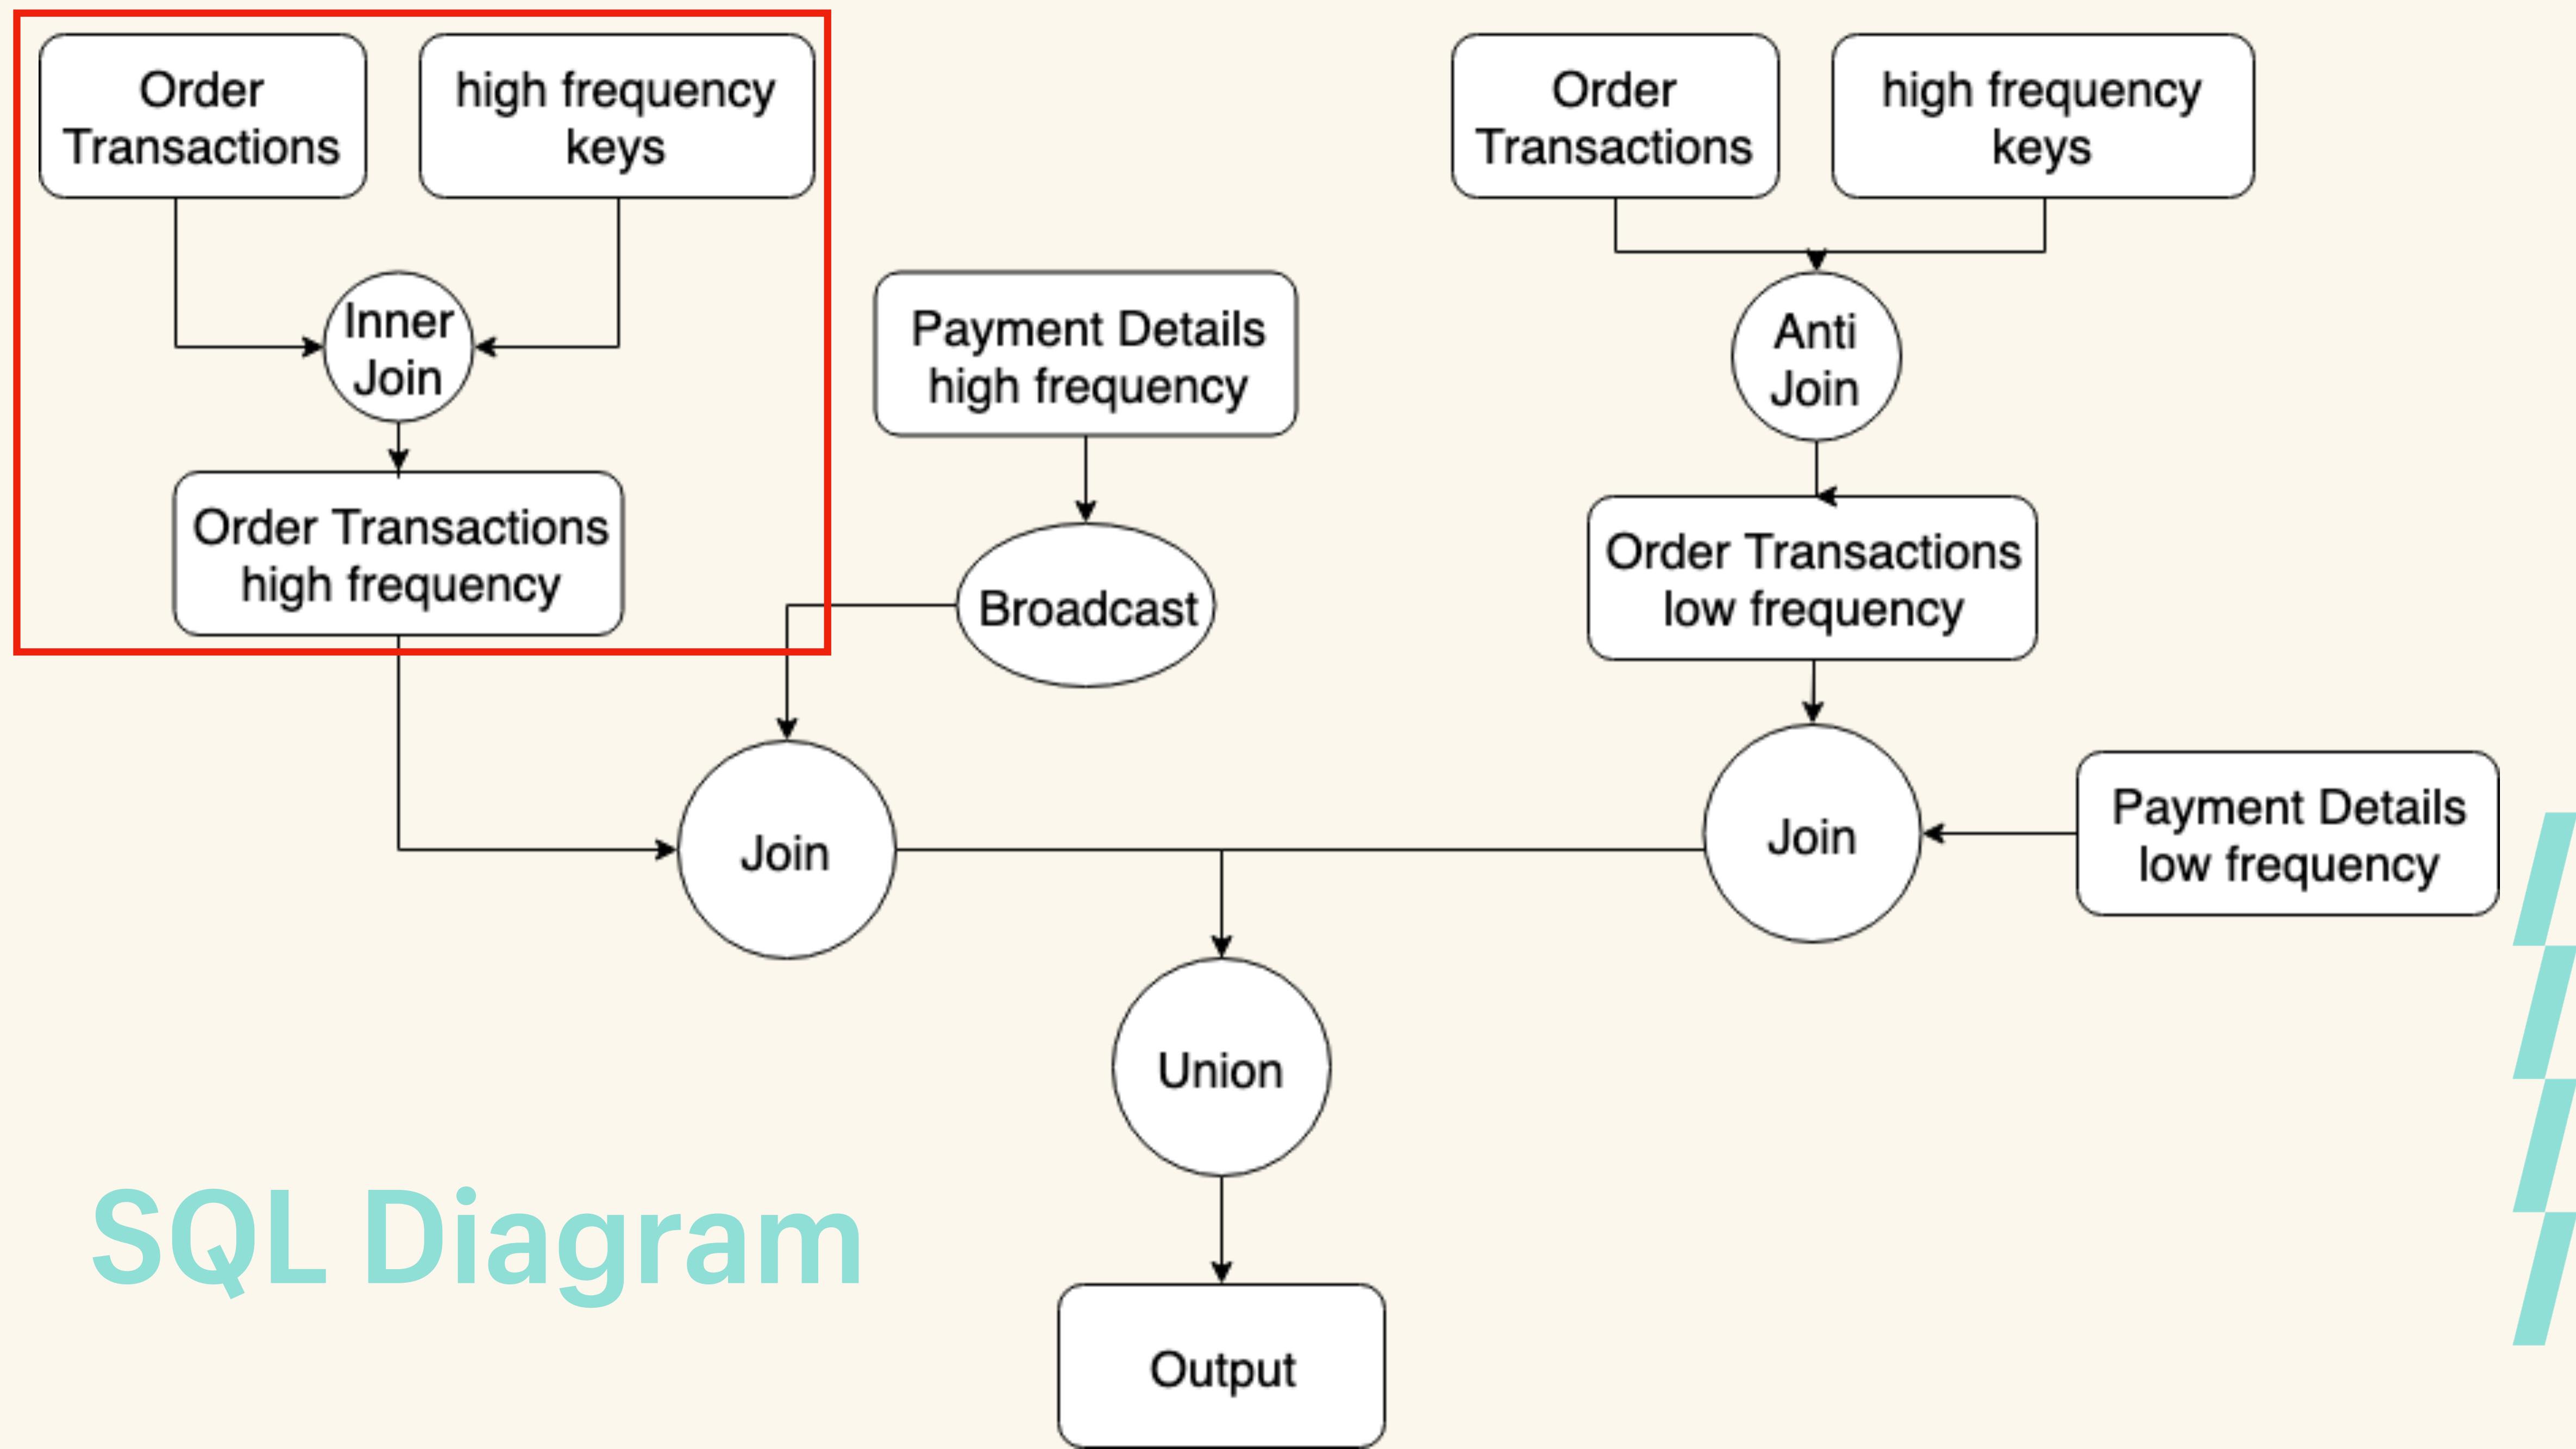
# Table Diagram



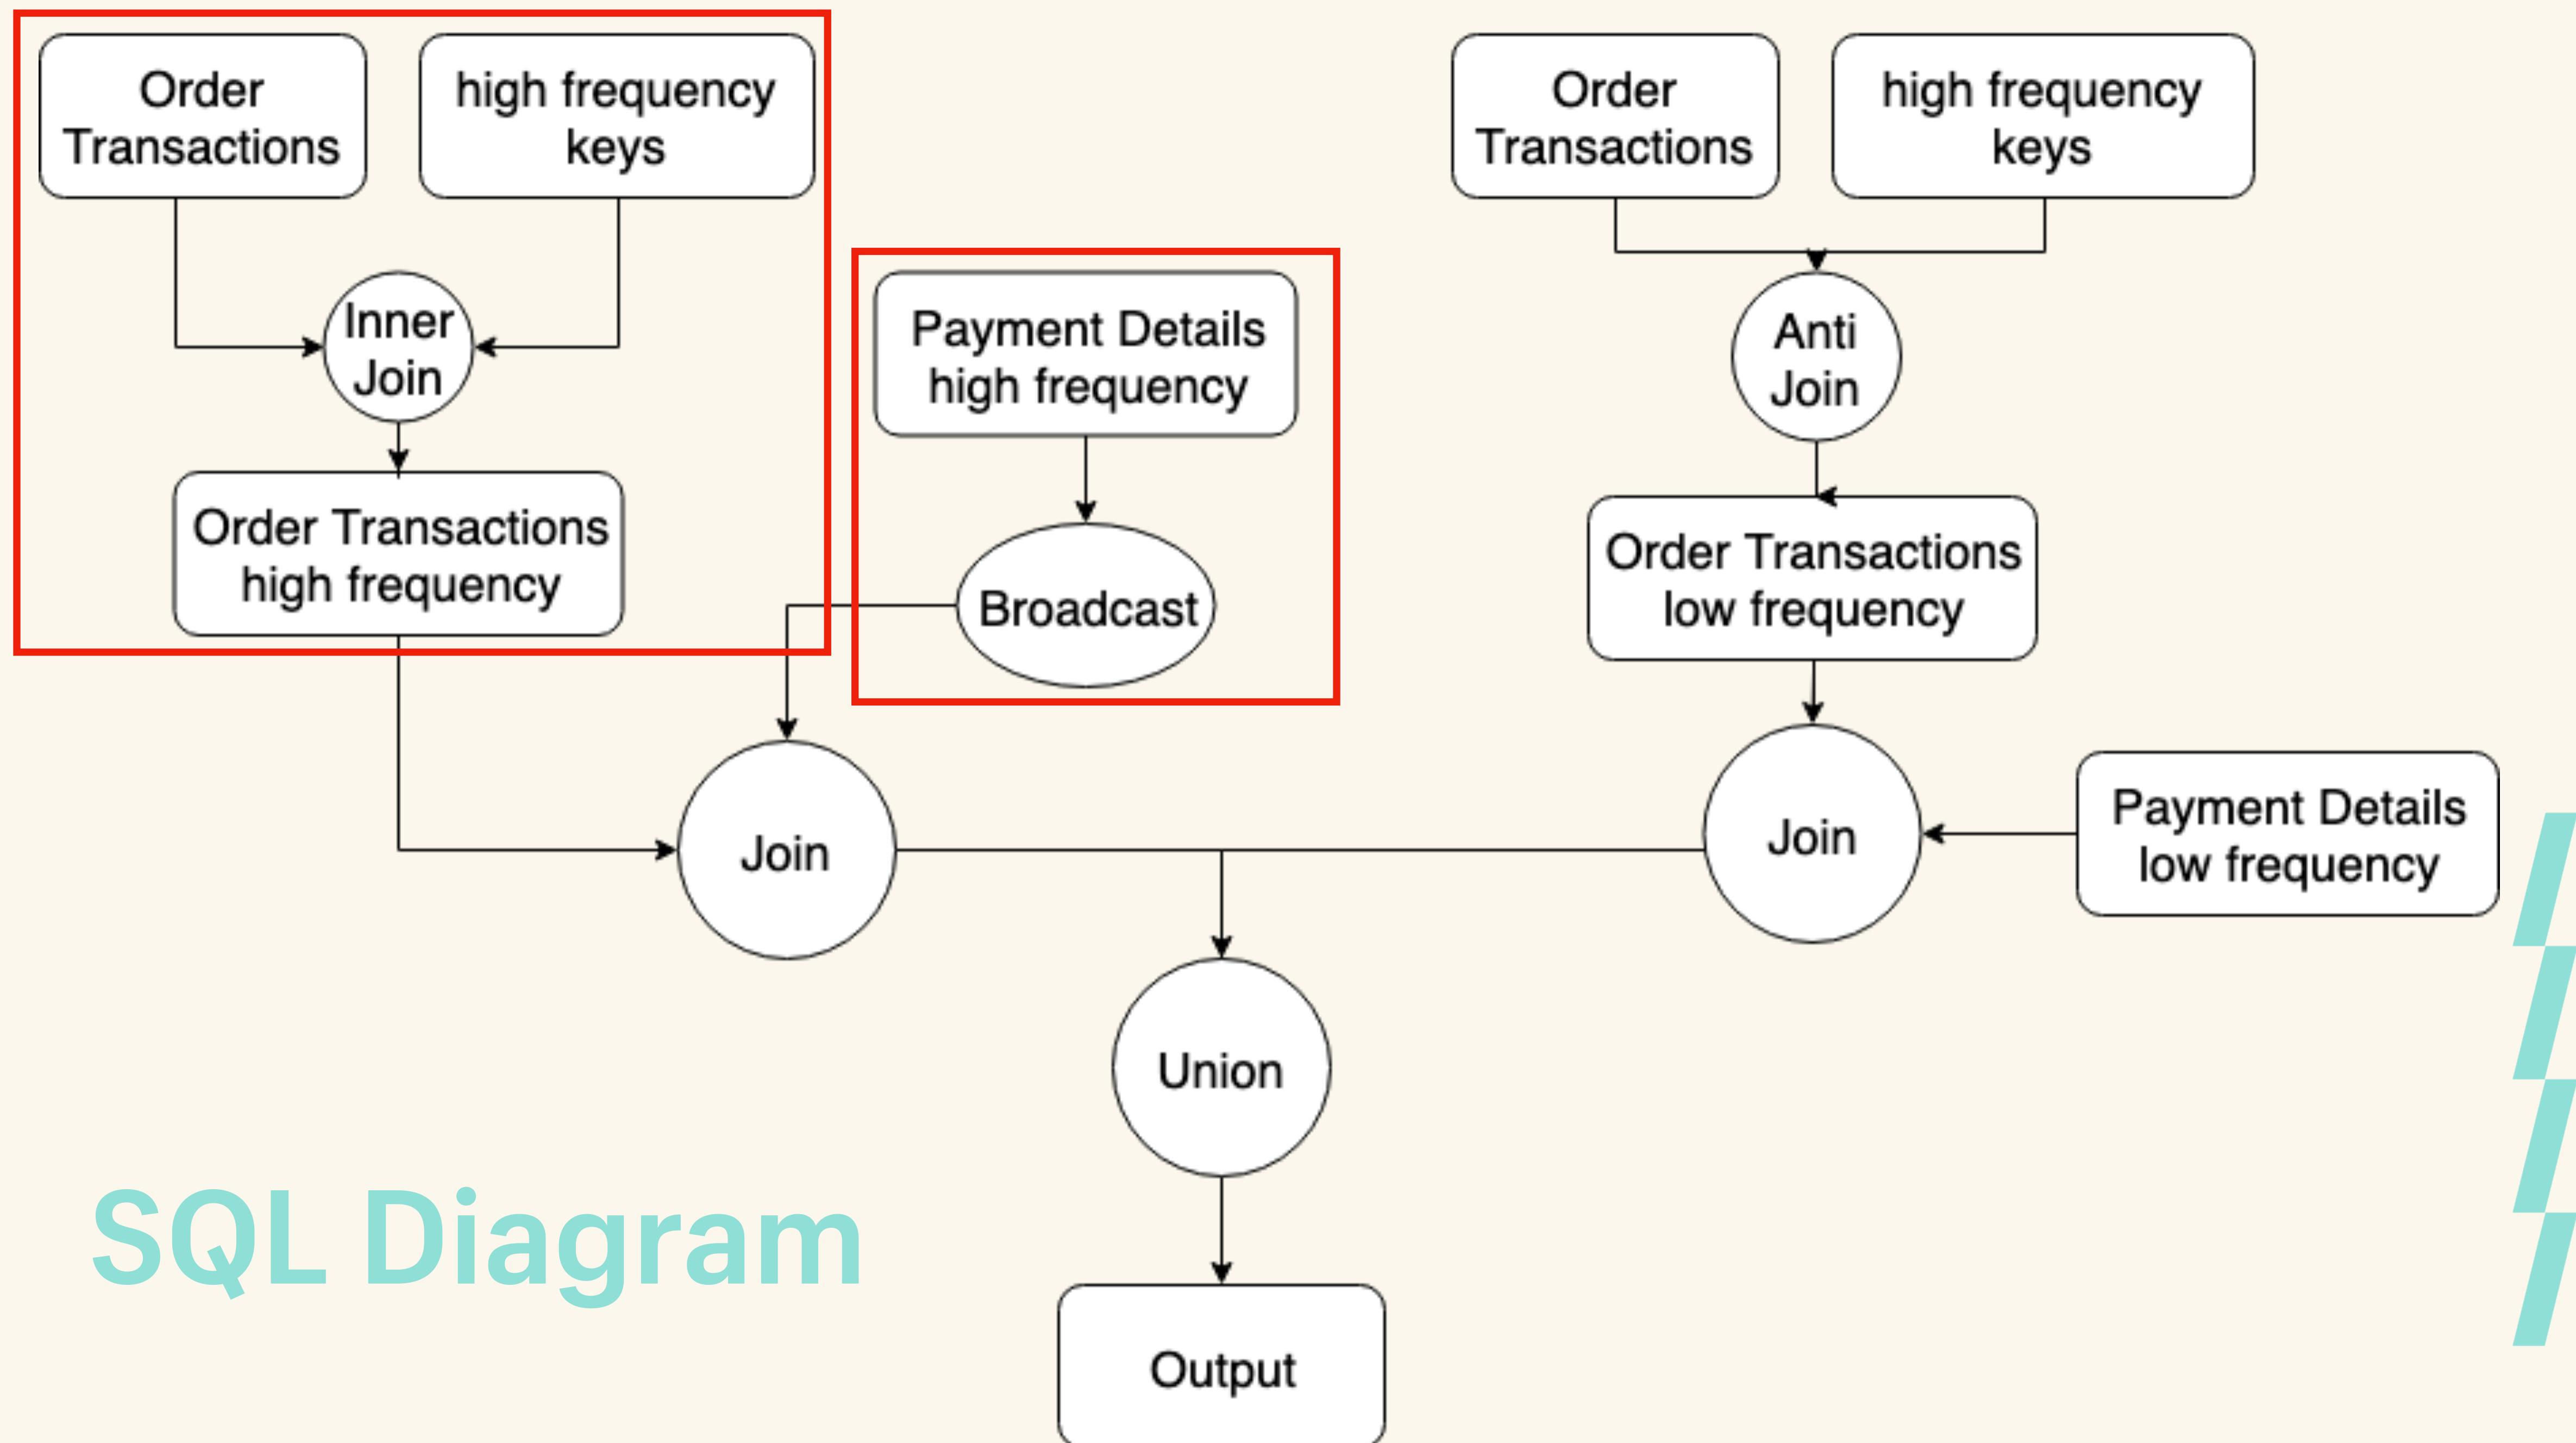
# SQL Diagram

# SQL Diagram

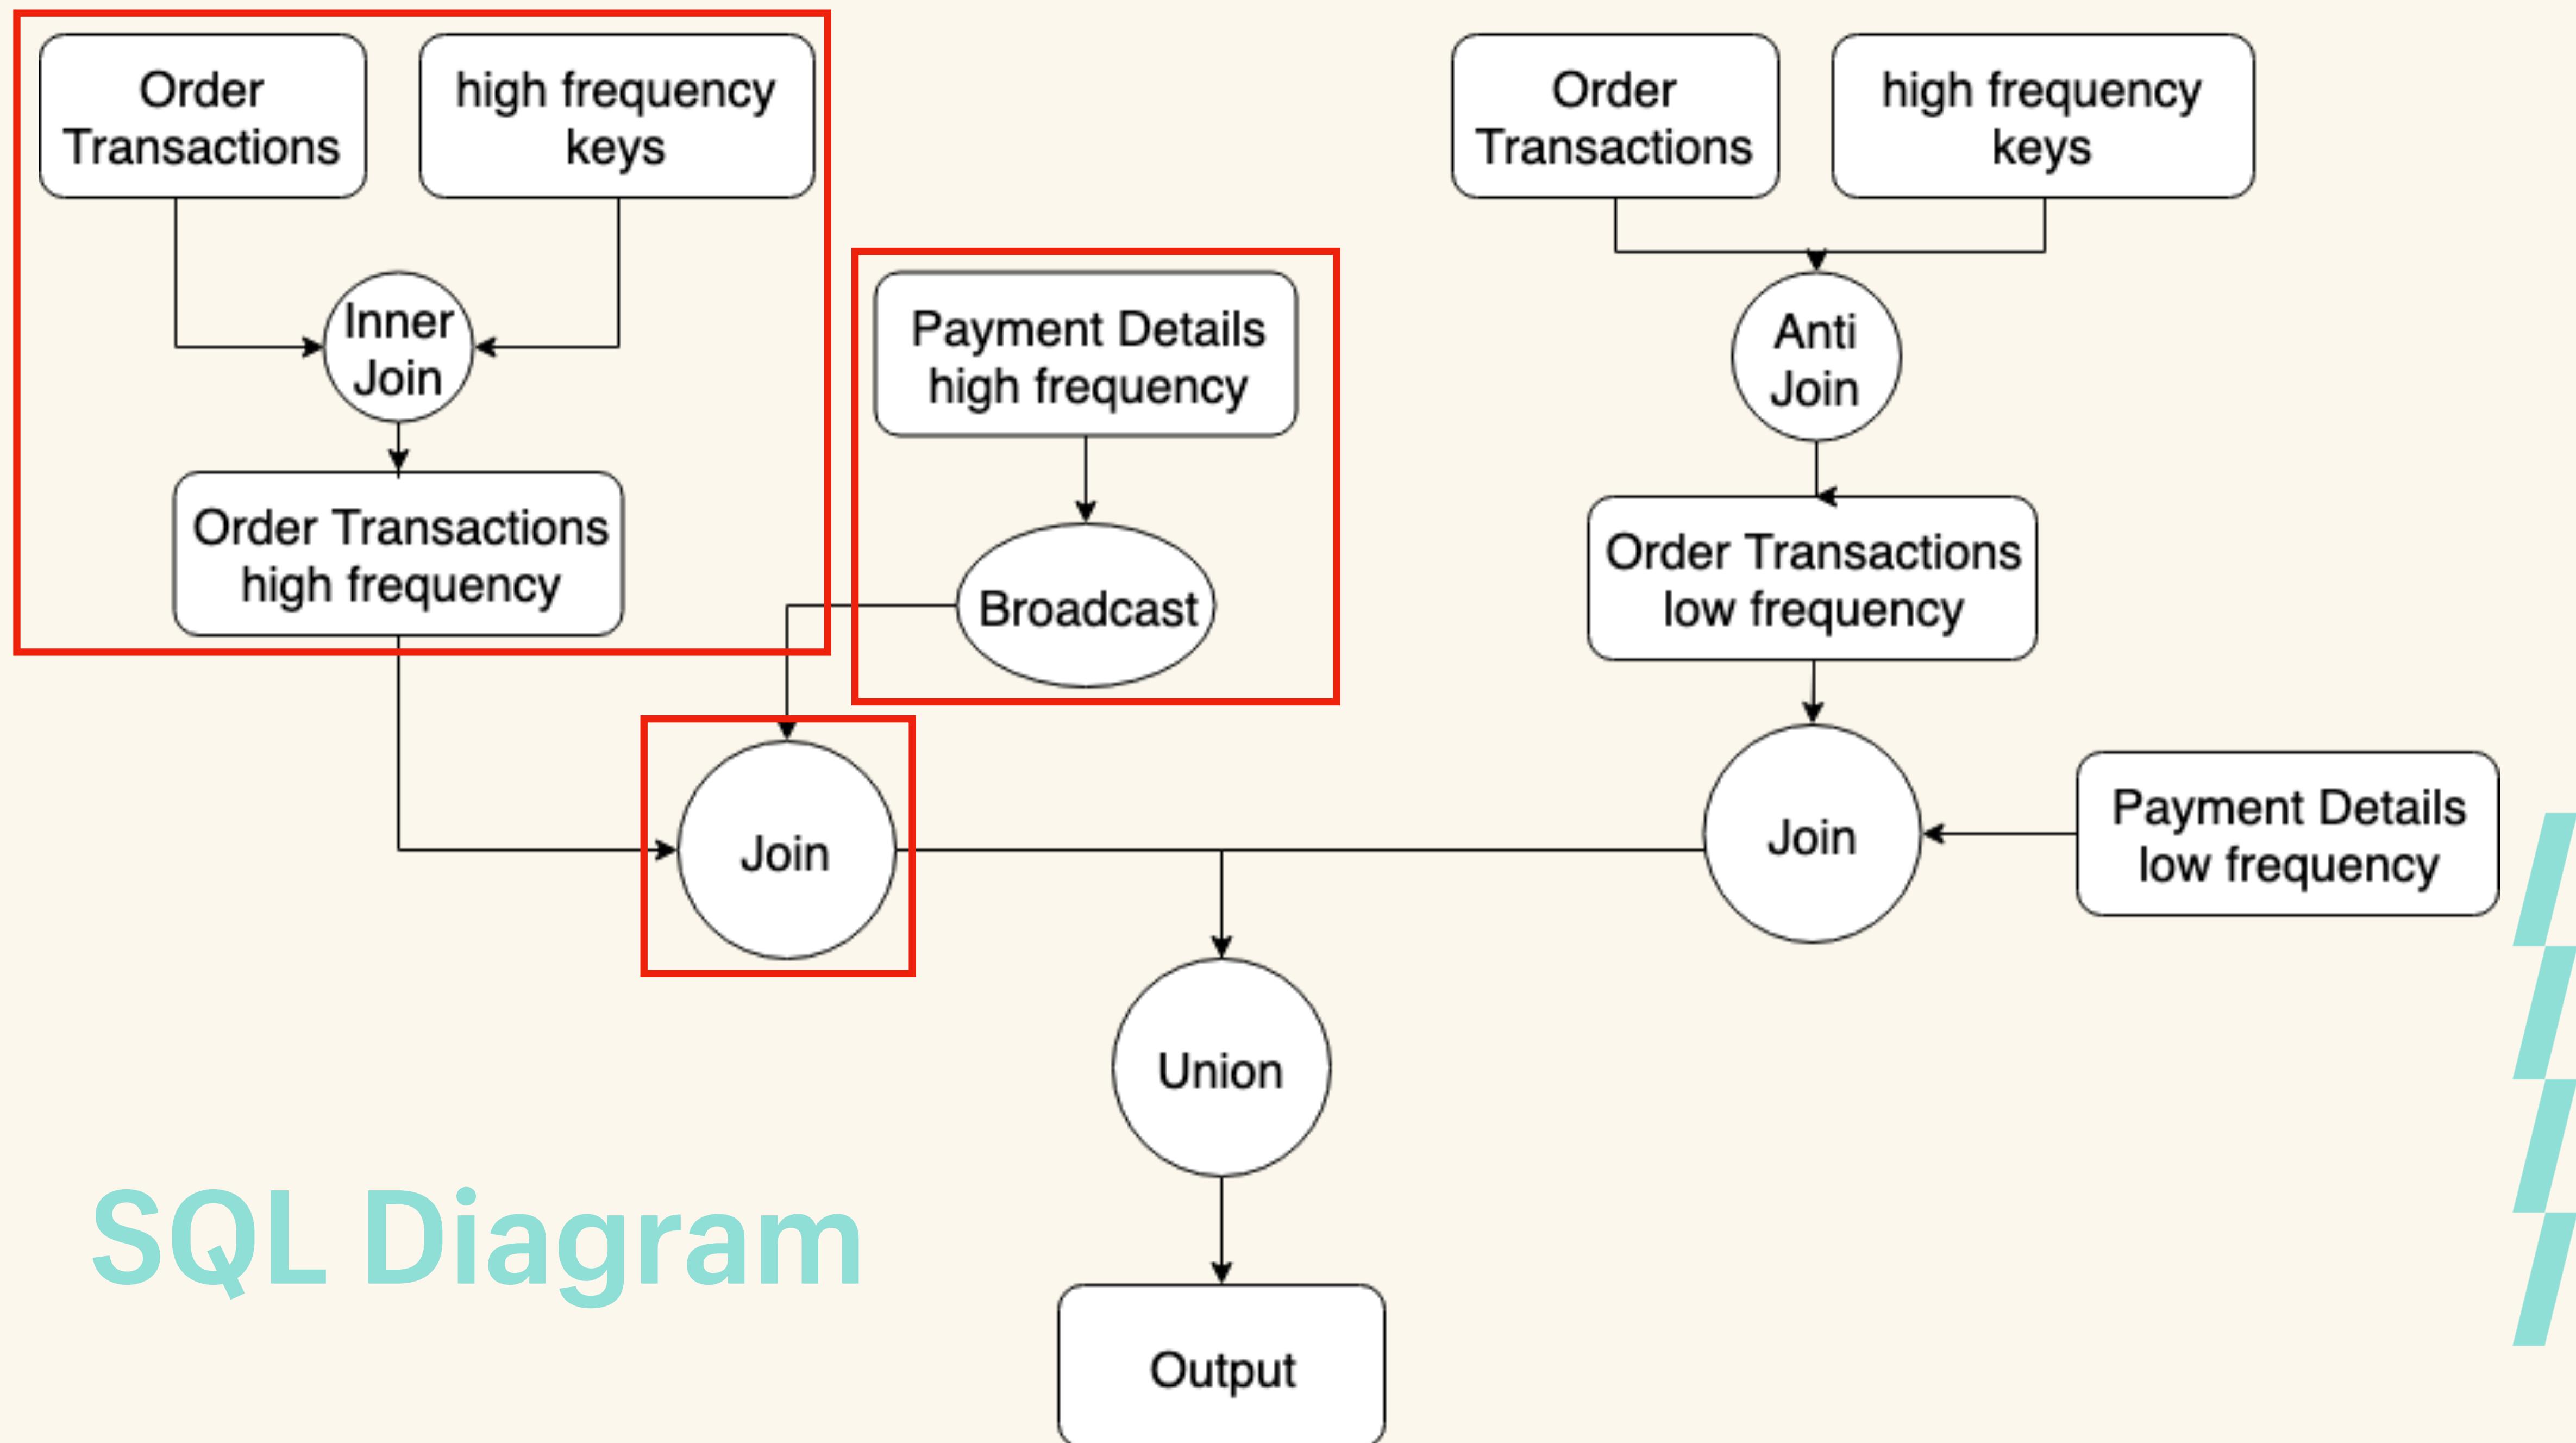




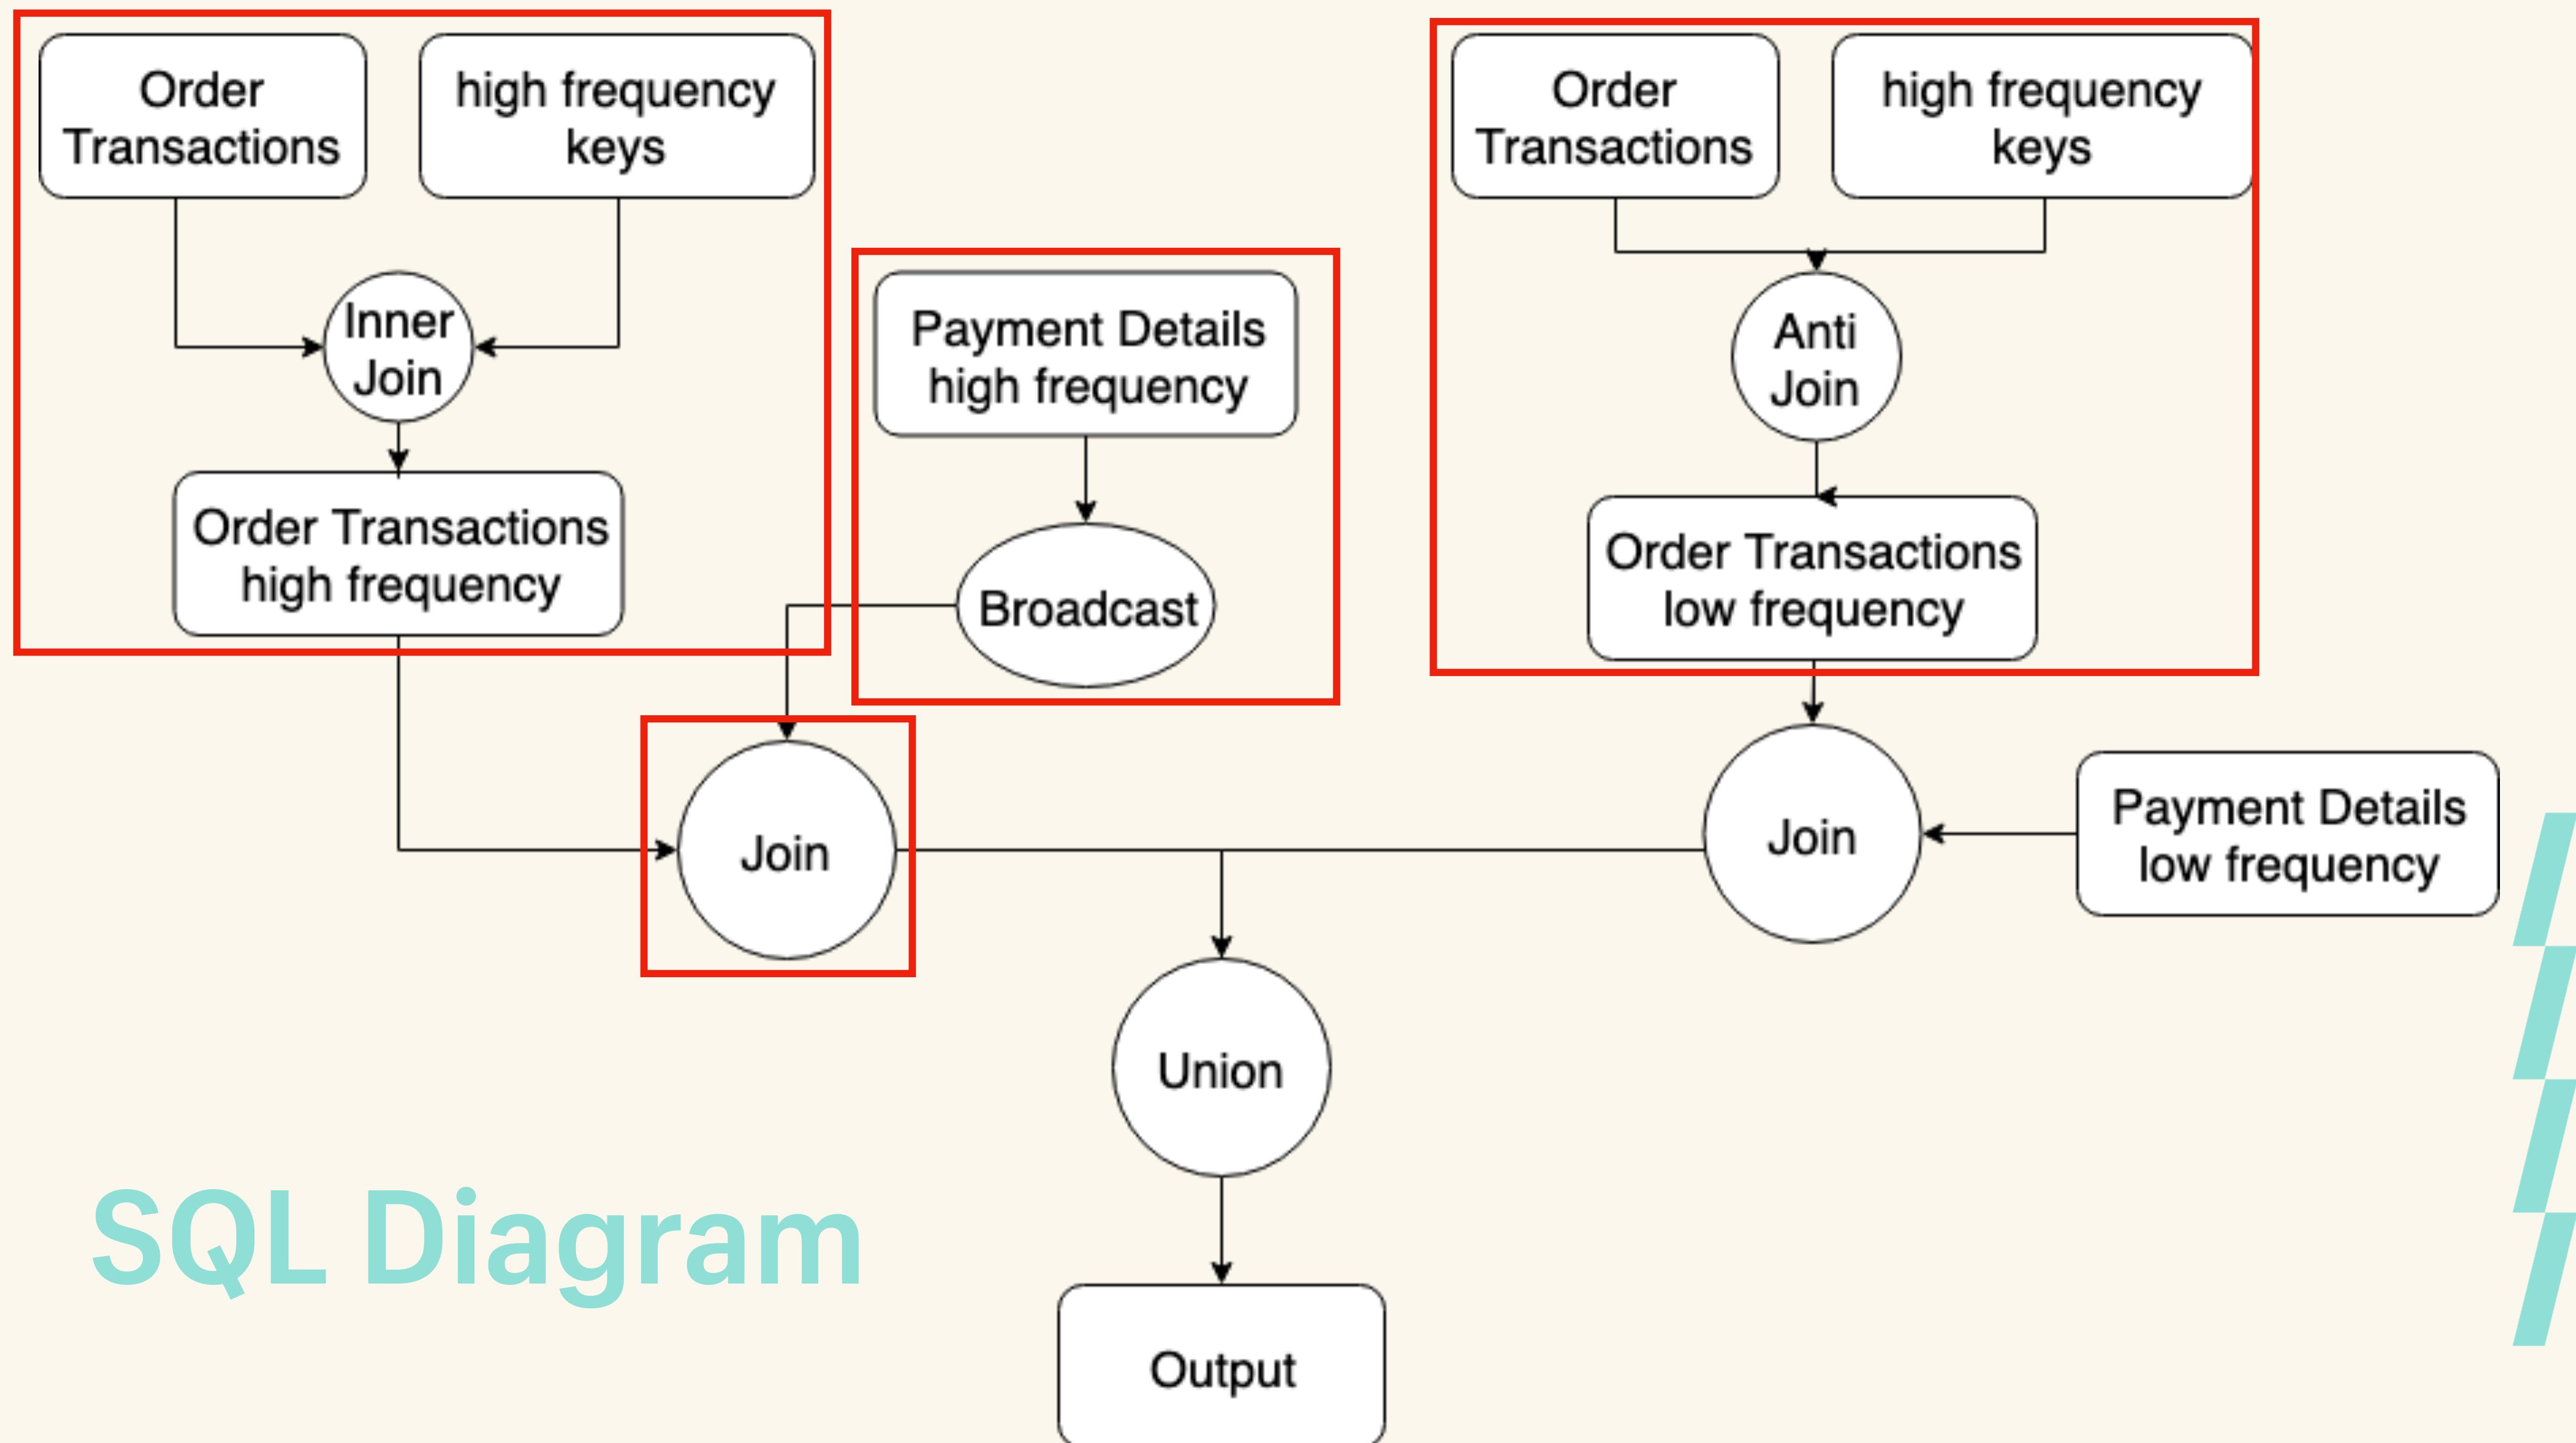
# SQL Diagram



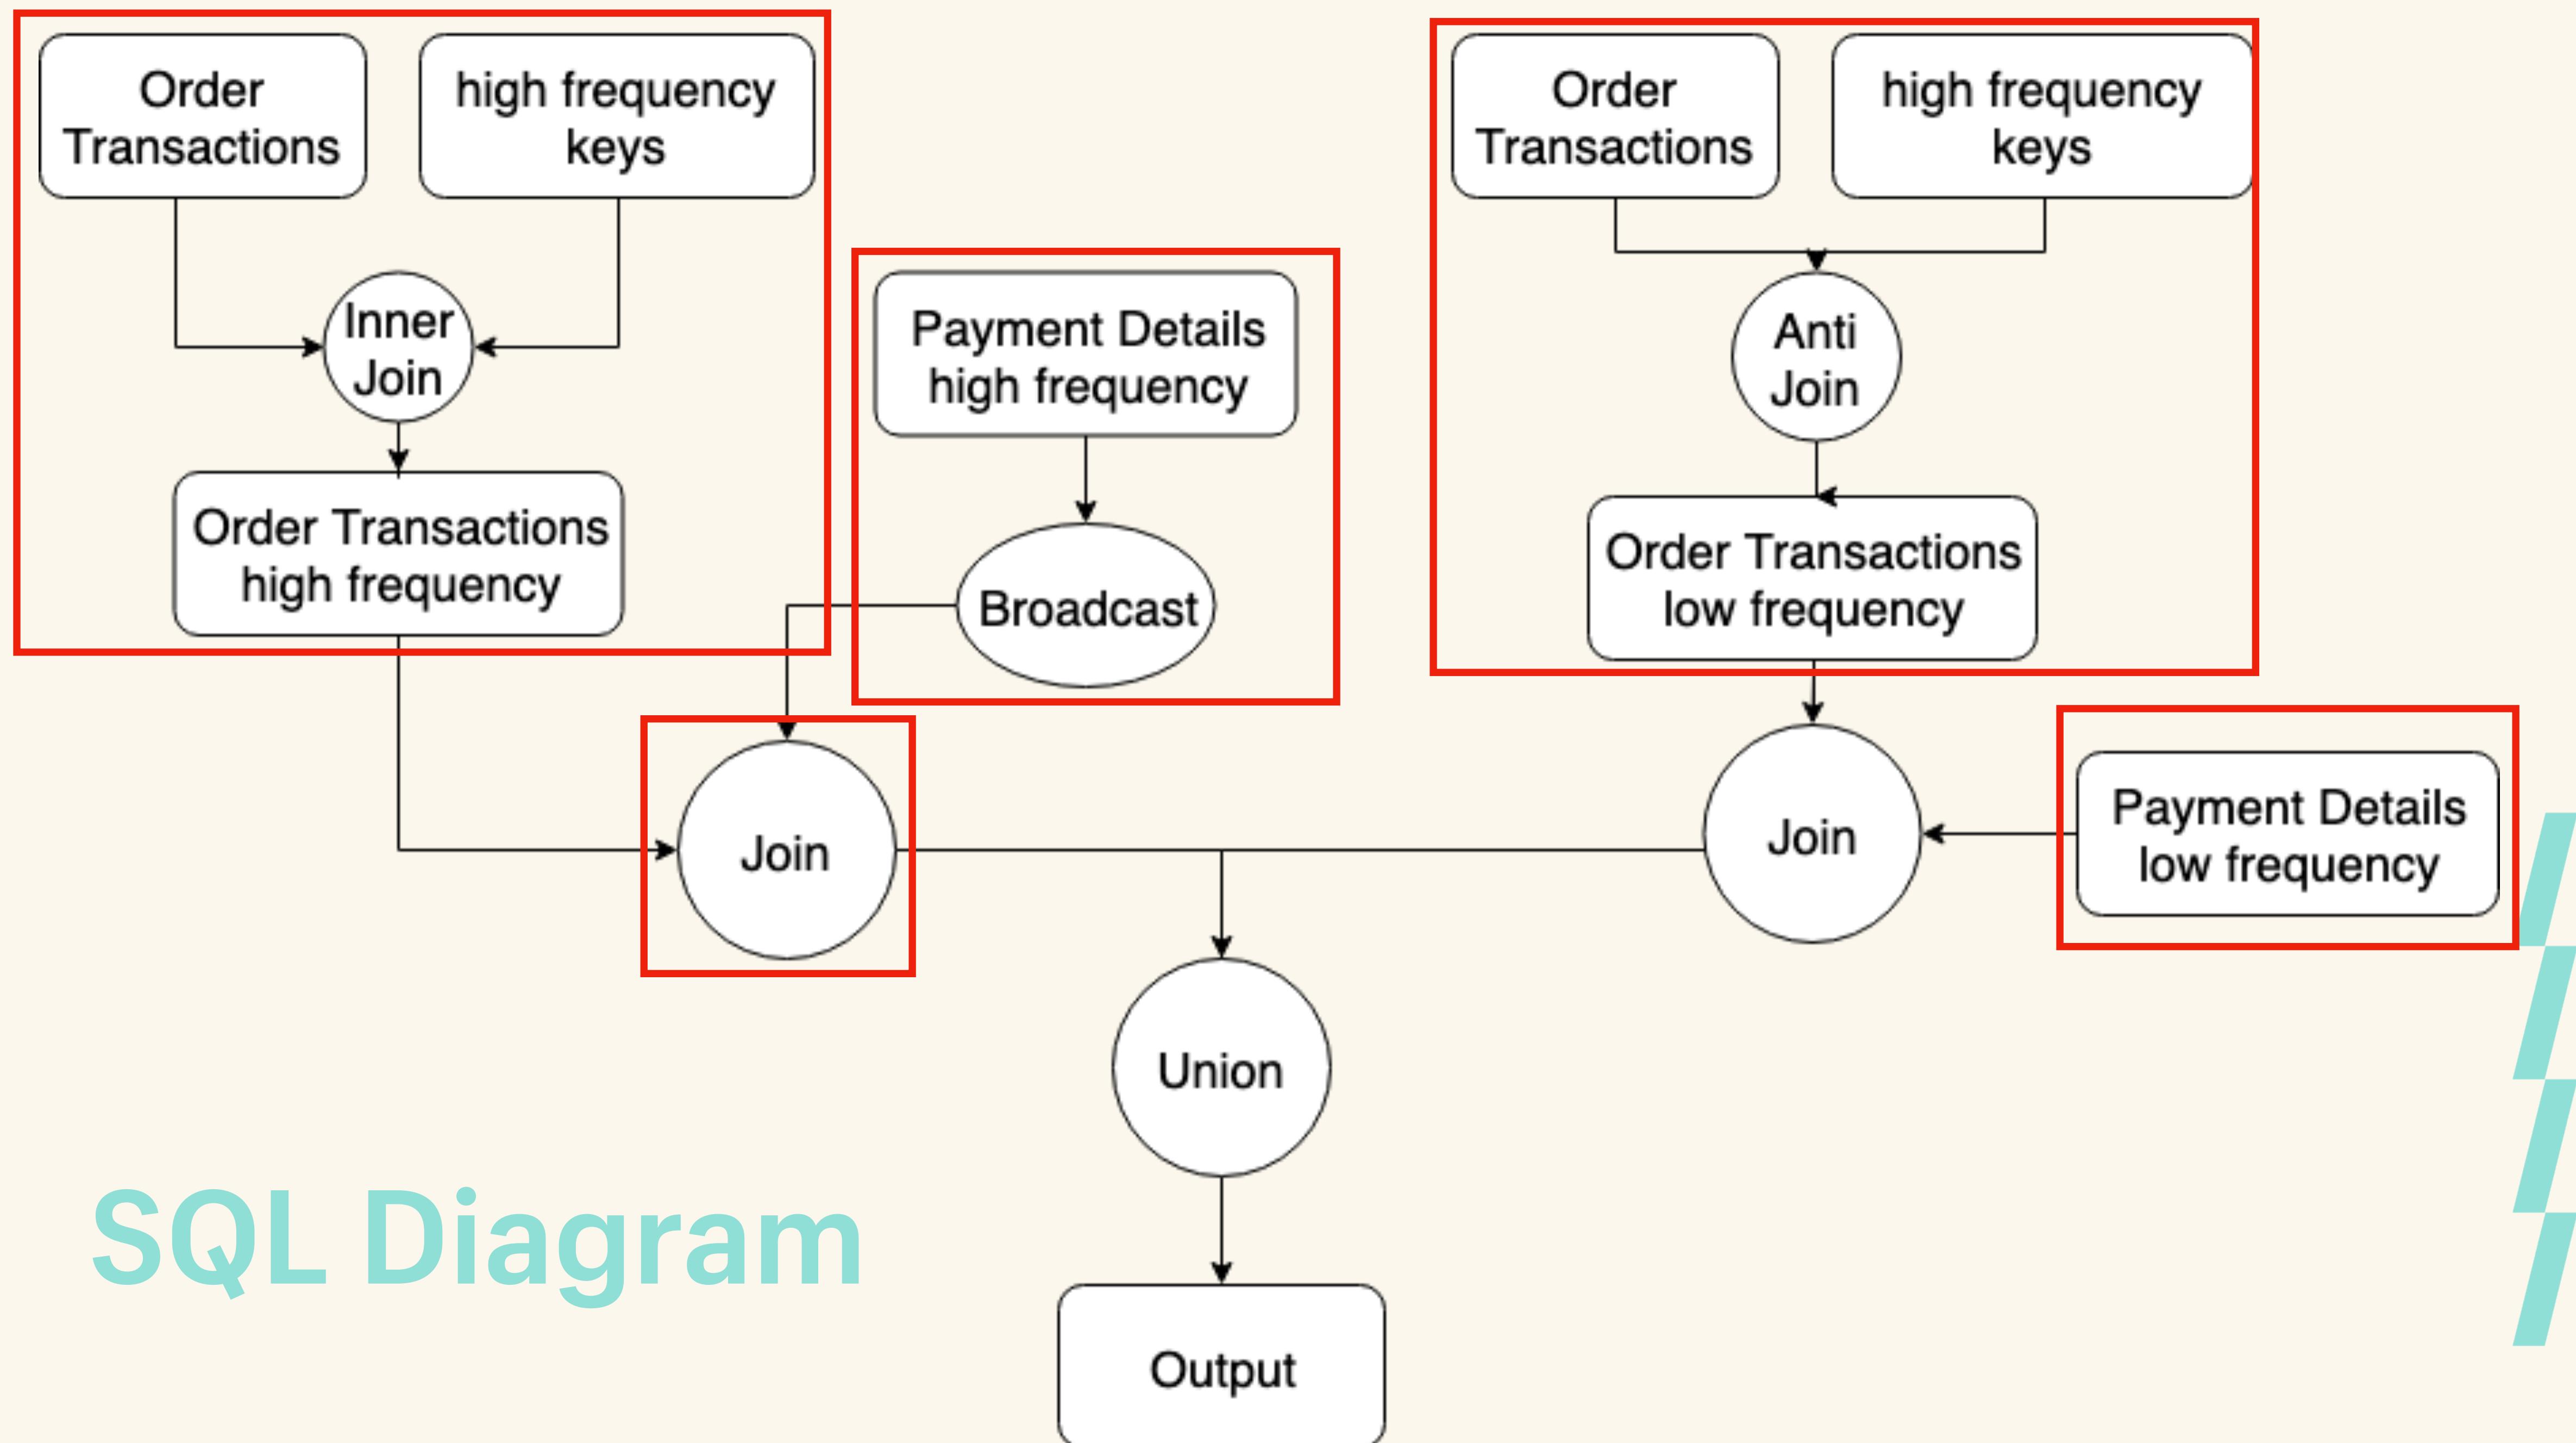
# SQL Diagram



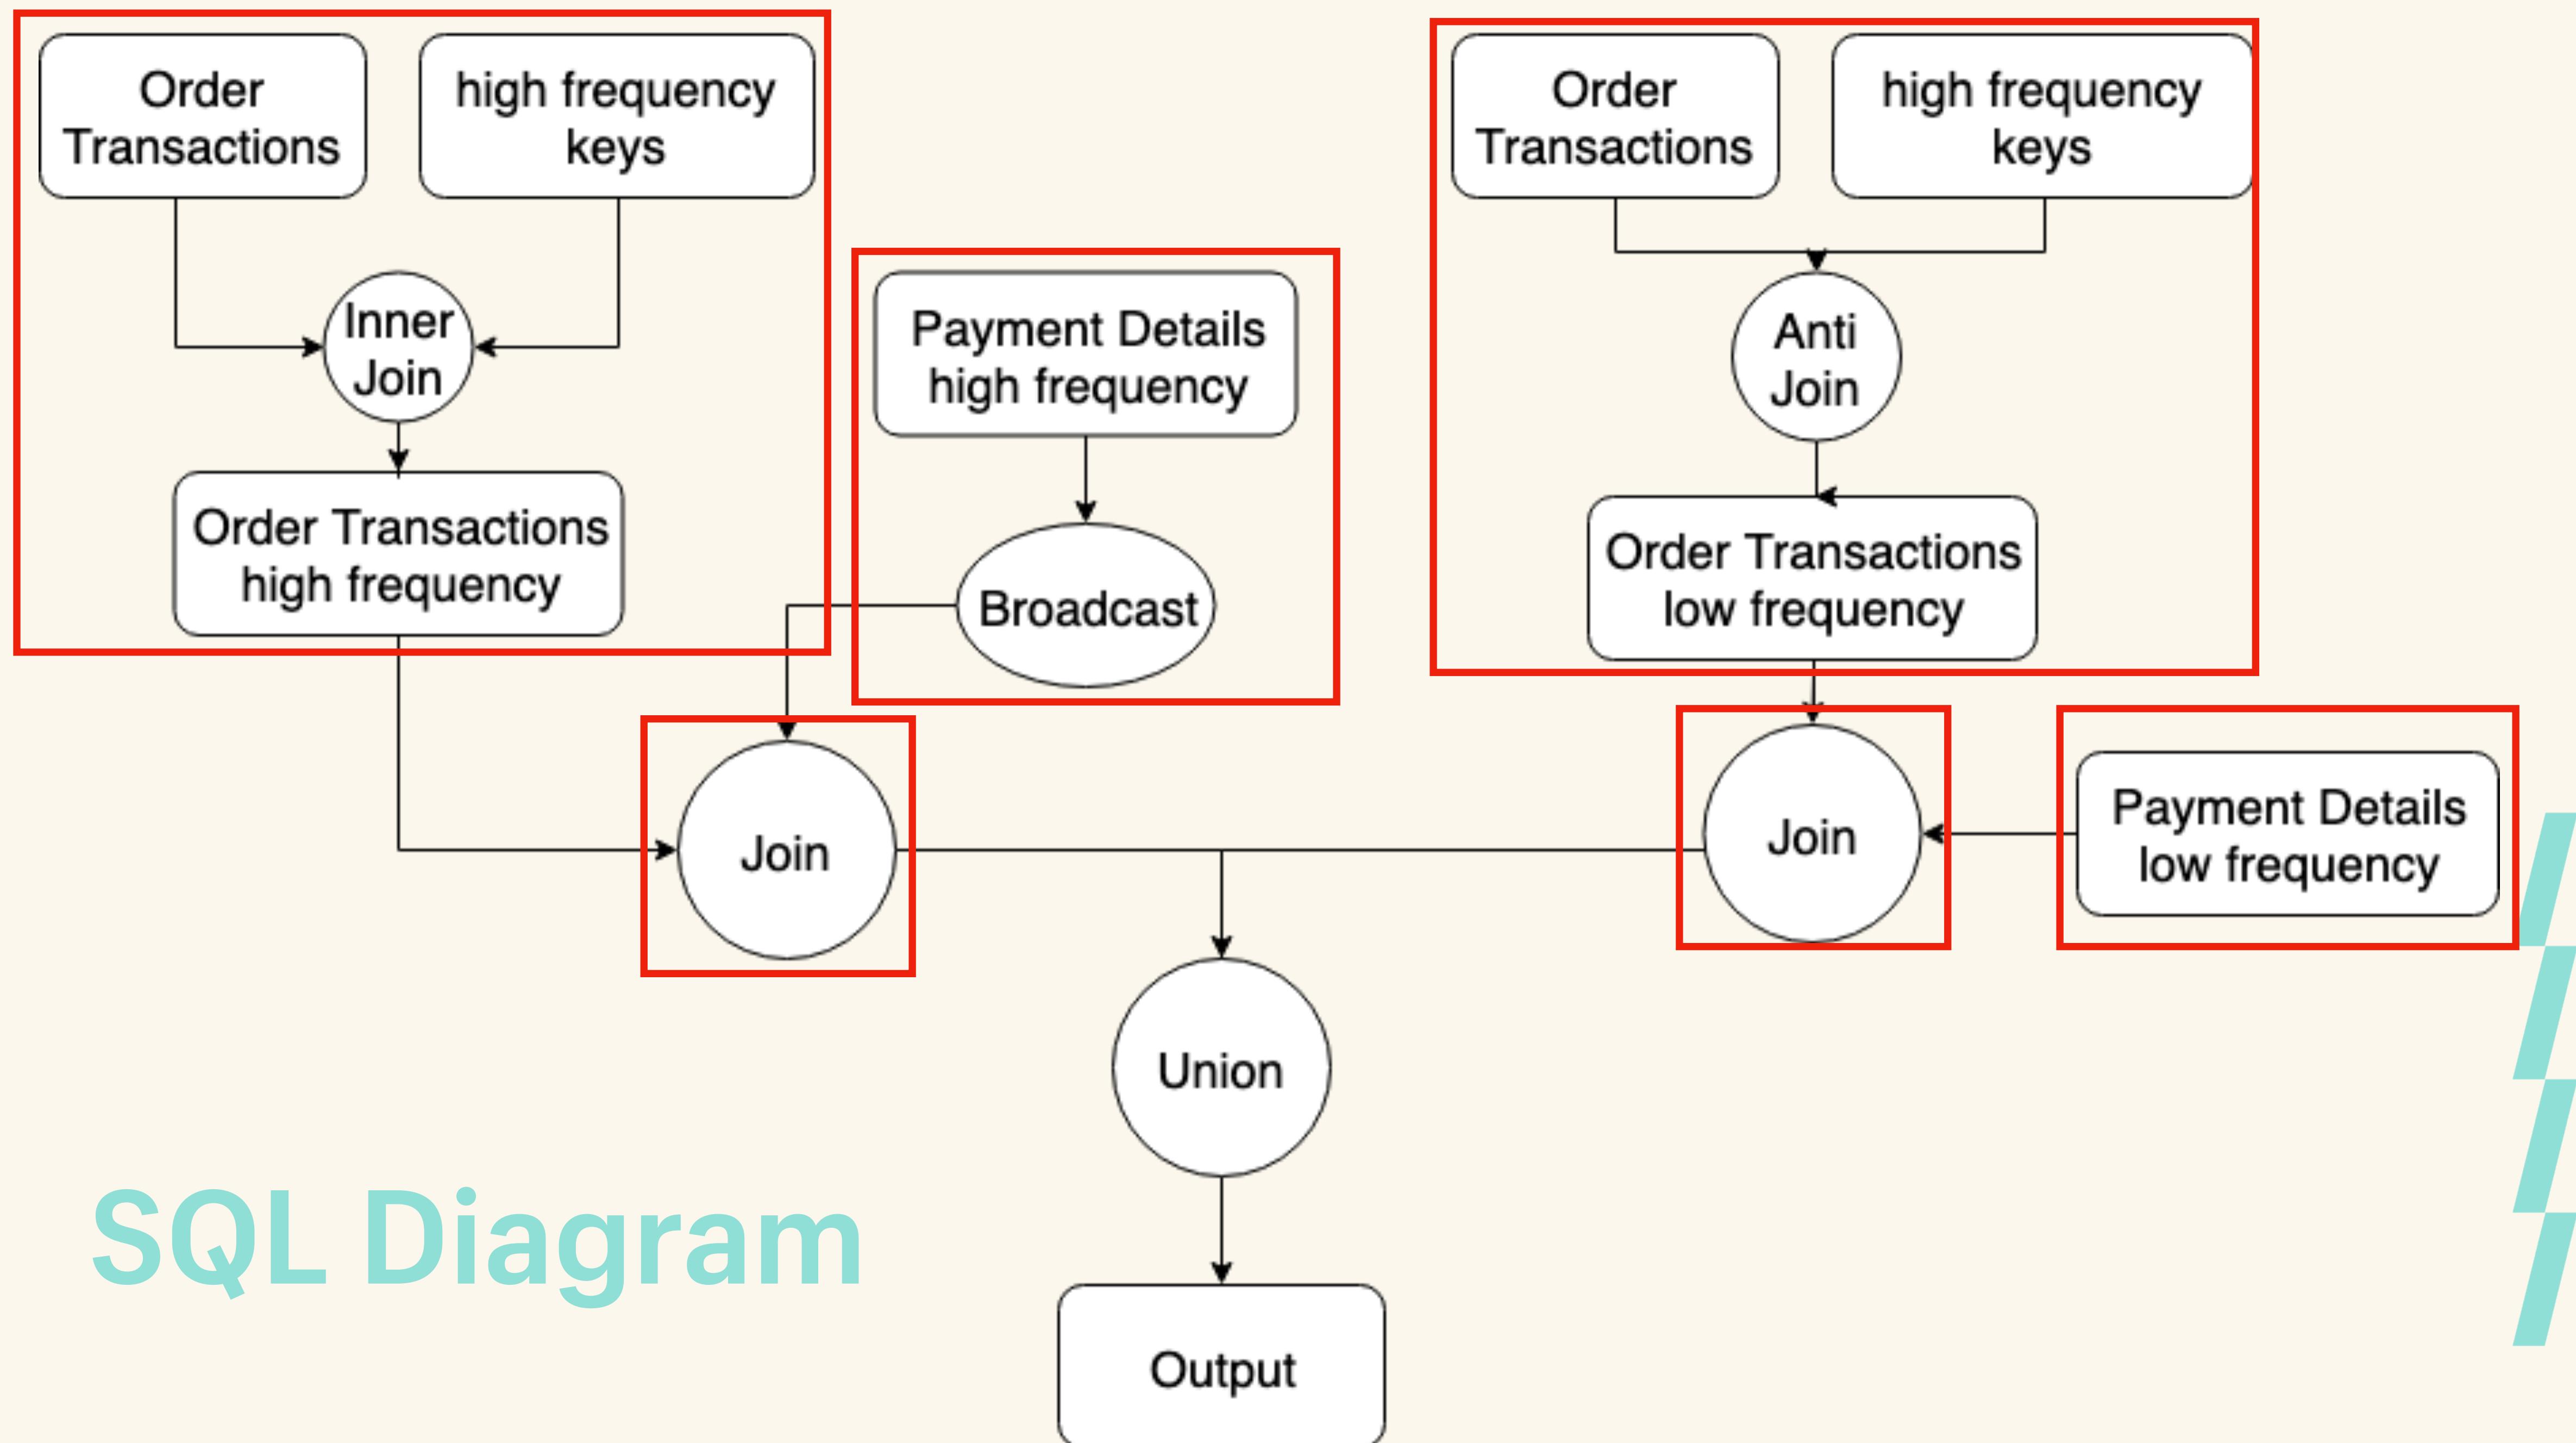
# SQL Diagram

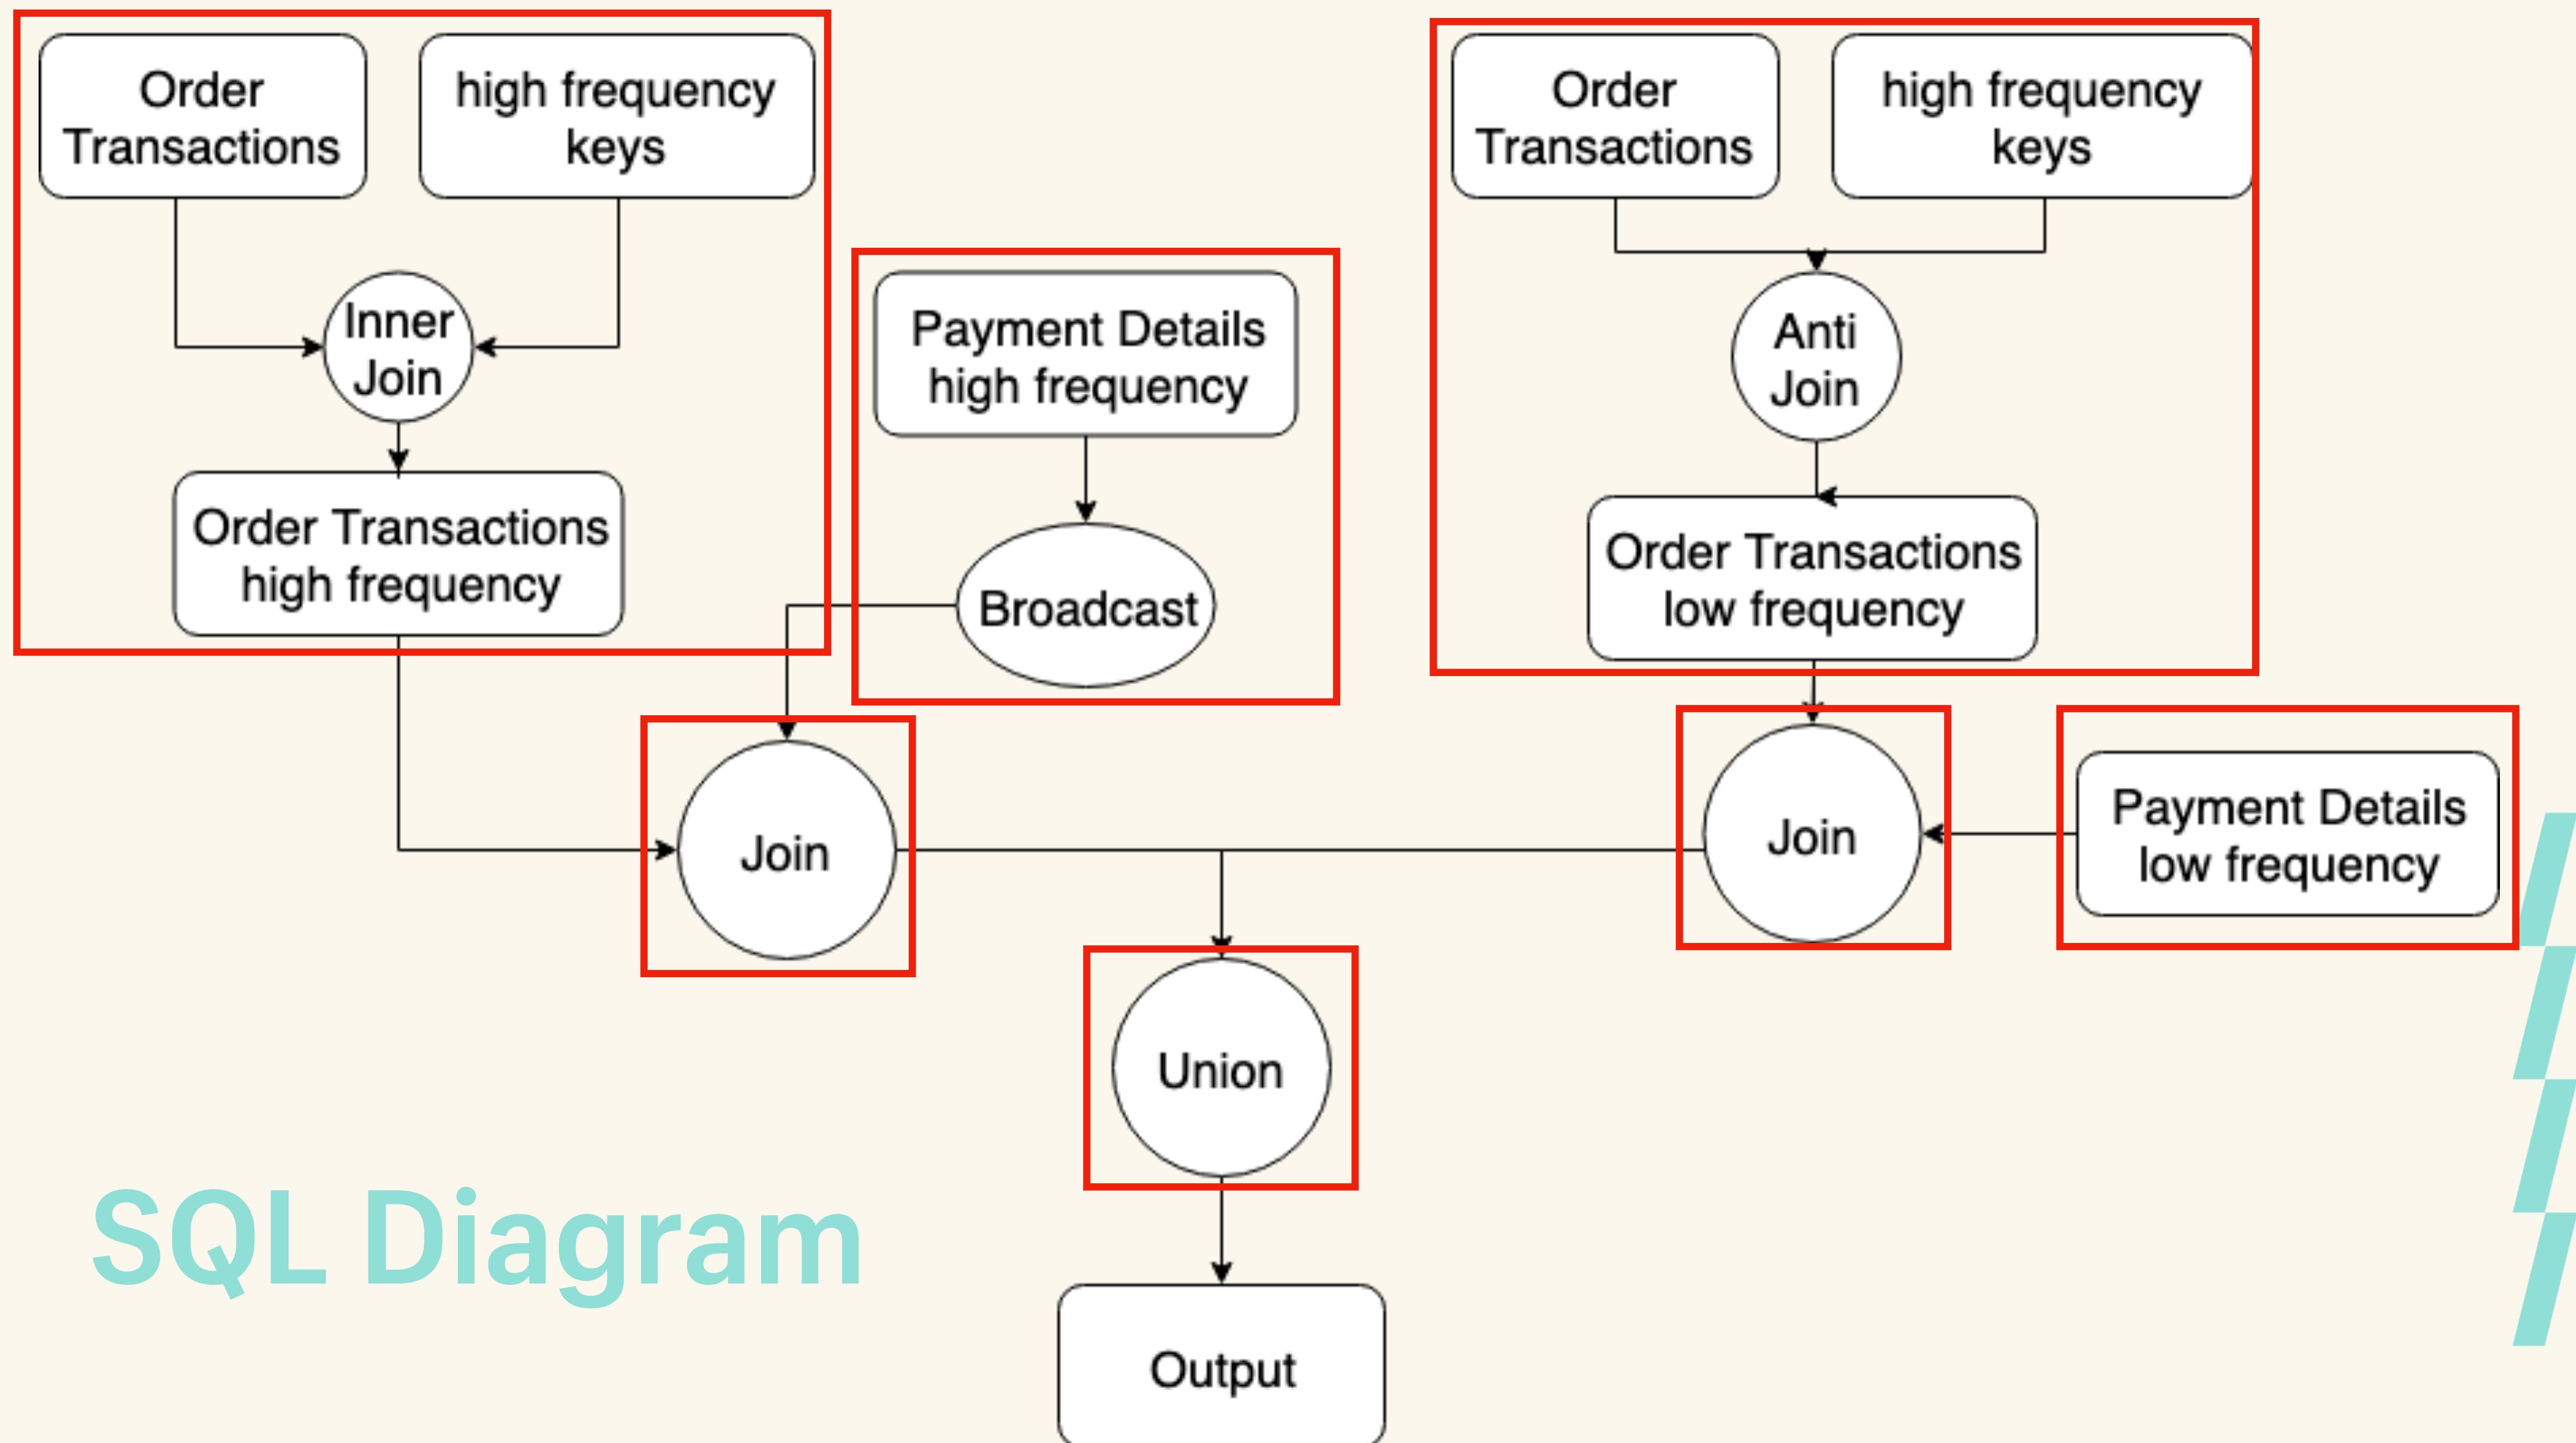


# SQL Diagram



# SQL Diagram





# SQL Diagram

# Downsides

- The complexity of the job increases as this is a more complex join.
- The broadcasting might timeout?



# Key Skew (cont'd)

- The techniques can be used together.
- ie. Largest skew key is null, then a non-null key.
  - Remove the nulls, then salt.



# SQL Tab Recap

1. Interpreted the overall spark DAG.
2. Found the spark plans from the details dropdown.
3. Mapped parts of the plan to the exact code line.
4. Analyzed aggregate level stats.

## Red Flags:

1. Long and wide DAGs.
2. Big variation in min/med/max times.

\*\* Evaluate plan in the order that they are listed.



# Join Optimization Recap

- Spot a join block that have widely varying task times.
- Run dev skew to analyze the distribution of keys.
- Pick the correct join optimization.
- Implement.



# **Section 4: Optimization 2 - Eliminating Data Spill**

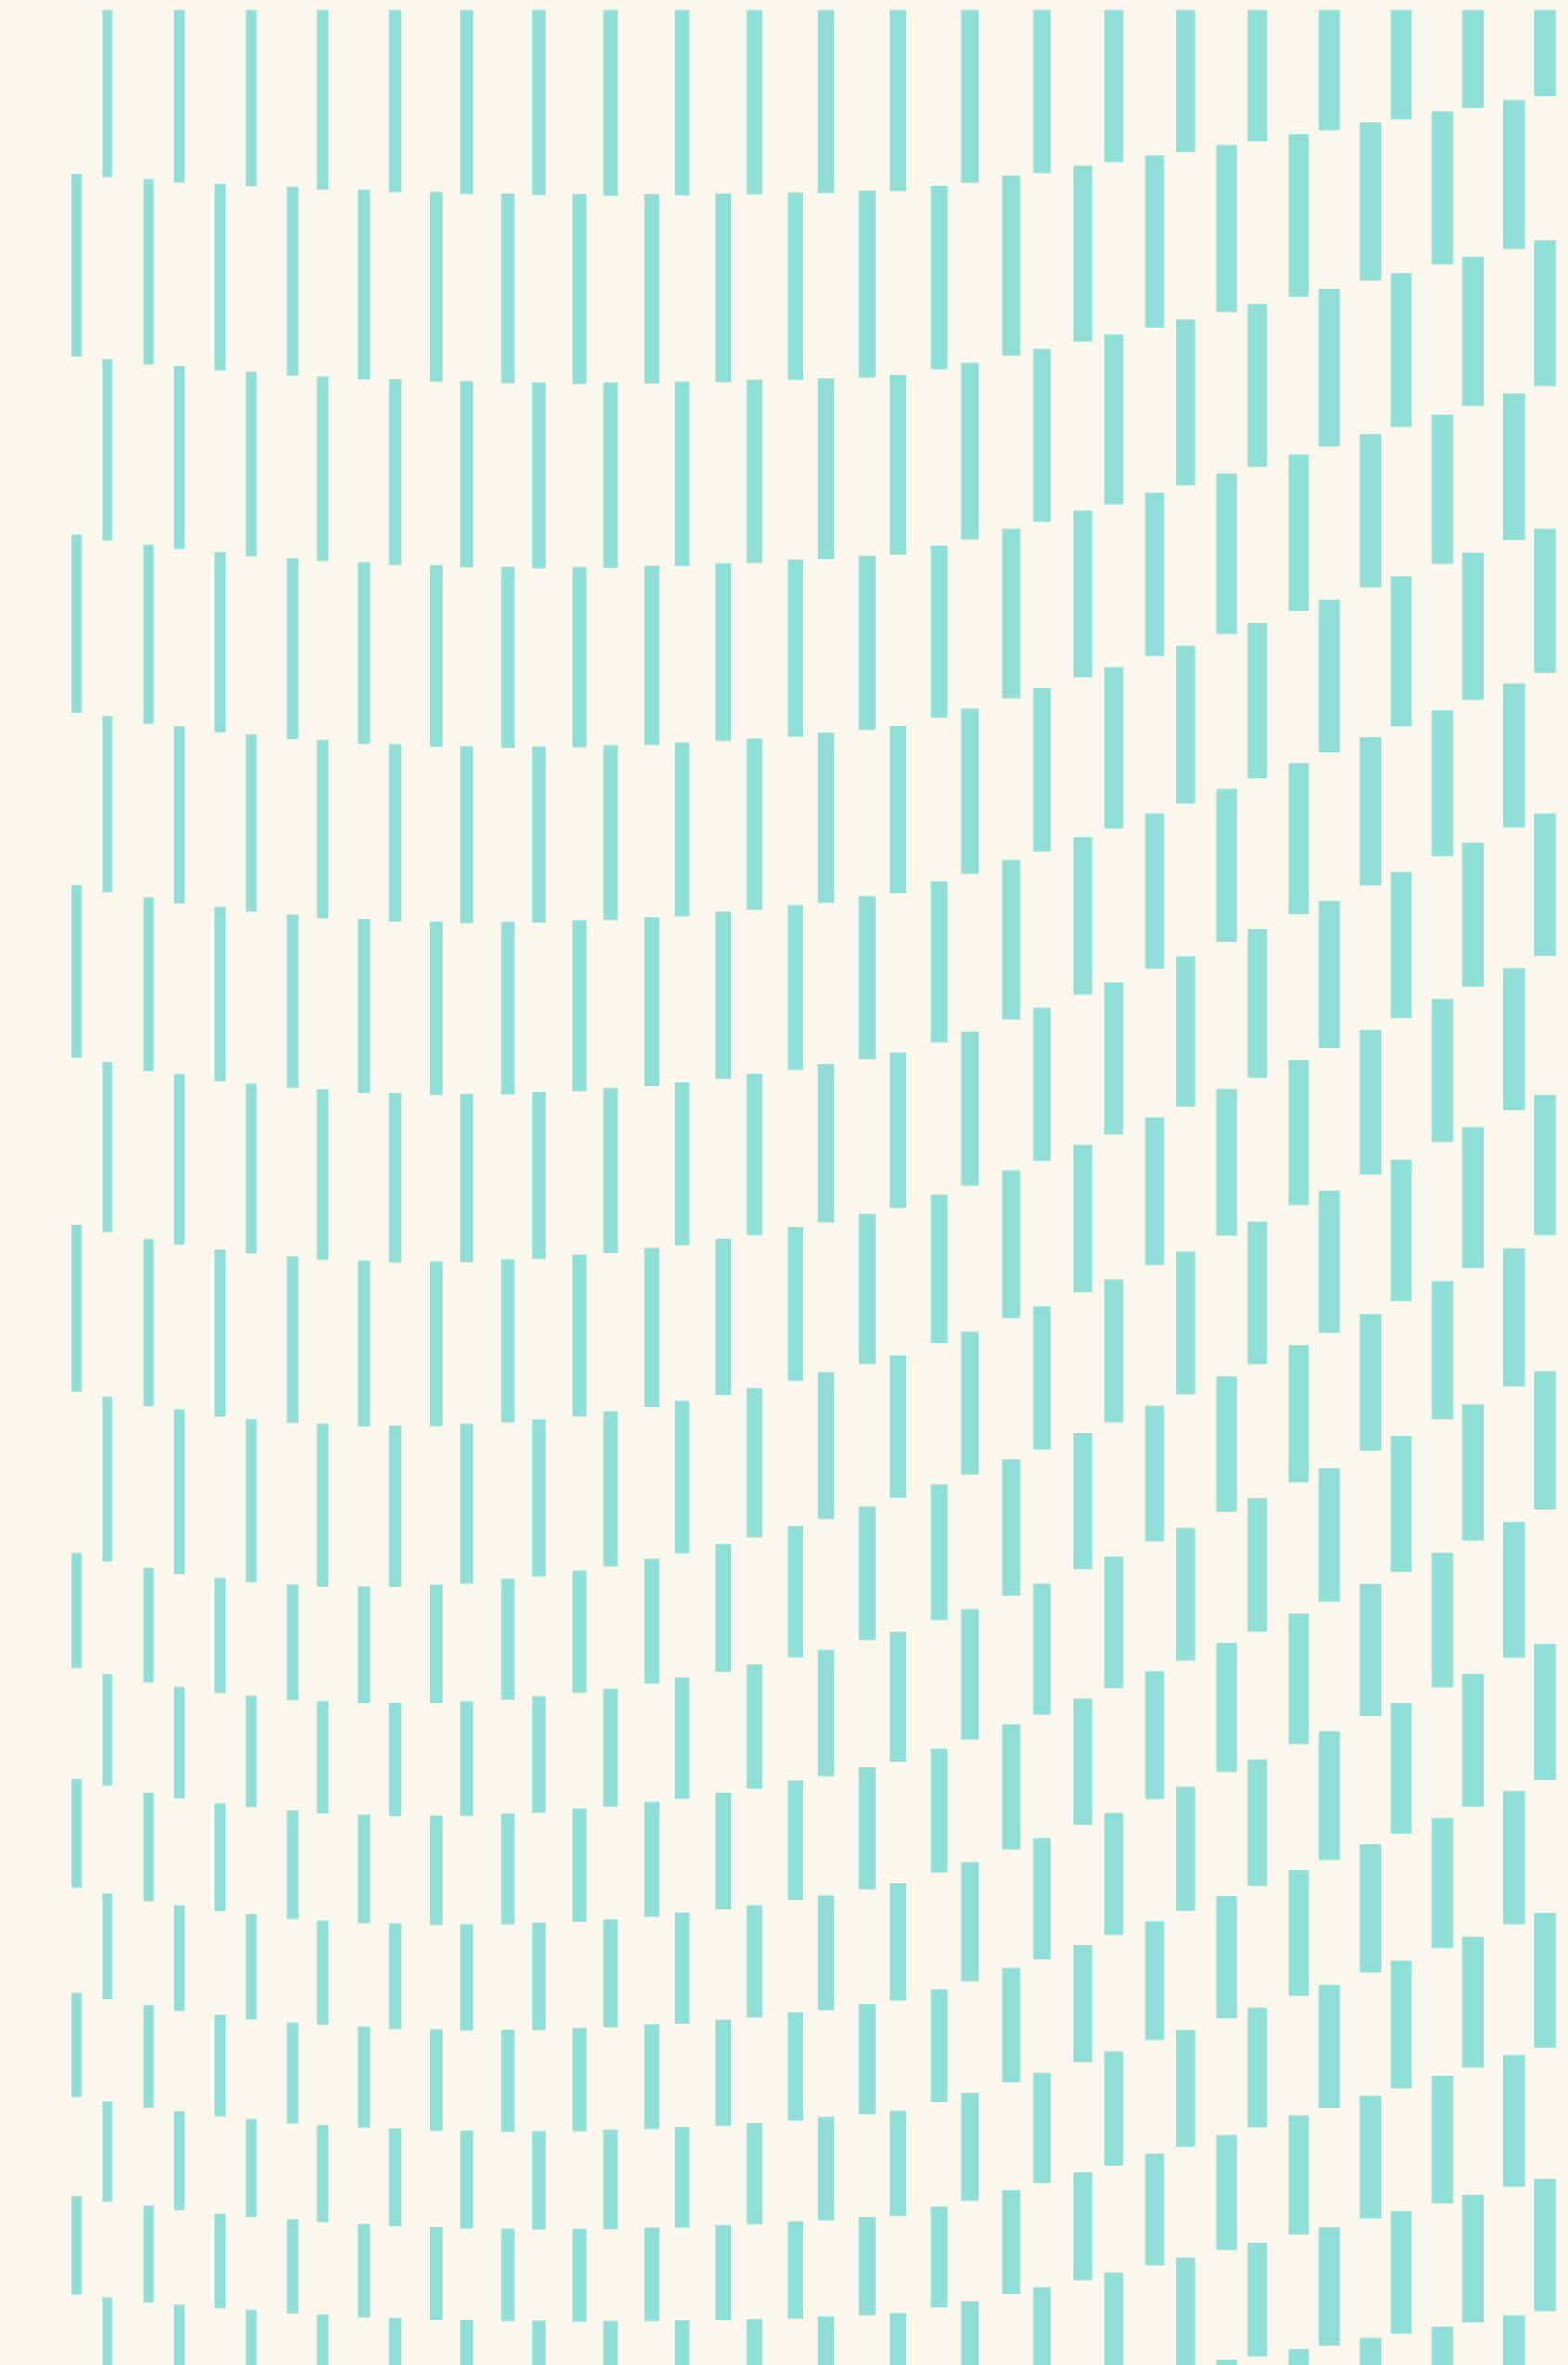
**Jobs, Stages, Tasks**

**Executors**

**Spill**

**Partitions / Tasks**

**Optimal Partition Calculations**



# Disclaimer

These screenshot are taken of the stages when we had the join skew.

They illustrate more things that you want to be aware of in your stage.

But not needed for the second optimization.



# Jobs, Stages, Tasks, Executors

- Every Spark “action” triggers a “job”.
- Every job contains multiple “stages”.
- New stages are created when a shuffle operation is required.
- A stage is a collection of transformations.
- A stage divides the work into a number of “tasks”.
- Tasks run in parallel on “executors”.

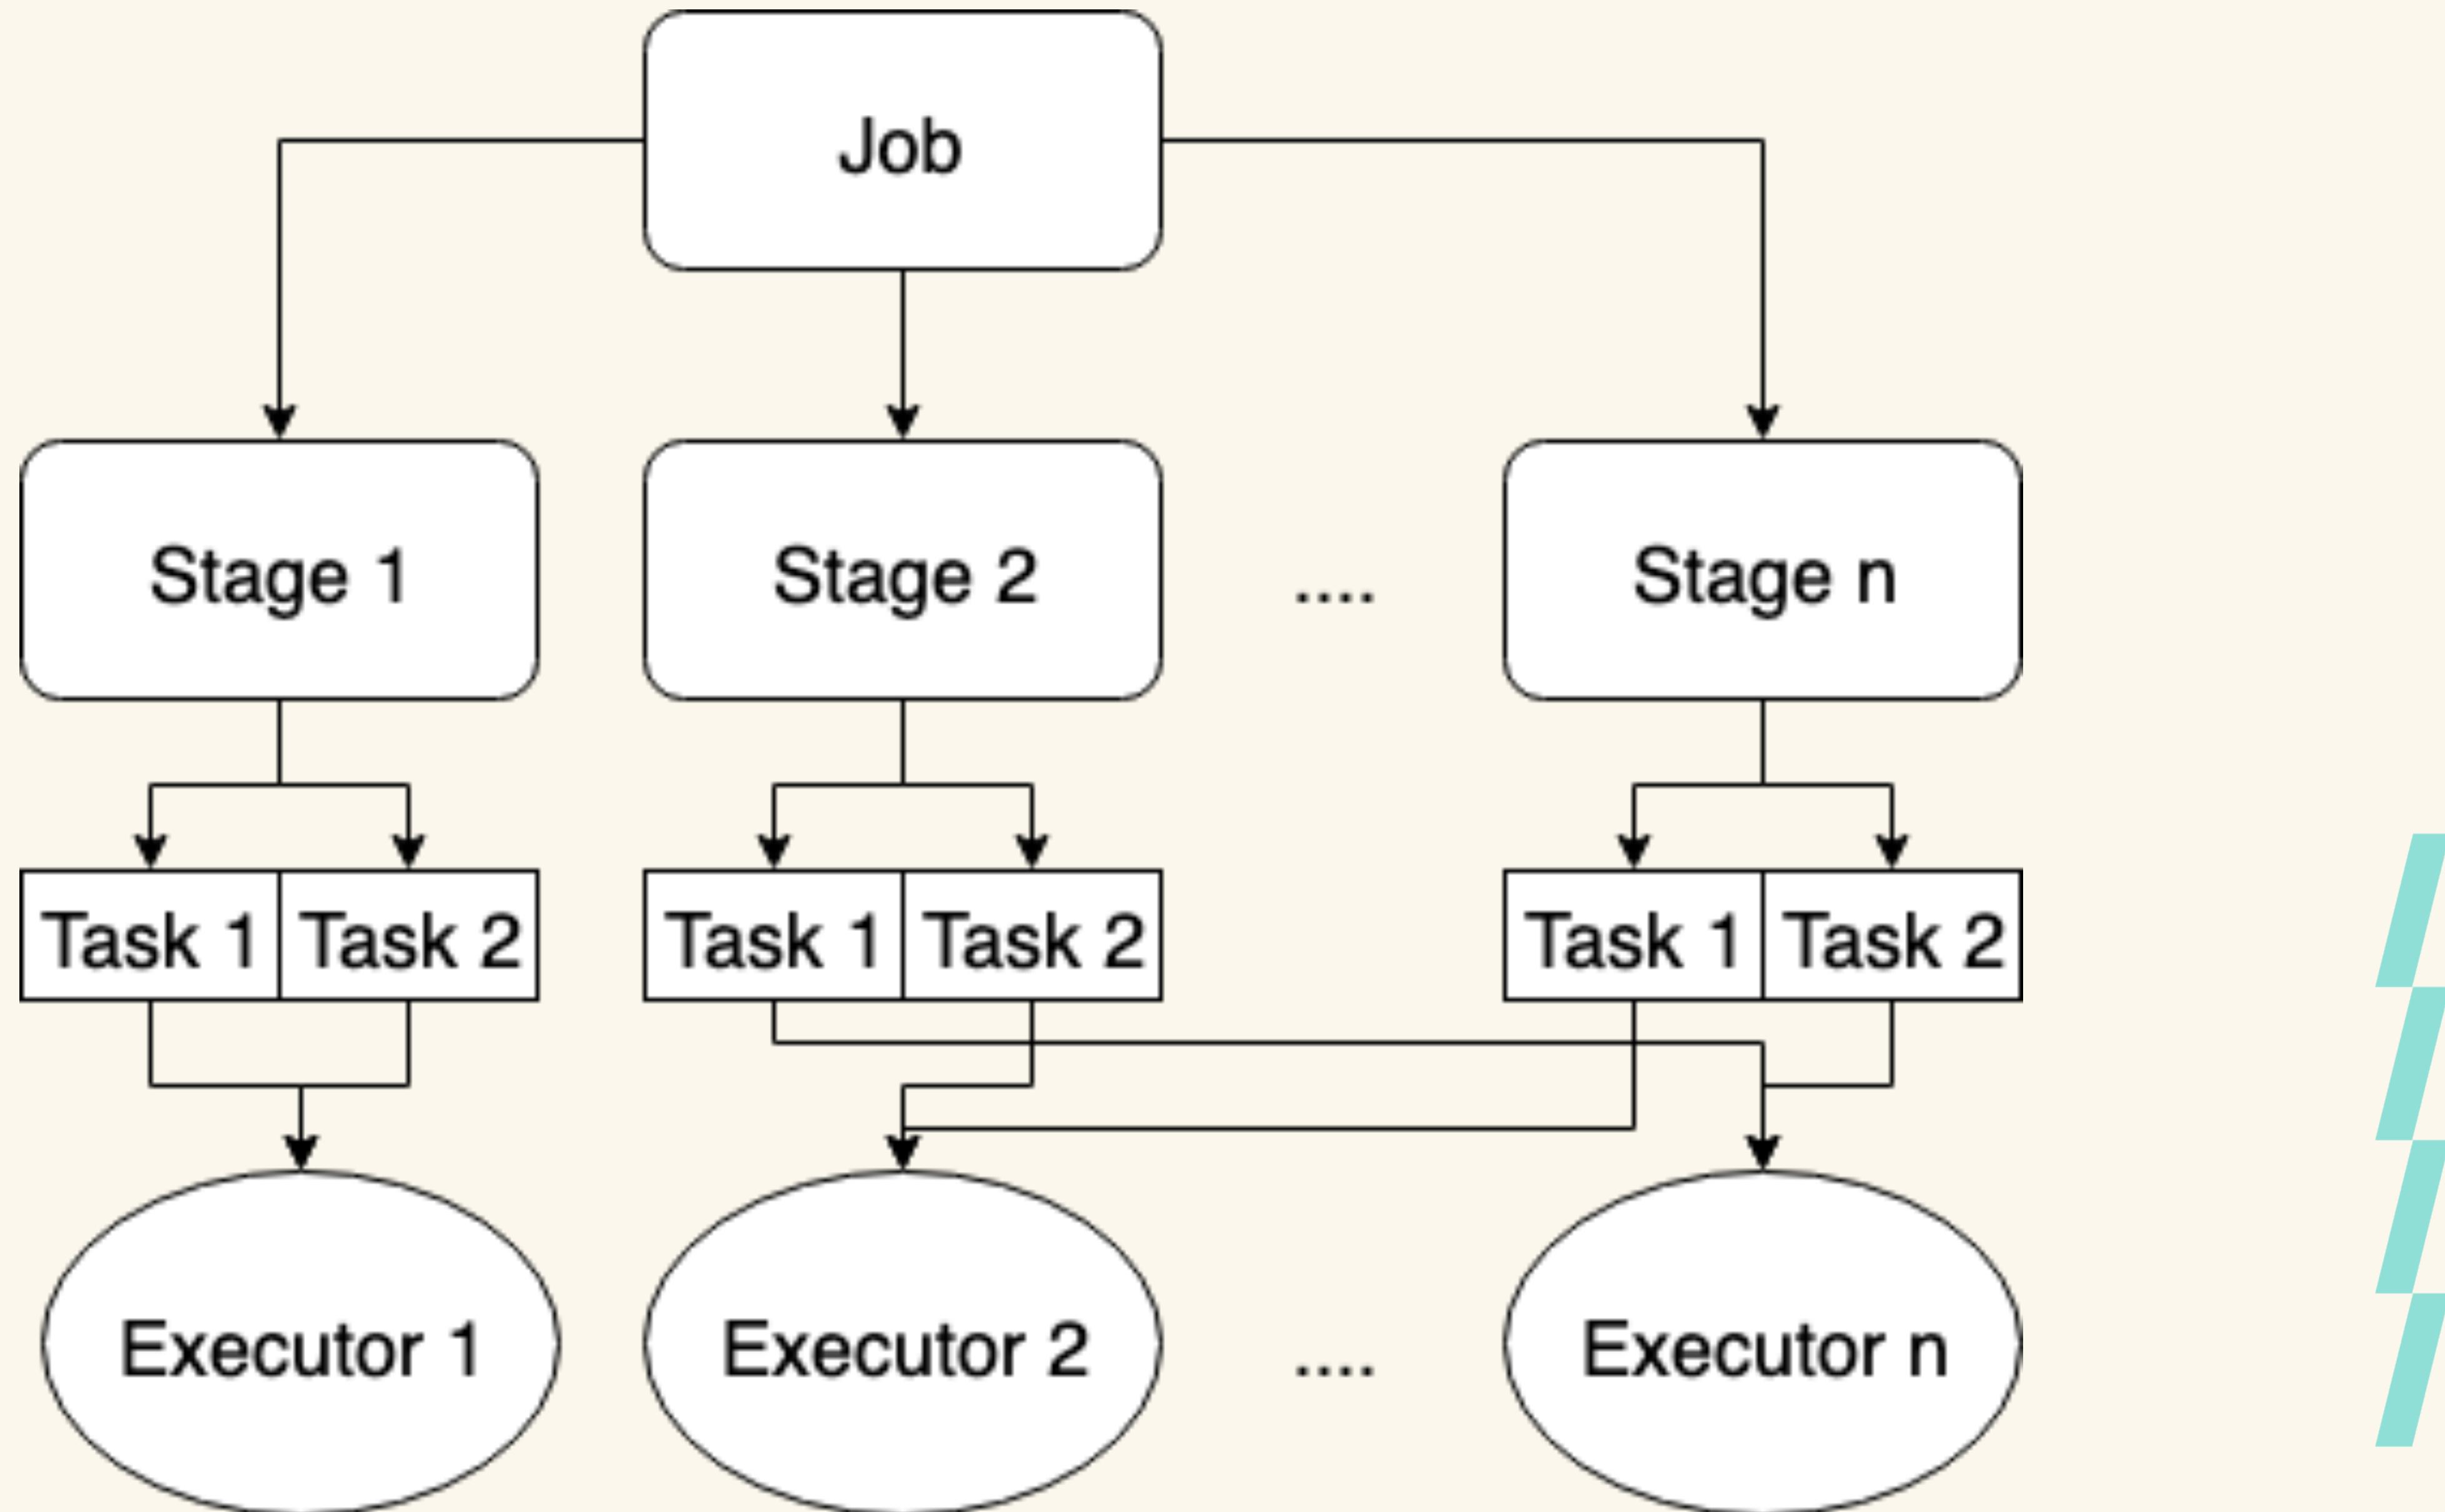


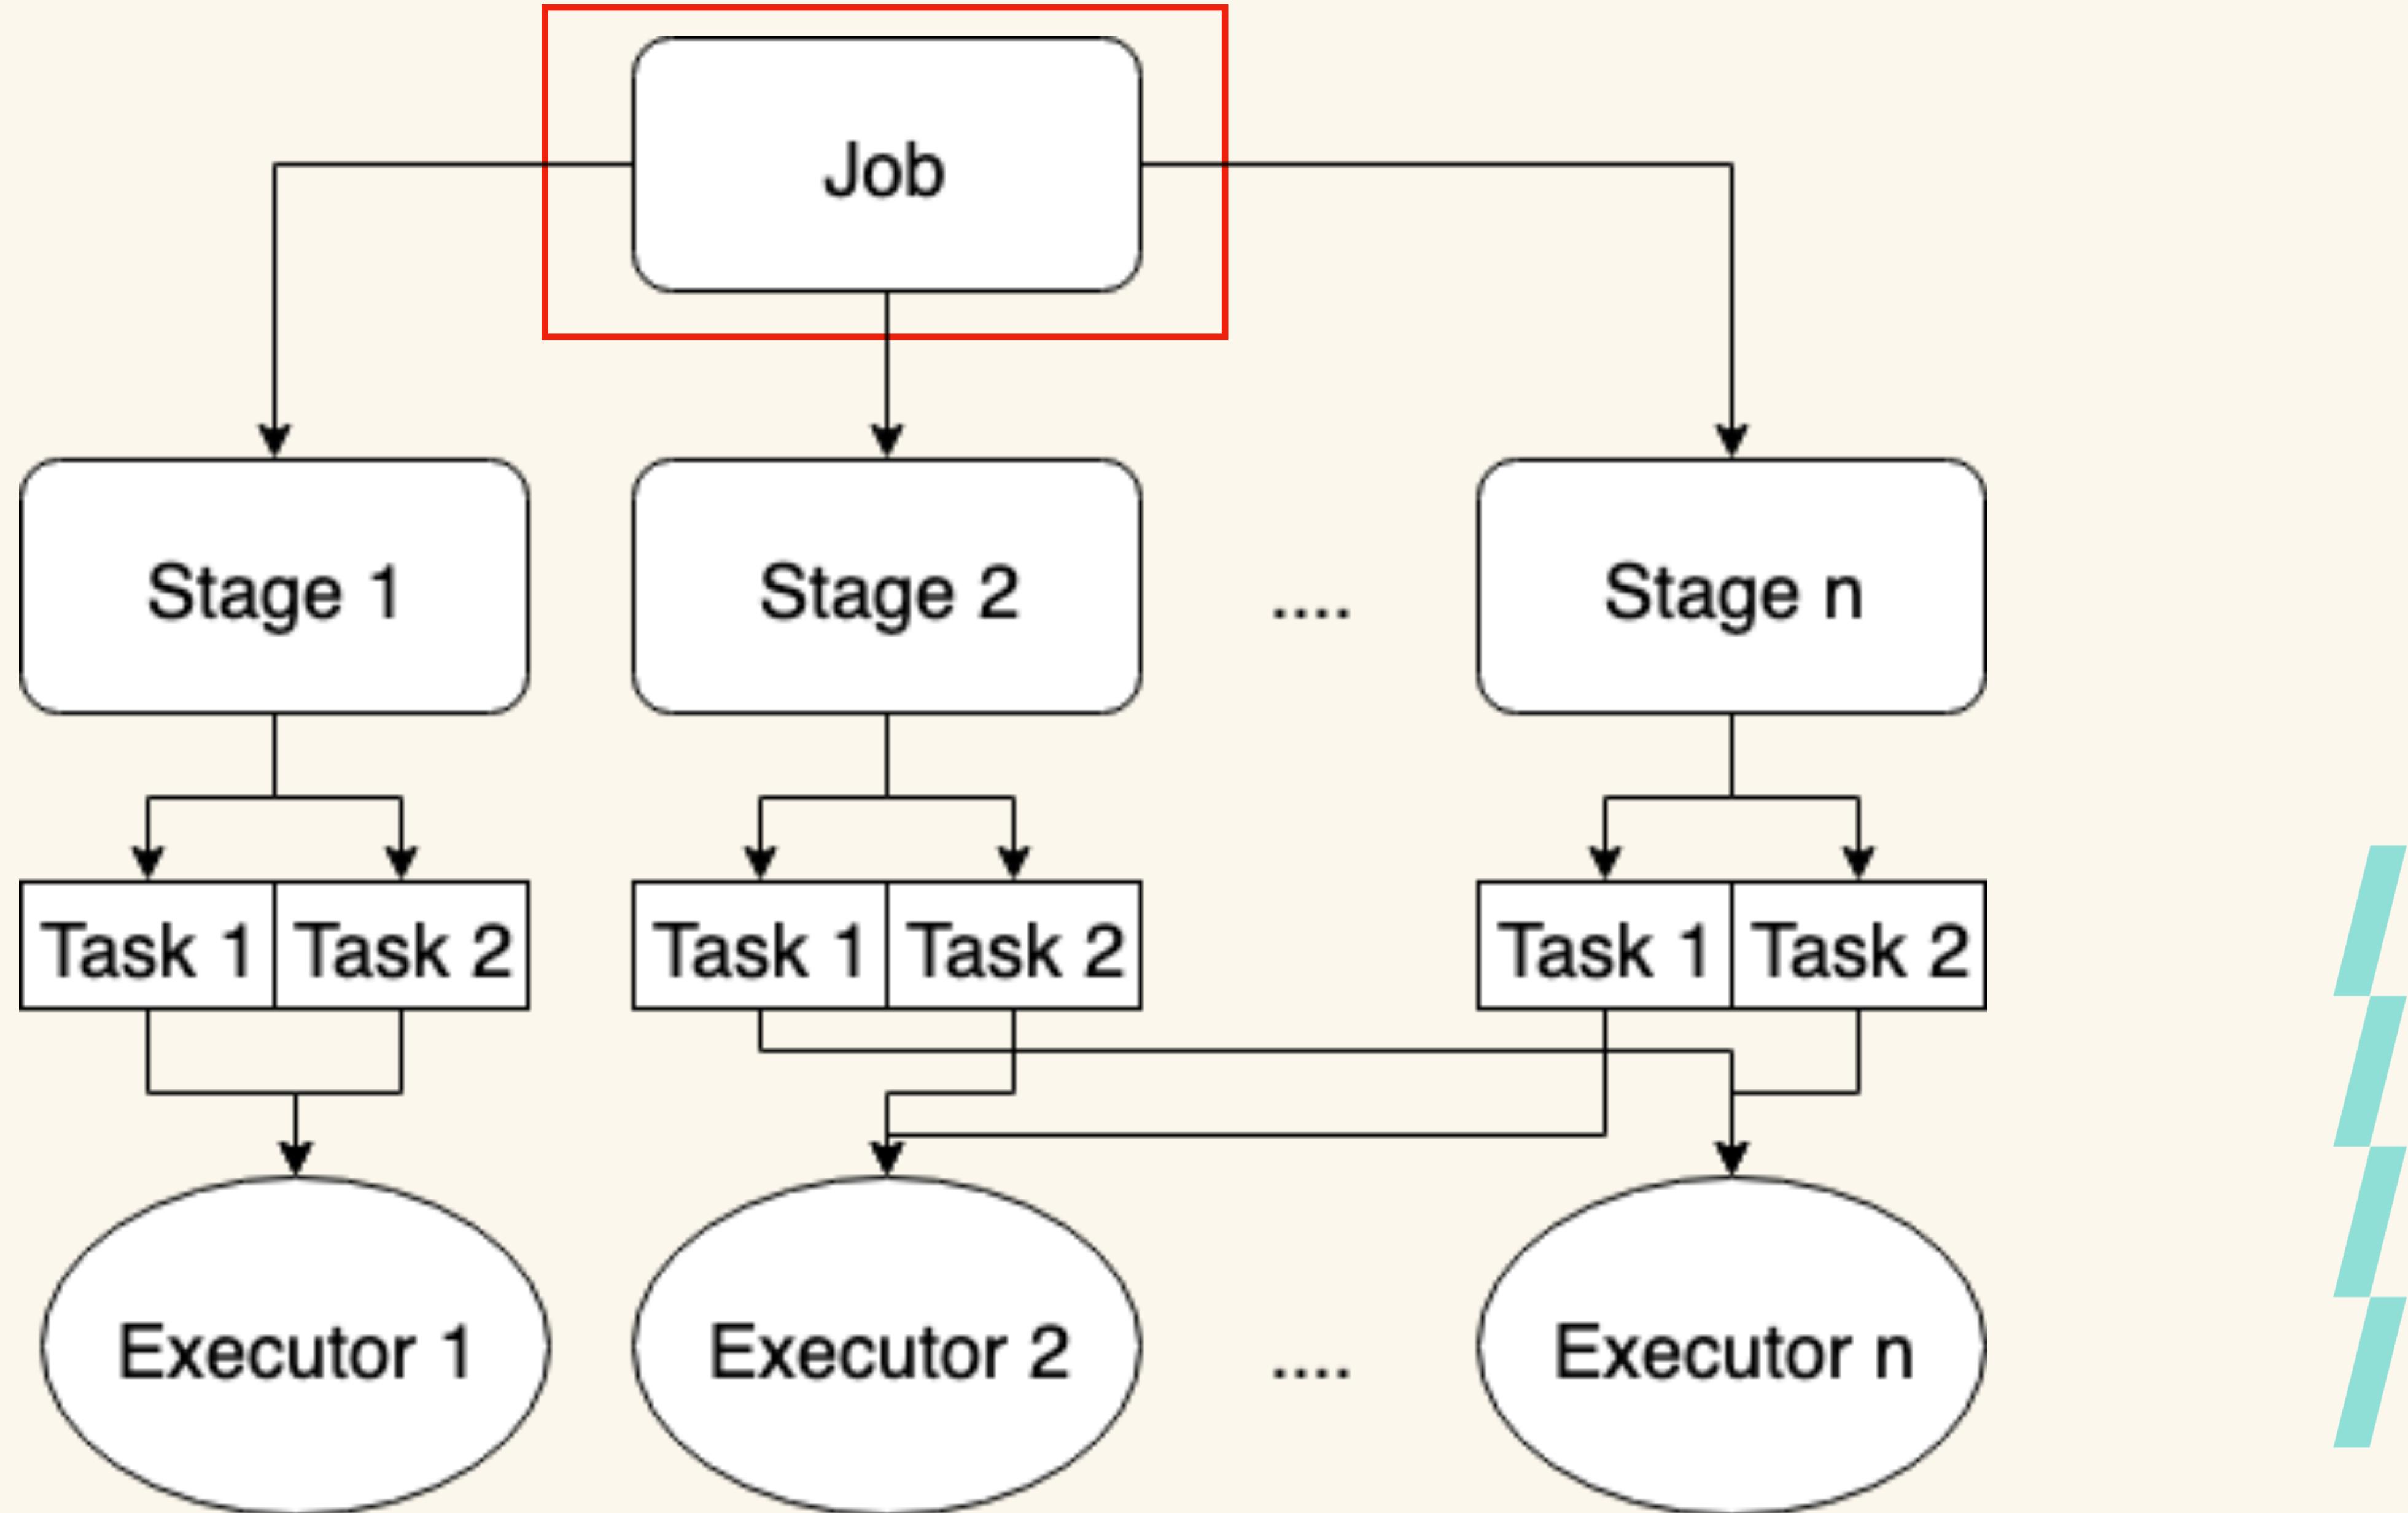
# Executors

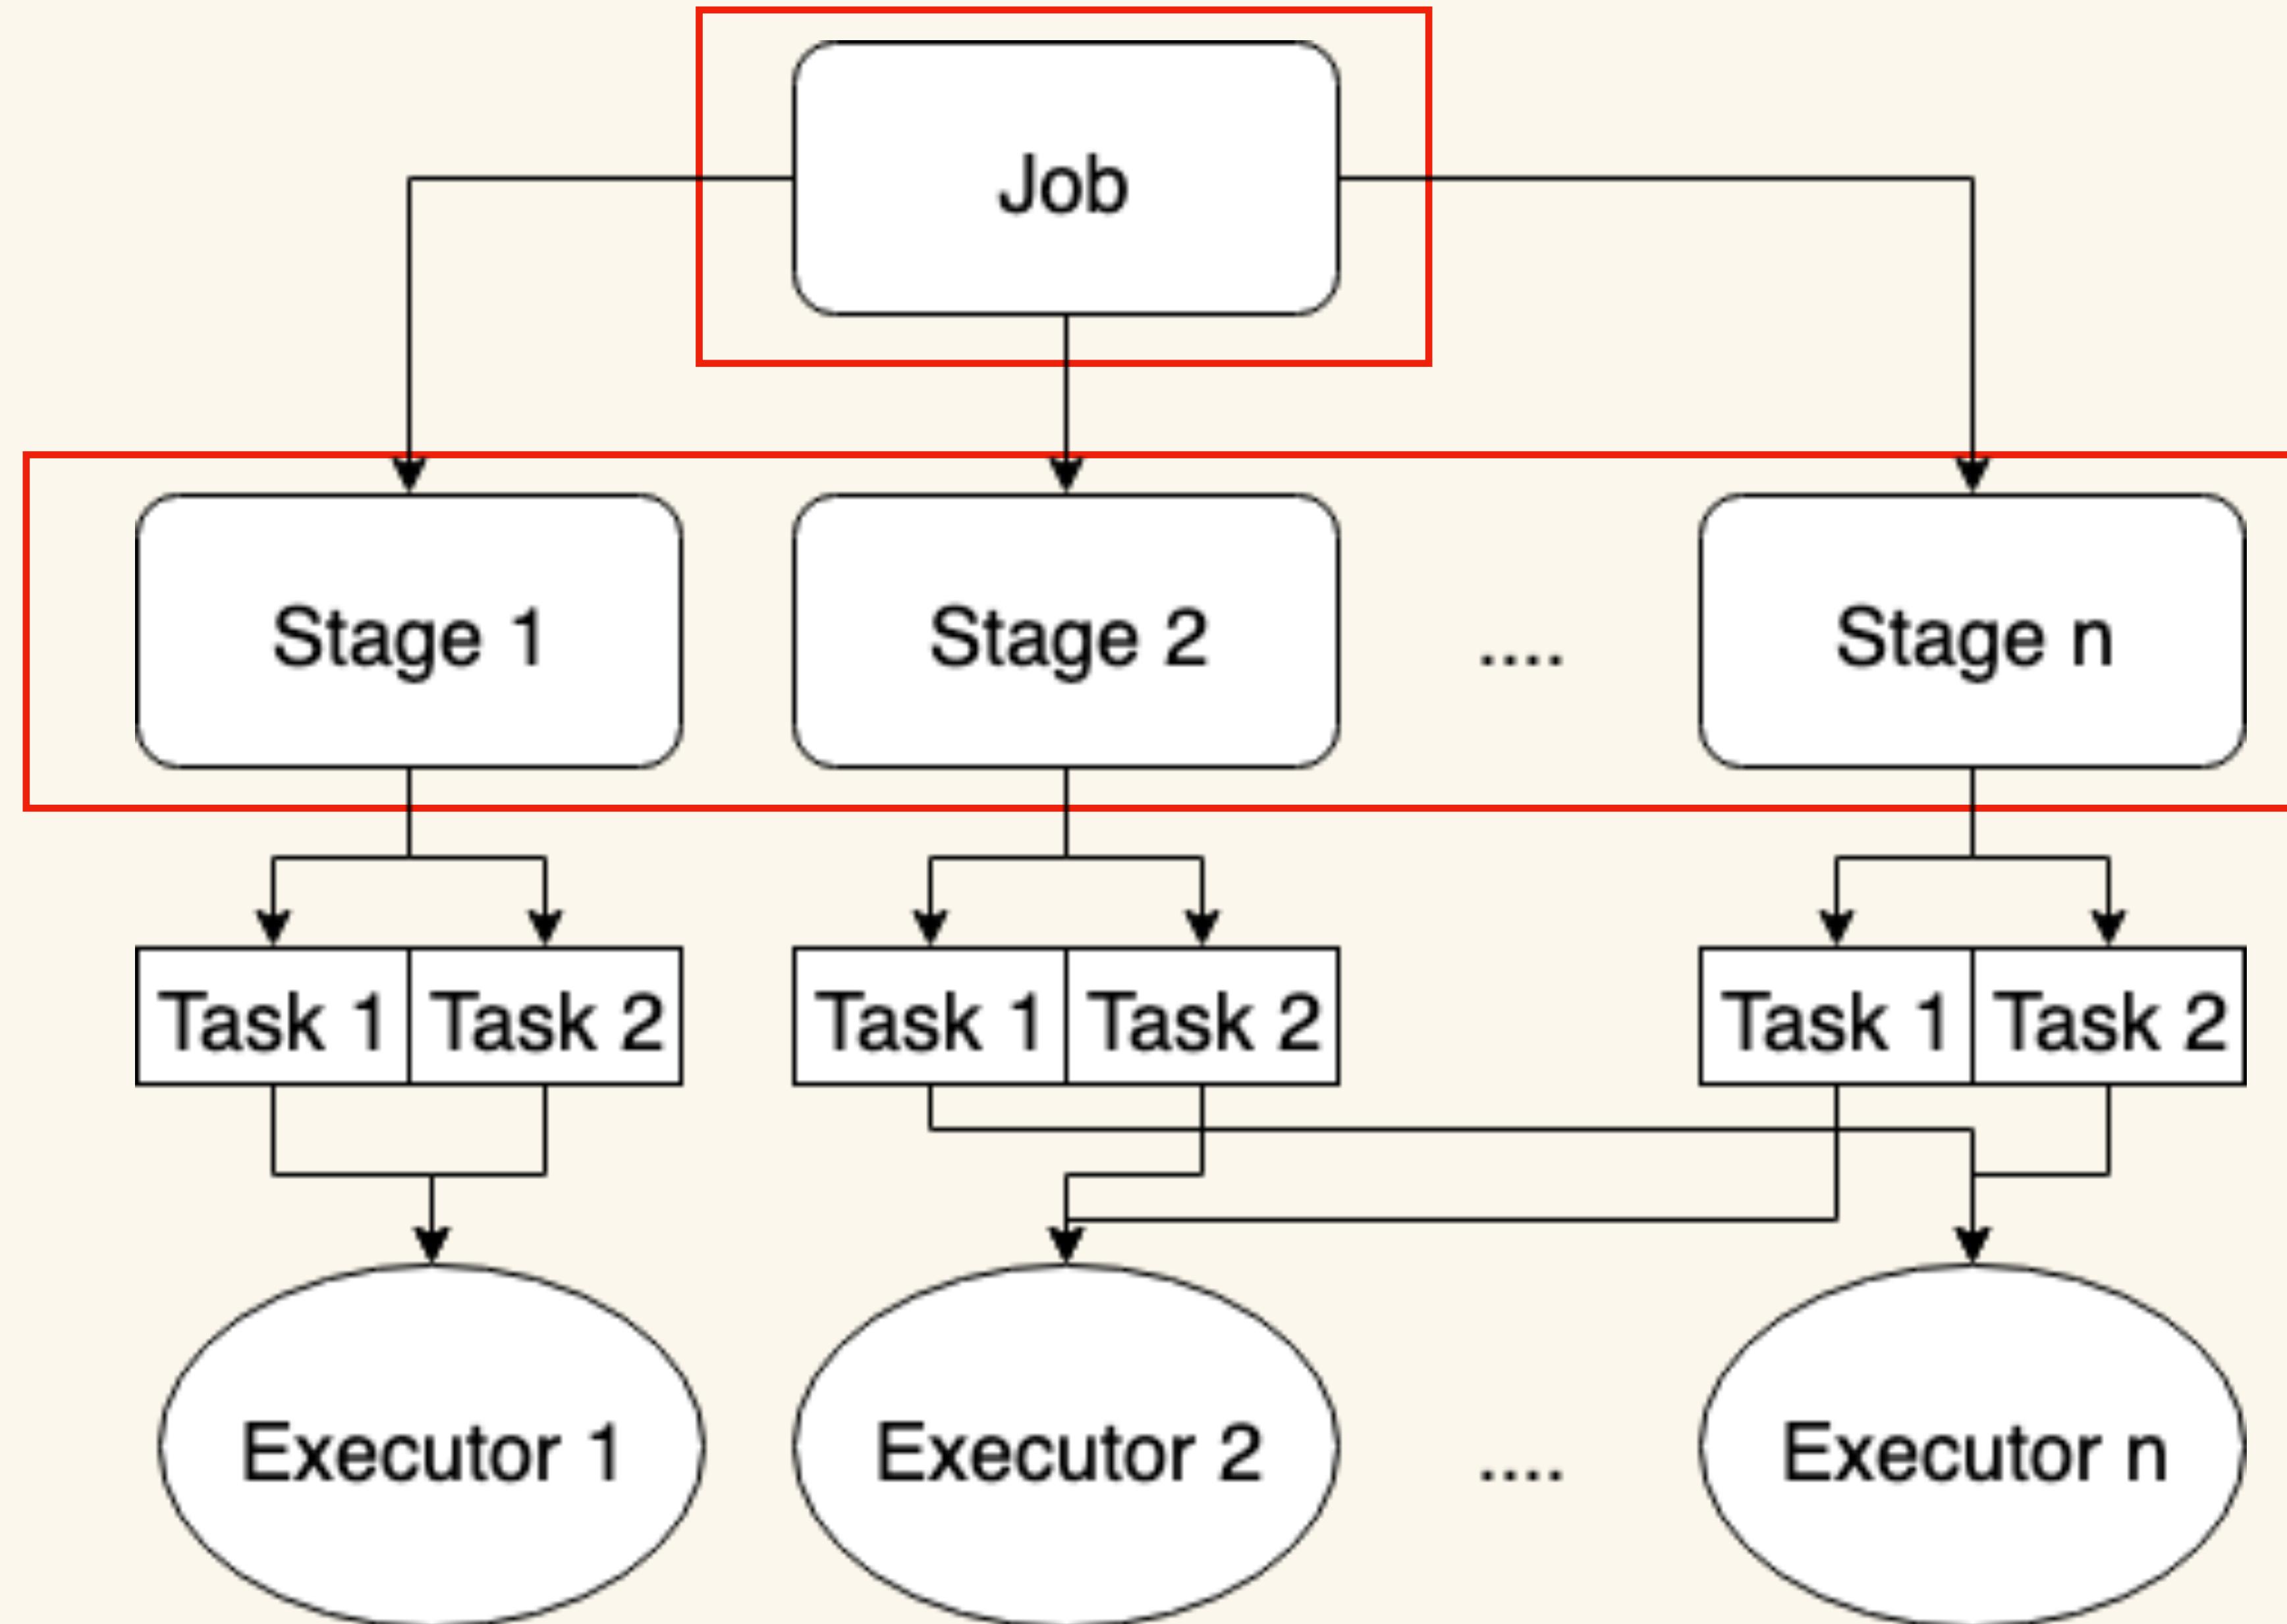
- Read Karl's discourse response to understand the architecture of Spark or go to Michael Style's talk.
- All we need to know is that executors are responsible for running tasks.
- The number of executors and memory per executor is set in the resource classes.

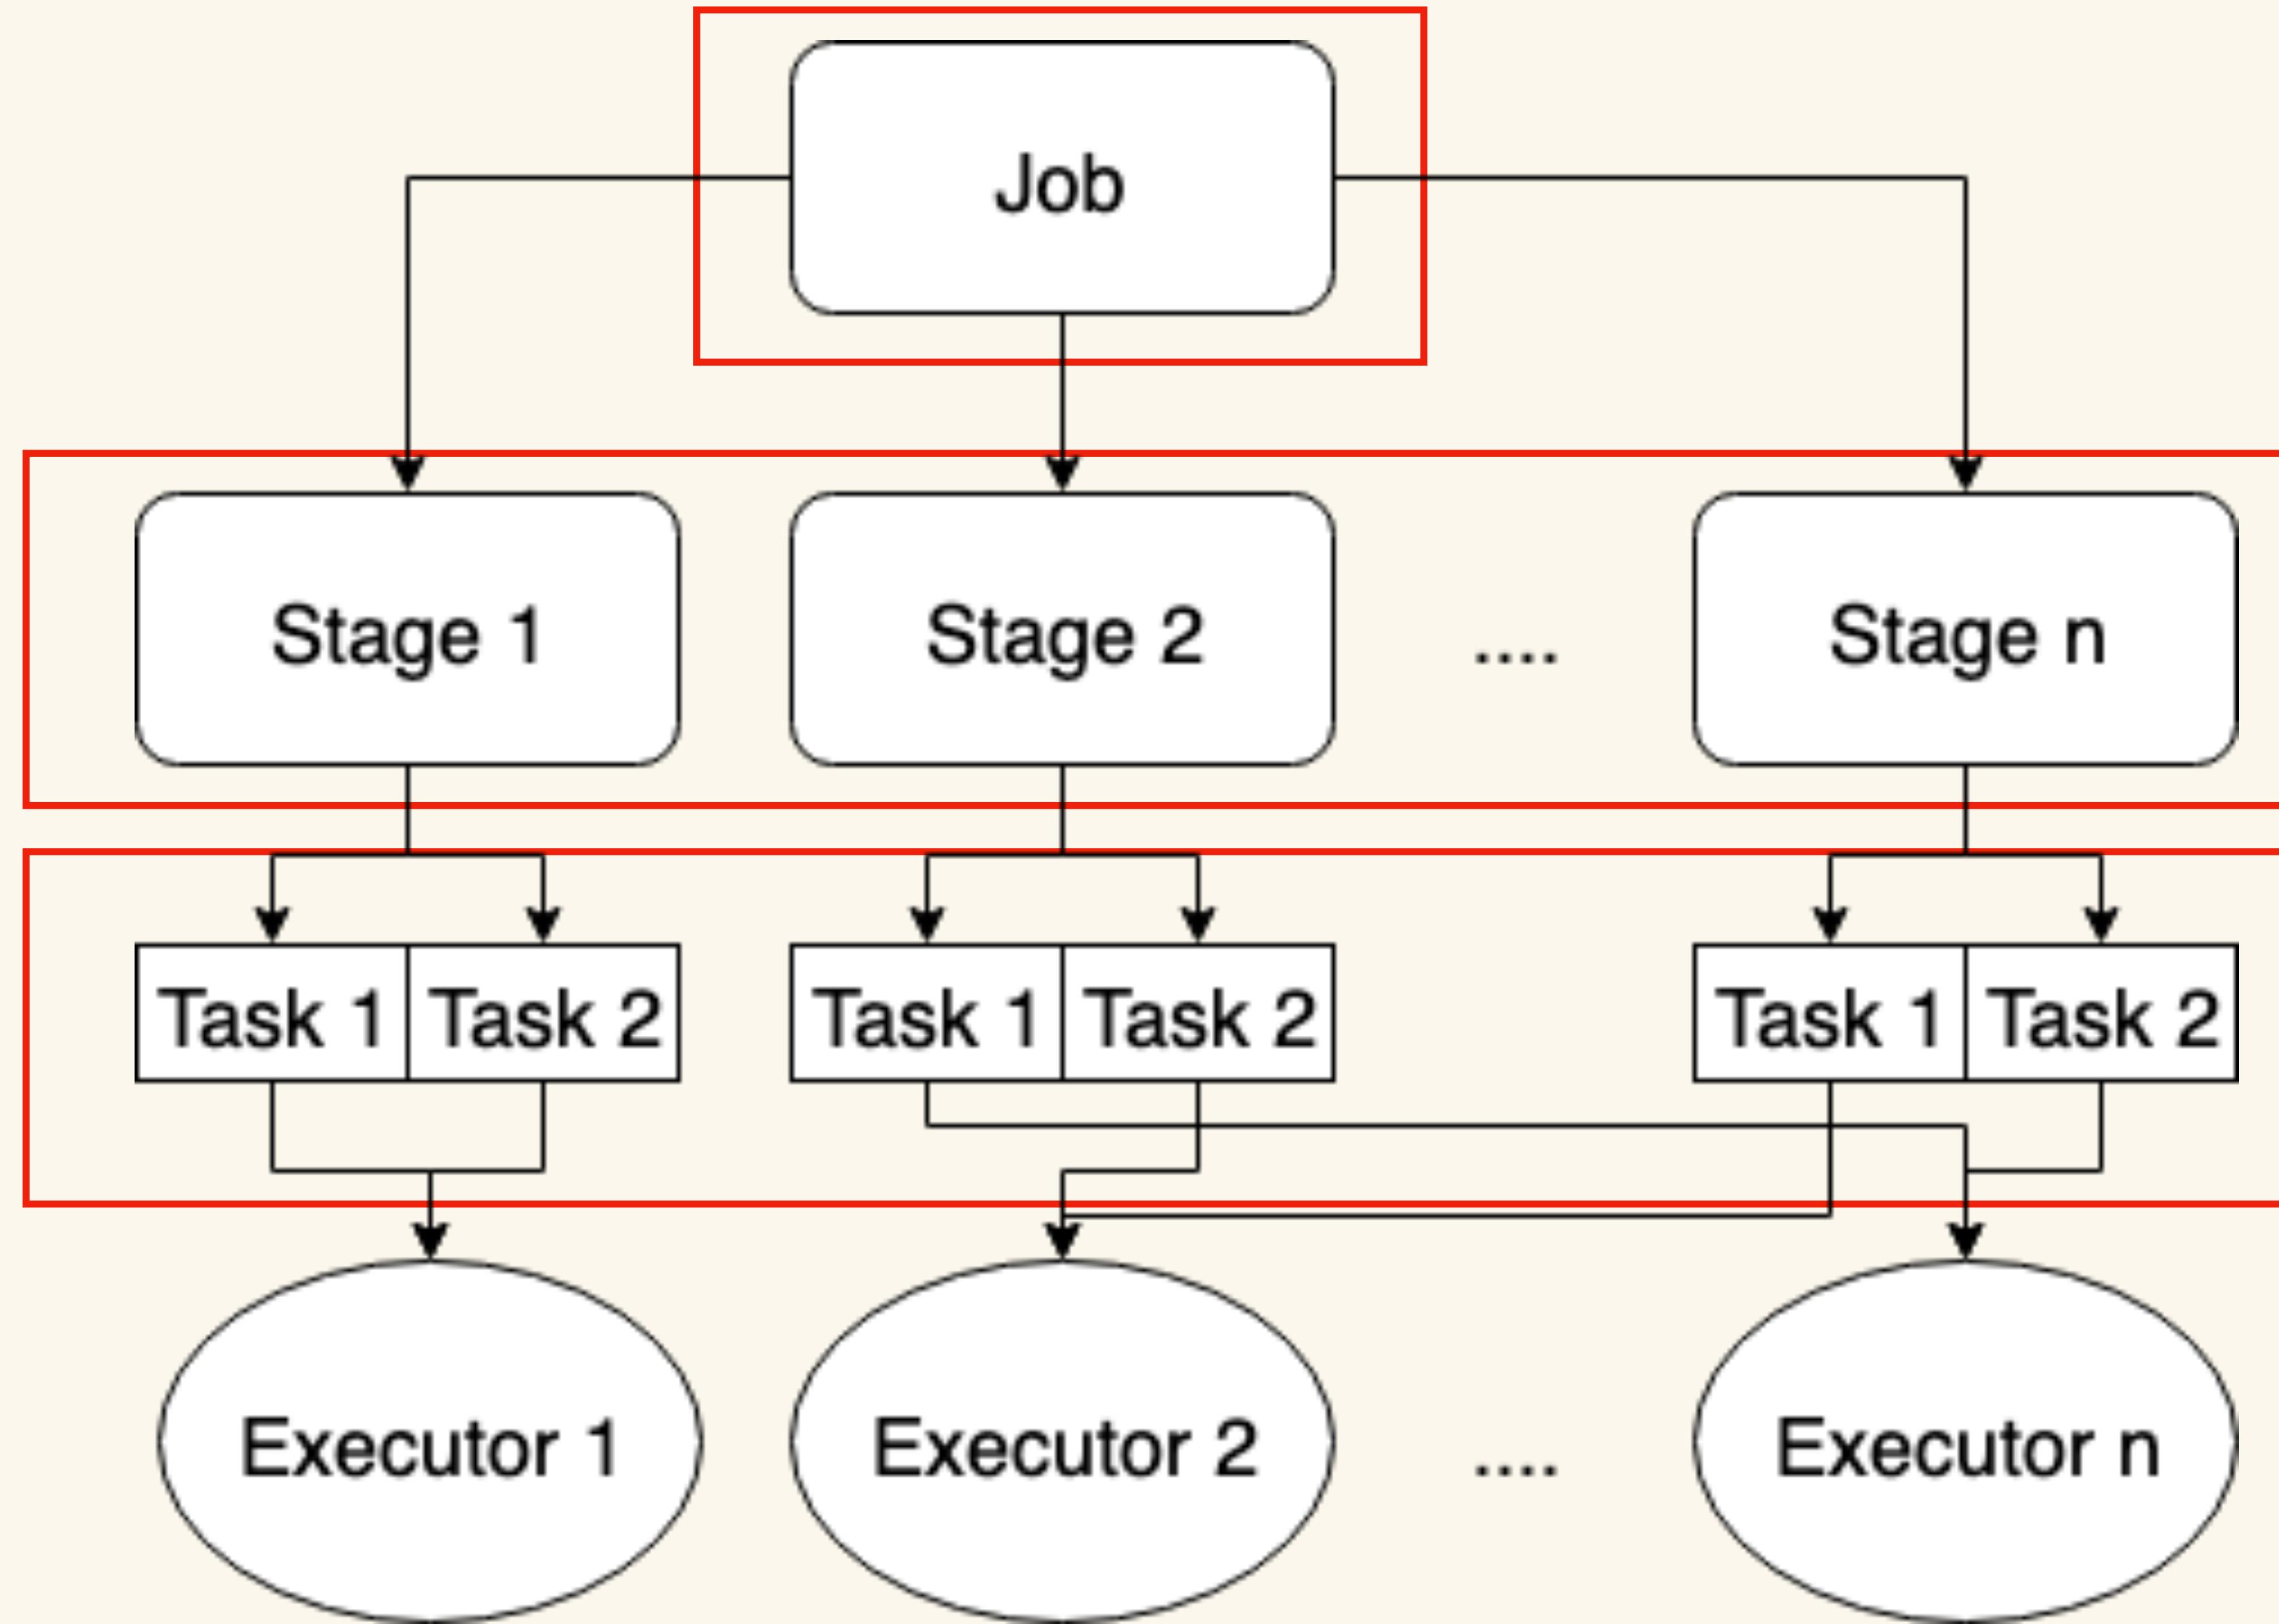
Karl's discourse response: <https://discourse.shopify.io/t/what-do-each-of-the-resource-class-arguments-mean/1920/3>

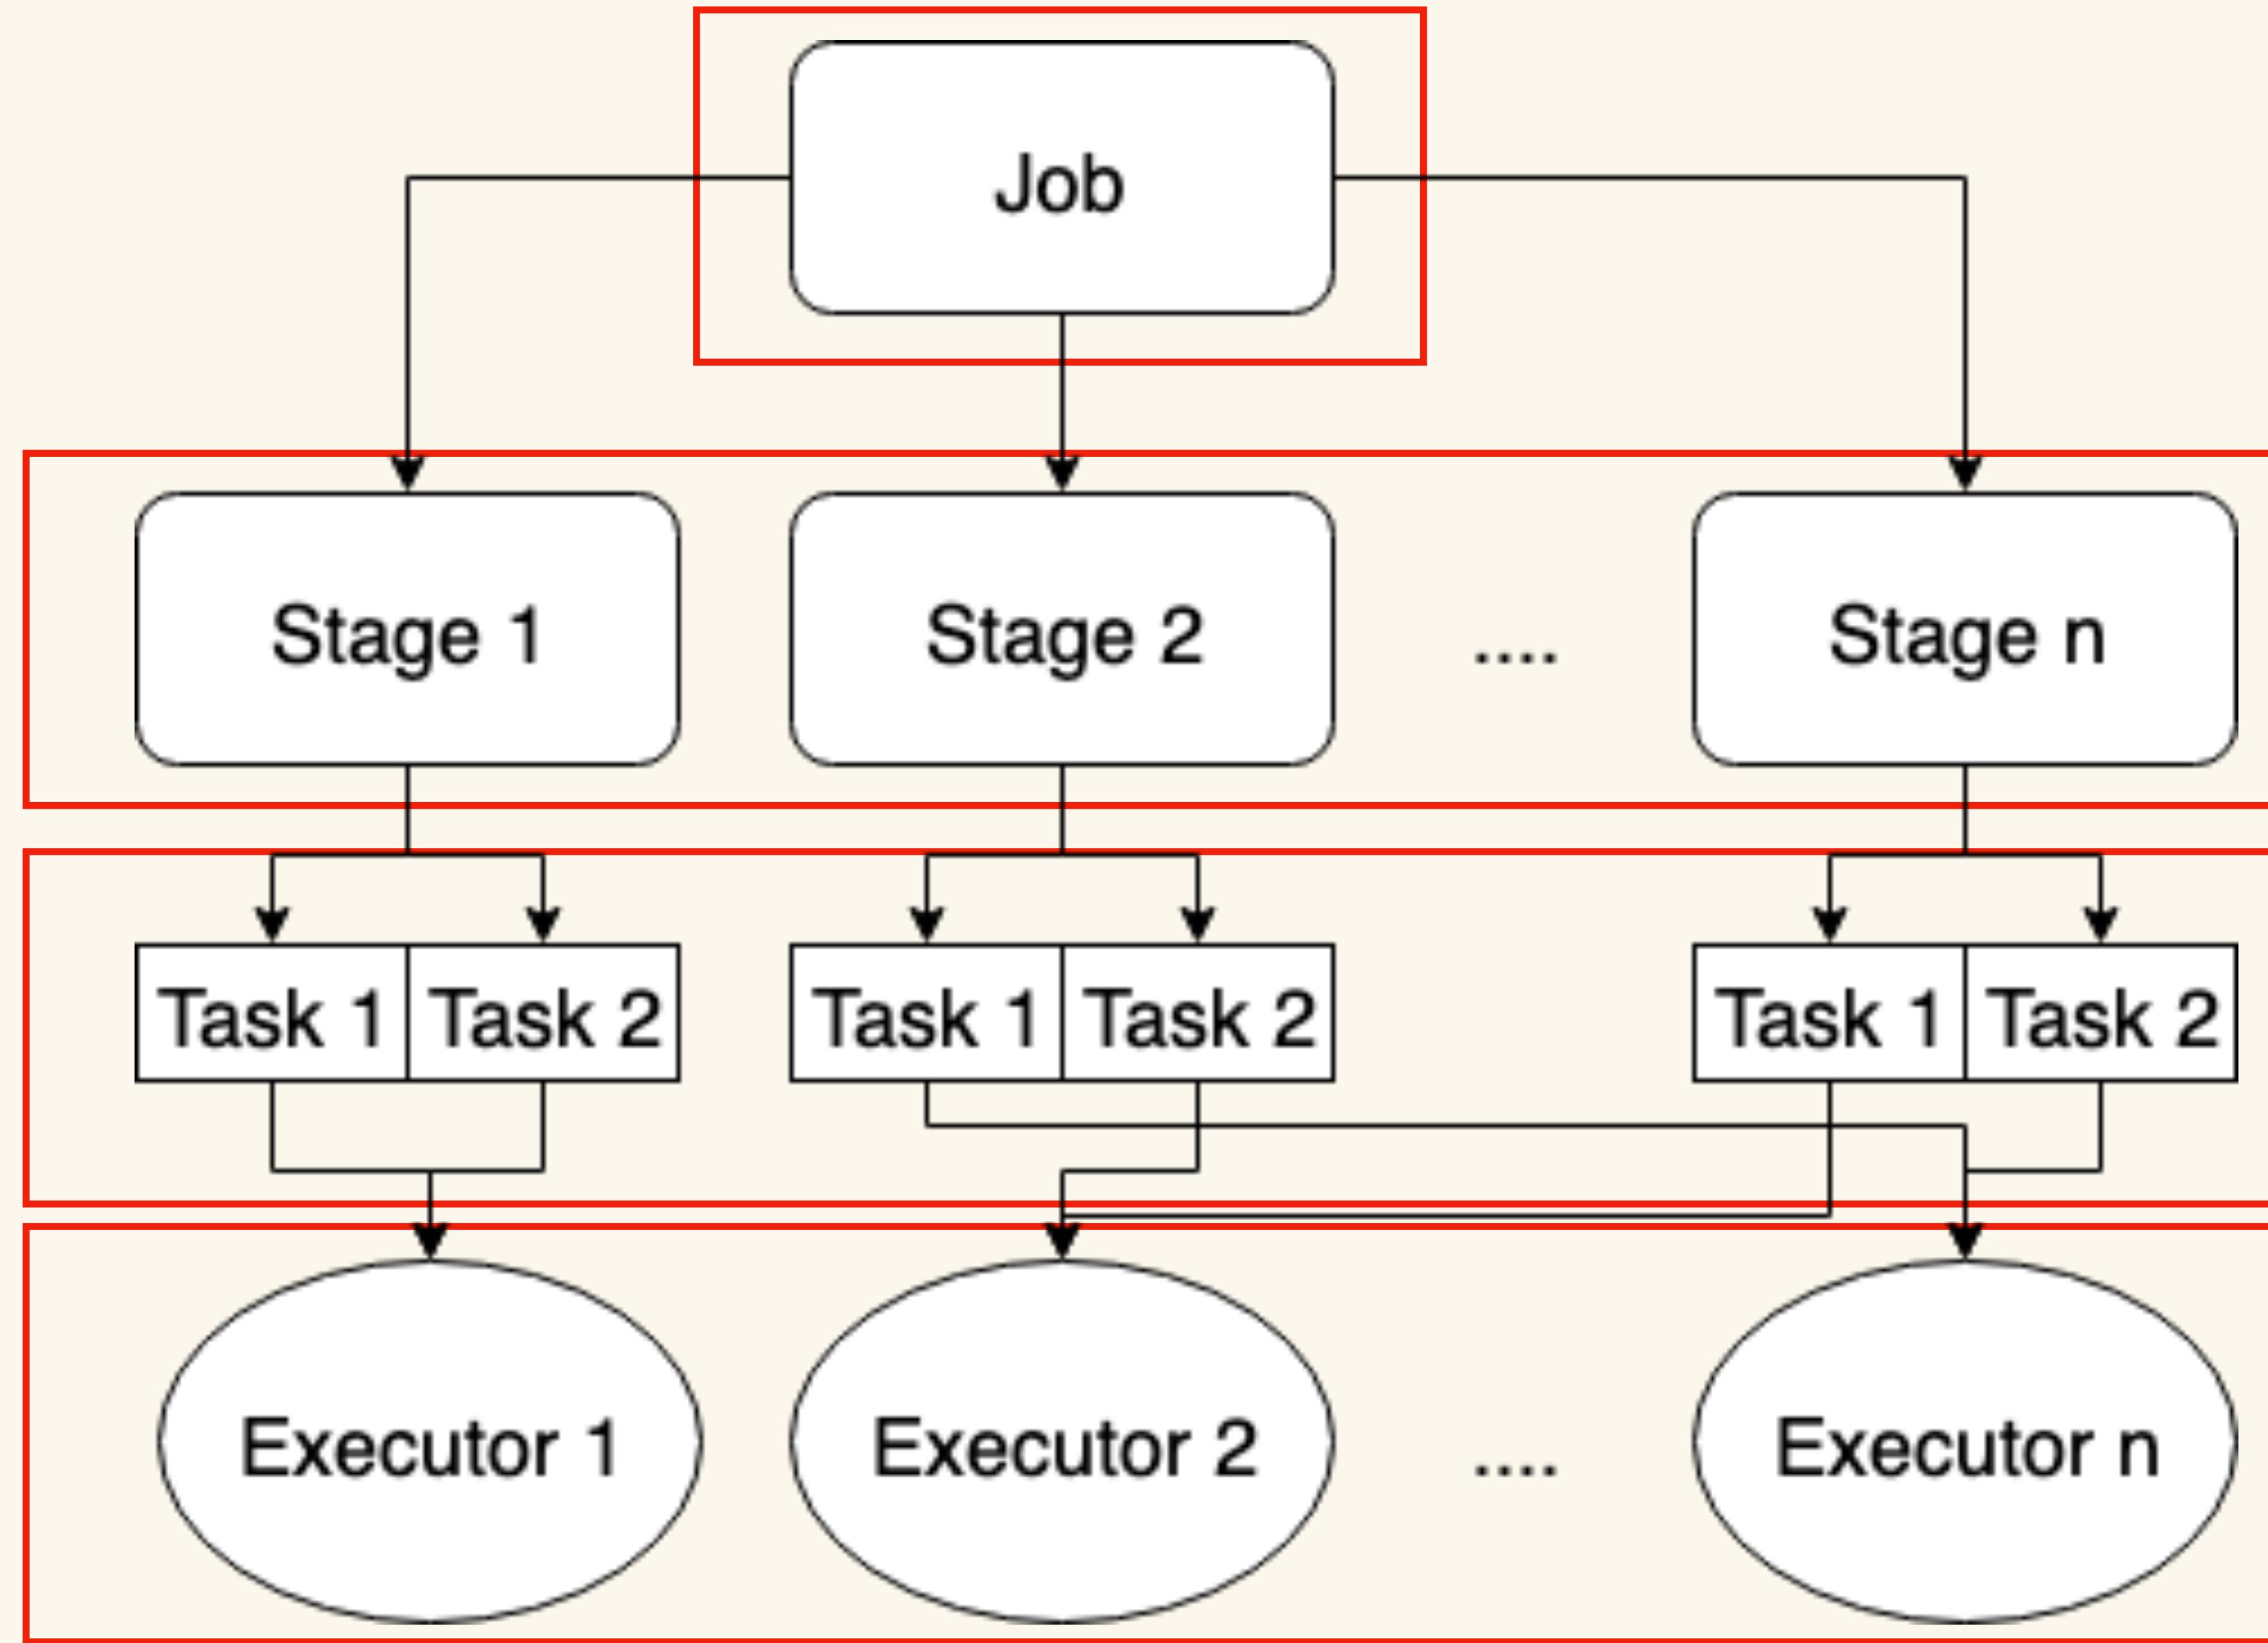












# Spill

- An executor has a set amount of memory, set at the start of the spark application.
- An executor is responsible for running tasks.
- A task is a unit of transformation(s) on a partition of data.
- If the partition of data exceeds the executor memory, it will spill the data from memory to disk.



# Spill

- Spilling is very bad, lots of I/O, serialization, etc. costs.
- If we spill too much we can potentially take the executors down.



# Jobs

- A job in a Spark application corresponds to a single “action” performed.
- A job consists of many stages.
- A new stage is created for every “shuffle operation”.

```
1 # Review of actions in Spark:  
2 df.write()  
3 df.head(n)  
4 df.take(n)  
5 df.collect()  
6 df.show()  
7 df.toPandas()
```

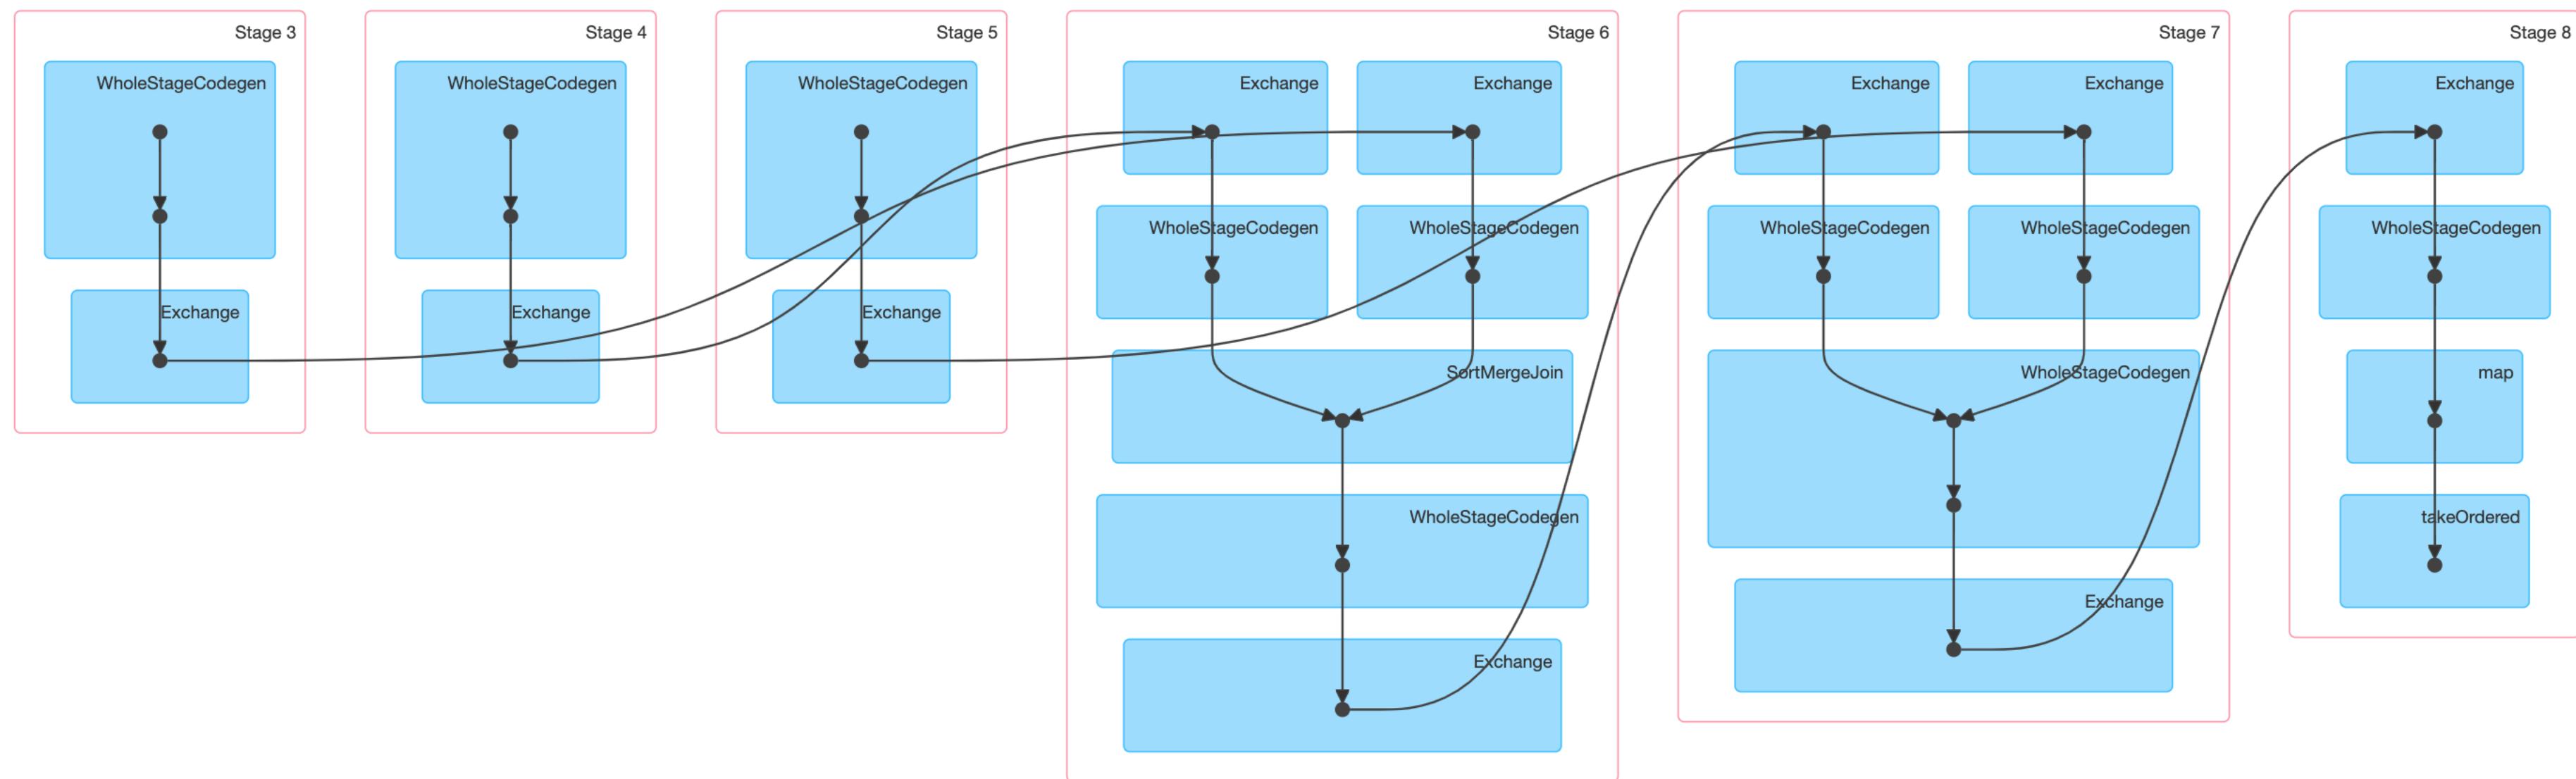
## Details for Job 3

Status: SUCCEEDED

Completed Stages: 6

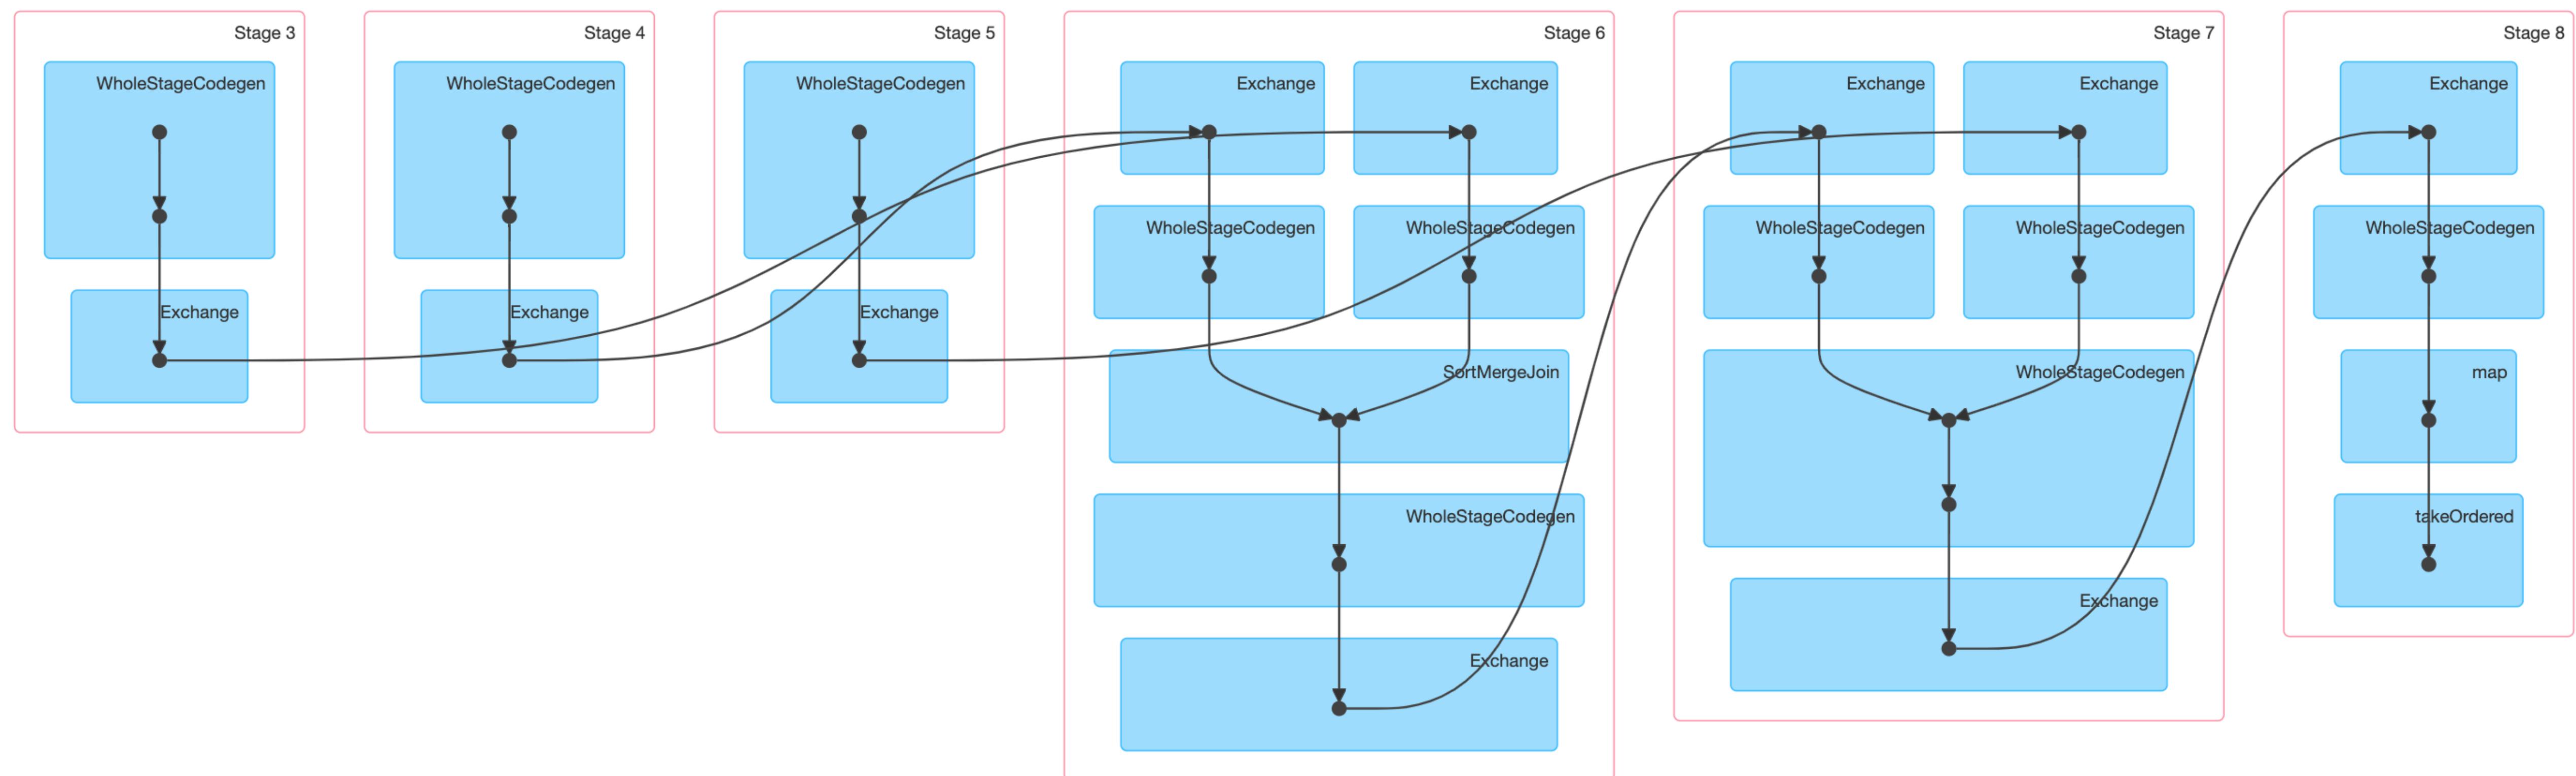
► Event Timeline

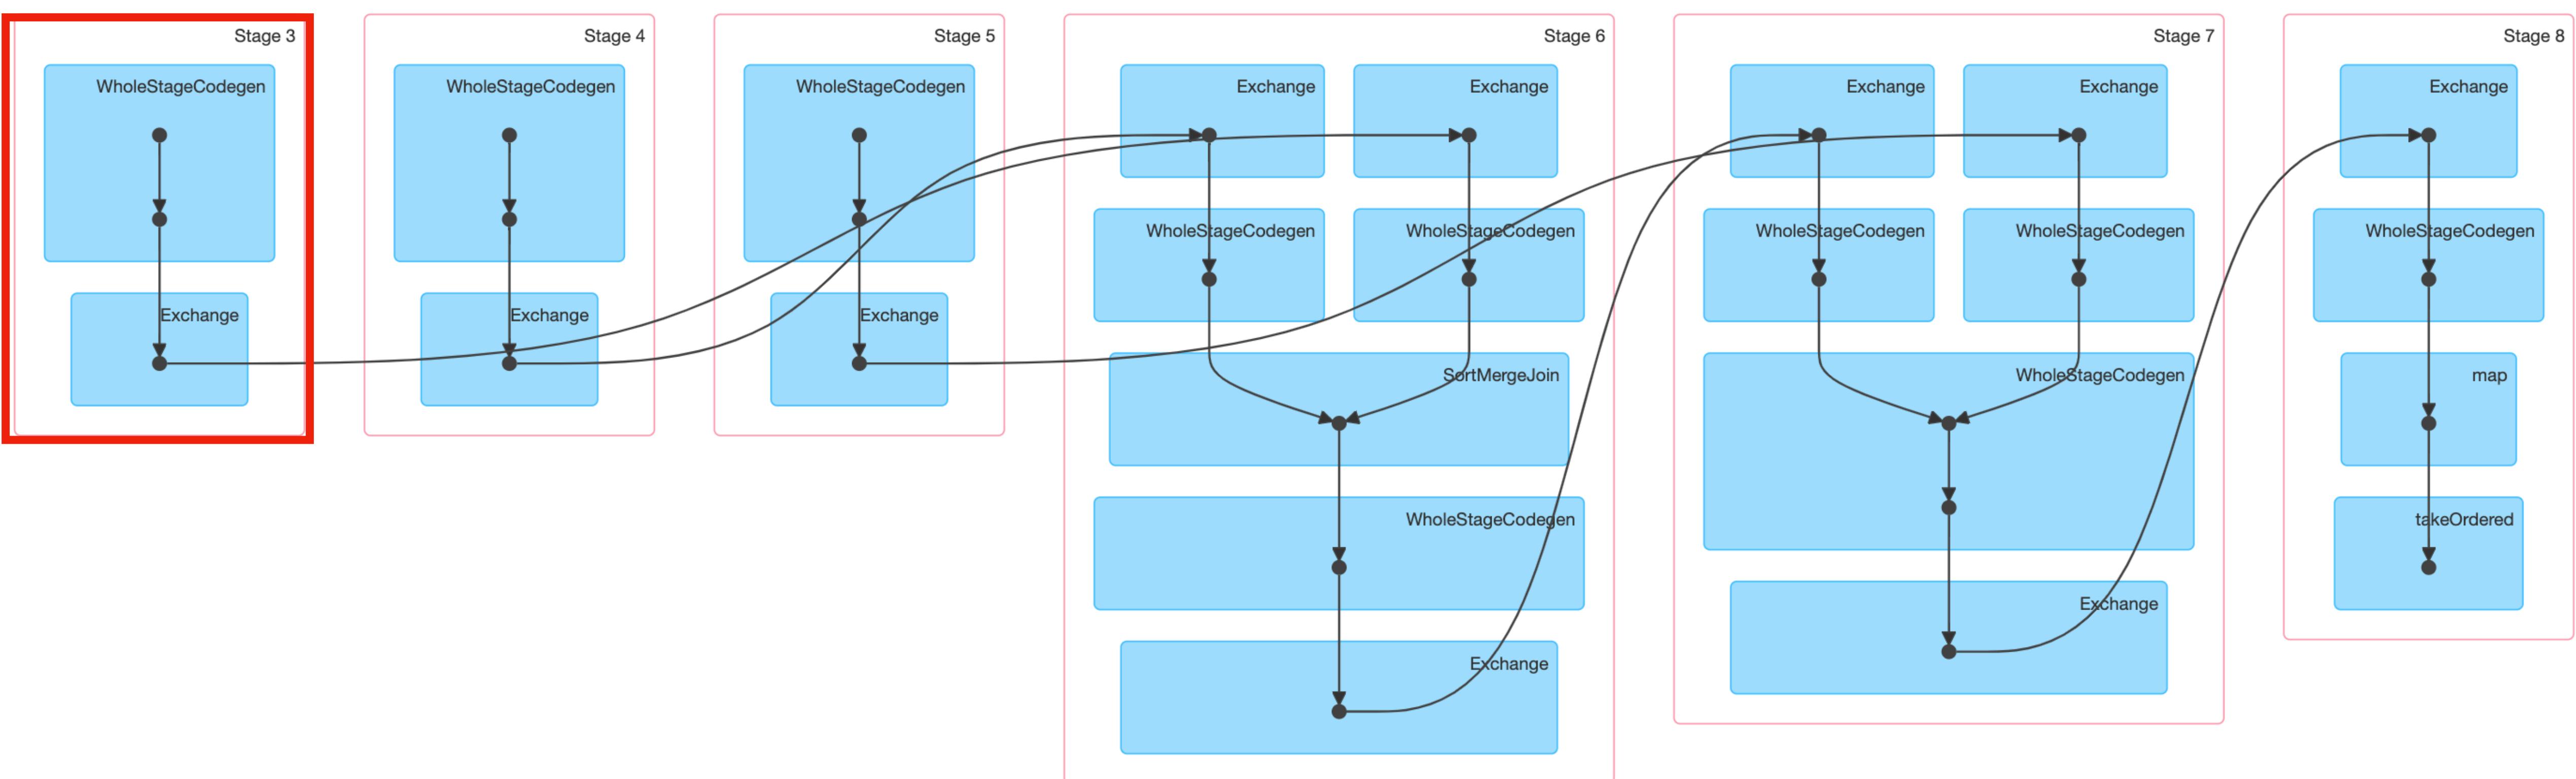
▼ DAG Visualization

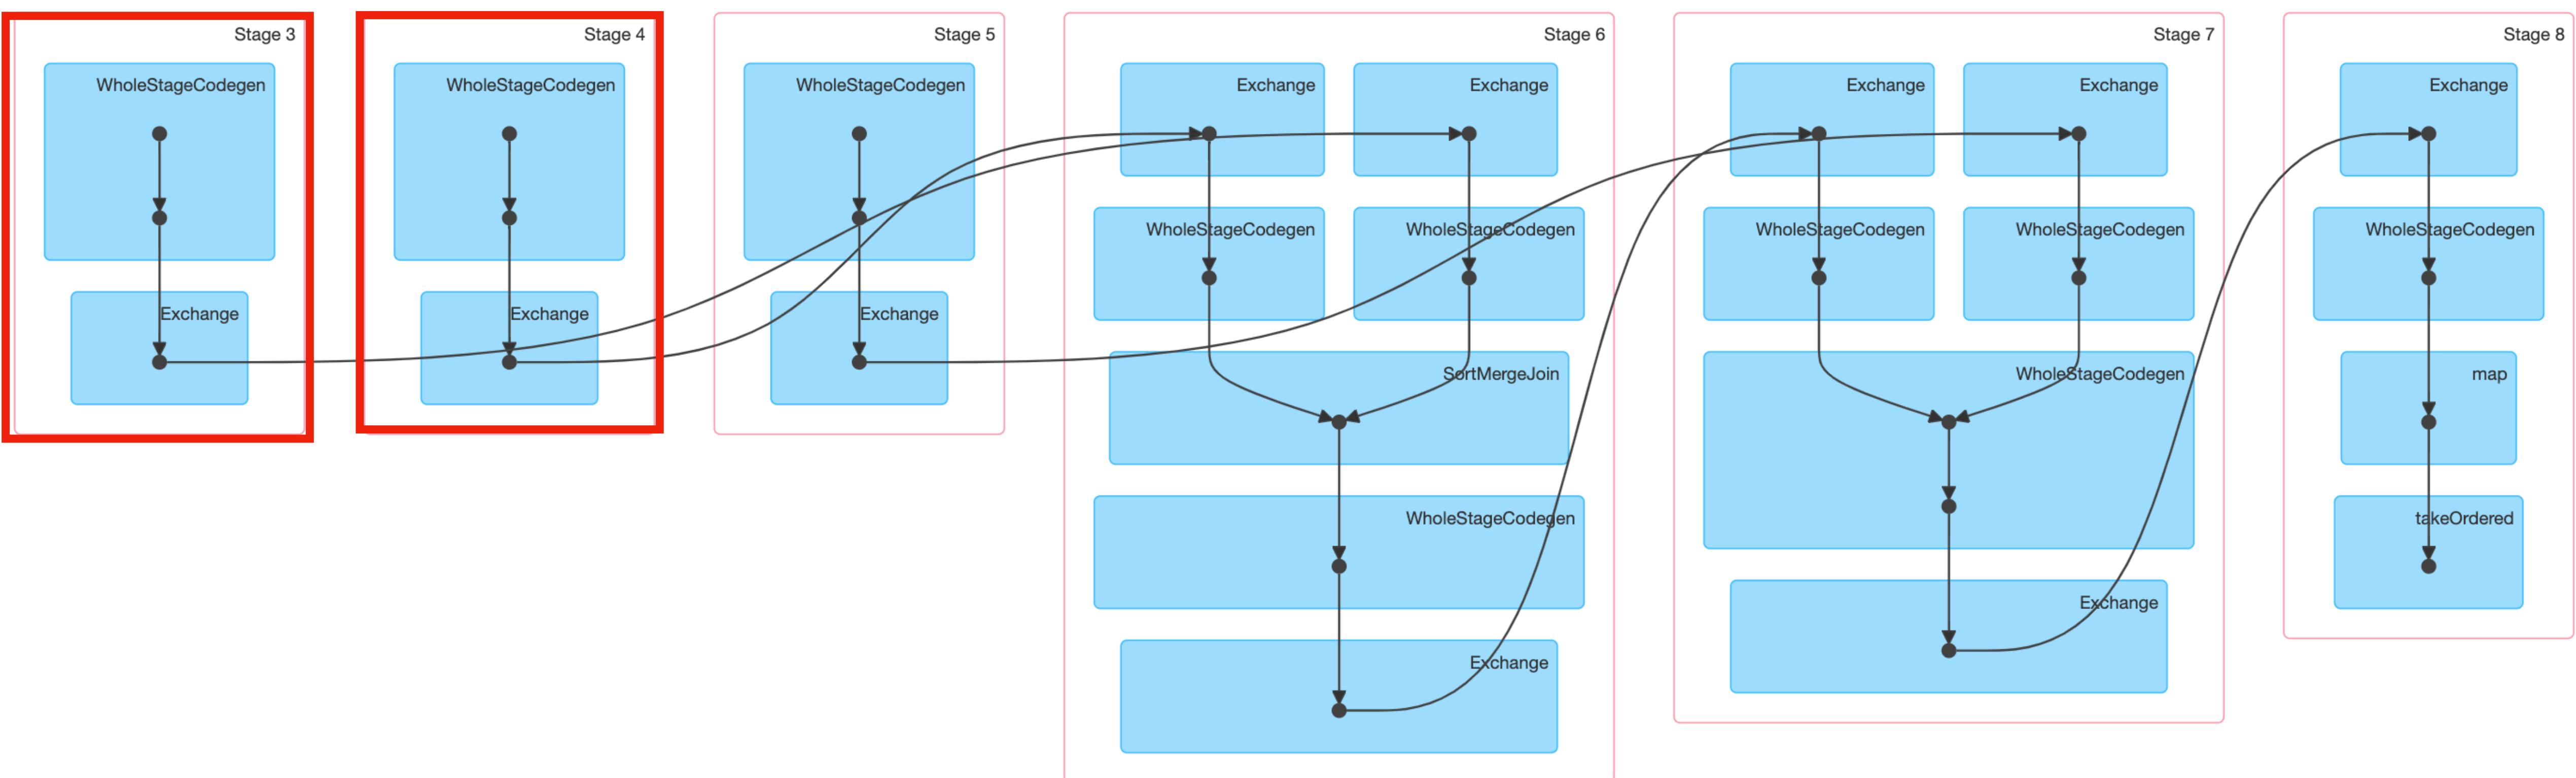


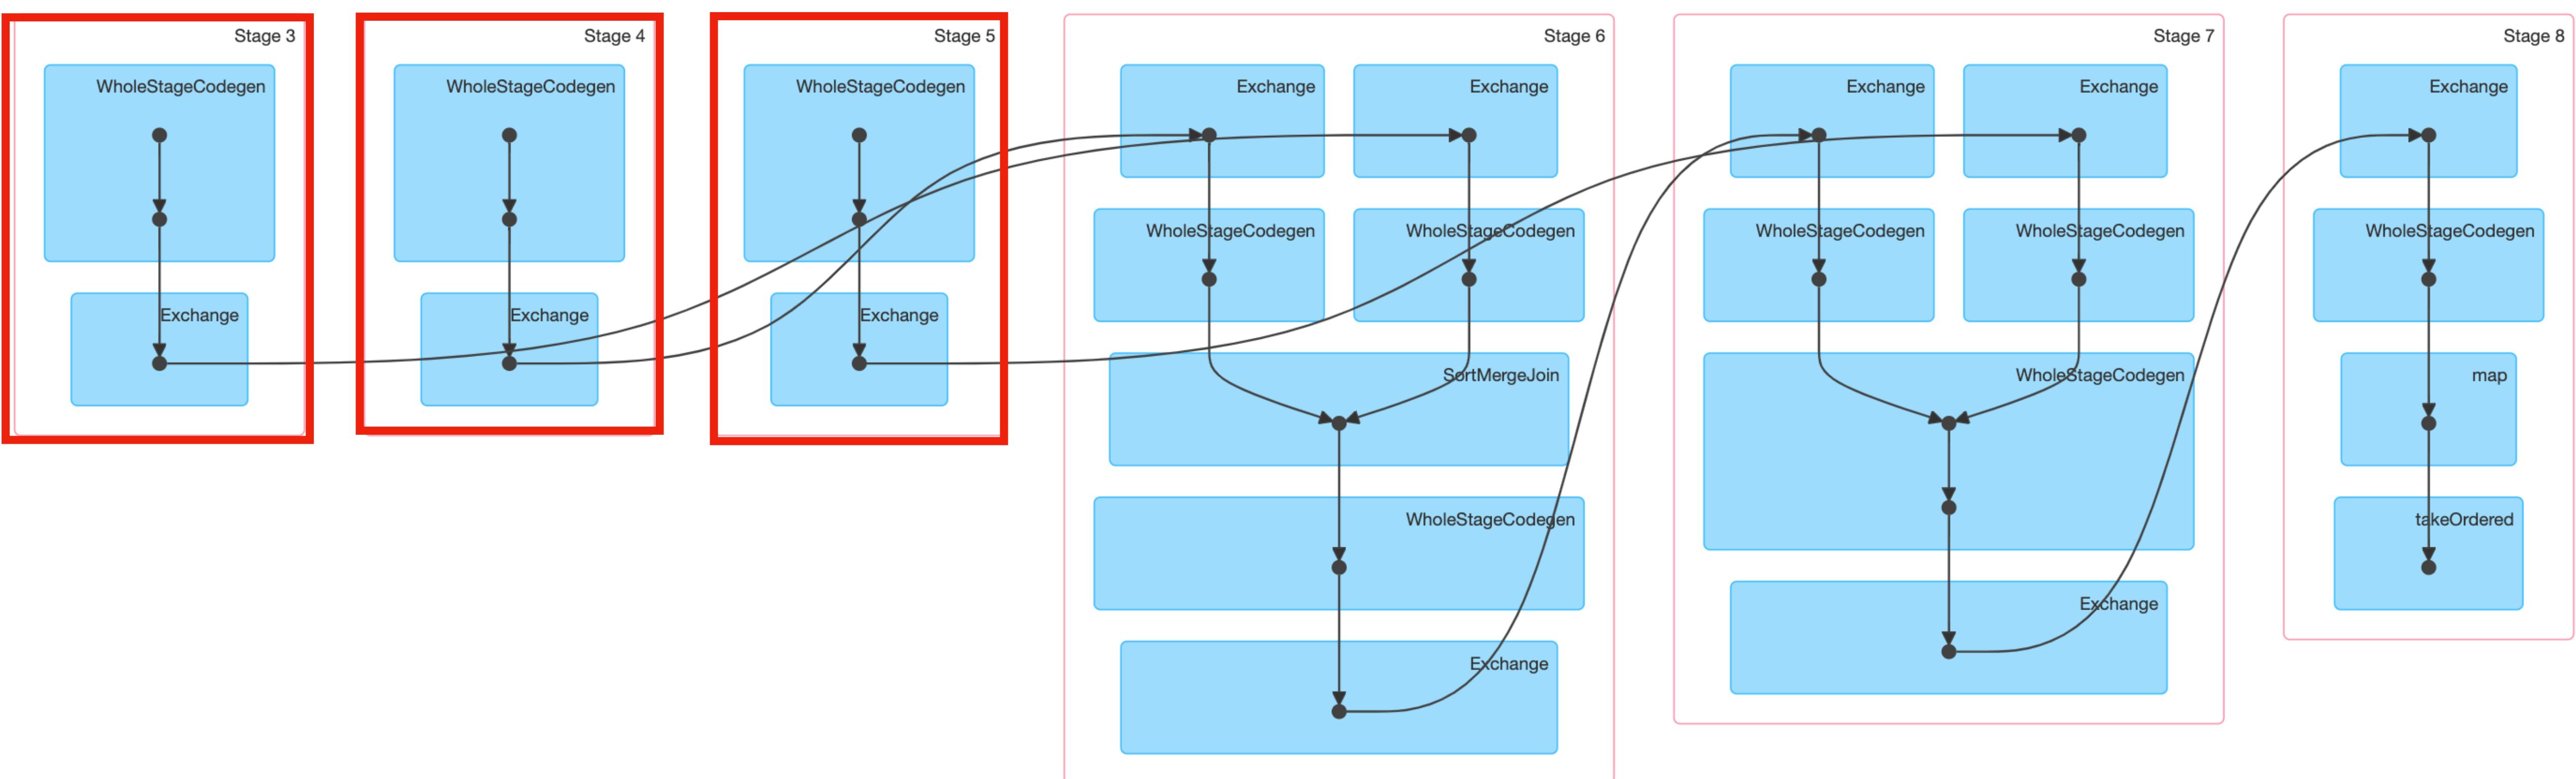
### Completed Stages (6)

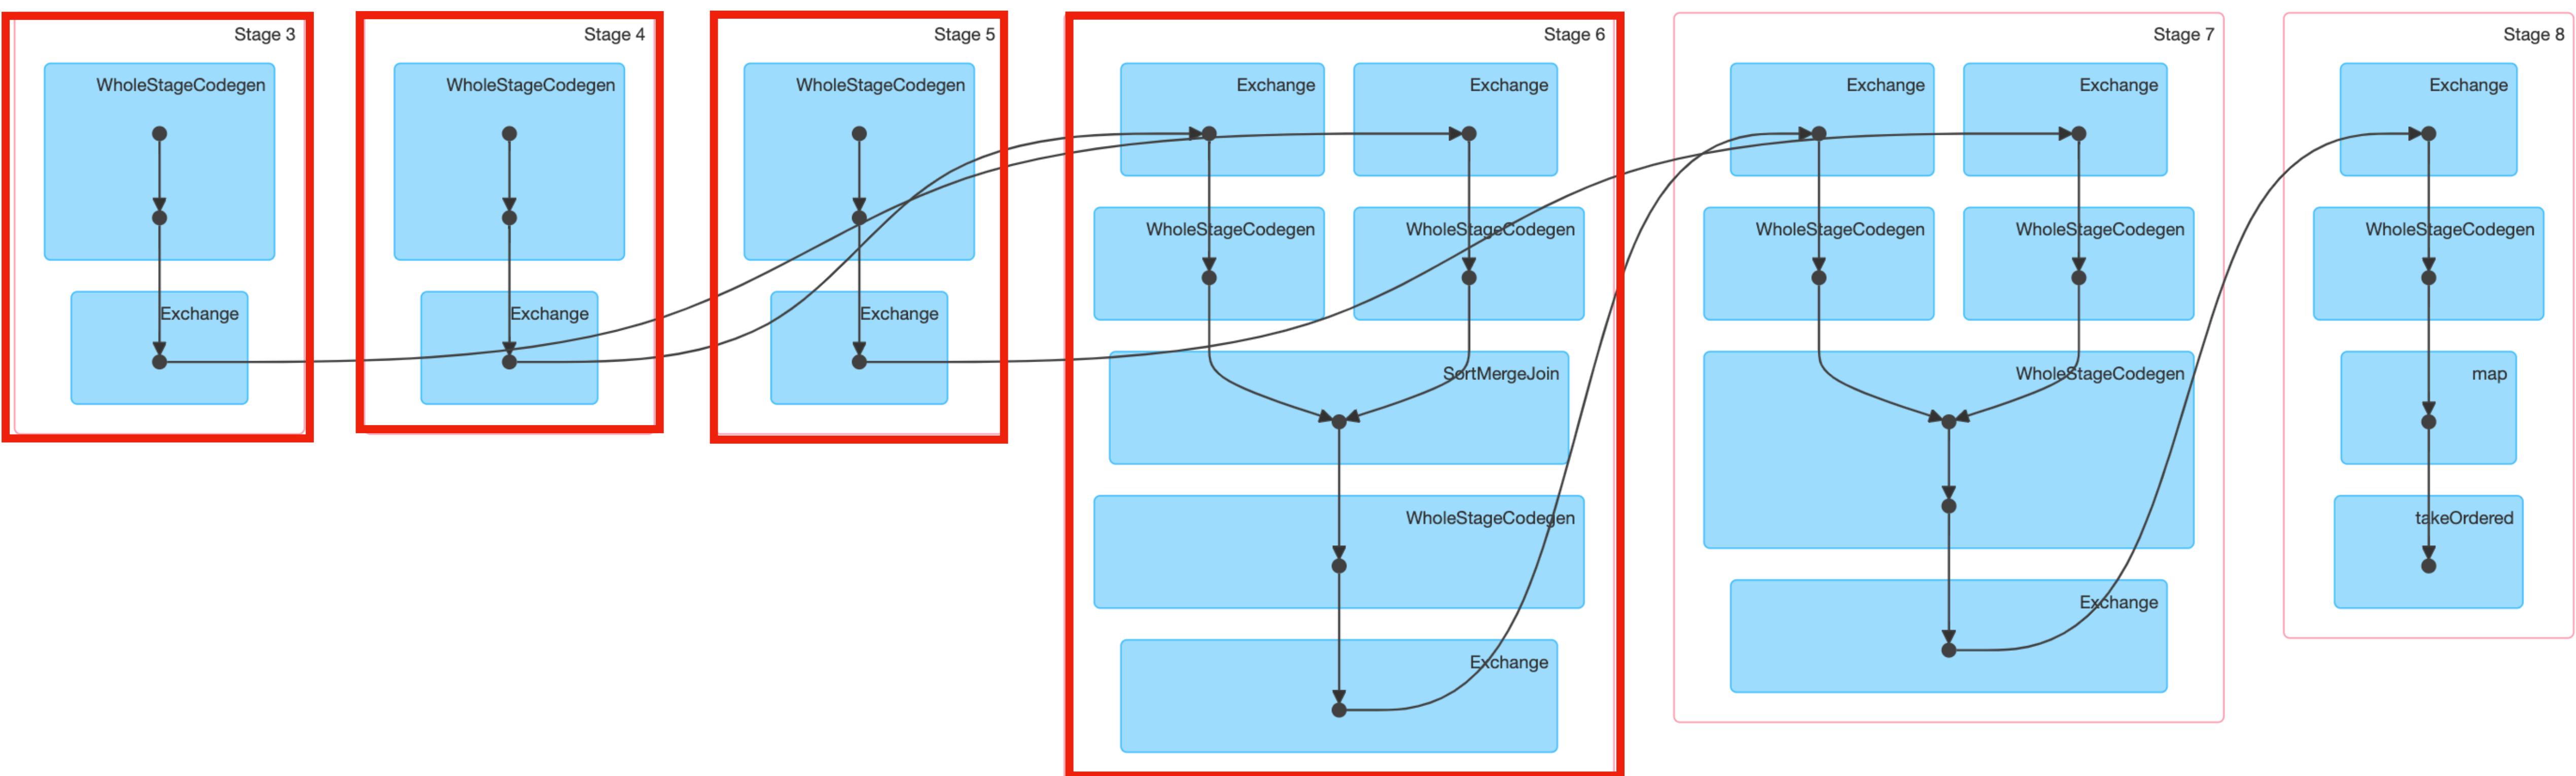
Stage Id ▾	Description	Submitted	Duration	Tasks: Succeeded/Total	Input	Output	Shuffle Read	Shuffle Write
8	showString at NativeMethodAccessorImpl.java:0	+details 2019/07/25 14:36:29	0.4 s	200/200			287.4 KB	
7	showString at NativeMethodAccessorImpl.java:0	+details 2019/07/25 14:34:55	1.6 min	200/200			25.1 GB	287.4 KB
6	showString at NativeMethodAccessorImpl.java:0	+details 2019/07/25 14:22:50	12 min	200/200 (1 killed: another attempt succeeded)			29.7 GB	18.6 GB
5	showString at NativeMethodAccessorImpl.java:0	+details 2019/07/25 14:19:18	3.9 min	1044/1044	16.0 GB			8.5 GB
4	showString at NativeMethodAccessorImpl.java:0	+details 2019/07/25 14:19:18	3.3 min	14629/14629	86.2 GB			25.1 GB
3	showString at NativeMethodAccessorImpl.java:0	+details 2019/07/25 14:19:18	12 s	413/413	213.4 MB			151.0 MB

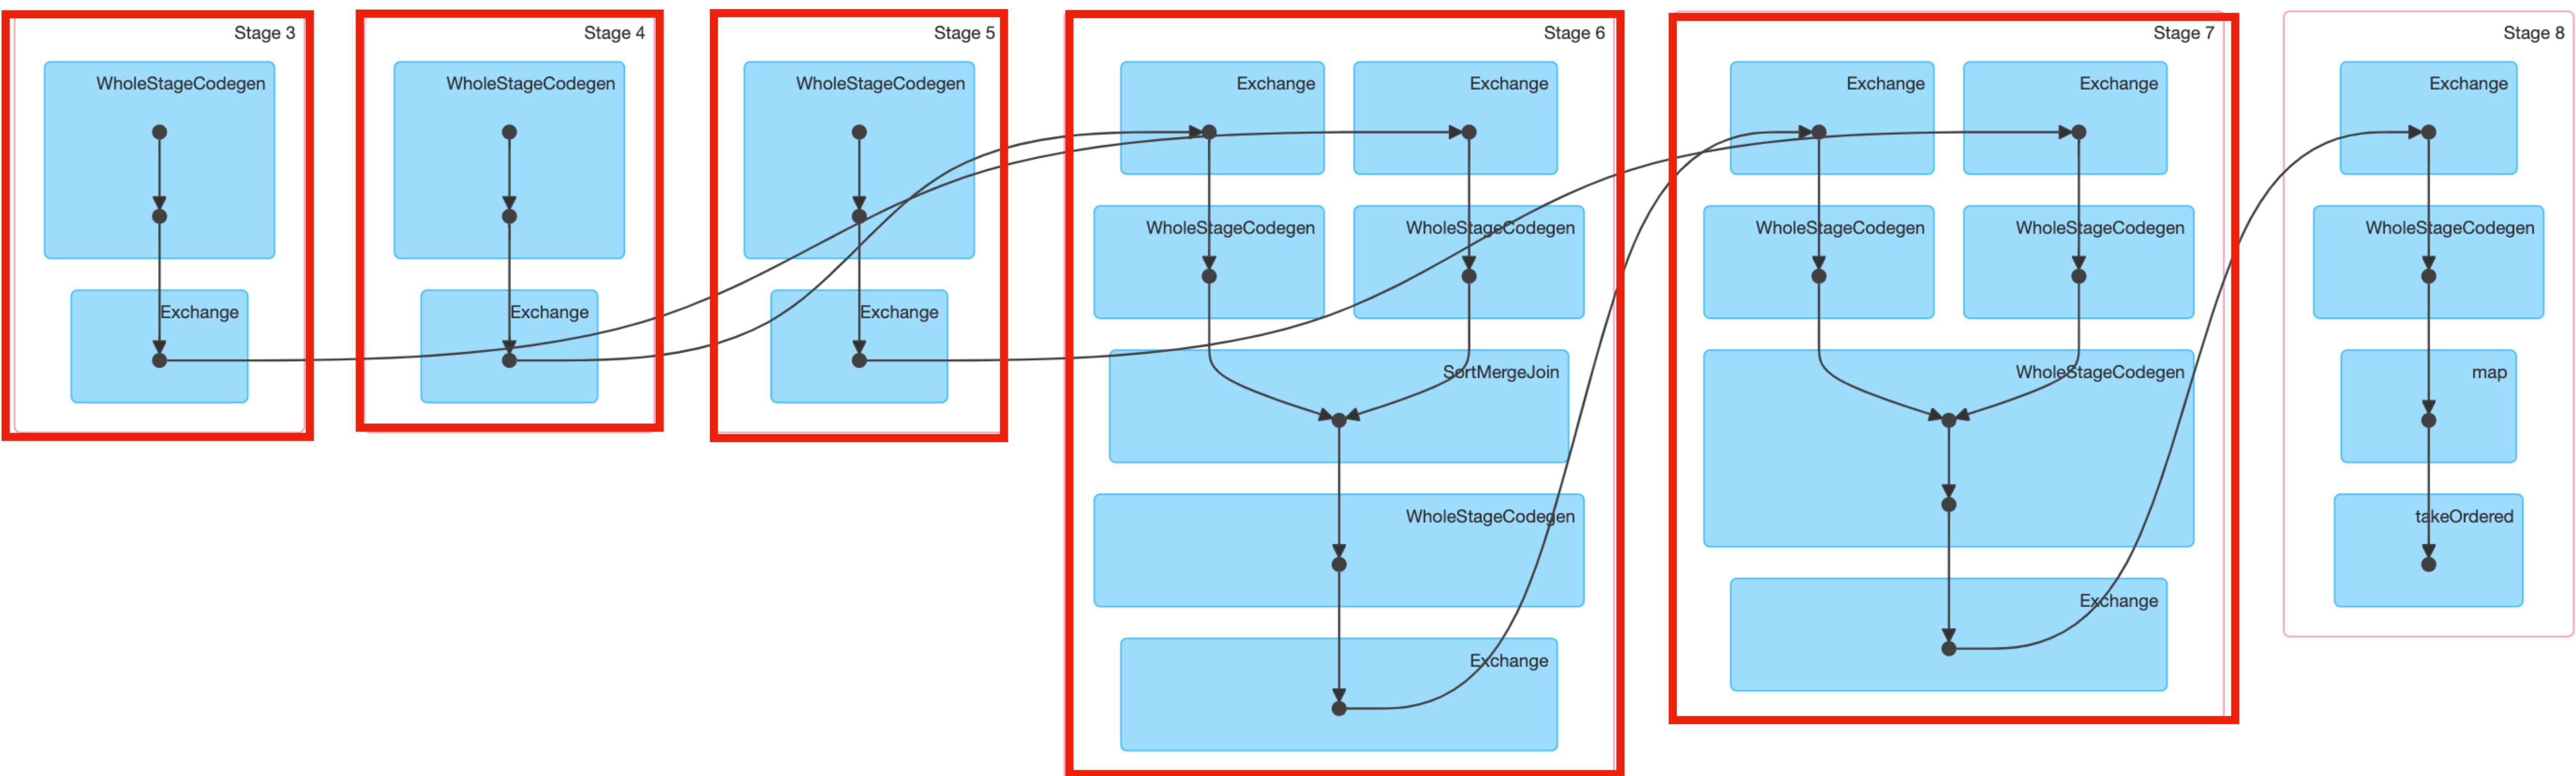


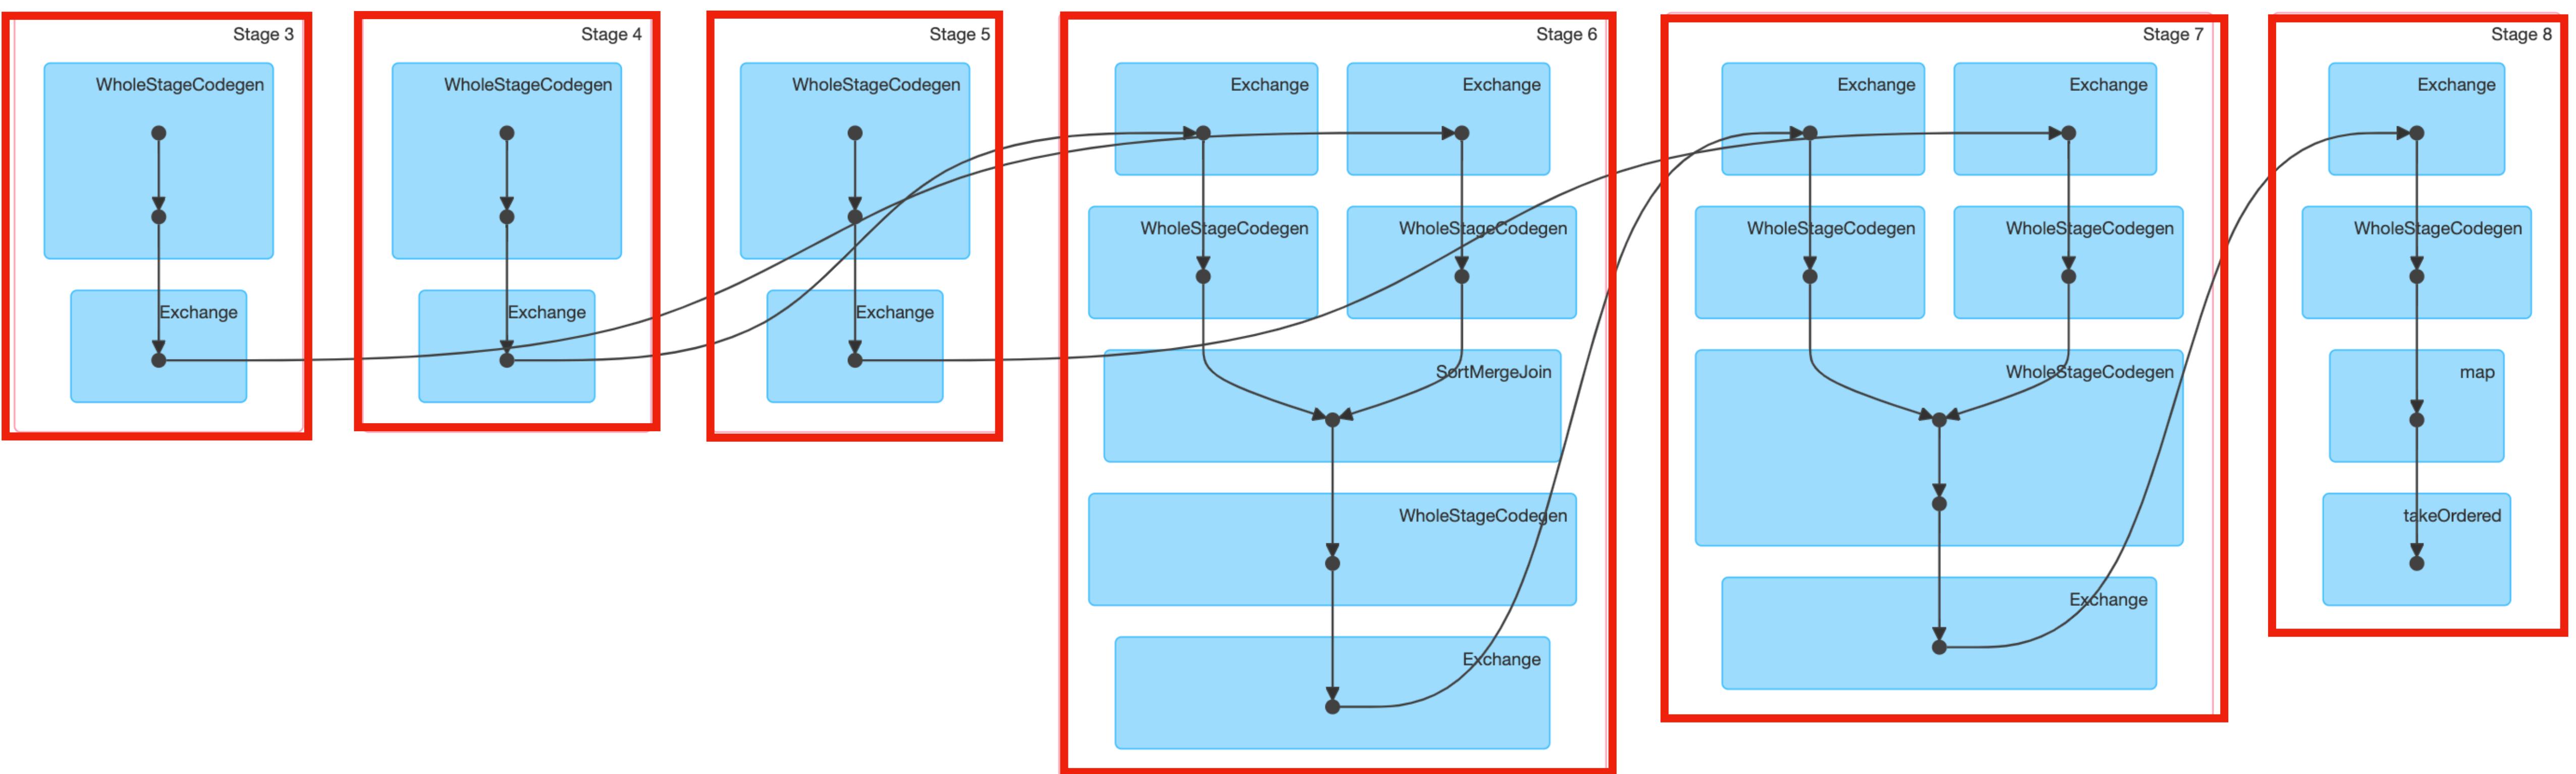


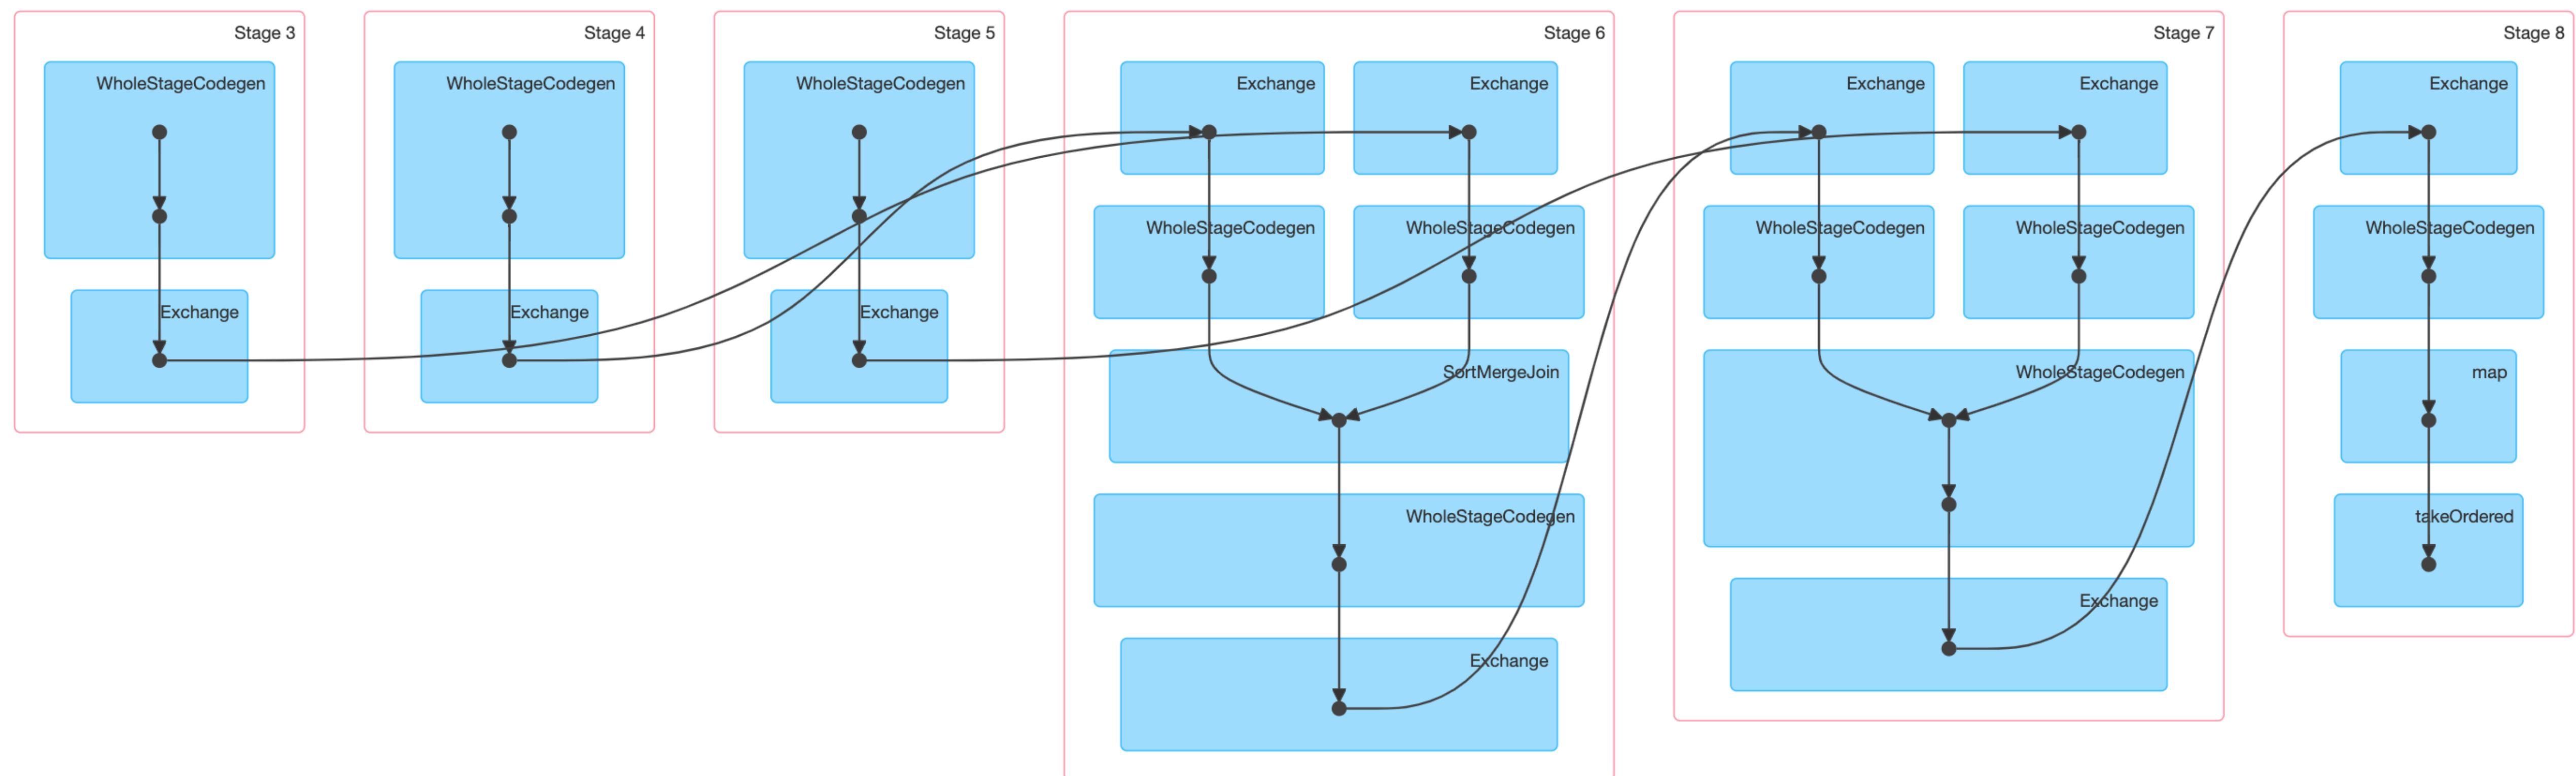


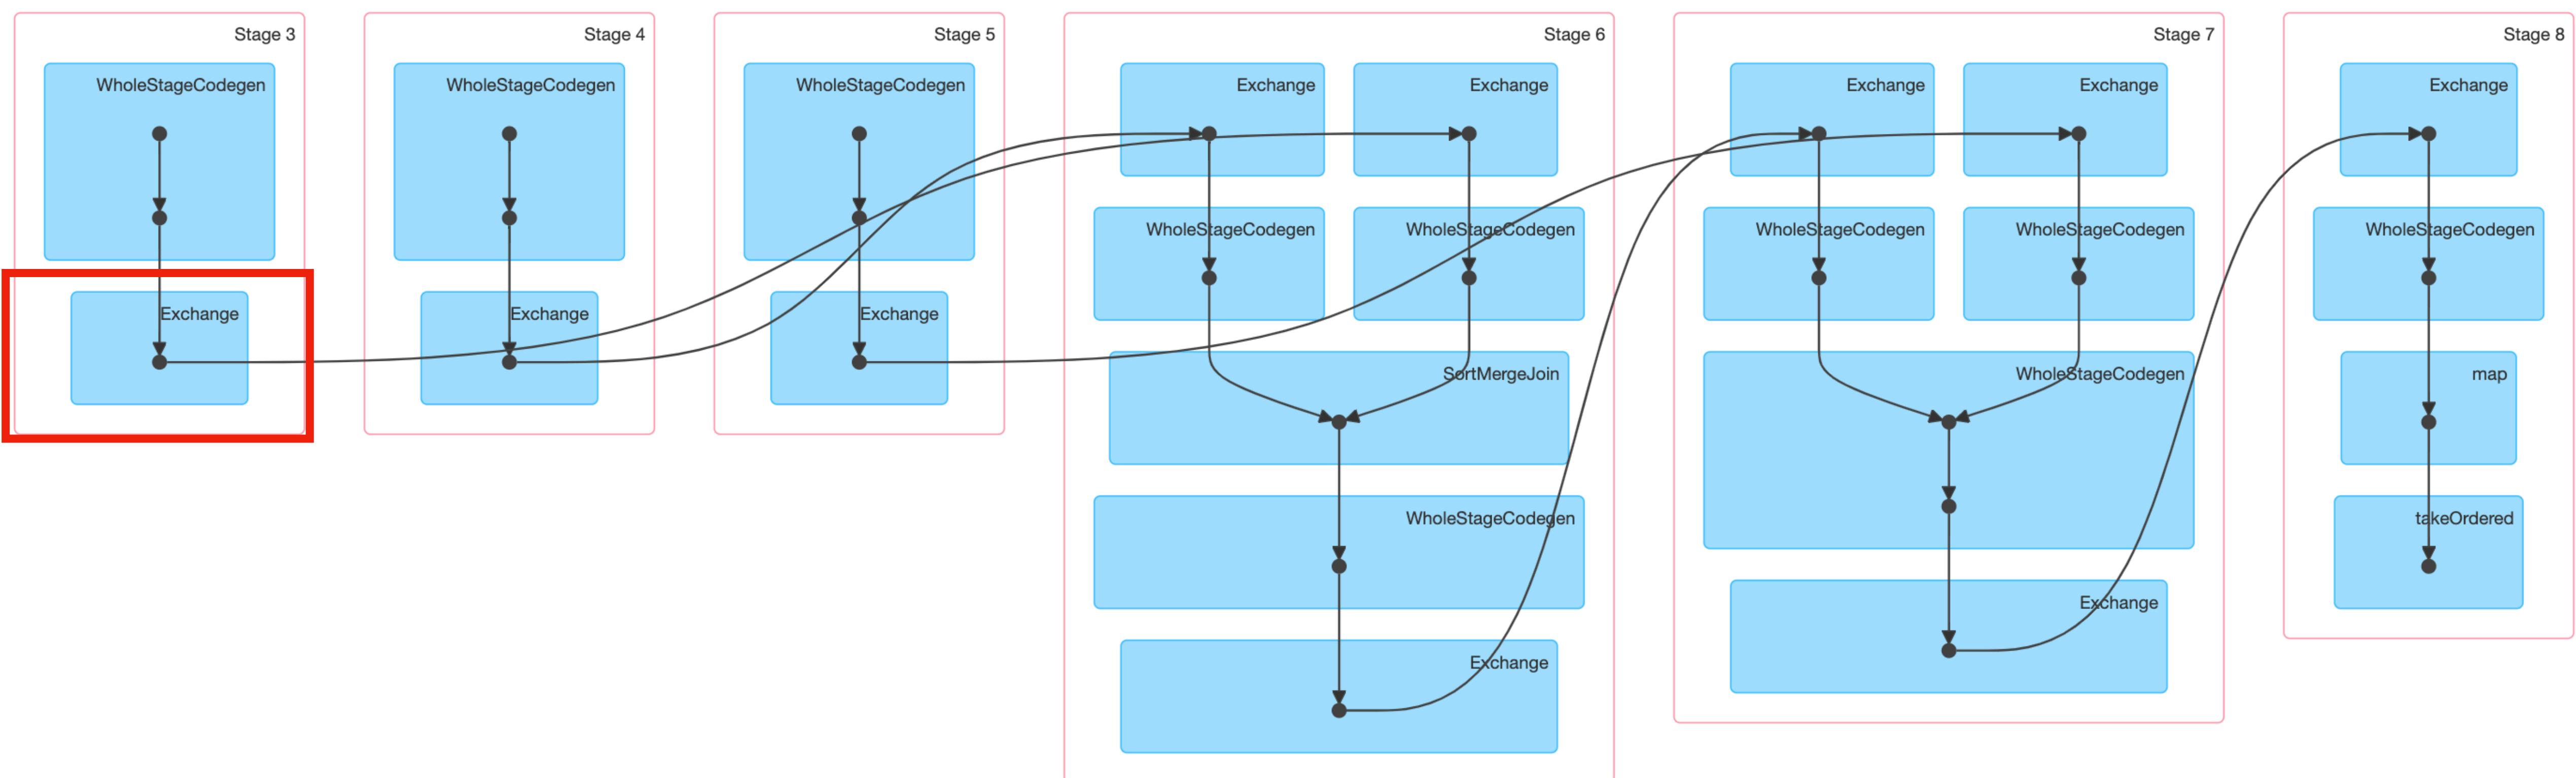


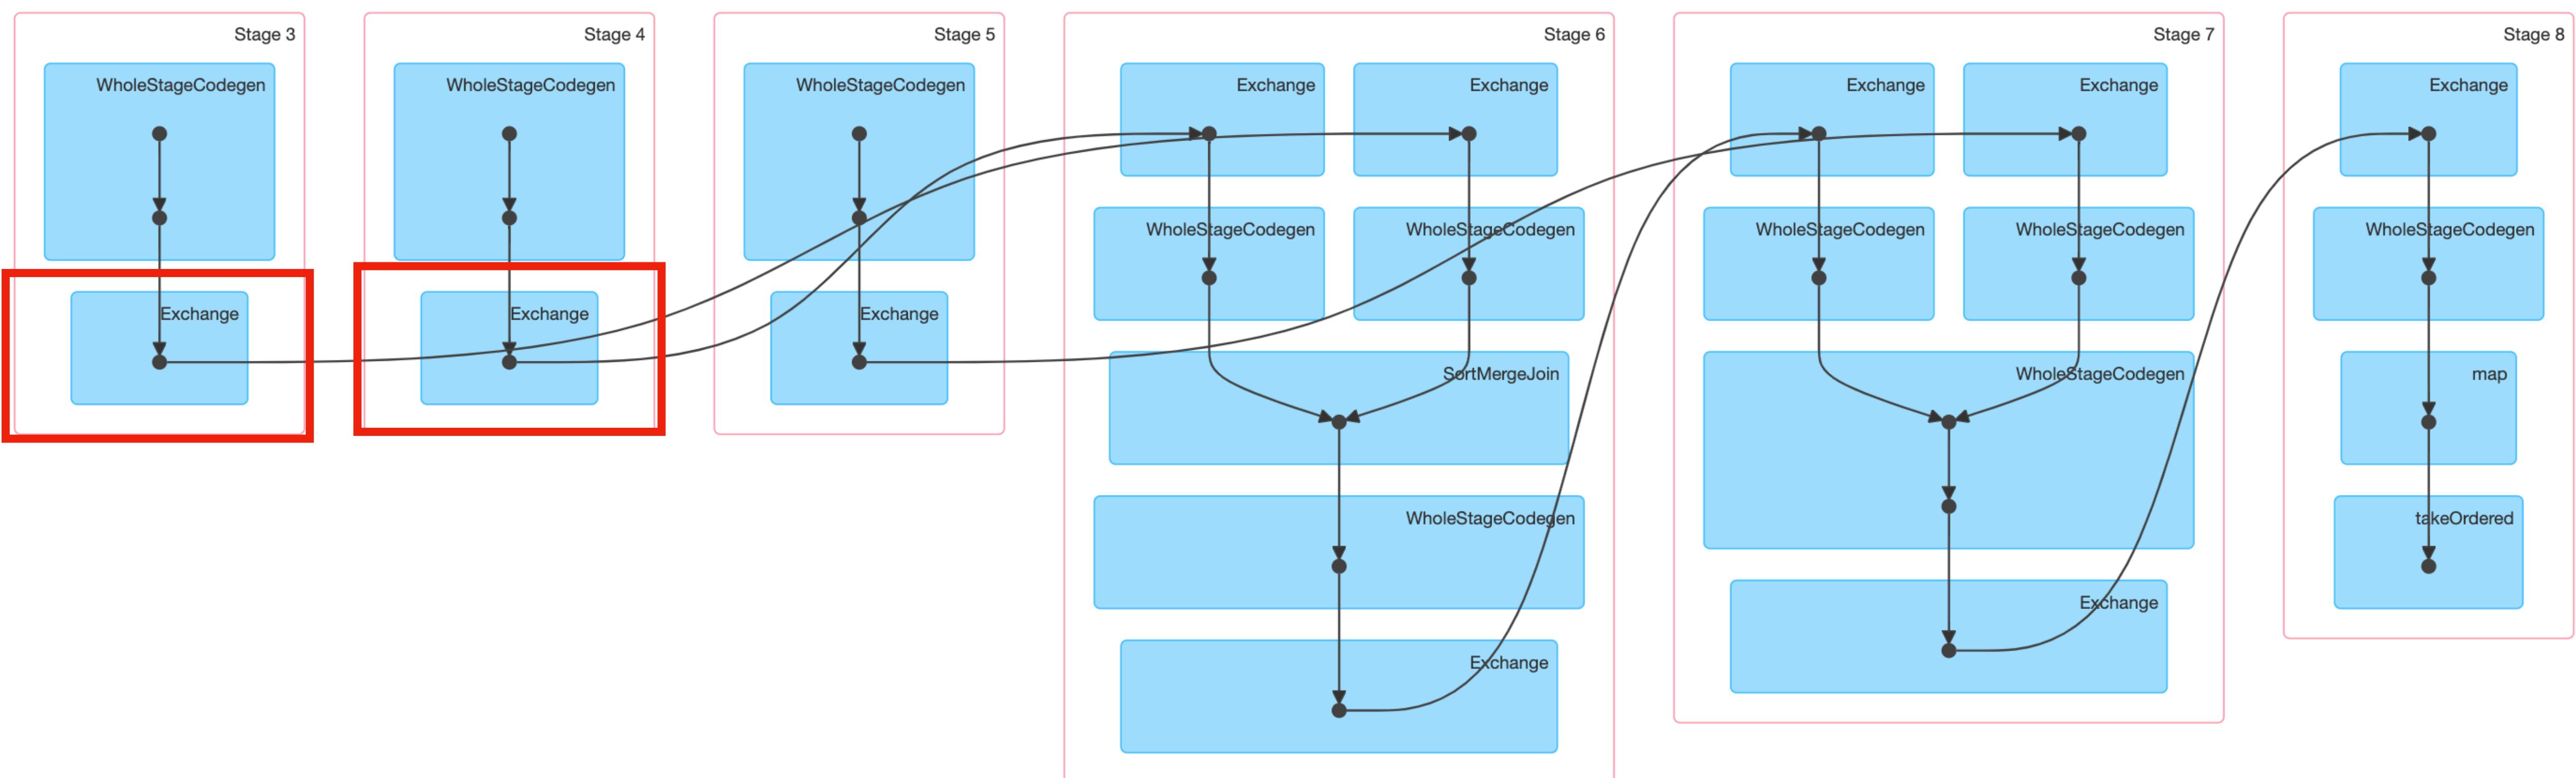


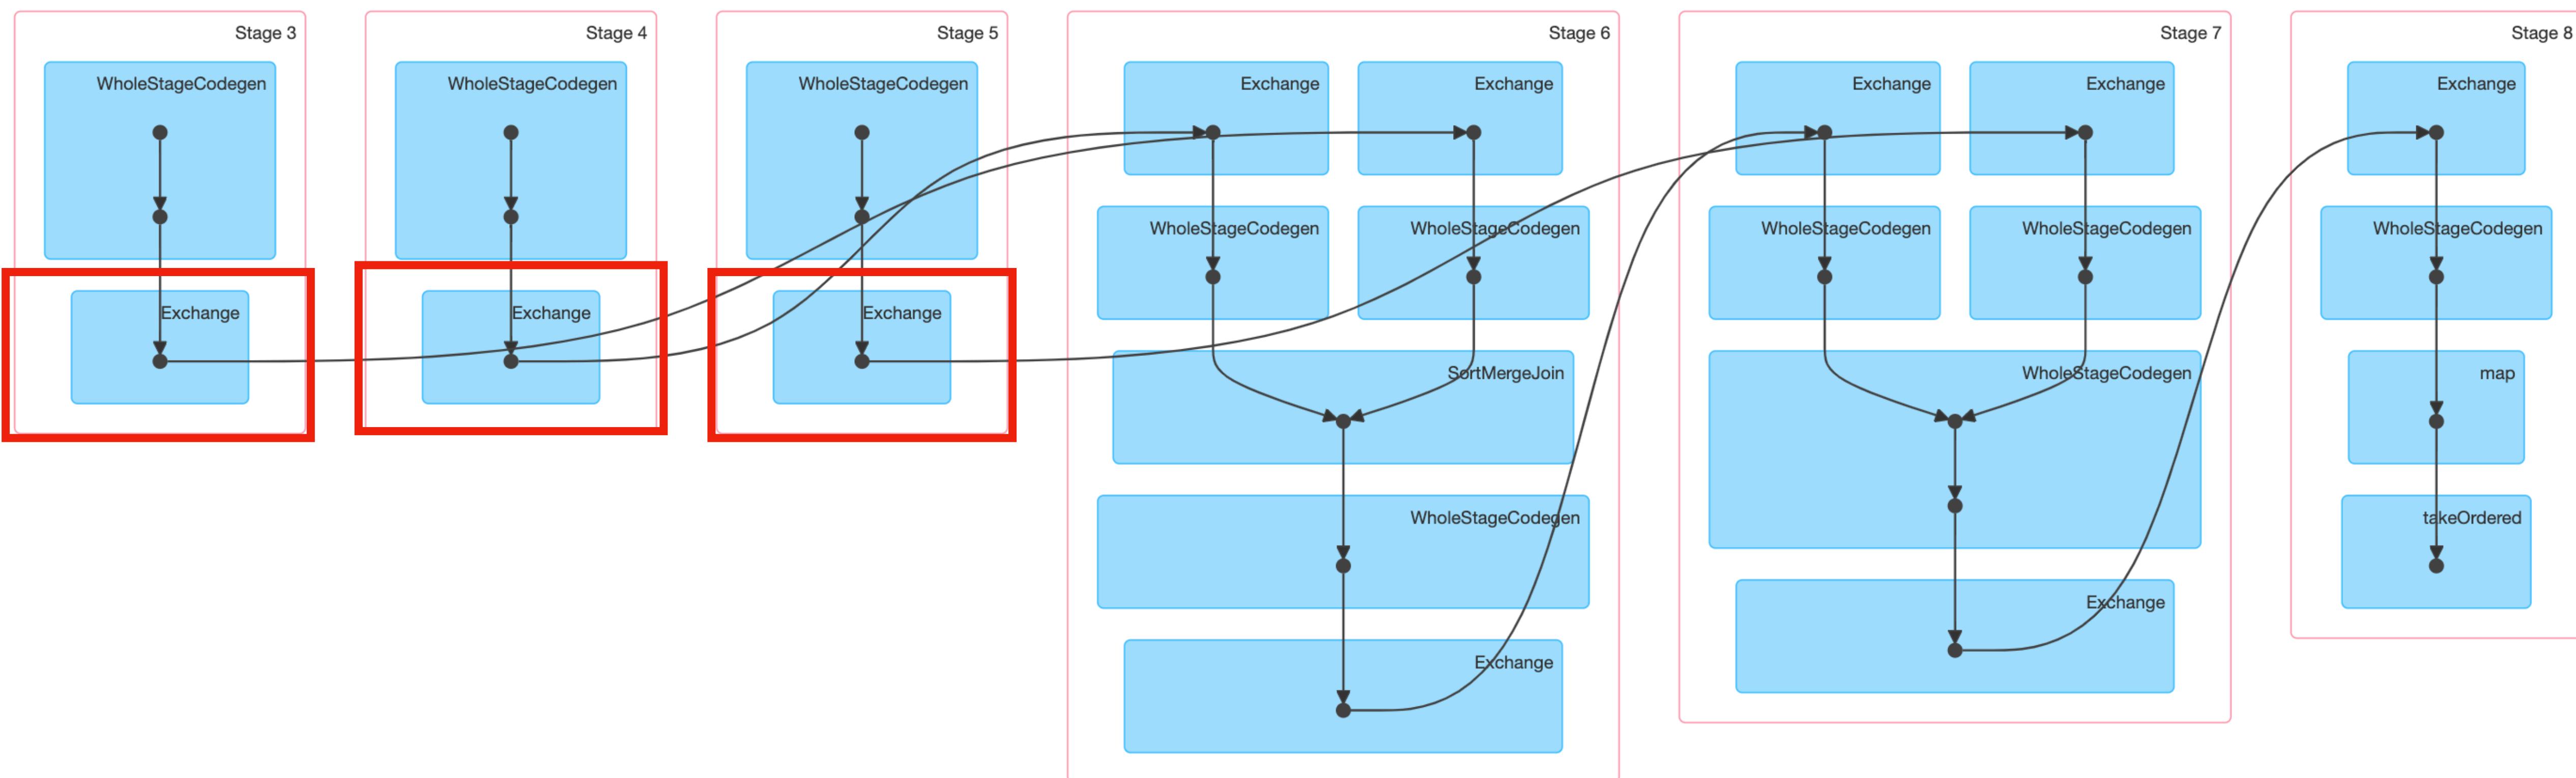


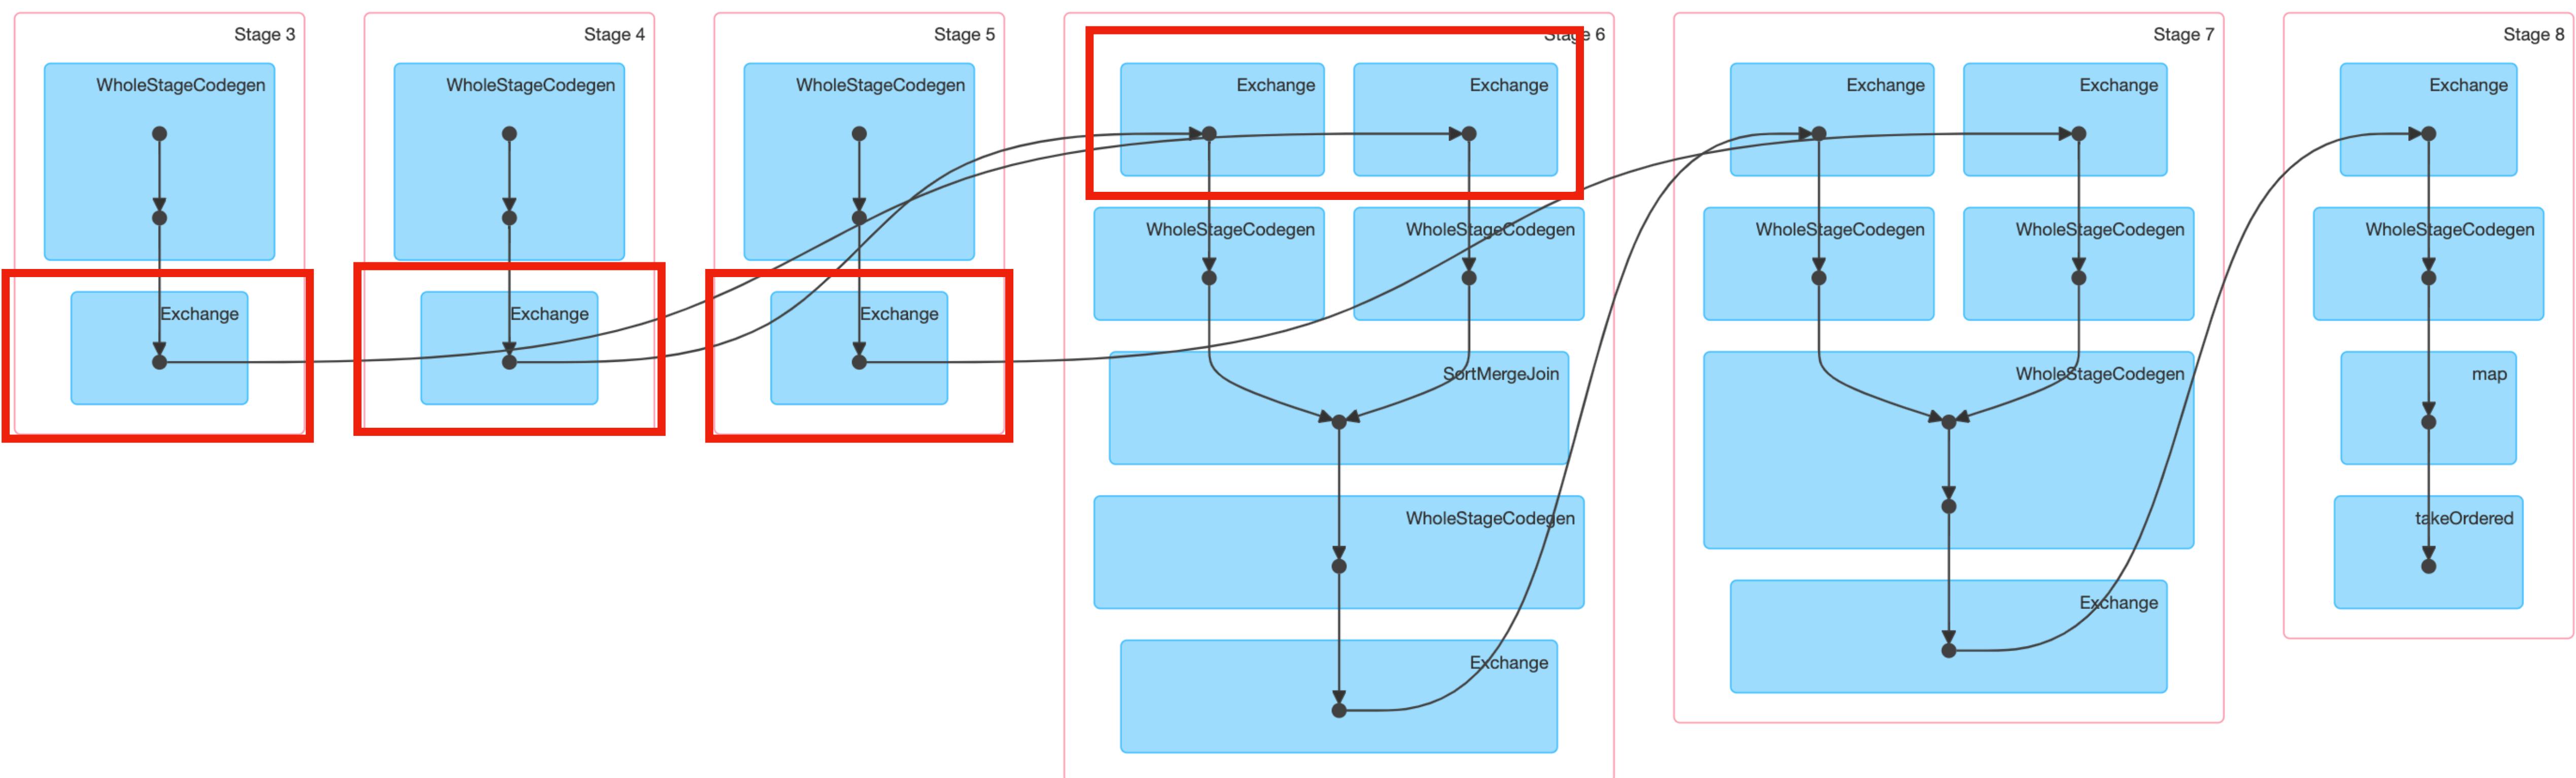


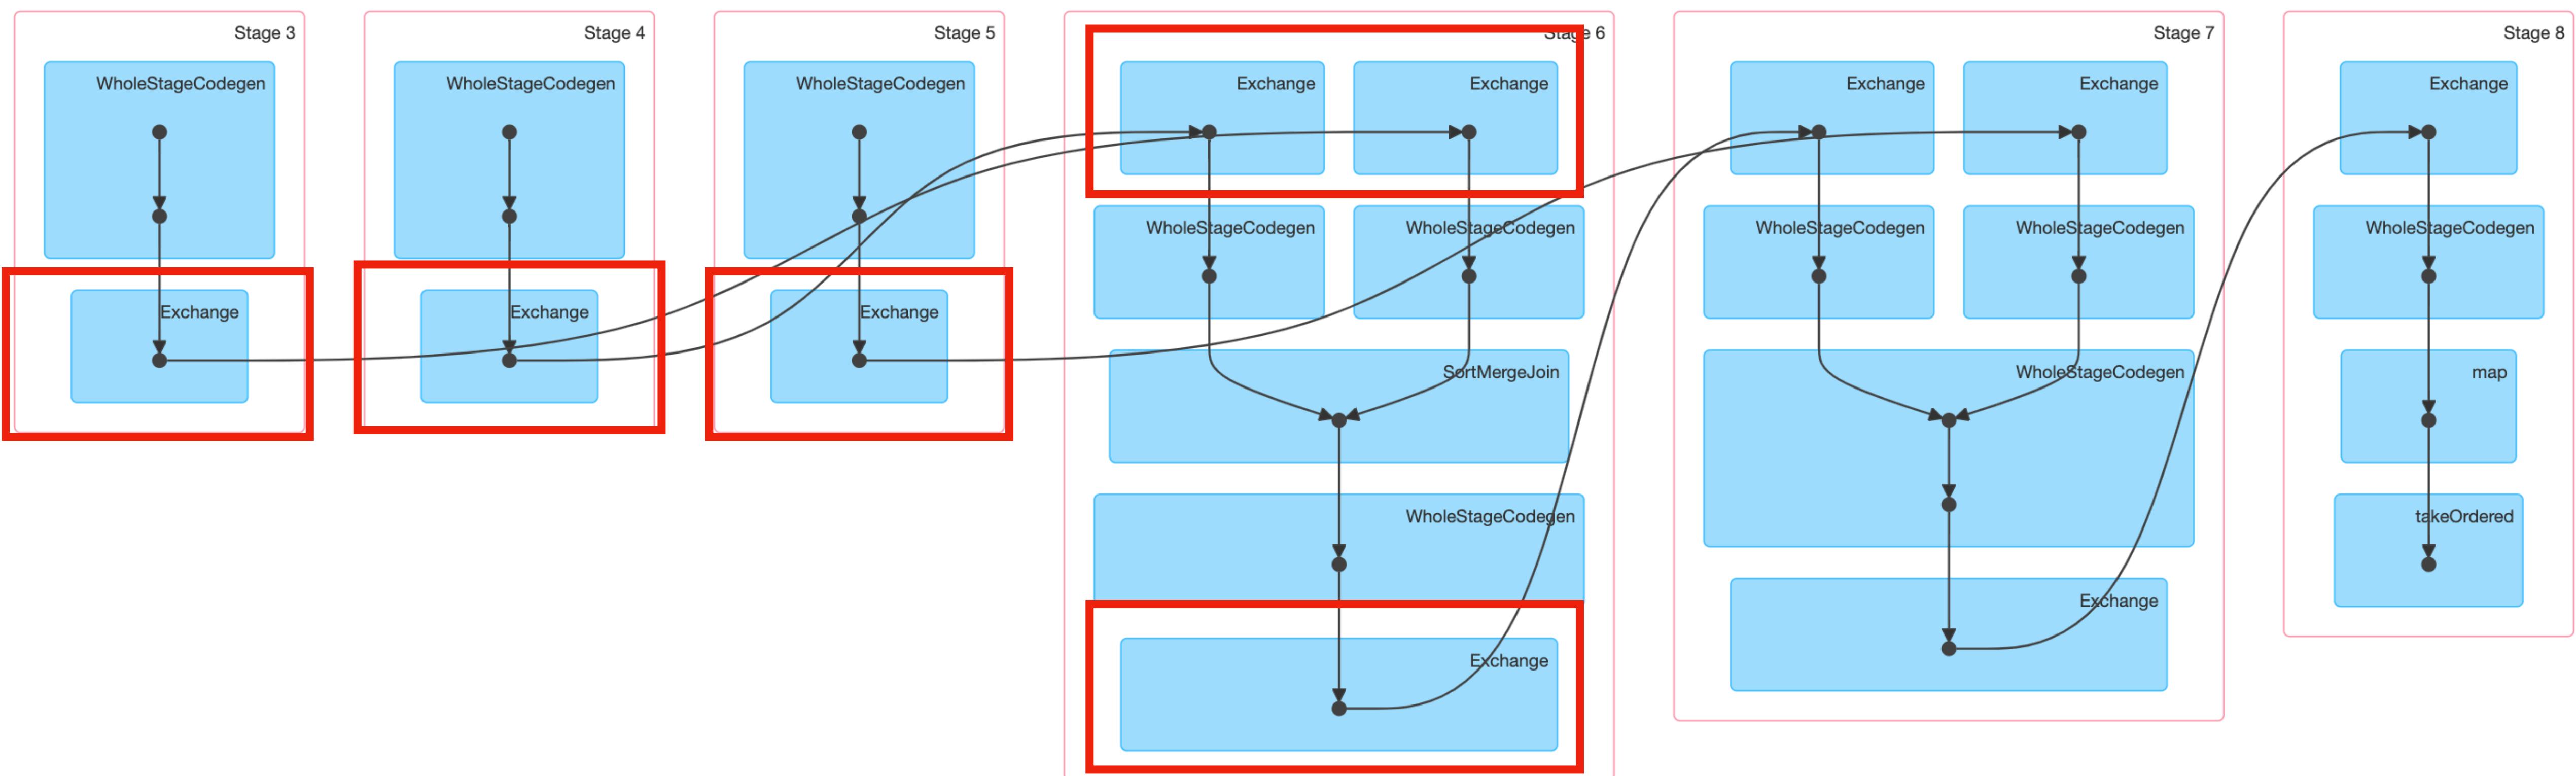


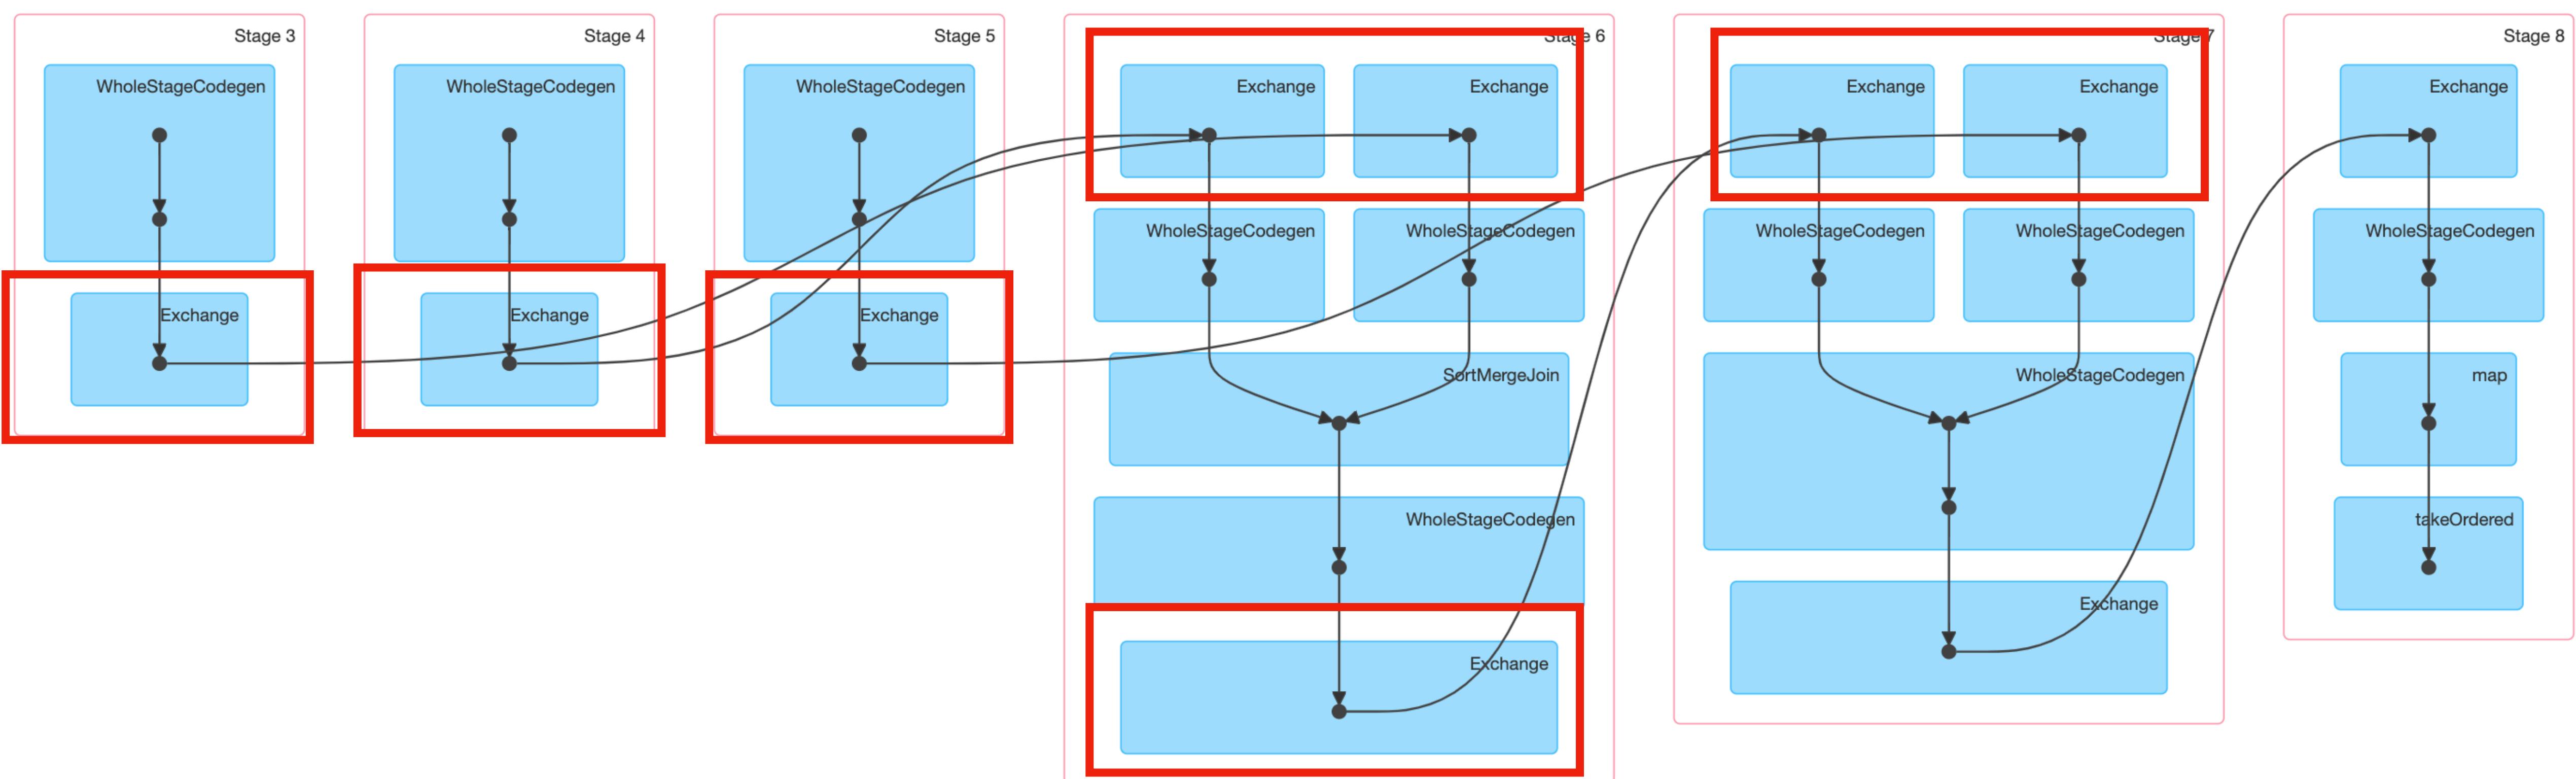


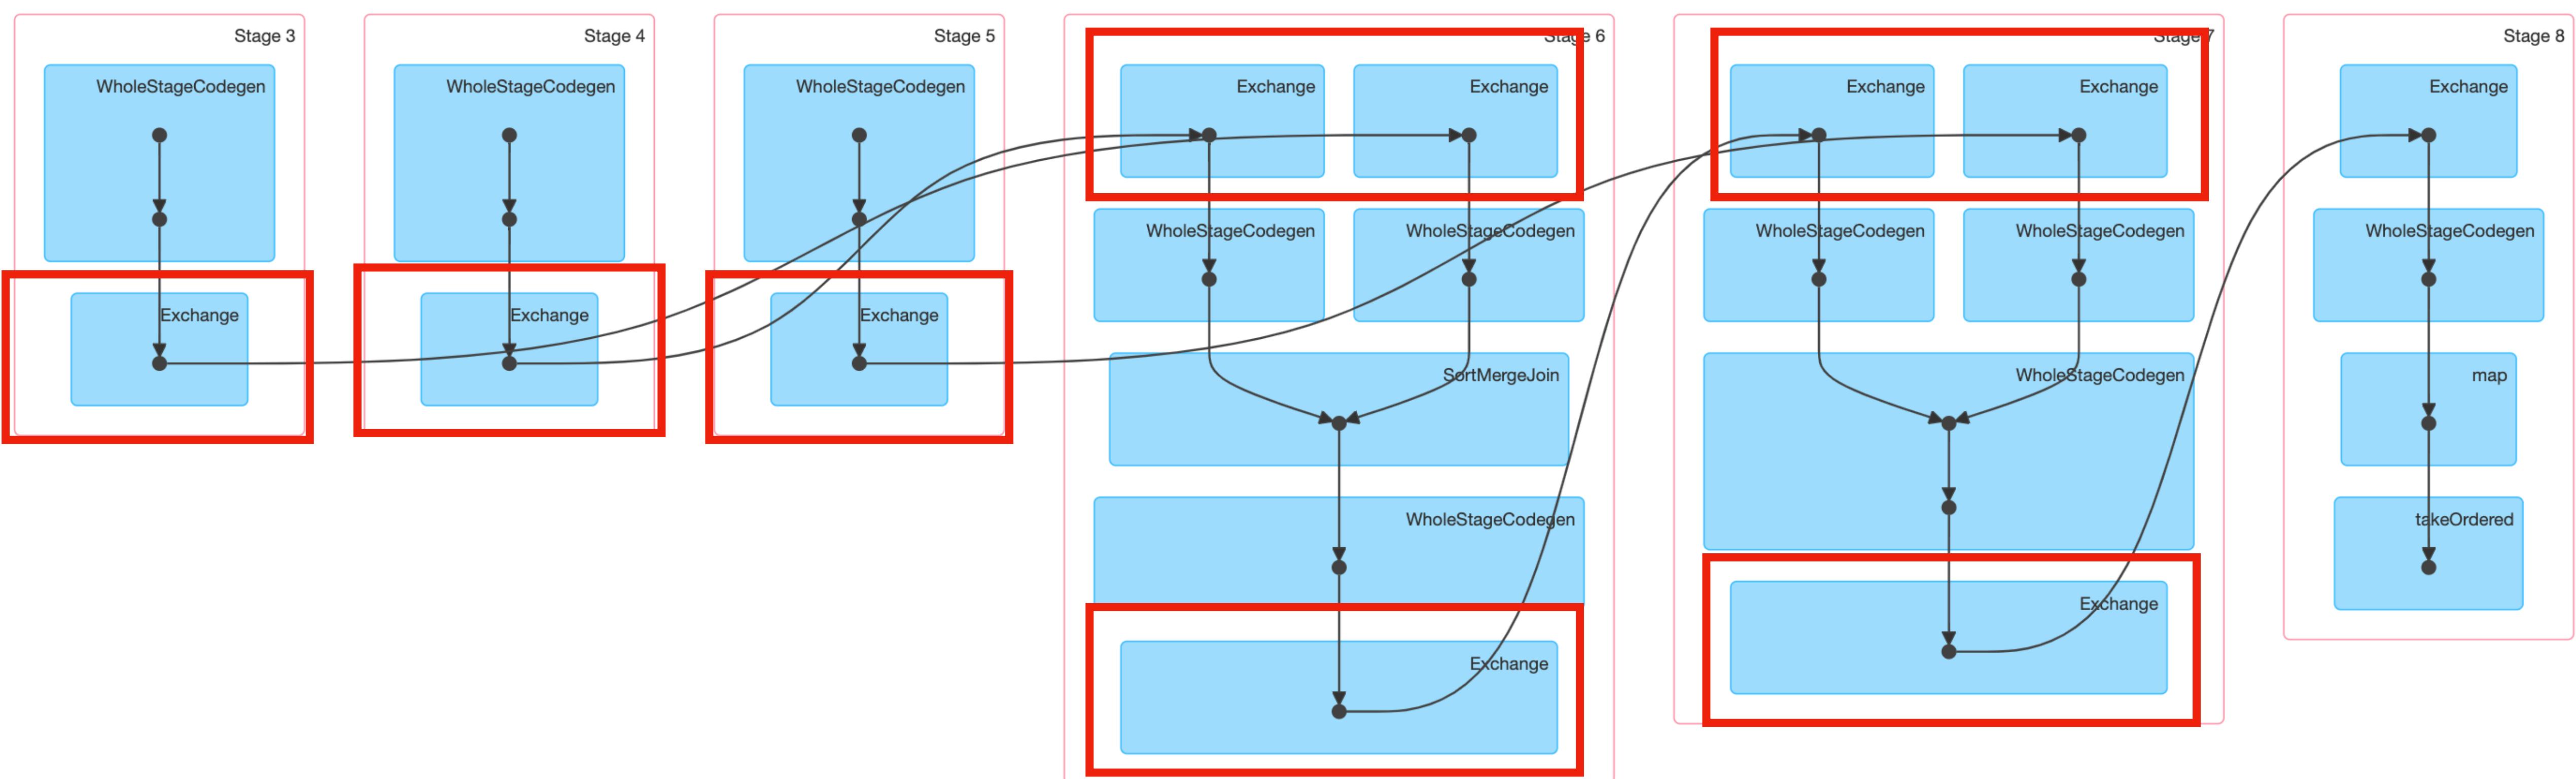


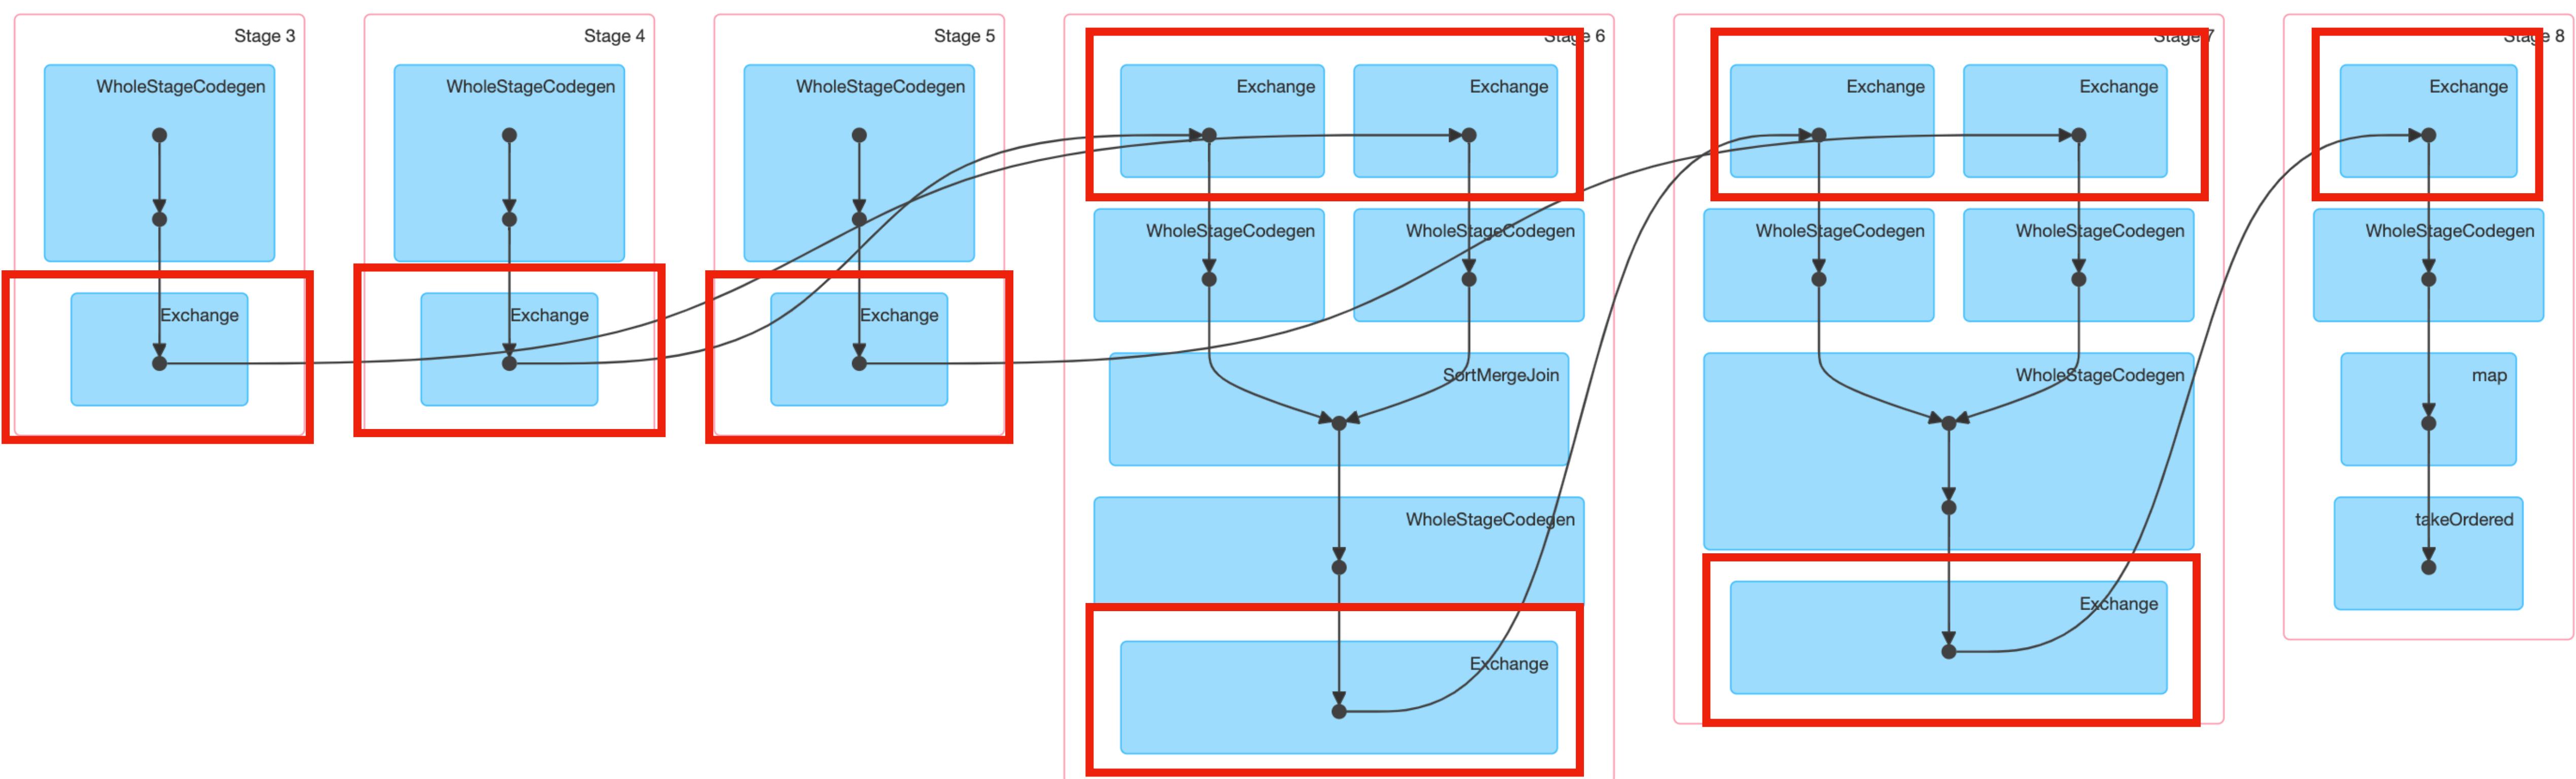












## Multiple stages of a stage:

- Pending
- Active
- Completed
- Failed

## RED FLAGS:

- Failed stages should be the first thing to check.

Completed Stages (6)	
Stage Id ▾	Description
8	showString at N
7	showString at N
6	showString at N
5	showString at N
4	showString at N
3	showString at N



Duration	Tasks: Succeeded/Total	Input	Output	Shuffle Read	Shuffle Write
0.4 s	200/200			287.4 KB	
1.6 min	200/200			25.1 GB	287.4 KB
12 min	200/200 (1 killed: another attempt succeeded)			29.7 GB	18.6 GB
3.9 min	1044/1044	16.0 GB			8.5 GB
3.3 min	14629/14629	86.2 GB			25.1 GB
12 s	413/413	213.4 MB			151.0 MB

## RED FLAGS:

- Stages with failed tasks.

`tasks.spark.task.maxFailures = 4`

Duration	Tasks: Succeeded/Total	Input	Output	Shuffle Read	Shuffle Write
0.4 s	200/200			287.4 KB	
1.6 min	200/200			25.1 GB	287.4 KB
12 min	200/200 (1 killed: another attempt succeeded)			29.7 GB	18.6 GB
3.9 min	1044/1044	16.0 GB			8.5 GB
3.3 min	14629/14629	86.2 GB			25.1 GB
12 s	413/413	213.4 MB			151.0 MB

## RED FLAGS:

- Stages with failed tasks.
- Long running stages.

Duration	Tasks: Succeeded/Total	Input	Output	Shuffle Read	Shuffle Write
0.4 s	200/200			287.4 KB	
1.6 min	200/200			25.1 GB	287.4 KB
12 min	200/200 (1 killed: another attempt succeeded)			29.7 GB	18.6 GB
3.9 min	1044/1044	16.0 GB			8.5 GB
3.3 min	14629/14629	86.2 GB			25.1 GB
12 s	413/413	213.4 MB			151.0 MB

## RED FLAGS:

- Stages with failed tasks.
- Long running stages.
- Lots of input data and data being shuffled around.

Duration	Tasks: Succeeded/Total	Input	Output	Shuffle Read	Shuffle Write
0.4 s	200/200			287.4 KB	
1.6 min	200/200			25.1 GB	287.4 KB
12 min	200/200 (1 killed: another attempt succeeded)			29.7 GB	18.6 GB
3.9 min	1044/1044	16.0 GB			8.5 GB
3.3 min	14629/14629	86.2 GB			25.1 GB
12 s	413/413	213.4 MB			151.0 MB

## RED FLAGS:

- Stages with failed tasks.
- Long running stages.
- Lots of input data and data being shuffled around.
- Low number of tasks.

Duration	Tasks: Succeeded/Total	Input	Output	Shuffle Read	Shuffle Write
0.4 s	200/200			287.4 KB	
1.6 min	200/200			25.1 GB	287.4 KB
12 min	200/200 (1 killed: another attempt succeeded)			29.7 GB	18.6 GB
3.9 min	1044/1044	16.0 GB			8.5 GB
3.3 min	14629/14629	86.2 GB			25.1 GB
12 s	413/413	213.4 MB			151.0 MB

## RED FLAGS:

- Stages with failed tasks.
- Long running stages.
- Lots of input data and data being shuffled around.
- Low number of tasks.

# Stages

- Is a group of transformations.
- Stages are separated by a shuffle of data:
  - join on x
  - group by x
  - window over x
  - etc.
- A **stage** consists of many **tasks**.

```
1 # Review of transformations in Spark:  
2 df.where(...)  
3 df.withColumn(...)  
4 df.withRenamedColumn(...)  
5  
6  
7
```

## Details for Stage 6 (Attempt 0)

**Total Time Across All Tasks:** 1.4 h

**Locality Level Summary:** Node local: 76; Rack local: 125

**Shuffle Read:** 29.7 GB / 3142989145

**Shuffle Write:** 18.6 GB / 2750912512

**Shuffle Spill (Memory):** 62.0 GB

**Shuffle Spill (Disk):** 15.1 GB

▶ [DAG Visualization](#)

▶ [Show Additional Metrics](#)

▶ [Event Timeline](#)

### Summary Metrics for 200 Completed Tasks

Metric	Min	25th percentile	Median	75th percentile	Max
Duration	0 ms	13 s	20 s	29 s	12 min
Scheduler Delay	2 ms	4 ms	5 ms	7 ms	7.5 min
Task Deserialization Time	0 ms	2 ms	3 ms	9 ms	0.2 s
GC Time	91 ms	0.4 s	4 s	8 s	27 s
Result Serialization Time	0 ms	0 ms	0 ms	0 ms	1 ms
Getting Result Time	0 ms	0 ms	0 ms	0 ms	0 ms
Peak Execution Memory	424.0 MB	520.0 MB	552.0 MB	552.0 MB	3.8 GB
Shuffle Read Size / Records	83.9 MB / 6948025	96.0 MB / 8116444	101.6 MB / 8666248	109.8 MB / 9503281	4.5 GB / 643068554
Shuffle Write Size / Records	45.9 MB / 6810020	53.6 MB / 7979098	57.1 MB / 8528805	62.7 MB / 9365457	4.6 GB / 642931319
Shuffle spill (memory)	0.0 B	0.0 B	264.0 MB	264.0 MB	16.4 GB
Shuffle spill (disk)	0.0 B	0.0 B	58.4 MB	59.2 MB	4.3 GB

### ▼ Aggregated Metrics by Executor

Executor ID ▲	Address	Task Time	Total Tasks	Failed Tasks	Killed Tasks	Succeeded Tasks	Shuffle Read Size / Records	Shuffle Write Size / Records	Shuffle Spill (Memory)	Shuffle Spill (Disk)	Blacklisted
1 stdout stderr	nm-yarn-adhoc-tbm3.yarn-adhoc.starscream-adhoc.yarn.shopifykloud.com:35955	25 min	76	0	1	75	12.5 GB / 1360445397	6.6 GB / 985547990	28.4 GB	6.9 GB	false
2 stdout stderr	nm-yarn-adhoc-5s0m.yarn-adhoc.starscream-adhoc.yarn.shopifykloud.com:37695	8.2 min	23	0	0	23	2.6 GB / 245329024	1618.2 MB / 242167060	4.9 GB	1109.6 MB	false
3 stdout stderr	nm-yarn-adhoc-zq40.yarn-adhoc.starscream-adhoc.yarn.shopifykloud.com:36223	7.1 min	21	0	0	21	2.0 GB / 177559979	1172.5 MB / 174673800	1584.0 MB	354.9 MB	false
4 stdout stderr	nm-yarn-adhoc-32lx.yarn-adhoc.starscream-adhoc.yarn.shopifykloud.com:35857	19 min	16	0	0	16	6.1 GB / 785470154	5.5 GB / 783269709	18.4 GB	4.7 GB	false
5 stdout stderr	nm-yarn-adhoc-32lx.yarn-adhoc.starscream-adhoc.yarn.shopifykloud.com:40177	8.0 min	16	0	0	16	1656.5 MB / 143606734	948.9 MB / 141406594	1848.0 MB	413.0 MB	false
6 stdout stderr	nm-yarn-adhoc-32lx.yarn-adhoc.starscream-adhoc.yarn.shopifykloud.com:35105	7.9 min	16	0	0	16	1623.8 MB / 139477928	918.1 MB / 137279959	2.6 GB	591.2 MB	false
7 stdout stderr	nm-yarn-adhoc-32lx.yarn-adhoc.starscream-adhoc.yarn.shopifykloud.com:36969	7.3 min	17	0	0	17	1730.3 MB / 149373384	986.8 MB / 147038438	1848.0 MB	414.5 MB	false
8 stdout stderr	nm-yarn-adhoc-32lx.yarn-adhoc.starscream-adhoc.yarn.shopifykloud.com:34453	7.9 min	16	0	0	16	1647.4 MB / 141726545	936.3 MB / 139528962	2.6 GB	591.7 MB	false

## Details for Stage 6 (Attempt 0)

**Total Time Across All Tasks:** 1.4 h

**Locality Level Summary:** Node local: 76; Rack local: 125

**Shuffle Read:** 29.7 GB / 3142989145

**Shuffle Write:** 18.6 GB / 2750912512

**Shuffle Spill (Memory):** 62.0 GB

**Shuffle Spill (Disk):** 15.1 GB

### RED FLAGS:

- Having spill in your stage.

# Stage Level Metrics

Summary Metrics for 200 Completed Tasks

Metric	Min	25th percentile	Median	75th percentile	Max
Duration	0 ms	13 s	20 s	29 s	12 min
Scheduler Delay	2 ms	4 ms	5 ms	7 ms	7.5 min
Task Deserialization Time	0 ms	2 ms	3 ms	9 ms	0.2 s
GC Time	91 ms	0.4 s	4 s	8 s	27 s
Result Serialization Time	0 ms	0 ms	0 ms	0 ms	1 ms
Getting Result Time	0 ms	0 ms	0 ms	0 ms	0 ms
Peak Execution Memory	424.0 MB	520.0 MB	552.0 MB	552.0 MB	3.8 GB
Shuffle Read Size / Records	83.9 MB / 6948025	96.0 MB / 8116444	101.6 MB / 8666248	109.8 MB / 9503281	4.5 GB / 643068554
Shuffle Write Size / Records	45.9 MB / 6810020	53.6 MB / 7979098	57.1 MB / 8528805	62.7 MB / 9365457	4.6 GB / 642931319
Shuffle spill (memory)	0.0 B	0.0 B	264.0 MB	264.0 MB	16.4 GB
Shuffle spill (disk)	0.0 B	0.0 B	58.4 MB	59.2 MB	4.3 GB

## RED FLAGS:

- Having spill in your stage.
- Big difference in the min and max time and memory.

# Stage Level Metrics

Summary Metrics for 200 Completed Tasks

Metric	Min	25th percentile	Median	75th percentile	Max
Duration	0 ms	13 s	20 s	29 s	12 min
Scheduler Delay	2 ms	4 ms	5 ms	7 ms	7.5 min
Task Deserialization Time	0 ms	2 ms	3 ms	9 ms	0.2 s
GC Time	91 ms	0.4 s	4 s	8 s	27 s
Result Serialization Time	0 ms	0 ms	0 ms	0 ms	1 ms
Getting Result Time	0 ms	0 ms	0 ms	0 ms	0 ms
Peak Execution Memory	424.0 MB	520.0 MB	552.0 MB	552.0 MB	3.8 GB
Shuffle Read Size / Records	83.9 MB / 6948025	96.0 MB / 8116444	101.6 MB / 8666248	109.8 MB / 9503281	4.5 GB / 643068554
Shuffle Write Size / Records	45.9 MB / 6810020	53.6 MB / 7979098	57.1 MB / 8528805	62.7 MB / 9365457	4.6 GB / 642931319
Shuffle spill (memory)	0.0 B	0.0 B	264.0 MB	264.0 MB	16.4 GB
Shuffle spill (disk)	0.0 B	0.0 B	58.4 MB	59.2 MB	4.3 GB

## RED FLAGS:

- Having spill in your stage.
- Big difference in the min and max time and memory.

# Stage Level Metrics

Summary Metrics for 200 Completed Tasks

Metric	Min	25th percentile	Median	75th percentile	Max
Duration	0 ms	13 s	20 s	29 s	12 min
Scheduler Delay	2 ms	4 ms	5 ms	7 ms	7.5 min
Task Deserialization Time	0 ms	2 ms	3 ms	9 ms	0.2 s
GC Time	91 ms	0.4 s	4 s	8 s	27 s
Result Serialization Time	0 ms	0 ms	0 ms	0 ms	1 ms
Getting Result Time	0 ms	0 ms	0 ms	0 ms	0 ms
Peak Execution Memory	424.0 MB	520.0 MB	552.0 MB	552.0 MB	3.8 GB
Shuffle Read Size / Records	83.9 MB / 6948025	96.0 MB / 8116444	101.6 MB / 8666248	109.8 MB / 9503281	4.5 GB / 643068554
Shuffle Write Size / Records	45.9 MB / 6810020	53.6 MB / 7979098	57.1 MB / 8528805	62.7 MB / 9365457	4.6 GB / 642931319
Shuffle spill (memory)	0.0 B	0.0 B	264.0 MB	264.0 MB	16.4 GB
Shuffle spill (disk)	0.0 B	0.0 B	58.4 MB	59.2 MB	4.3 GB

## RED FLAGS:

- Having spill in your stage.
- Big difference in the min and max time and memory.

# Executor Level Metrics

## ▼ Aggregated Metrics by Executor

Executor ID ▲	Address	Task Time	Total Tasks	Failed Tasks	Killed Tasks	Succeeded Tasks
1 stdout stderr	nm-yarn-adhoc-tbm3.yarn-adhoc.starscream-adhoc.yarn.shopifykloud.com:35955	25 min	76	0	1	75
2 stdout stderr	nm-yarn-adhoc-5s0m.yarn-adhoc.starscream-adhoc.yarn.shopifykloud.com:37695	8.2 min	23	0	0	23
3 stdout stderr	nm-yarn-adhoc-zq40.yarn-adhoc.starscream-adhoc.yarn.shopifykloud.com:36223	7.1 min	21	0	0	21
4 stdout stderr	nm-yarn-adhoc-32lx.yarn-adhoc.starscream-adhoc.yarn.shopifykloud.com:35857	19 min	16	0	0	16
5 stdout stderr	nm-yarn-adhoc-32lx.yarn-adhoc.starscream-adhoc.yarn.shopifykloud.com:40177	8.0 min	16	0	0	16
6 stdout stderr	nm-yarn-adhoc-32lx.yarn-adhoc.starscream-adhoc.yarn.shopifykloud.com:35105	7.9 min	16	0	0	16
7 stdout stderr	nm-yarn-adhoc-32lx.yarn-adhoc.starscream-adhoc.yarn.shopifykloud.com:36969	7.3 min	17	0	0	17
8 stdout stderr	nm-yarn-adhoc-32lx.yarn-adhoc.starscream-adhoc.yarn.shopifykloud.com:34453	7.9 min	16	0	0	16

## RED FLAGS:

- Having spill in your stage.
- Big difference in the min and max time and memory.
- Uneven distribution of work per executor.

# Executor Level Metrics

## ▼ Aggregated Metrics by Executor

Executor ID ▲	Address	Task Time	Total Tasks	Failed Tasks	Killed Tasks	Succeeded Tasks
1 stdout stderr	nm-yarn-adhoc-tbm3.yarn-adhoc.starscream-adhoc.yarn.shopifykloud.com:35955	25 min	76	0	1	75
2 stdout stderr	nm-yarn-adhoc-5s0m.yarn-adhoc.starscream-adhoc.yarn.shopifykloud.com:37695	8.2 min	23	0	0	23
3 stdout stderr	nm-yarn-adhoc-zq40.yarn-adhoc.starscream-adhoc.yarn.shopifykloud.com:36223	7.1 min	21	0	0	21
4 stdout stderr	nm-yarn-adhoc-32lx.yarn-adhoc.starscream-adhoc.yarn.shopifykloud.com:35857	19 min	16	0	0	16
5 stdout stderr	nm-yarn-adhoc-32lx.yarn-adhoc.starscream-adhoc.yarn.shopifykloud.com:40177	8.0 min	16	0	0	16
6 stdout stderr	nm-yarn-adhoc-32lx.yarn-adhoc.starscream-adhoc.yarn.shopifykloud.com:35105	7.9 min	16	0	0	16
7 stdout stderr	nm-yarn-adhoc-32lx.yarn-adhoc.starscream-adhoc.yarn.shopifykloud.com:36969	7.3 min	17	0	0	17
8 stdout stderr	nm-yarn-adhoc-32lx.yarn-adhoc.starscream-adhoc.yarn.shopifykloud.com:34453	7.9 min	16	0	0	16

## RED FLAGS:

- Having spill in your stage.
- Big difference in the min and max time and memory.
- Uneven distribution of work per executor.

# Executor Level Metrics

## ▼ Aggregated Metrics by Executor

Executor ID ▲	Address	Task Time	Total Tasks	Failed Tasks	Killed Tasks	Succeeded Tasks
1 stdout stderr	nm-yarn-adhoc-tbm3.yarn-adhoc.starscream-adhoc.yarn.shopifykloud.com:35955	25 min	76	0	1	75
2 stdout stderr	nm-yarn-adhoc-5s0m.yarn-adhoc.starscream-adhoc.yarn.shopifykloud.com:37695	8.2 min	23	0	0	23
3 stdout stderr	nm-yarn-adhoc-zq40.yarn-adhoc.starscream-adhoc.yarn.shopifykloud.com:36223	7.1 min	21	0	0	21
4 stdout stderr	nm-yarn-adhoc-32lx.yarn-adhoc.starscream-adhoc.yarn.shopifykloud.com:35857	19 min	16	0	0	16
5 stdout stderr	nm-yarn-adhoc-32lx.yarn-adhoc.starscream-adhoc.yarn.shopifykloud.com:40177	8.0 min	16	0	0	16
6 stdout stderr	nm-yarn-adhoc-32lx.yarn-adhoc.starscream-adhoc.yarn.shopifykloud.com:35105	7.9 min	16	0	0	16
7 stdout stderr	nm-yarn-adhoc-32lx.yarn-adhoc.starscream-adhoc.yarn.shopifykloud.com:36969	7.3 min	17	0	0	17
8 stdout stderr	nm-yarn-adhoc-32lx.yarn-adhoc.starscream-adhoc.yarn.shopifykloud.com:34453	7.9 min	16	0	0	16

## RED FLAGS:

- Having spill in your stage.
- Big difference in the min and max time and memory.
- Uneven distribution of work per executor.

# Executor Level Metrics

## ▼ Aggregated Metrics by Executor

Executor ID ▲	Address	Task Time	Total Tasks	Failed Tasks	Killed Tasks	Succeeded Tasks
1 stdout stderr	nm-yarn-adhoc-tbm3.yarn-adhoc.starscream-adhoc.yarn.shopifykloud.com:35955	25 min	76	0	1	75
2 stdout stderr	nm-yarn-adhoc-5s0m.yarn-adhoc.starscream-adhoc.yarn.shopifykloud.com:37695	8.2 min	23	0	0	23
3 stdout stderr	nm-yarn-adhoc-zq40.yarn-adhoc.starscream-adhoc.yarn.shopifykloud.com:36223	7.1 min	21	0	0	21
4 stdout stderr	nm-yarn-adhoc-32lx.yarn-adhoc.starscream-adhoc.yarn.shopifykloud.com:35857	19 min	16	0	0	16
5 stdout stderr	nm-yarn-adhoc-32lx.yarn-adhoc.starscream-adhoc.yarn.shopifykloud.com:40177	8.0 min	16	0	0	16
6 stdout stderr	nm-yarn-adhoc-32lx.yarn-adhoc.starscream-adhoc.yarn.shopifykloud.com:35105	7.9 min	16	0	0	16
7 stdout stderr	nm-yarn-adhoc-32lx.yarn-adhoc.starscream-adhoc.yarn.shopifykloud.com:36969	7.3 min	17	0	0	17
8 stdout stderr	nm-yarn-adhoc-32lx.yarn-adhoc.starscream-adhoc.yarn.shopifykloud.com:34453	7.9 min	16	0	0	16

## RED FLAGS:

- Having spill in your stage.
- Big difference in the min and max time and memory.
- Uneven distribution of work per executor.

# Task Level Metrics

## Tasks (201)

Index ▲	ID	Attempt	Status	Locality Level	Executor ID	Host	Launch Time	Duration	Scheduler Delay	Task Deserialization Time	GC Time	Result Serialization Time	Getting Result Time	Peak Execution Memory	Shuffle Read Size / Records	Write Time	Shuffle Write Size / Records	Shuffle Spill (Memory)	Shuffle Spill (Disk)	Errors
0	16089	0	SUCCESS	RACK_LOCAL	4	nm-yarn-adhoc-32lx.yarn-adhoc.starscream-adhoc.yarn.shopifykcloud.com	2019/07/25 14:23:12	32 s	6 ms	0.1 s	14 s	0 ms	0 ms	520.0 MB	90.3 MB / 7515616	0.8 s	50.3 MB / 7378423	0.0 B	0.0 B	
1	16090	0	SUCCESS	RACK_LOCAL	7	nm-yarn-adhoc-32lx.yarn-adhoc.starscream-adhoc.yarn.shopifykcloud.com	2019/07/25 14:23:12	30 s	2 ms	0.1 s	15 s	0 ms	0 ms	520.0 MB	93.3 MB / 8253612	0.7 s	53.8 MB / 8116264	0.0 B	0.0 B	
2	16091	0	SUCCESS	NODE_LOCAL	1	nm-yarn-adhoc-tbm3.yarn-adhoc.starscream-adhoc.yarn.shopifykcloud.com	2019/07/25 14:23:12	26 s	4 ms	0.1 s	5 s	0 ms	0 ms	552.0 MB	122.2 MB / 11025909	0.3 s	73.8 MB / 10888717	264.0 MB	59.4 MB	
3	16092	0	SUCCESS	NODE_LOCAL	1	nm-yarn-adhoc-tbm3.yarn-adhoc.starscream-adhoc.yarn.shopifykcloud.com	2019/07/25 14:23:12	23 s	7 ms	95 ms	4 s	0 ms	0 ms	552.0 MB	117.1 MB / 10221419	0.3 s	66.9 MB / 10083829	264.0 MB	59.0 MB	
4	16093	0	SUCCESS	NODE_LOCAL	1	nm-yarn-adhoc-tbm3.yarn-adhoc.starscream-adhoc.yarn.shopifykcloud.com	2019/07/25 14:23:12	20 s	3 ms	89 ms	3 s	0 ms	0 ms	552.0 MB	101.3 MB / 8654774	0.8 s	57.1 MB / 8517328	264.0 MB	59.1 MB	
5	16094	0	SUCCESS	NODE_LOCAL	1	nm-yarn-adhoc-tbm3.yarn-adhoc.starscream-adhoc.yarn.shopifykcloud.com	2019/07/25 14:23:12	21 s	4 ms	48 ms	3 s	0 ms	0 ms	552.0 MB	122.1 MB / 10710569	0.3 s	71.1 MB / 10573267	264.0 MB	59.4 MB	
6	16095	0	SUCCESS	NODE_LOCAL	1	nm-yarn-adhoc-tbm3.yarn-adhoc.starscream-adhoc.yarn.shopifykcloud.com	2019/07/25 14:23:12	17 s	4 ms	2 ms	3 s	0 ms	0 ms	520.0 MB	93.6 MB / 7715981	0.2 s	51.5 MB / 7578841	0.0 B	0.0 B	

Duration	Scheduler Delay	Task Deserialization Time	GC Time	Result Serialization Time	Getting Result Time	Peak Execution Memory	Shuffle Read Size / Records
32 s	6 ms	0.1 s	14 s	0 ms	0 ms	520.0 MB	90.3 MB / 7515616
30 s	2 ms	0.1 s	15 s	0 ms	0 ms	520.0 MB	93.3 MB / 8253612
26 s	4 ms	0.1 s	5 s	0 ms	0 ms	552.0 MB	122.2 MB / 11025909
23 s	7 ms	95 ms	4 s	0 ms	0 ms	552.0 MB	117.1 MB / 10221419
20 s	3 ms	89 ms	3 s	0 ms	0 ms	552.0 MB	101.3 MB / 8654774
21 s	4 ms	48 ms	3 s	0 ms	0 ms	552.0 MB	122.1 MB / 10710569

## RED FLAGS:

- Having spill in your stage.
- Big difference in the min and max time and memory.
- Uneven distribution of work per executor.
- Tasks with more than 128MB of Shuffle Read data.

Duration	Scheduler Delay	Task Deserialization Time	GC Time	Result Serialization Time	Getting Result Time	Peak Execution Memory	Shuffle Read Size / Records
32 s	6 ms	0.1 s	14 s	0 ms	0 ms	520.0 MB	90.3 MB / 7515616
30 s	2 ms	0.1 s	15 s	0 ms	0 ms	520.0 MB	93.3 MB / 8253612
26 s	4 ms	0.1 s	5 s	0 ms	0 ms	552.0 MB	122.2 MB / 11025909
23 s	7 ms	95 ms	4 s	0 ms	0 ms	552.0 MB	117.1 MB / 10221419
20 s	3 ms	89 ms	3 s	0 ms	0 ms	552.0 MB	101.3 MB / 8654774
21 s	4 ms	48 ms	3 s	0 ms	0 ms	552.0 MB	122.1 MB / 10710569

## RED FLAGS:

- Having spill in your stage.
- Big difference in the min and max time and memory.
- Uneven distribution of work per executor.
- Tasks with more than 128MB of Shuffle Read data.

Duration	Scheduler Delay	Task Deserialization Time	GC Time	Result Serialization Time	Getting Result Time	Peak Execution Memory	Shuffle Read Size / Records
32 s	6 ms	0.1 s	14 s	0 ms	0 ms	520.0 MB	90.3 MB / 7515616
30 s	2 ms	0.1 s	15 s	0 ms	0 ms	520.0 MB	93.3 MB / 8253612
26 s	4 ms	0.1 s	5 s	0 ms	0 ms	552.0 MB	122.2 MB / 11025909
23 s	7 ms	95 ms	4 s	0 ms	0 ms	552.0 MB	117.1 MB / 10221419
20 s	3 ms	89 ms	3 s	0 ms	0 ms	552.0 MB	101.3 MB / 8654774
21 s	4 ms	48 ms	3 s	0 ms	0 ms	552.0 MB	122.1 MB / 10710569

## RED FLAGS:

- Having spill in your stage.
- Big difference in the min and max time and memory.
- Uneven distribution of work per executor.
- Tasks with more than 128MB of Shuffle Read data.

# After Skew Fix

Summary Metrics for 200 Completed Tasks

Metric	Min	25th percentile	Median	75th percentile	Max
Duration	14 s	16 s	20 s	25 s	45 s
Scheduler Delay	1 ms	3 ms	4 ms	5 ms	16 ms
Task Deserialization Time	1 ms	2 ms	3 ms	16 ms	35 ms
GC Time	0.1 s	0.8 s	2 s	3 s	14 s
Result Serialization Time	0 ms	0 ms	0 ms	0 ms	1 ms
Getting Result Time	0 ms				
Peak Execution Memory	0.0 B				
Shuffle Read Size / Records	128.6 MB / 17139979	128.7 MB / 17149116	128.7 MB / 17151913	128.7 MB / 17155092	128.8 MB / 17164081
Shuffle Write Size / Records	1403.0 B / 19	1477.0 B / 20	1478.0 B / 20	1479.0 B / 20	1558.0 B / 21
Shuffle spill (memory)	192.0 MB	384.0 MB	384.0 MB	384.0 MB	736.0 MB
Shuffle spill (disk)	66.0 MB	97.9 MB	98.0 MB	98.0 MB	98.0 MB

Varying task times and data volume has been fixed!



# After Skew Fix

Summary Metrics for 200 Completed Tasks

Metric	Min	25th percentile	Median	75th percentile	Max
Duration	14 s	16 s	20 s	25 s	45 s
Scheduler Delay	1 ms	3 ms	4 ms	5 ms	16 ms
Task Deserialization Time	1 ms	2 ms	3 ms	16 ms	35 ms
GC Time	0.1 s	0.8 s	2 s	3 s	14 s
Result Serialization Time	0 ms	0 ms	0 ms	0 ms	1 ms
Getting Result Time	0 ms				
Peak Execution Memory	0.0 B				
Shuffle Read Size / Records	128.6 MB / 17139979	128.7 MB / 17149116	128.7 MB / 17151913	128.7 MB / 17155092	128.8 MB / 17164081
Shuffle Write Size / Records	1403.0 B / 19	1477.0 B / 20	1478.0 B / 20	1479.0 B / 20	1558.0 B / 21
Shuffle spill (memory)	192.0 MB	384.0 MB	384.0 MB	384.0 MB	736.0 MB
Shuffle spill (disk)	66.0 MB	97.9 MB	98.0 MB	98.0 MB	98.0 MB

Varying task times and data volume has been fixed!



# After Skew Fix

Summary Metrics for 200 Completed Tasks

Metric	Min	25th percentile	Median	75th percentile	Max
Duration	14 s	16 s	20 s	25 s	45 s
Scheduler Delay	1 ms	3 ms	4 ms	5 ms	16 ms
Task Deserialization Time	1 ms	2 ms	3 ms	16 ms	35 ms
GC Time	0.1 s	0.8 s	2 s	3 s	14 s
Result Serialization Time	0 ms	0 ms	0 ms	0 ms	1 ms
Getting Result Time	0 ms				
Peak Execution Memory	0.0 B				
Shuffle Read Size / Records	128.6 MB / 17139979	128.7 MB / 17149116	128.7 MB / 17151913	128.7 MB / 17155092	128.8 MB / 17164081
Shuffle Write Size / Records	1403.0 B / 19	1477.0 B / 20	1478.0 B / 20	1479.0 B / 20	1558.0 B / 21
Shuffle spill (memory)	192.0 MB	384.0 MB	384.0 MB	384.0 MB	736.0 MB
Shuffle spill (disk)	66.0 MB	97.9 MB	98.0 MB	98.0 MB	98.0 MB

Varying task times and data volume has been fixed!



# After Skew Fix

## Details for Stage 13 (Attempt 0)

**Total Time Across All Tasks:** 1.2 h

**Locality Level Summary:** Process local: 200

**Shuffle Read:** 25.1 GB / 3430413322

**Shuffle Write:** 287.4 KB / 3982

**Shuffle Spill (Memory):** 73.6 GB

**Shuffle Spill (Disk):** 18.7 GB

Still some spill of data though :(.

# Goal 2

Reducing the shuffle spilt to disk.

Motivation:

“When the records destined for these aggregation operations do not easily fit in memory, some mayhem can ensue. First, holding many records in these data structures puts pressure on garbage collection, which can lead to pauses down the line. Second, when the records do not fit in memory, Spark will spill them to disk, which causes disk I/O and sorting. This overhead during large shuffles is probably the number one cause of job stalls I have seen at Cloudera customers.”

Reference: <https://blog.cloudera.com/blog/2015/03/how-to-tune-your-apache-spark-jobs-part-2/>

# Tasks

- The smallest unit of transformation(s) applied on a unit of data.
- Tasks run on executors within a stage.
- Number of tasks per stage is determined by various settings.



# Partitions & Tasks

Number of tasks for a given stage can be:

`spark.sql.shuffle.partitions = 200`

1. If partition number is not tweaked:
  - If stage is reading input data:
    - Number of tasks  $\approx$  total number of “blocks” in all the files.
  - Else:
    - Default value of 200 partitions / stage.
2. If partition number is tweaked:
  - Partitions passed to the repartition / coalesce function.
  - Min(cardinality of a dataset, number of partitions argument).



# Partitions & Tasks

## Adding More Partitions / Tasks:

- Smaller manageable amounts of data to process at a time.
- Data (hopefully) fit in memory on the executor.

## DISCLAIMER:

- More tasks also means more load on the driver.

## RED FLAGS:

- Lots of shuffle spill.
- Driver OOMs → too many partitions / tasks.

Review:

Tasks in a stage are processed sequentially by executors.



# Choosing an Optimal Number of Tasks

For the most optimal setup, you want to dynamically change the number of partitions on a stage level, as it will differ stage by stage.

But for simplicity sacks we will find an overarching number for the entire job.

Steps:

1. Find the non read input stage with shuffle spill data.
2. Rule of thumb, ideal task data size is **128MB**.
3. Add `spark.sql.shuffle.partitions: ideal_number` to schedule file.



# Calculating Ideal Task Number

(1) Memory available for a given task:

$$\frac{\text{Java memory per executor} * \text{margin of error}}{\text{cores per executor}}$$

$$\text{Margin of error} = 0.8 * 0.2 = 0.16$$

(2) In memory size of shuffle data:

$$\frac{\text{Shuffle spill (memory)} * \text{Shuffle write}}{\text{Shuffle spill (disk)}}$$

reference: <https://blog.cloudera.com/blog/2015/03/how-to-tune-your-apache-spark-jobs-part-2/>

# Calculating Ideal Task Number

Putting the 2 equations together:

Shuffle spill (memory) \* Shuffle write \* cores per executor  
Shuffle spill (disk) \* java memory per executor \* margin of error

reference: <https://blog.cloudera.com/blog/2015/03/how-to-tune-your-apache-spark-jobs-part-2/>



# All else Fails...

but.... This equation is not bullet proof..

If this equation gives a smaller estimate than the current number of tasks, then we can only guess and check. Increase the number of tasks by 1.5x incrementally until no shuffle spill is observed.



# After Skew Fix

## Details for Stage 13 (Attempt 0)

**Total Time Across All Tasks:** 1.2 h

**Locality Level Summary:** Process local: 200

**Shuffle Read:** 25.1 GB / 3430413322

**Shuffle Write:** 287.4 KB / 3982

**Shuffle Spill (Memory):** 73.6 GB

**Shuffle Spill (Disk):** 18.7 GB

Notebooks resource class:

- 8 executors
- 12GB java memory

# After Skew Fix

Shuffle spill (memory) \* Shuffle write \* cores per executor  
Shuffle spill (disk) \* java memory per executor \* margin of error

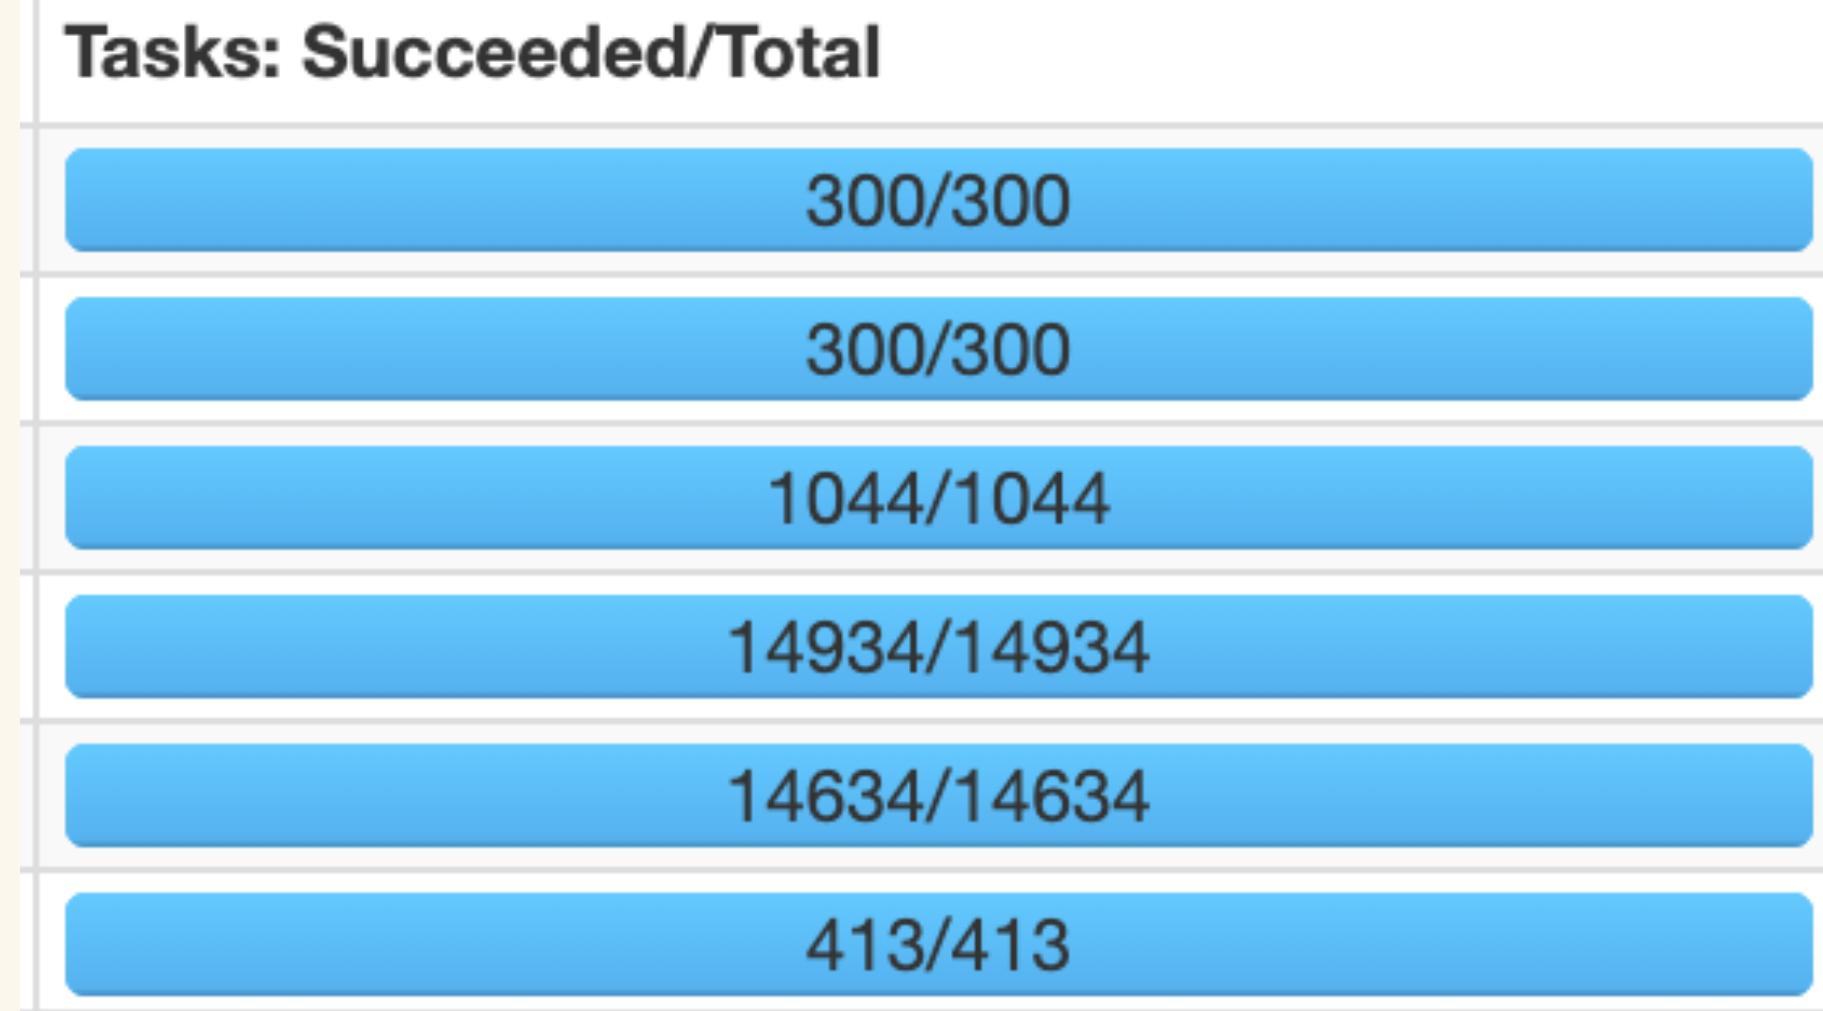
$$\frac{(73.6 \text{ GB} * 287.4 \text{ KB} * 8)}{(18.7 \text{ GB} * 12 \text{ GB} * 0.8)}$$

0.00094263101 partitions...

So let's try  $200 * 1.5 = 300$  partitions.



# Results



## Details for Stage 19 (Attempt 0)

**Total Time Across All Tasks:** 54 min

**Locality Level Summary:** Process local: 300

**Shuffle Read:** 25.0 GB / 3435988435

**Shuffle Write:** 425.5 KB / 5896

▶ [DAG Visualization](#)



# **Section 5: Recap**

# Recap

1. Understood how a Spark DataFrames application got executed.
2. Evaluated the Spark application by analyzing the SparkUI.
3. Identified bottlenecks.
4. Implemented optimizations to address the bottlenecks.



# Questions?

Thank You

